



GAM401

Riot Games: Standardizing Application Deployments Using Amazon ECS and Terraform

Adam Rozumek

November 28, 2016

RIOT GAMES | GAM 401

**Standardizing Application
Deployments Using Amazon ECS
and Terraform**



INTRODUCTIONS

WHO AM I?



ADAM ROZUMEK
SYSTEMS ENGINEER

WHAT TO EXPECT

AMAZON ECS. TERRAFORM. VIDEO GAMES.

- 1 ECS CONSOLIDATION WINS & LESSONS
Adapting existing deployments & infrastructure
- 2 TERRAFORM ENCAPSULATION STRATS
Object Oriented Development Operations
- 3 MULTI-GAME MODULAR SERVICE DESIGN
A different kind of scaling



WHAT TO EXPECT

AMAZON ECS. TERRAFORM. VIDEO GAMES.

- 1 ECS CONSOLIDATION WINS & LESSONS
Adapting existing deployments & infrastructure
- 2 TERRAFORM ENCAPSULATION STRATS
Object Oriented Development Operations
- 3 MULTI-GAME MODULAR SERVICE DESIGN
A different kind of scaling



WHAT TO EXPECT

AMAZON ECS. TERRAFORM. VIDEO GAMES.

- 1 ECS CONSOLIDATION WINS & LESSONS
Adapting existing deployments & infrastructure
- 2 TERRAFORM ENCAPSULATION STRATS
Object Oriented Development Operations
- 3 MULTI-GAME MODULAR SERVICE DESIGN
A different kind of scaling



WHAT TO EXPECT

AMAZON ECS. TERRAFORM. VIDEO GAMES.

- 1 ECS CONSOLIDATION WINS & LESSONS
Adapting existing deployments & infrastructure
- 2 TERRAFORM ENCAPSULATION STRATS
Object Oriented Development Operations
- 3 MULTI-GAME MODULAR SERVICE DESIGN
A different kind of scaling





LEAGUE OF
LEGENDS



2016 LEAGUE OF LEGENDS STATS



MORE THAN
100 MILLION

MONTHLY ACTIVE
PLAYERS



MORE THAN
27 MILLION

DAILY ACTIVE
PLAYERS



MORE THAN
7.5 MILLION

PEAK CONCURRENT
PLAYERS

DATA PRODUCTS & SERVICES

OUR MISSION

Empower teams at Riot to make timely, data-informed products by maintaining a scalable and reliable data platform



PROBLEM



PROBLEM

SCALING TOTAL OWNERSHIP ON AWS

Total ownership

We want to **empower** developers to:

- Provision their own infrastructure
- Execute their own deployments
- Monitor their own metrics



PROBLEM SCALING TOTAL OWNERSHIP ON AWS

Resource attribution

- Who owns these EBS volumes?
- What applications depend on these security groups?
- Can these AMIs be deleted?



A screenshot of an AWS EBS volume list. The search bar shows "Name : Not tagged". Below it, the list displays 1 to 50 of 1,491 volumes. The columns are labeled "Name", "Volume ID", "Size", and "Volume Type".



PROBLEM SCALING TOTAL OWNERSHIP ON AWS

Security



Custom TCP Rule	TCP	50000	sg-2a598753 (vertica_bogdan_vertica_internal_traffic)
Custom TCP Rule	TCP	5434	sg-2a598753 (vertica_bogdan_vertica_internal_traffic)
Custom ICMP Rule	Traceroute	N/A	sg-2a598753 (vertica_bogdan_vertica_internal_traffic)
DNS (UDP)	UDP	53	10.93.0.0/22
Custom TCP Rule	TCP	4803 - 4805	sg-2a598753 (vertica_bogdan_vertica_internal_traffic)
Custom TCP Rule	TCP	544	sg-2a598753 (vertica_bogdan_vertica_internal_traffic)
Custom ICMP Rule	Echo Reply	N/A	sg-2a598753 (vertica_bogdan_vertica_internal_traffic)
All traffic	All	All	0.0.0.0/0

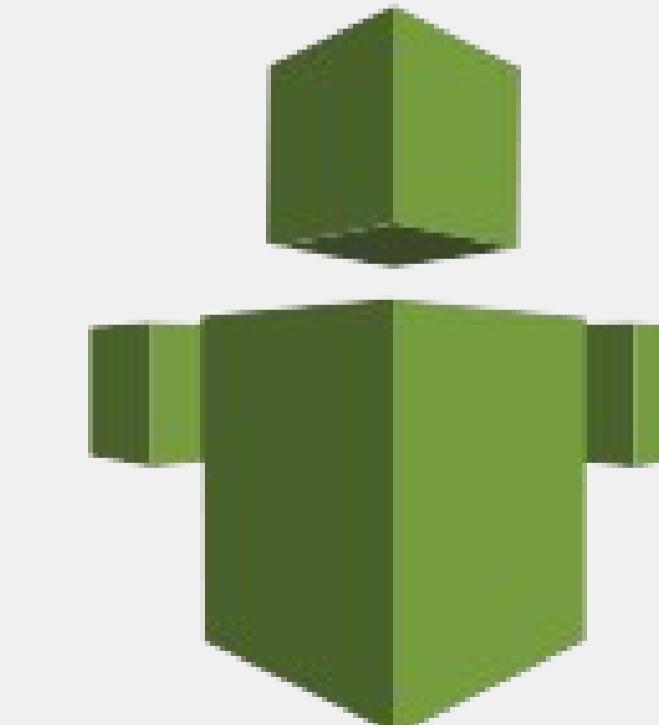
PROBLEM SCALING TOTAL OWNERSHIP ON AWS

Security

- Auditing is important, but it's reactive
- Operational time sink



Security Monkey



AWS Trusted Advisor

PROBLEM

SCALING TOTAL OWNERSHIP ON AWS

Enforcing conventions

- Common syntax for tags
- Simplify reservations
- Standardize networks

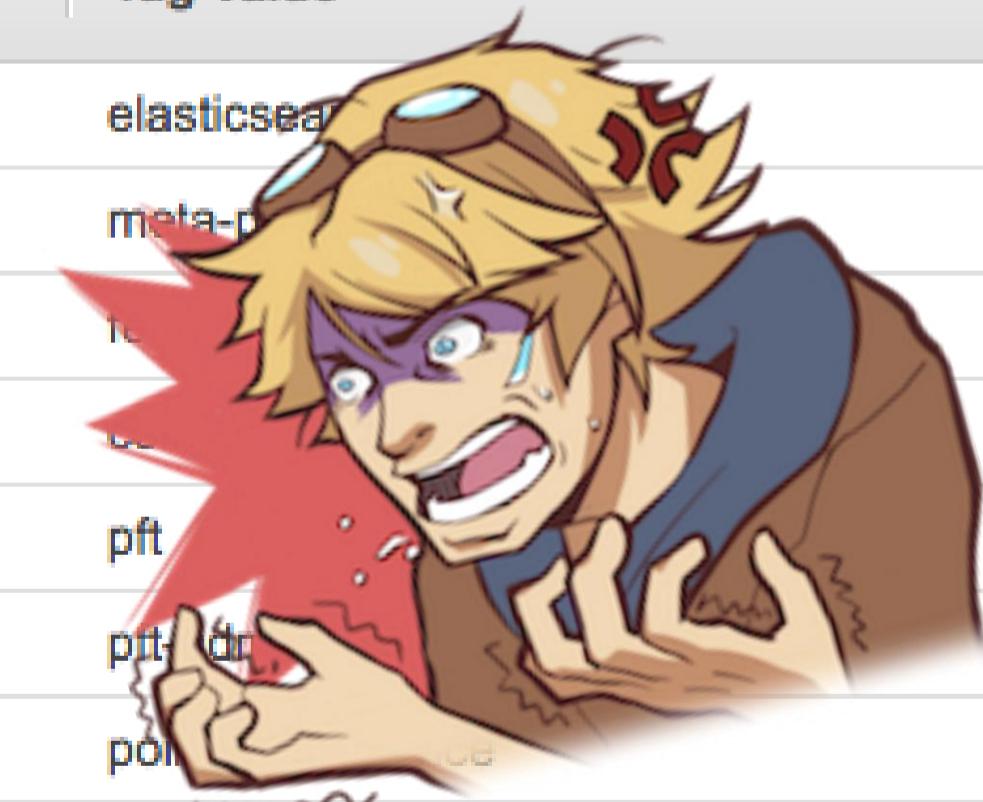
PROBLEM

SCALING TOTAL OWNERSHIP ON AWS

Enforcing conventions

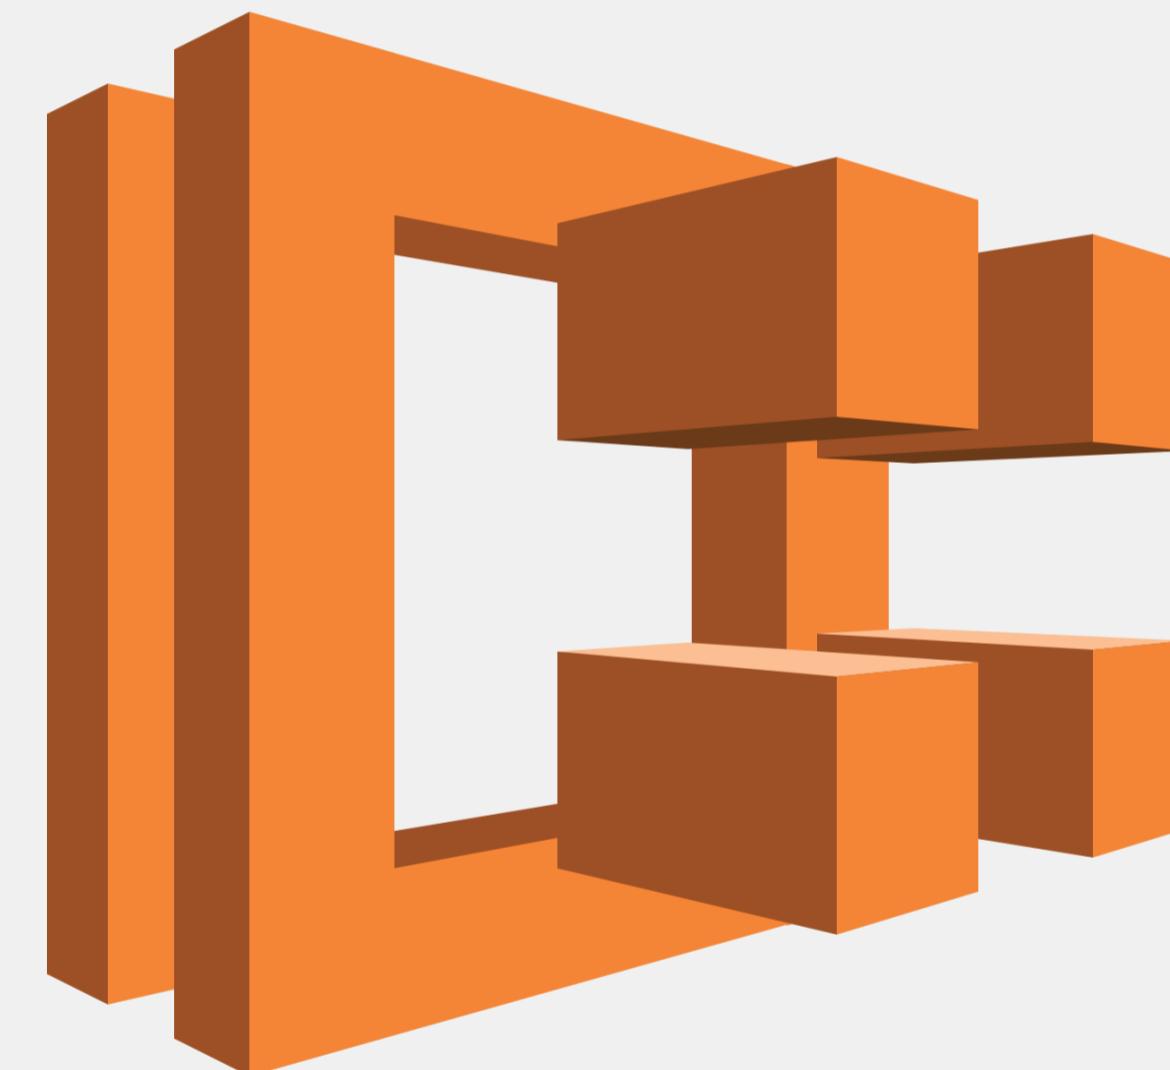
- Common syntax for tags
- Simplify reservations
- Standardize networks

	Tag Key	Tag Value	Total
Manage Tag	App	elasticsearch	33
Manage Tag	App	meta-data	10
Manage Tag	app	riot	1
Manage Tag	app	riot	1
Manage Tag	App_Name	pft	4
Manage Tag	App_Name	pft-dr	4
Manage Tag	application	polo	51
Manage Tag	AppName	legs-prod-proxy-west2	9



PROBLEM

SCALING TOTAL OWNERSHIP ON AWS



Amazon EC2
Container Service



Terraform



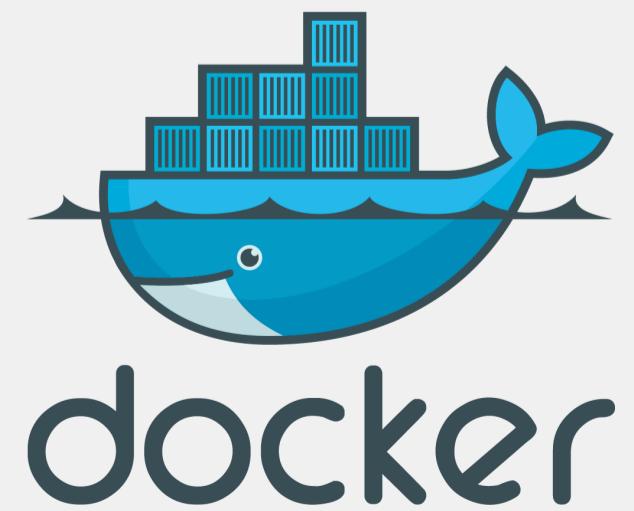
CONTAINERS



CONTAINERS

STANDARDIZED APPLICATION UNITS... ON AWS

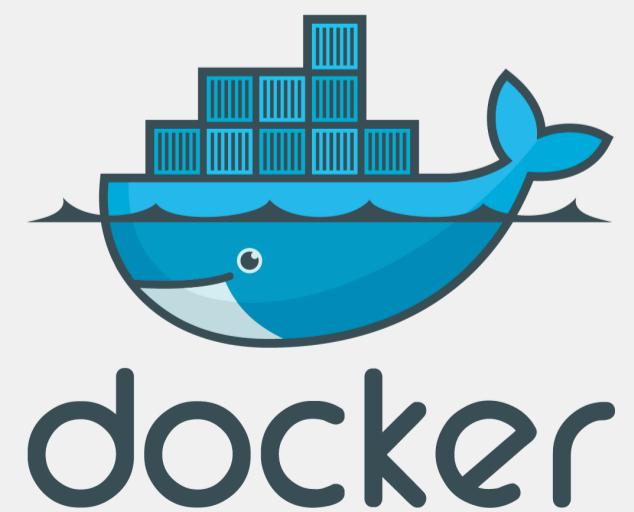
- Dockerfiles capture application dependencies



CONTAINERS

STANDARDIZED APPLICATION UNITS... ON AWS

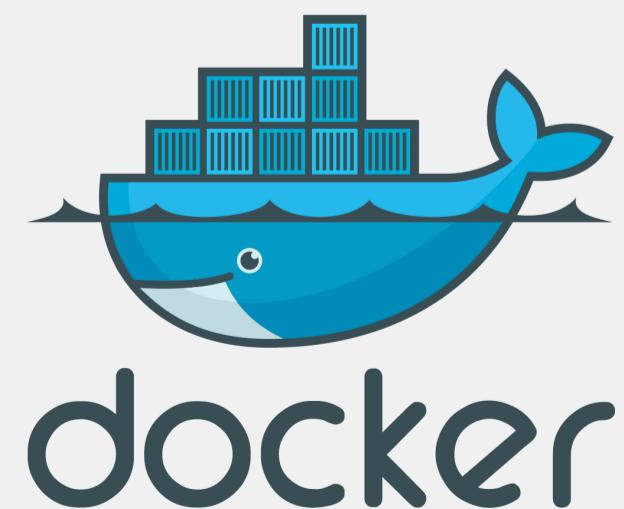
- Dockerfiles capture application dependencies
- Common use cases have great community support



CONTAINERS

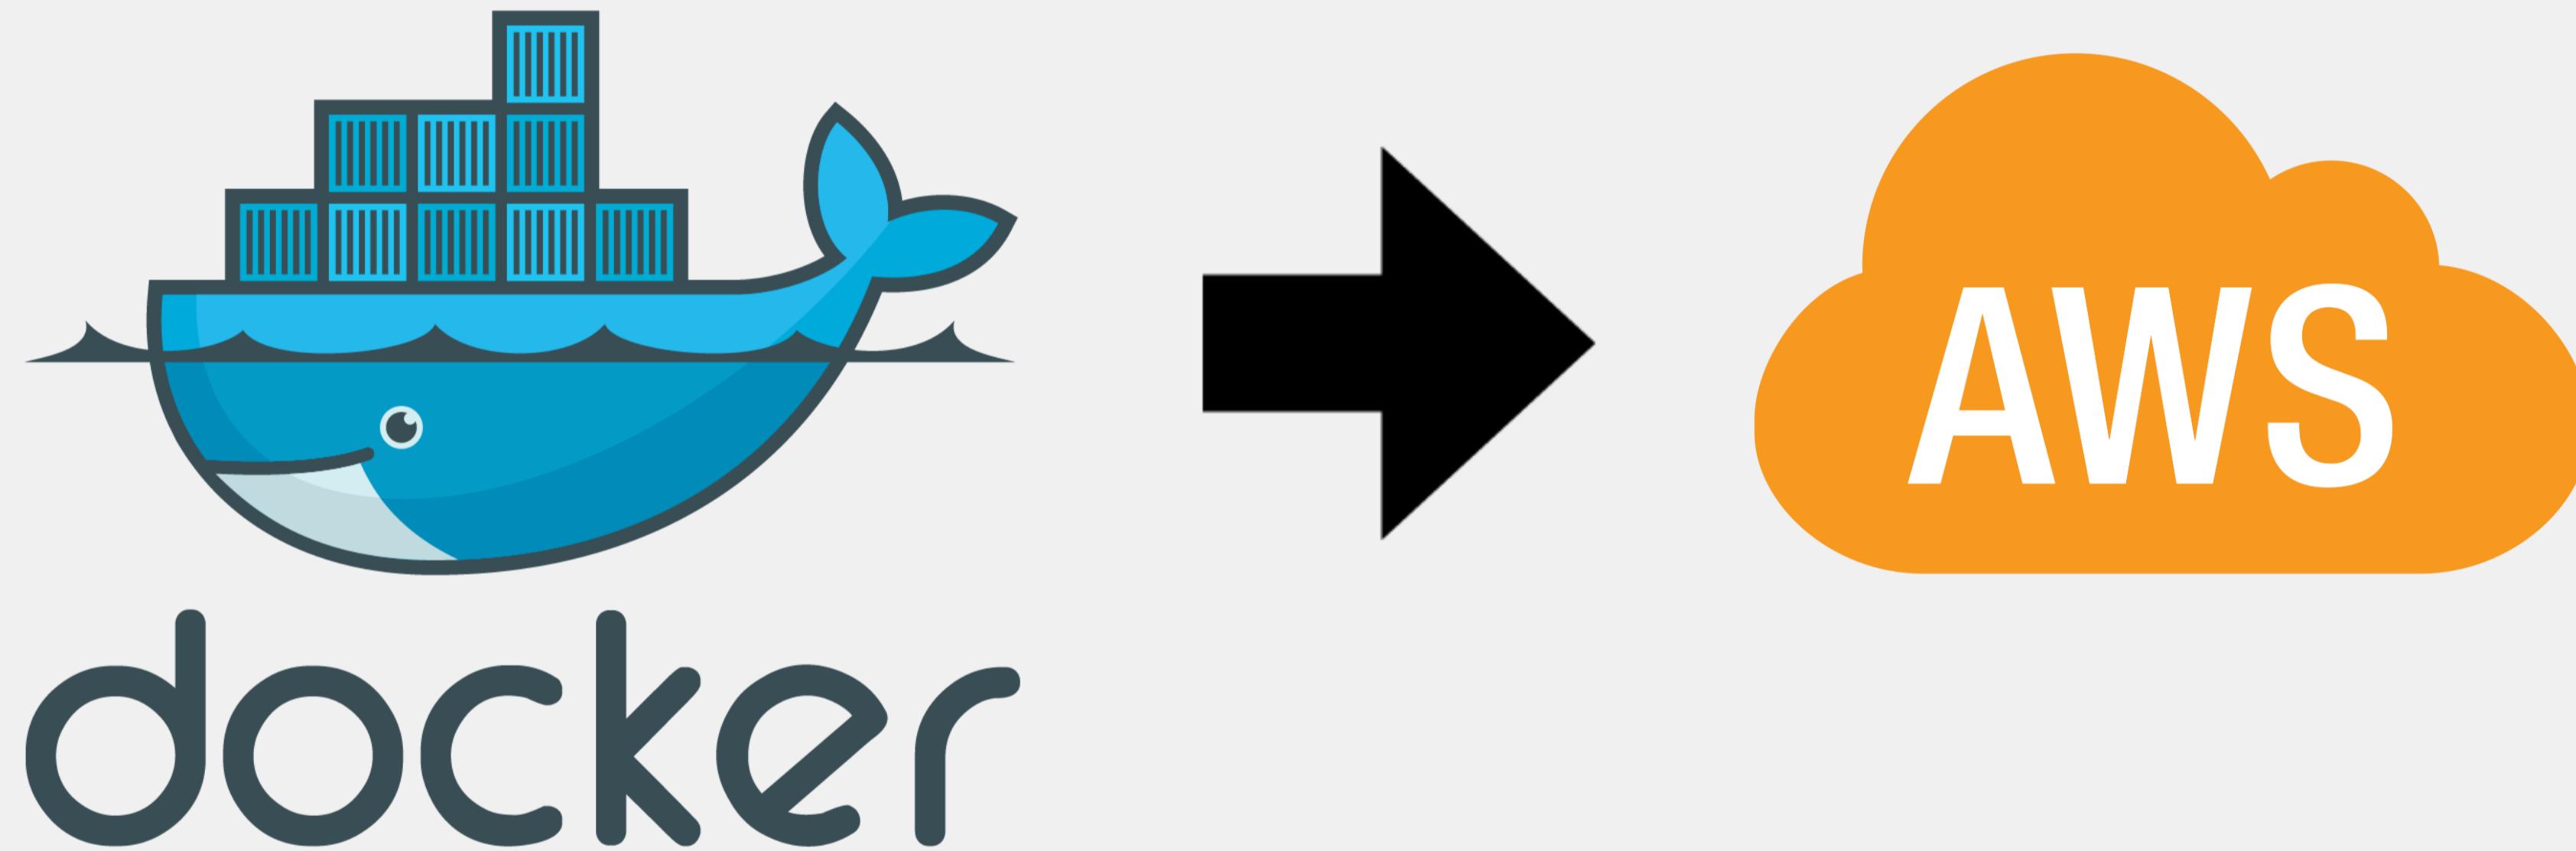
STANDARDIZED APPLICATION UNITS... ON AWS

- Dockerfiles capture application dependencies
- Common use cases have great community support
- Profit from our own engineering community



CONTAINERS

STANDARDIZED APPLICATION UNITS... ON AWS



CONTAINERS

STANDARDIZED APPLICATION UNITS... ON AWS

Embrace the abstraction

- Plan for failure at all levels
- Avoid manual intervention whenever possible
- It's ephemeral all the way down



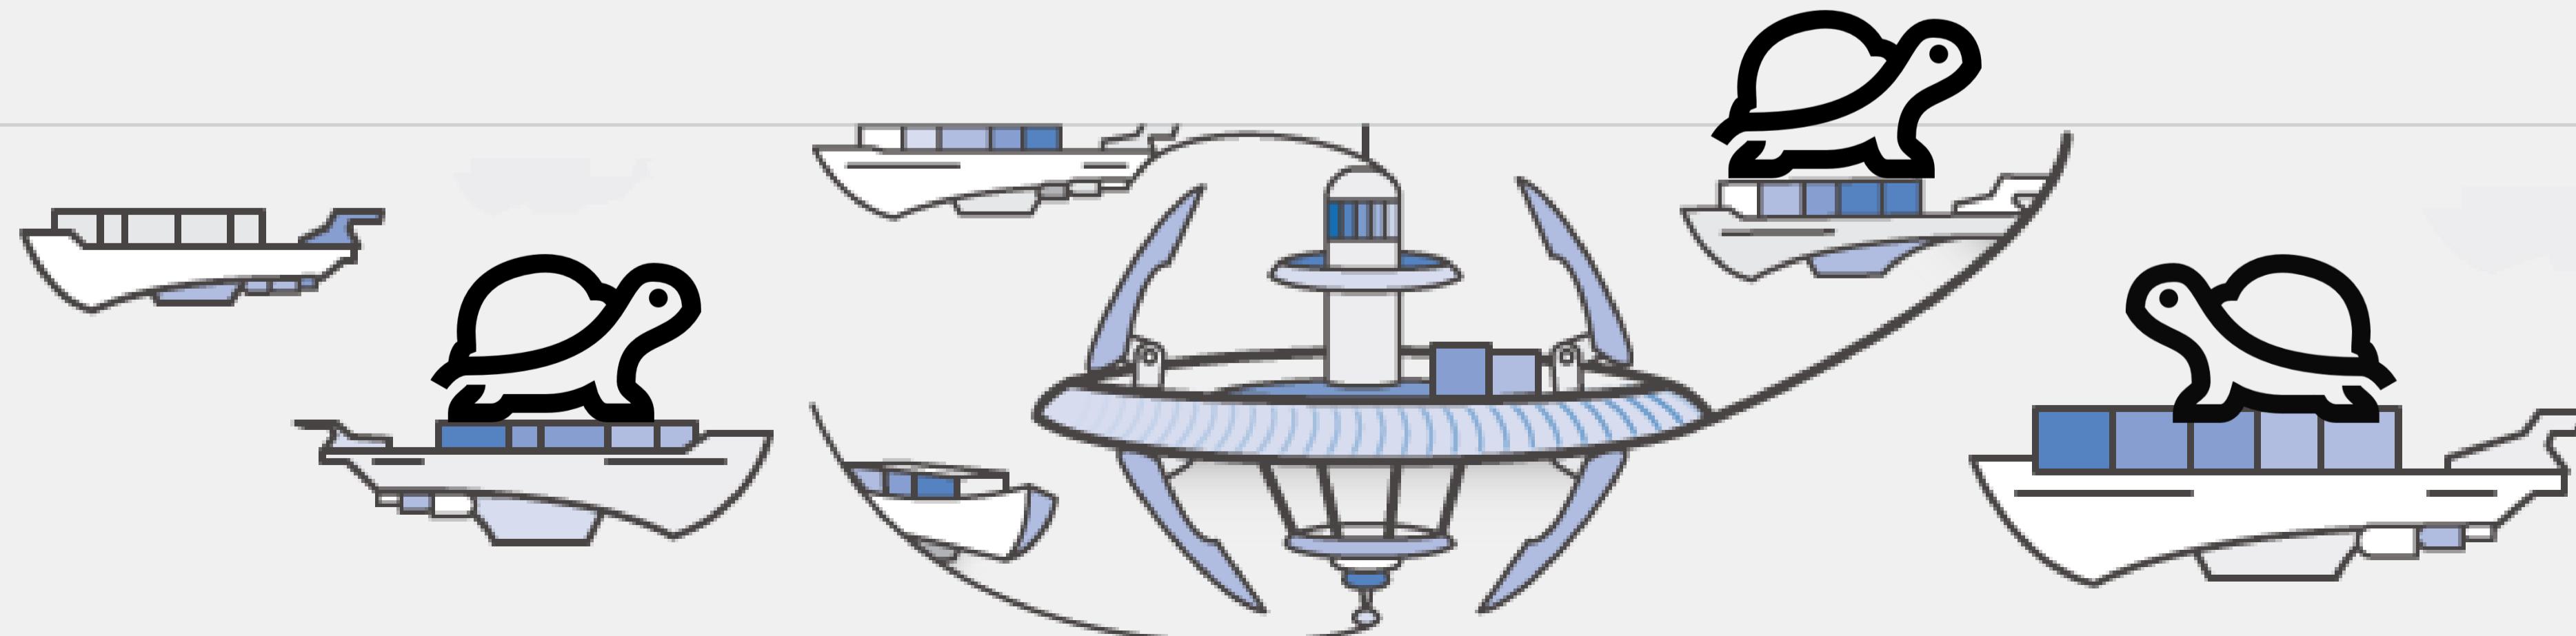
CONTAINERS

STANDARDIZED APPLICATION UNITS... ON AWS

Scheduling is hard

We need to:

- Quickly and “fairly” run tasks
- Prevent resource conflicts
- Provide reasonable fault tolerance

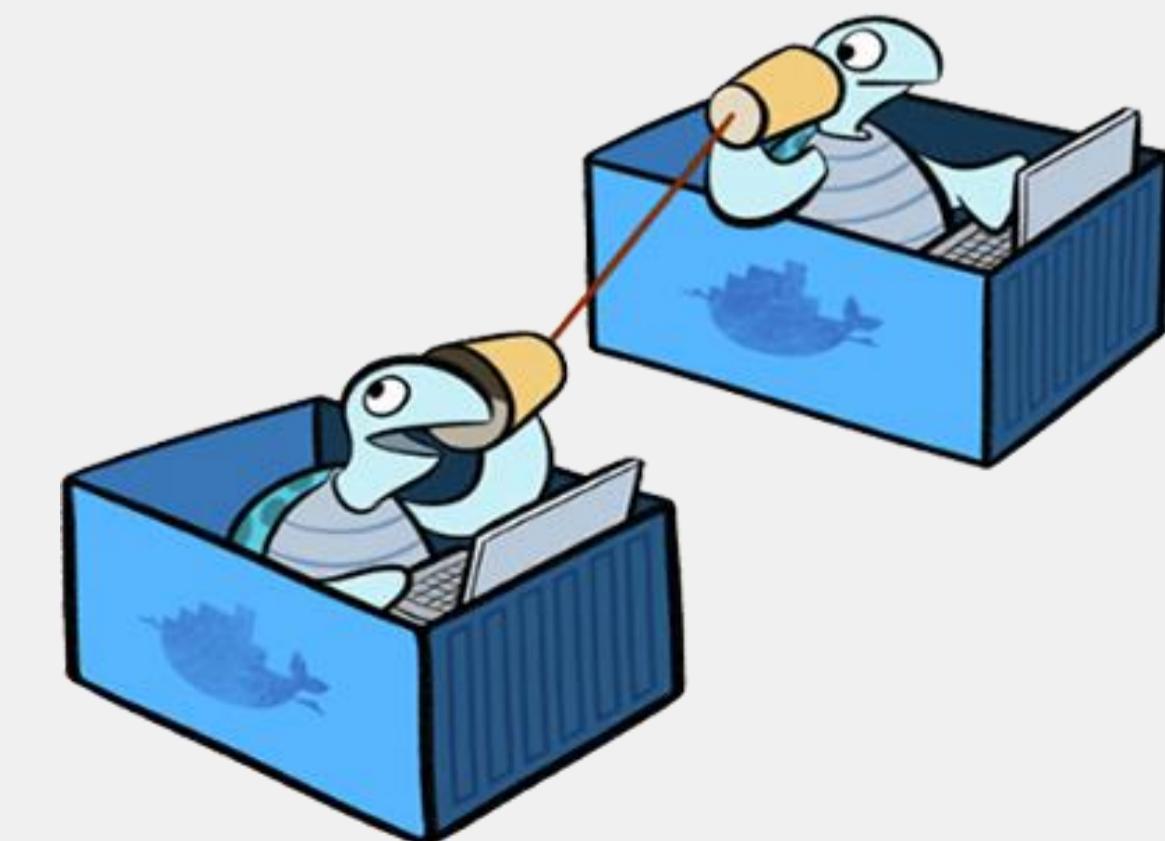


CONTAINERS

STANDARDIZED APPLICATION UNITS... ON AWS

We need AWS hardware

- Enable total ownership of the AWS resources backing our containers
- Avoid the security, resource attribution, and convention degradation pitfalls





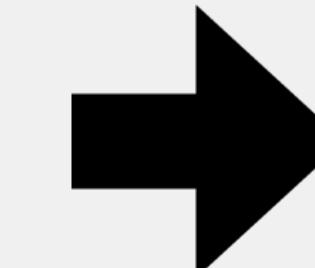
HISTORY

CONTAINERS

STANDARDIZED APPLICATION UNITS... ON AWS

First iteration (2013)

- Traditional ASGs
- AMIs configured to launch a single Docker container
- Managed by Netflix's Asgard



CONTAINERS

STANDARDIZED APPLICATION UNITS... ON AWS

First iteration (2013)

- Easy to adopt
- Familiar AWS concepts



CONTAINERS

STANDARDIZED APPLICATION UNITS... ON AWS

First iteration (2013)

- No increased resource utilization
- No elasticity improvements

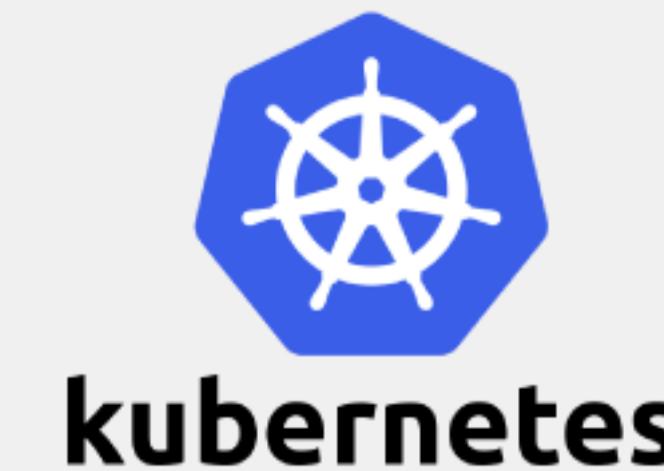


CONTAINERS

STANDARDIZED APPLICATION UNITS... ON AWS

Second iteration (2013)

- Third party container orchestration
- Production scale had unexpected challenges

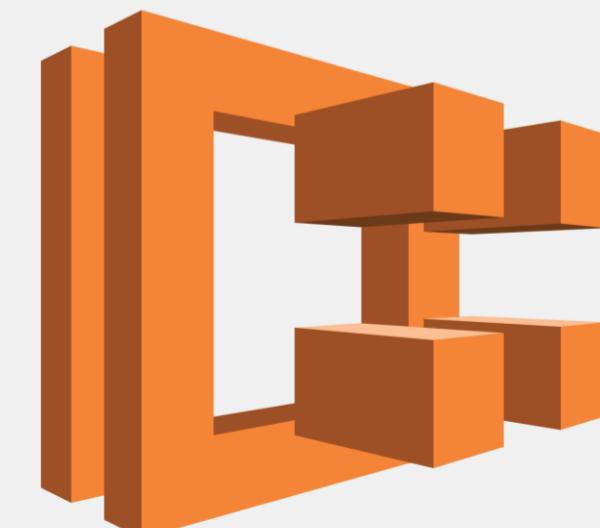


CONTAINERS

STANDARDIZED APPLICATION UNITS... ON ECS!

Third iteration: Amazon EC2 Container Service

- ECS AMI provides all necessary software
- Designed with AWS integration in mind
- Free!

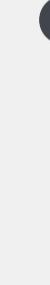


AMAZON EC2 CONTAINER SERVICE

KEY FEATURE TIMELINE

ECS ANNOUNCED

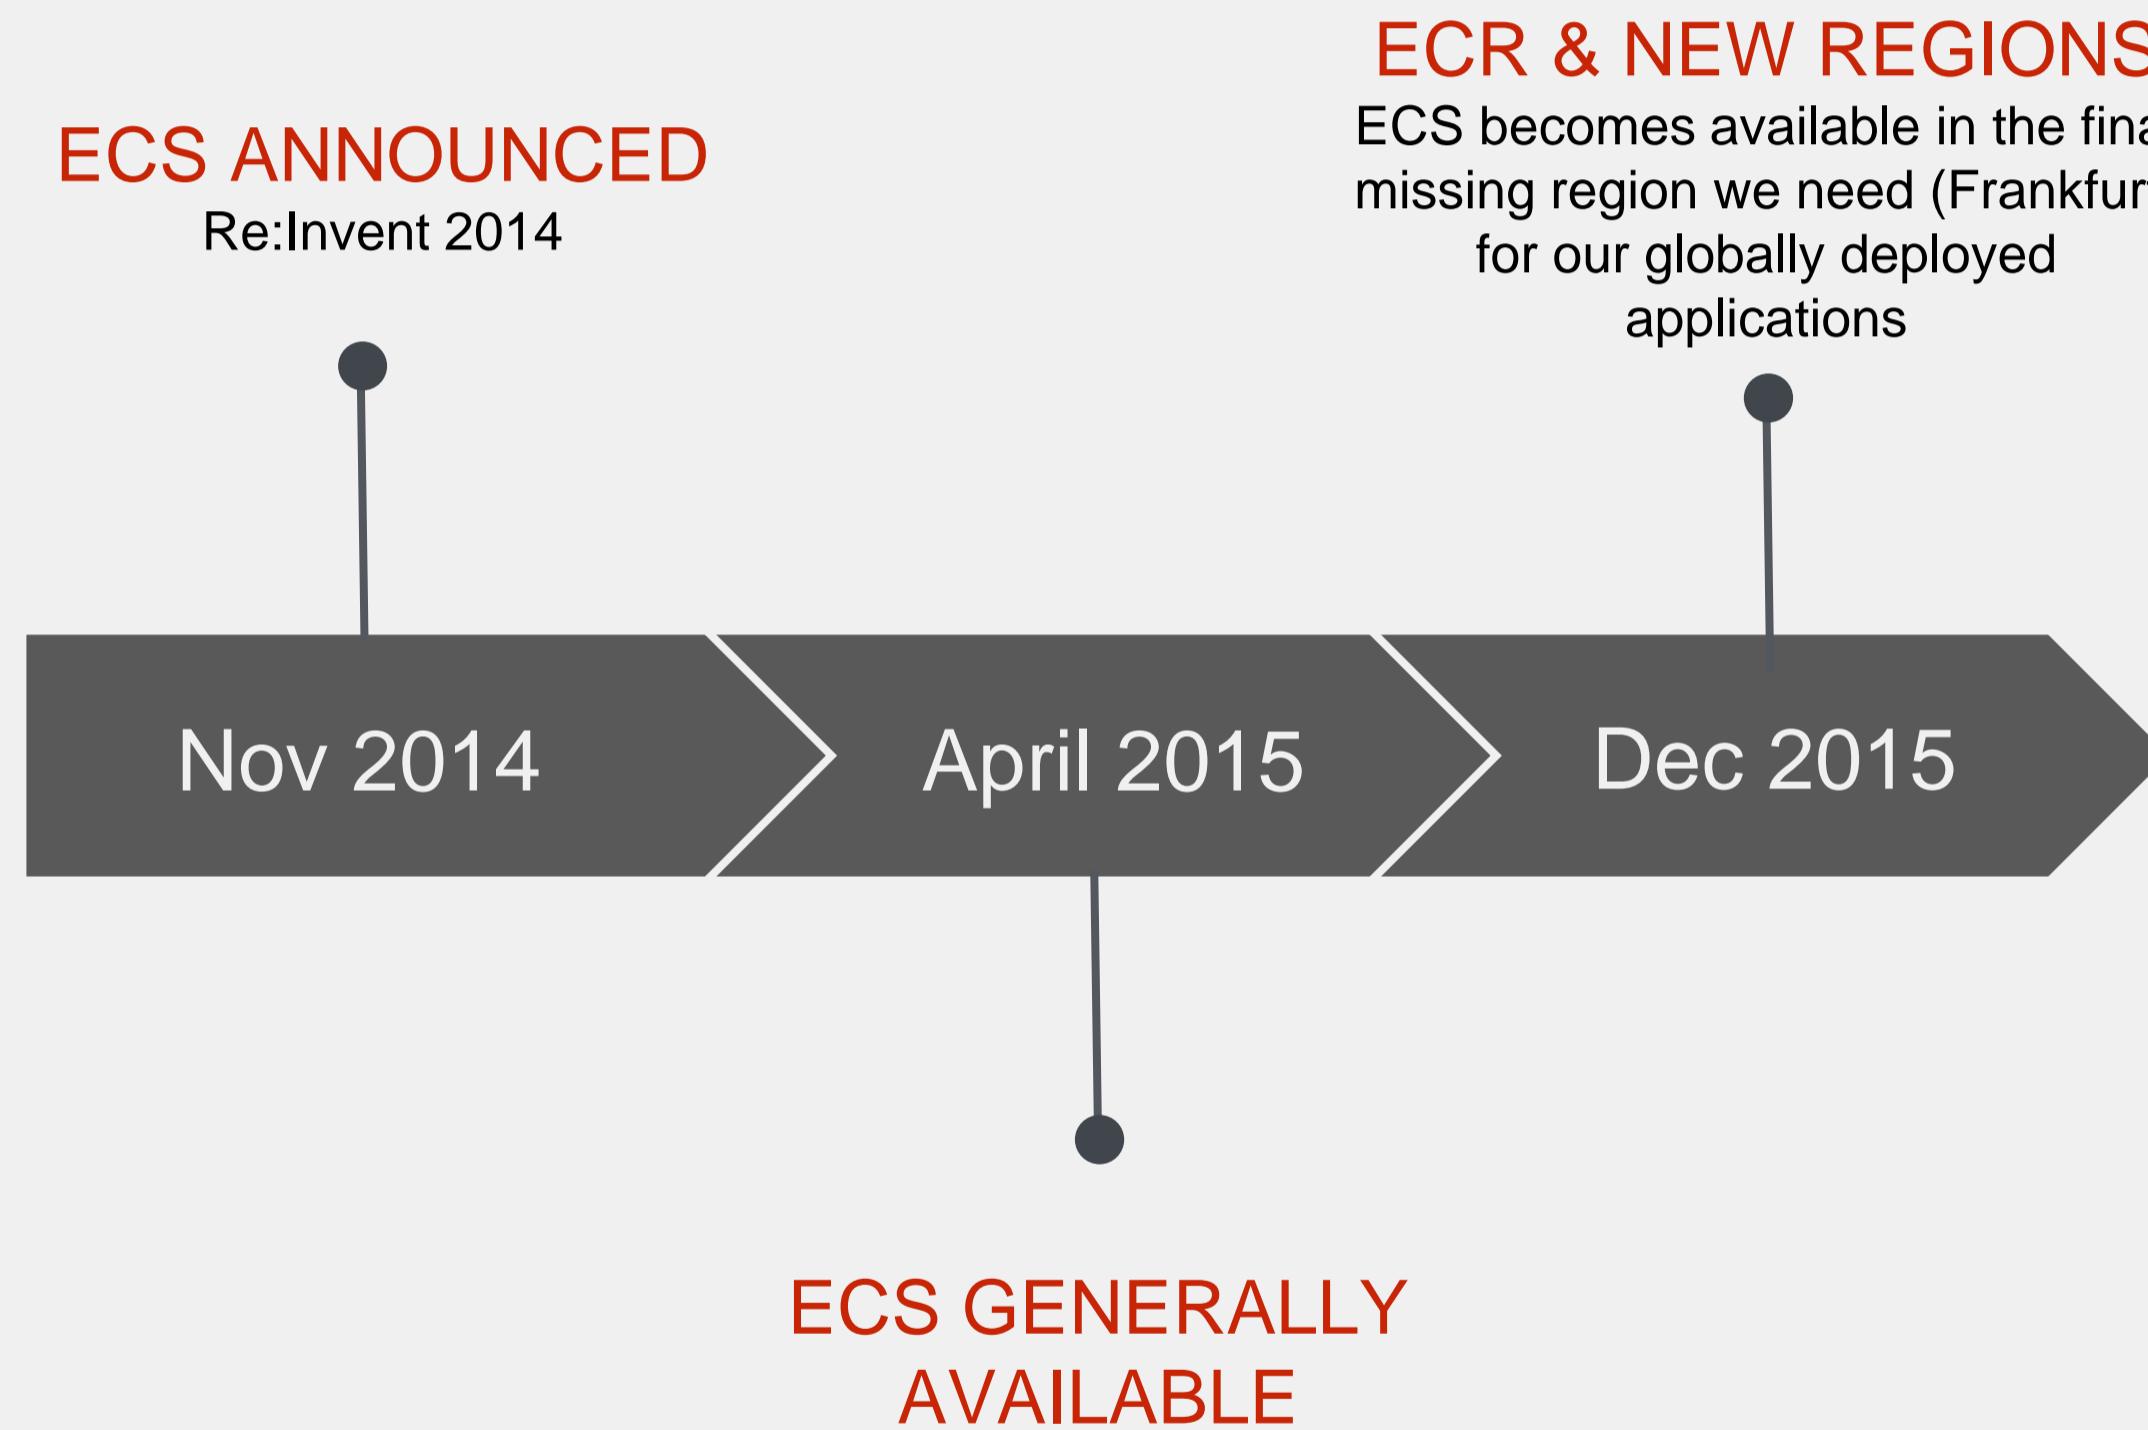
Re:Invent 2014



Nov 2014

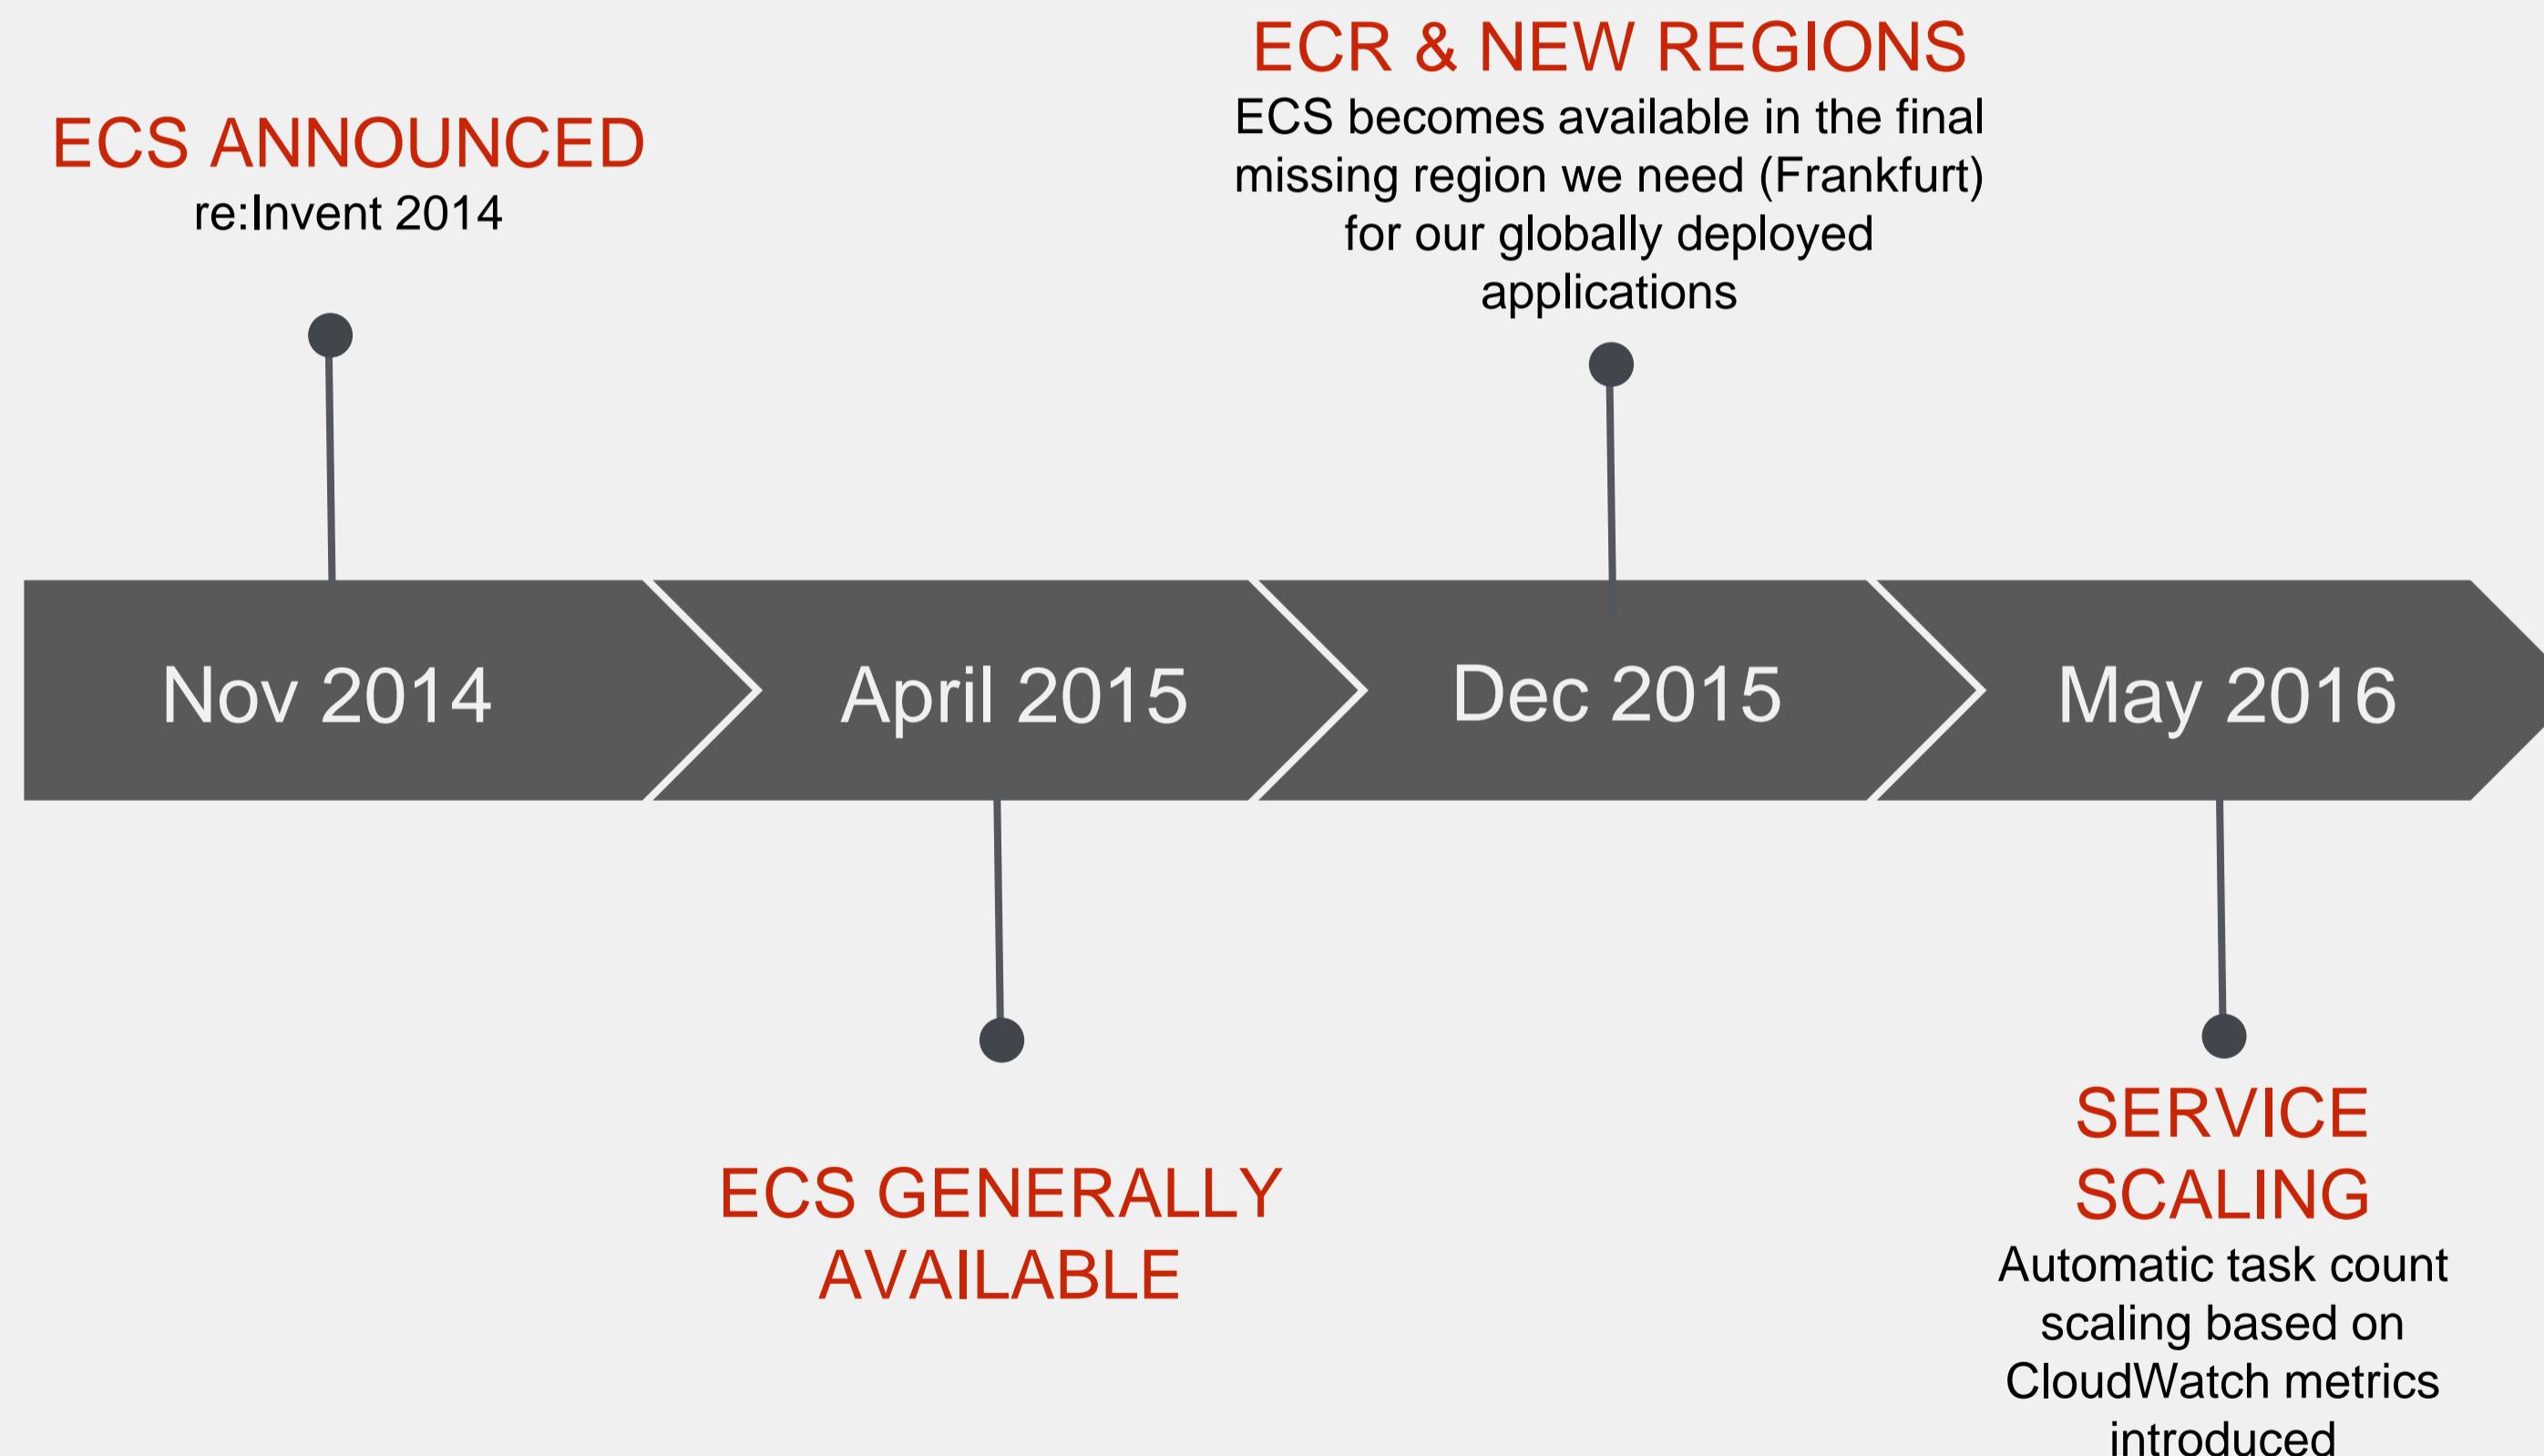
AMAZON EC2 CONTAINER SERVICE

KEY FEATURE TIMELINE



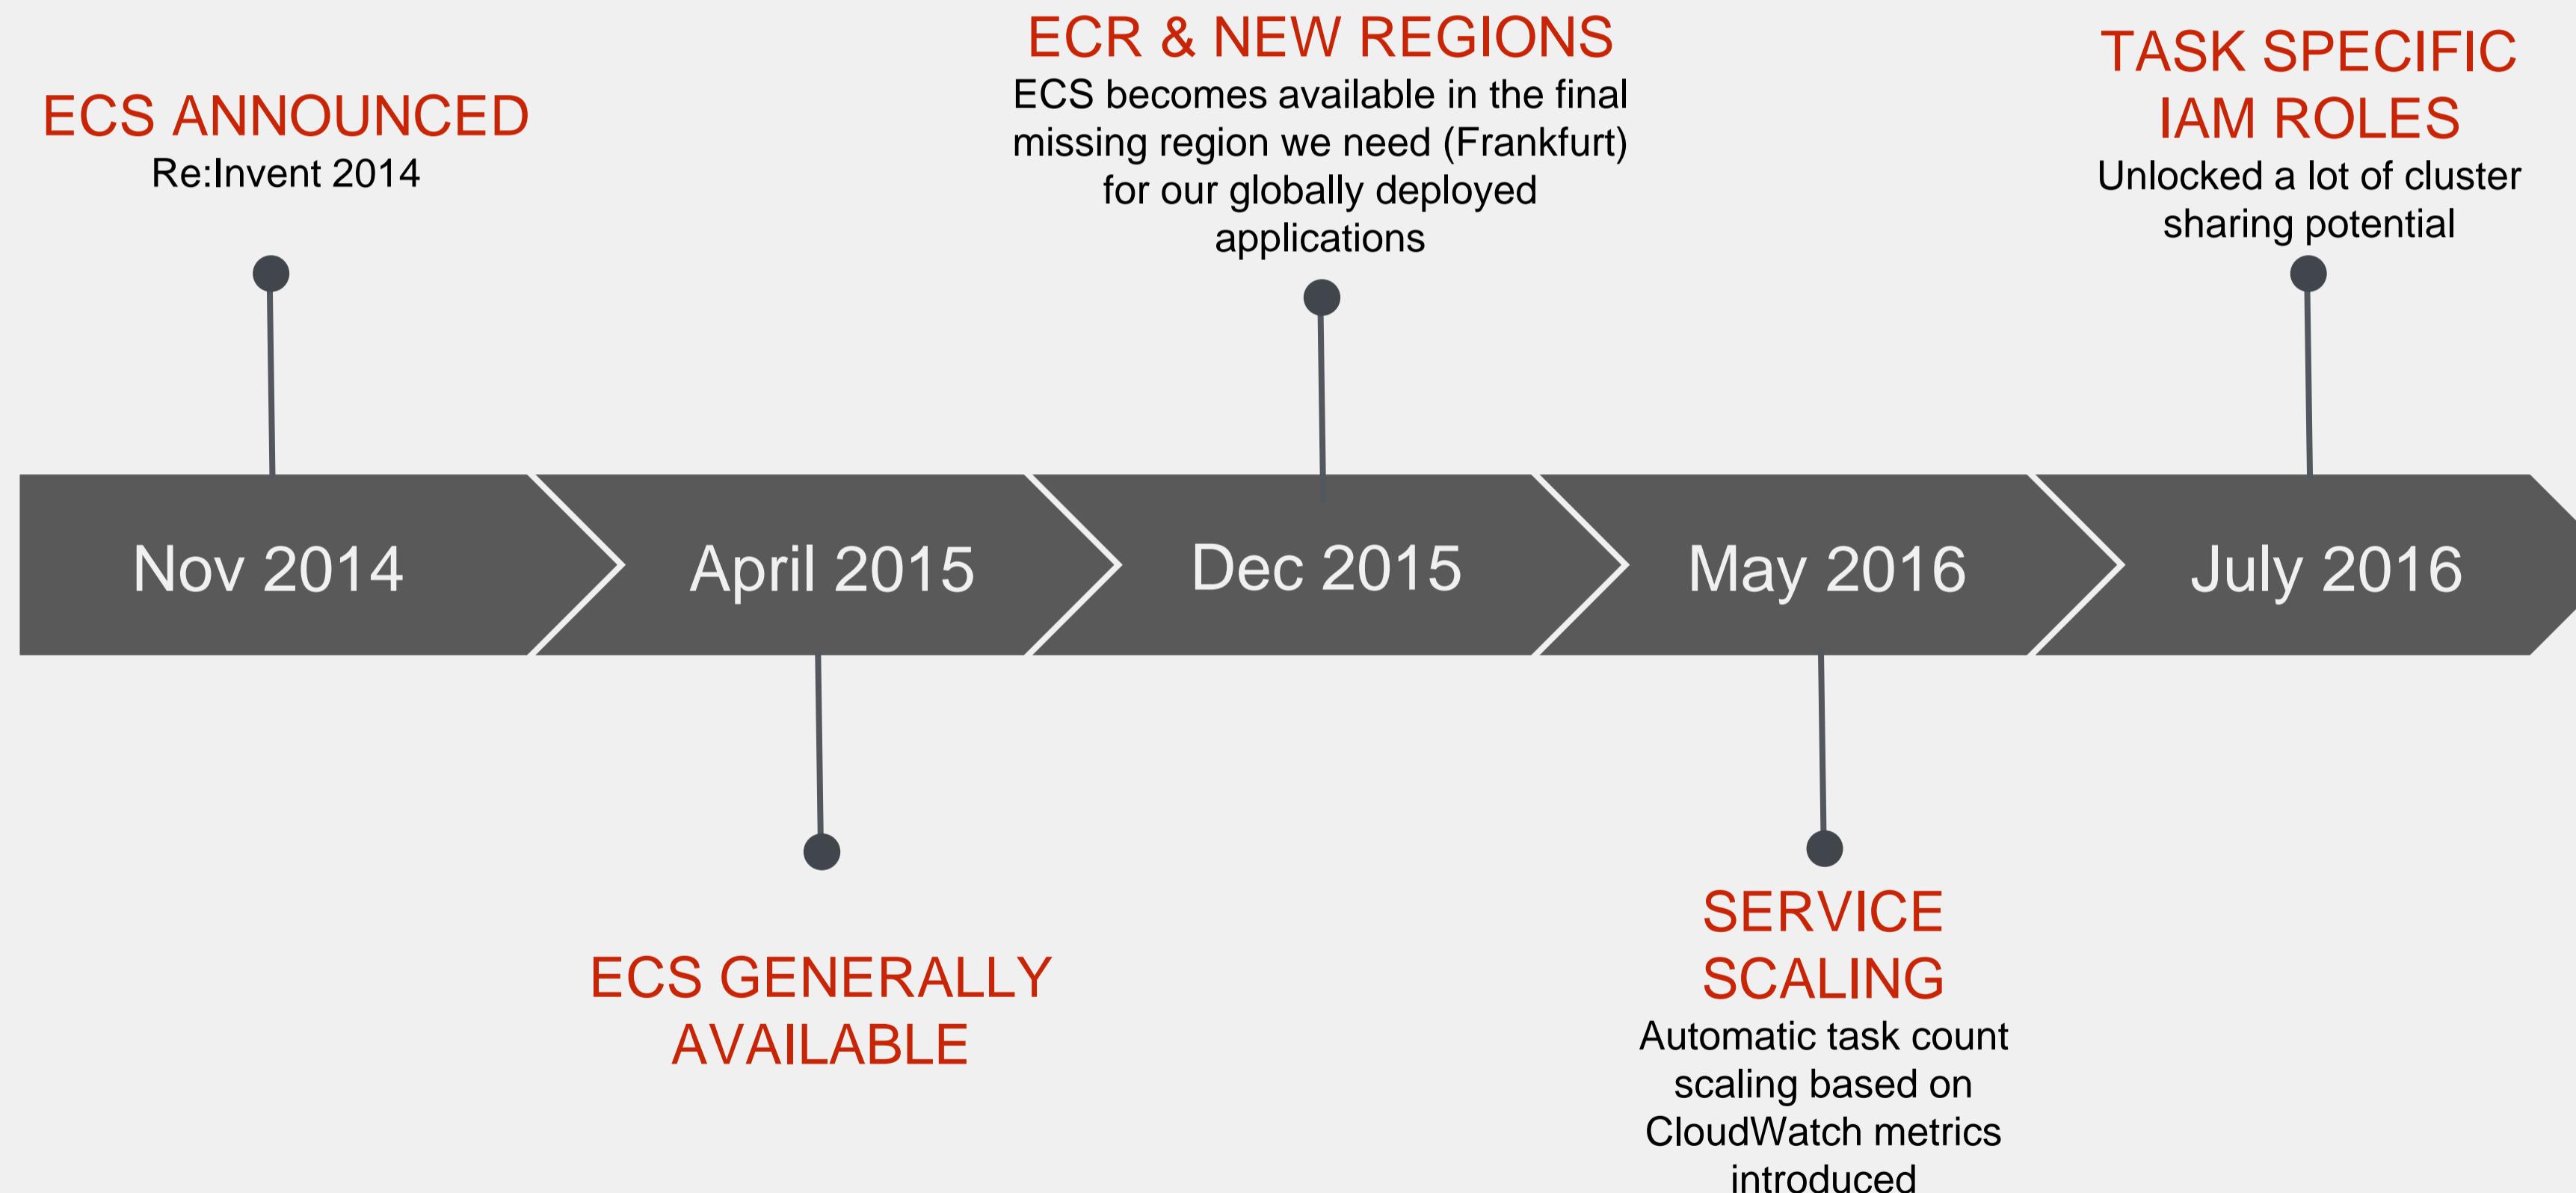
AMAZON EC2 CONTAINER SERVICE

KEY FEATURE TIMELINE



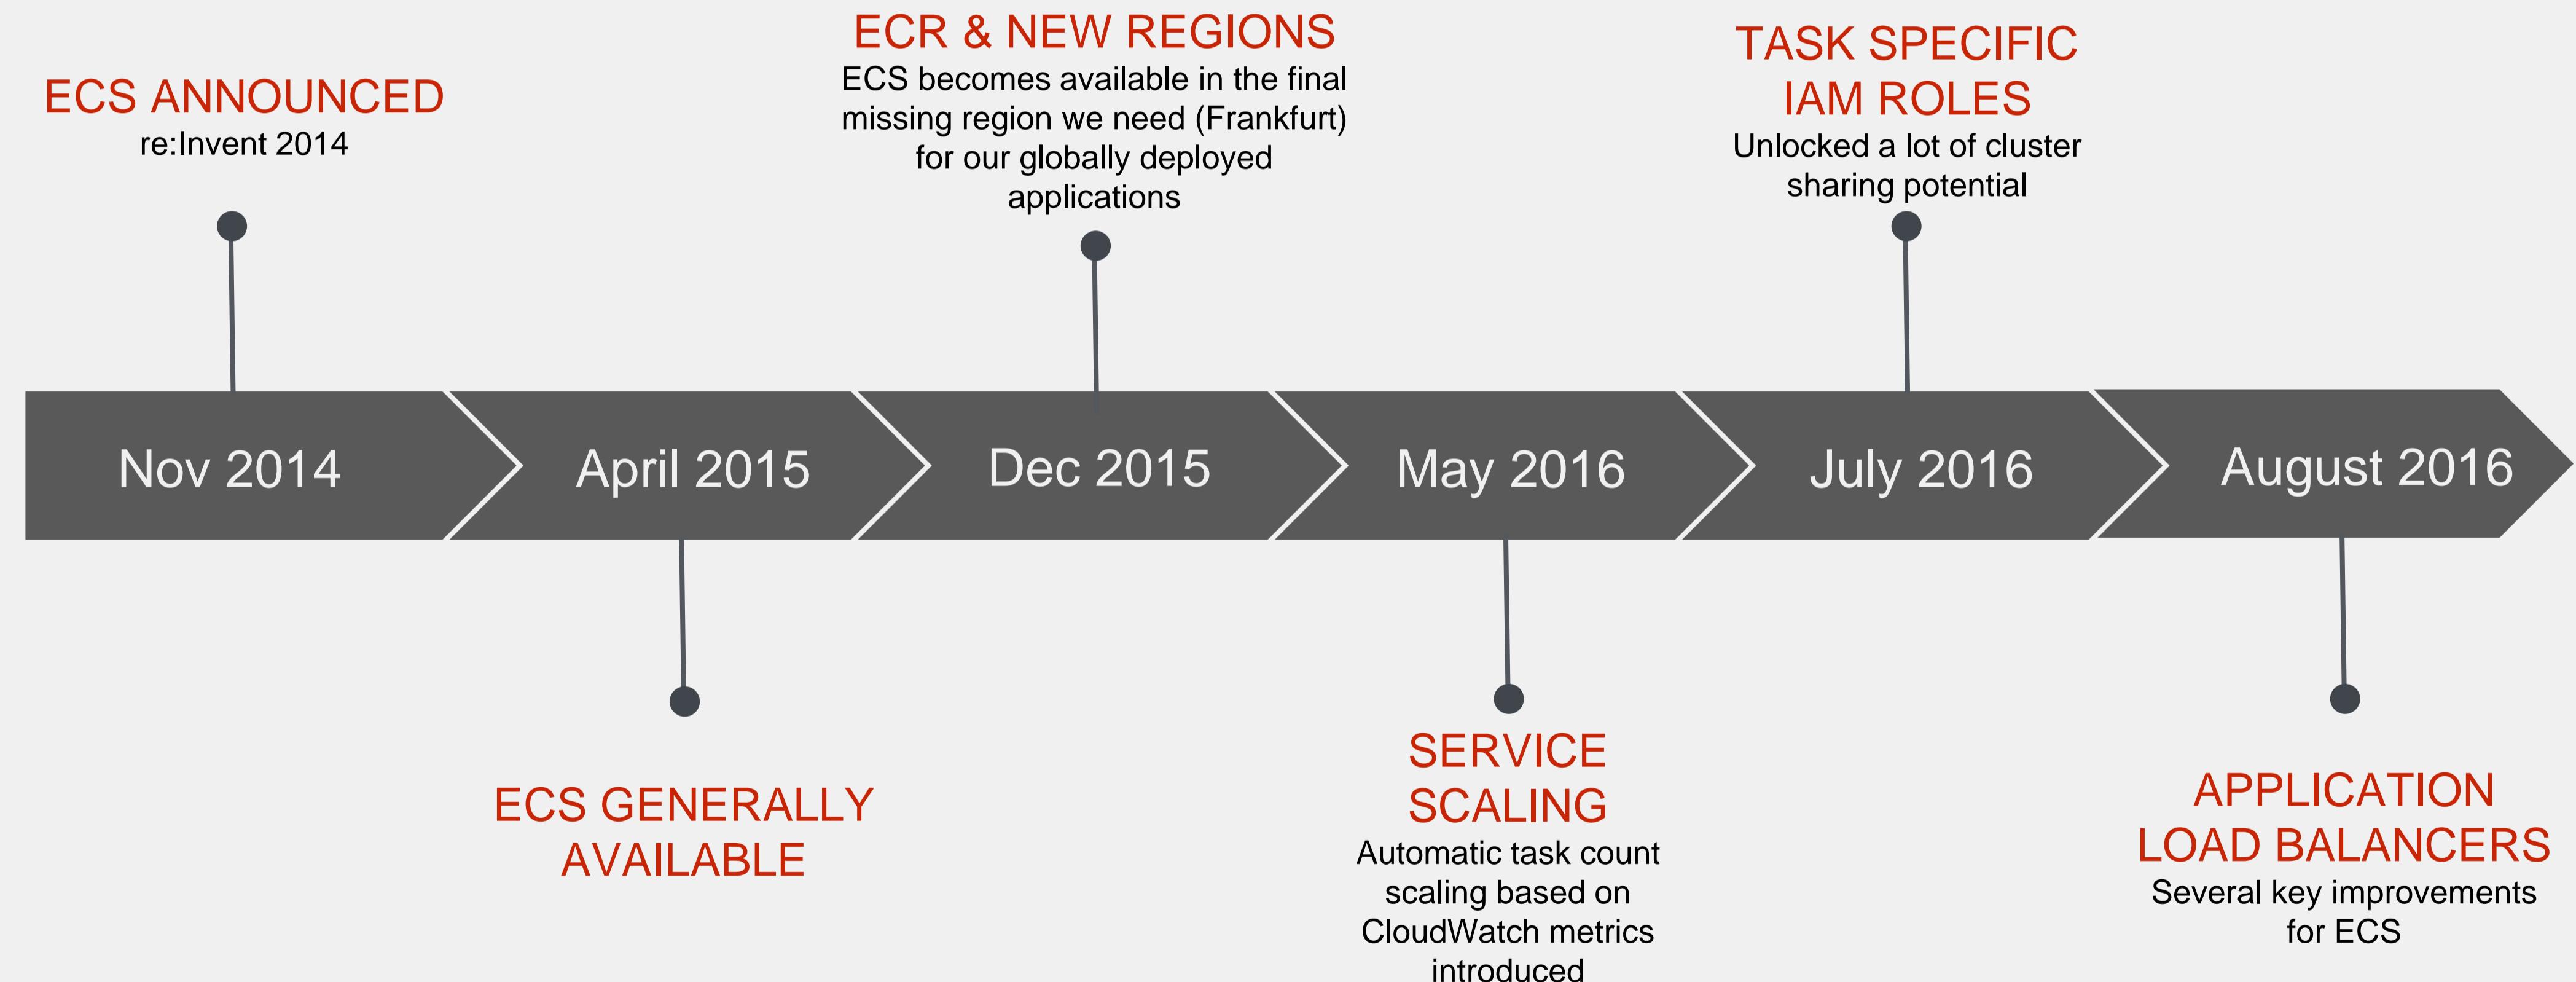
AMAZON EC2 CONTAINER SERVICE

KEY FEATURE TIMELINE



AMAZON EC2 CONTAINER SERVICE

KEY FEATURE TIMELINE



BEYOND ECS

INFRASTRUCTURE AS CODE

- At scale, orchestrating infrastructure in a consistent, reproducible way is key



BEYOND ECS

INFRASTRUCTURE AS CODE

- At scale, orchestrating infrastructure in a consistent, reproducible way is key

Total ownership





TERRAFORM

TERRAFORM

INFRASTRUCTURE AS OBJECT-ORIENTED CODE

- Intelligent & flexible flow control



TERRAFORM

INFRASTRUCTURE AS OBJECT-ORIENTED CODE

- Intelligent & flexible flow control
- Configuration drift managed



TERRAFORM

INFRASTRUCTURE AS OBJECT-ORIENTED CODE

- Intelligent & flexible flow control
- Configuration drift managed
- AWS resources documented



TERRAFORM

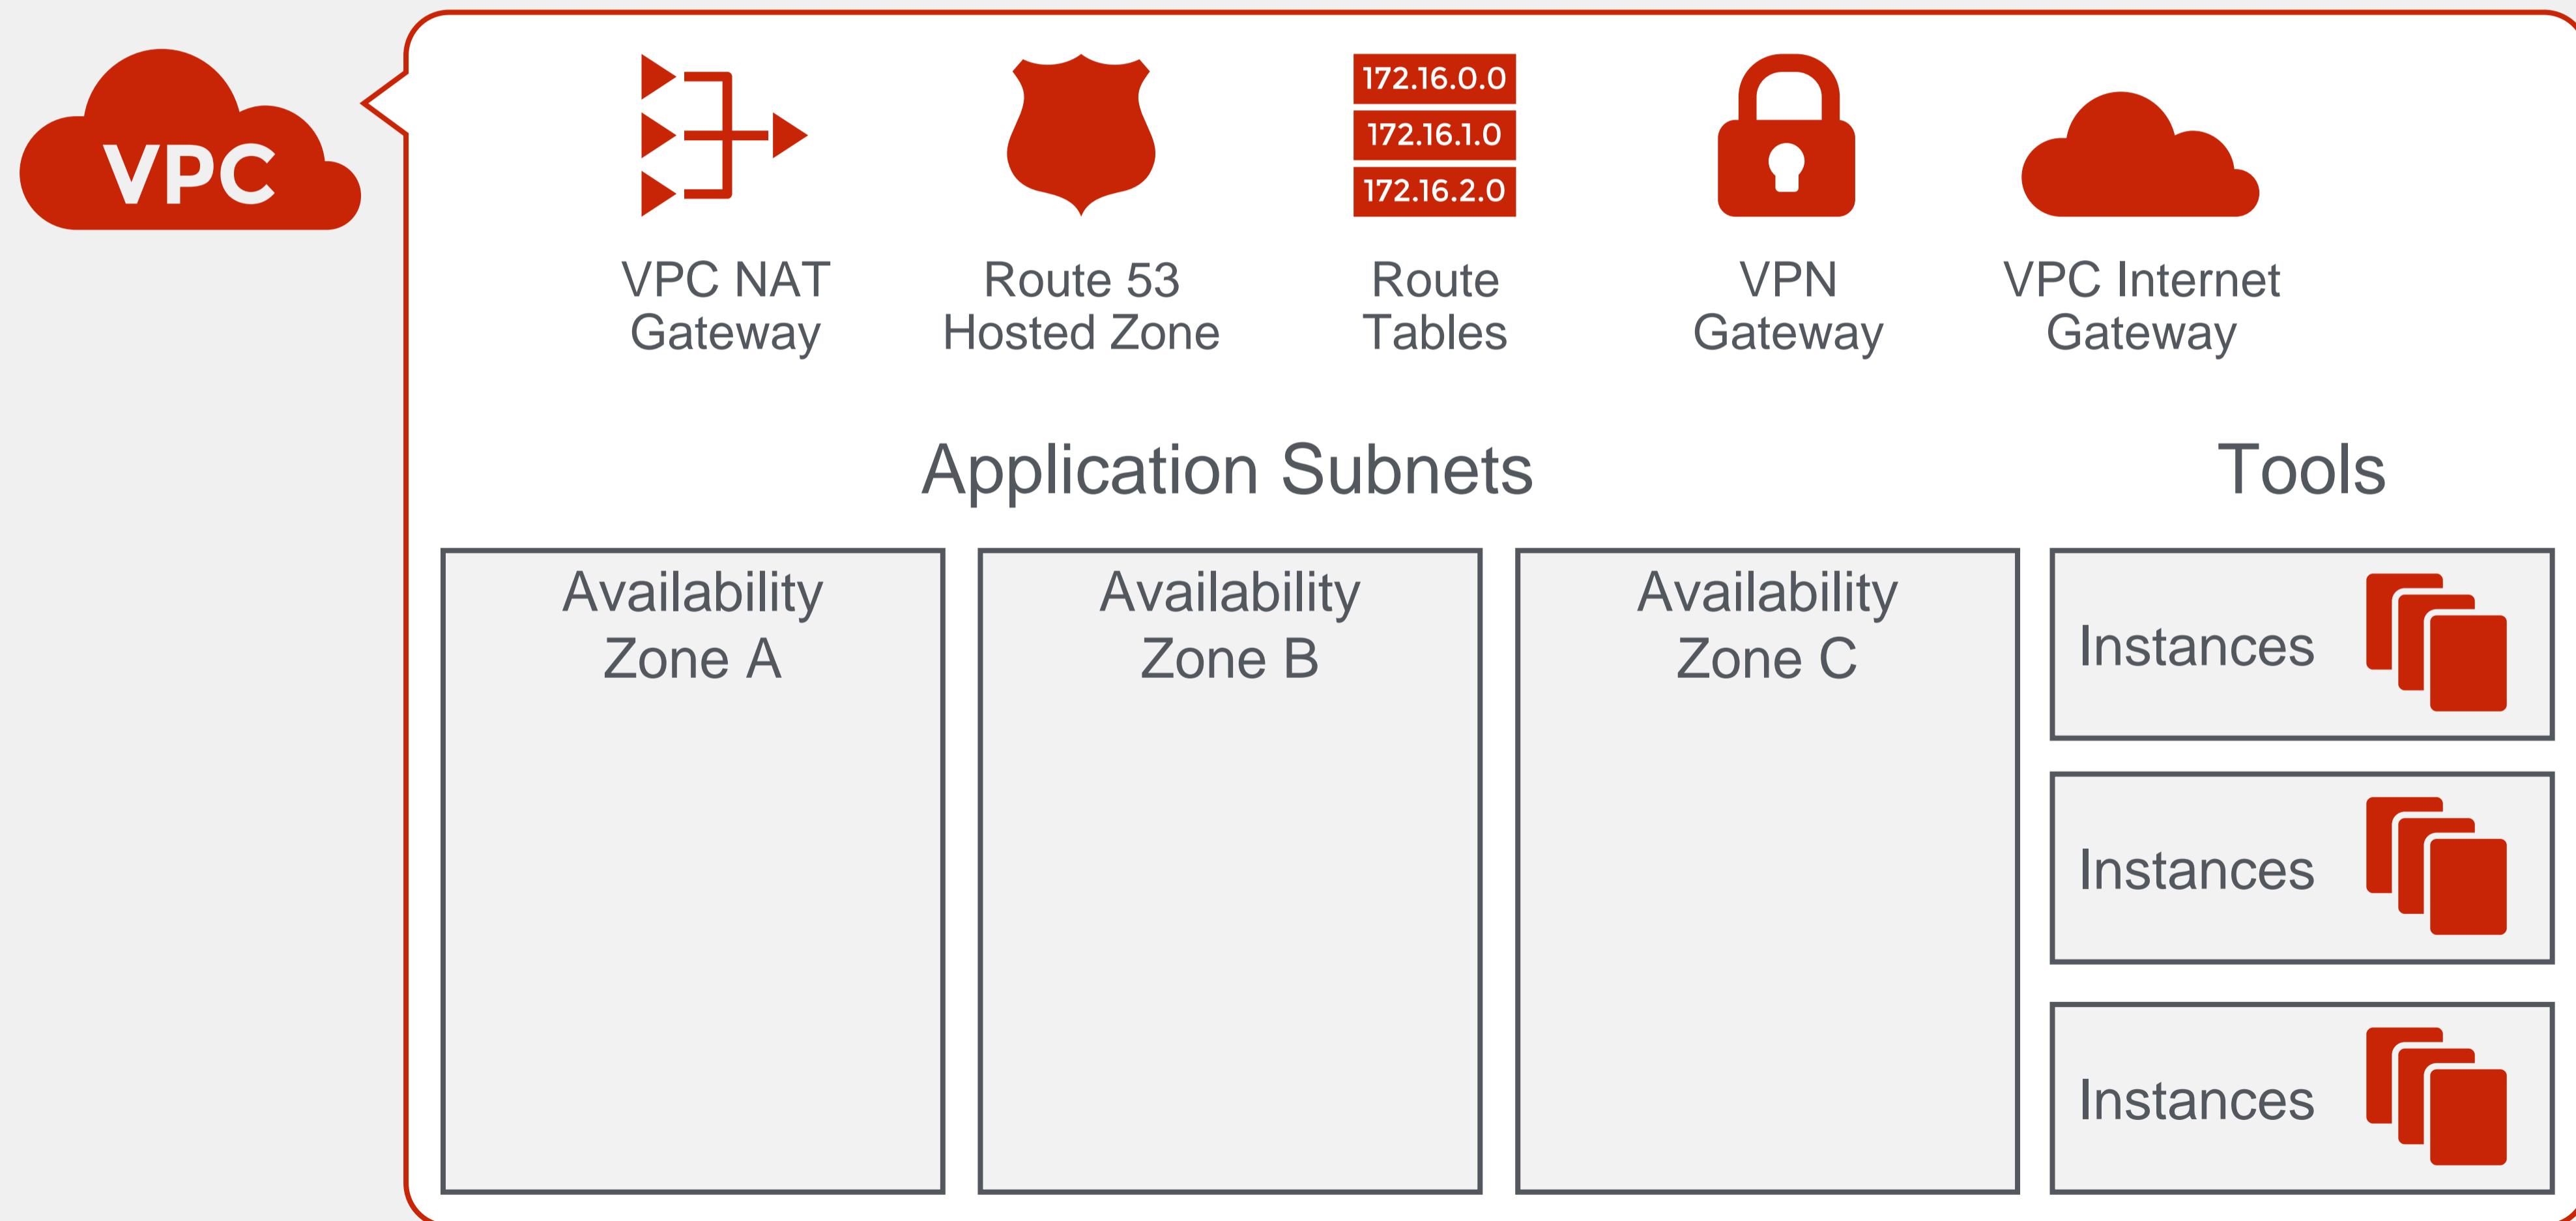
INFRASTRUCTURE AS OBJECT-ORIENTED CODE

- Intelligent & flexible flow control
- Configuration drift managed
- AWS resources documented
- Share common use cases



TERRAFORM

INFRASTRUCTURE AS OBJECT-ORIENTED CODE



TERRAFORM

INFRASTRUCTURE AS OBJECT-ORIENTED CODE

Simplified infrastructure definitions

```
# Your AWS account ID
account_id = "012345678901"

# Configure the AWS region you are in.
region = "us-east-1"

# Set the name of your INITIATIVE here
account_name = "rpg"

# This is the name of your VPC
vpc_name = "sandbox_vpc"

# Google Group Distribution for your team
owner = "myteamalias@riotgames.com"

# REQUIRED resource attribution tag
accounting = "territory.initiative.project"

# You should update this to the subnet that is allotted to you
vpc_cidr = "10.181.180.0/22"
azs = "us-east-1a,us-east-1c"

# USER CONFIGURED. This is an example of subnet distribution
frontend_subnets = "10.180.161.0/27,10.180.161.32/27"
backend_subnets = "10.180.161.128/27,10.180.161.160/27"
frontend_dev_subnets = "10.180.162.0/27,10.180.162.32/27"
backend_dev_subnets = "10.180.162.128/27,10.180.162.160/27"
```

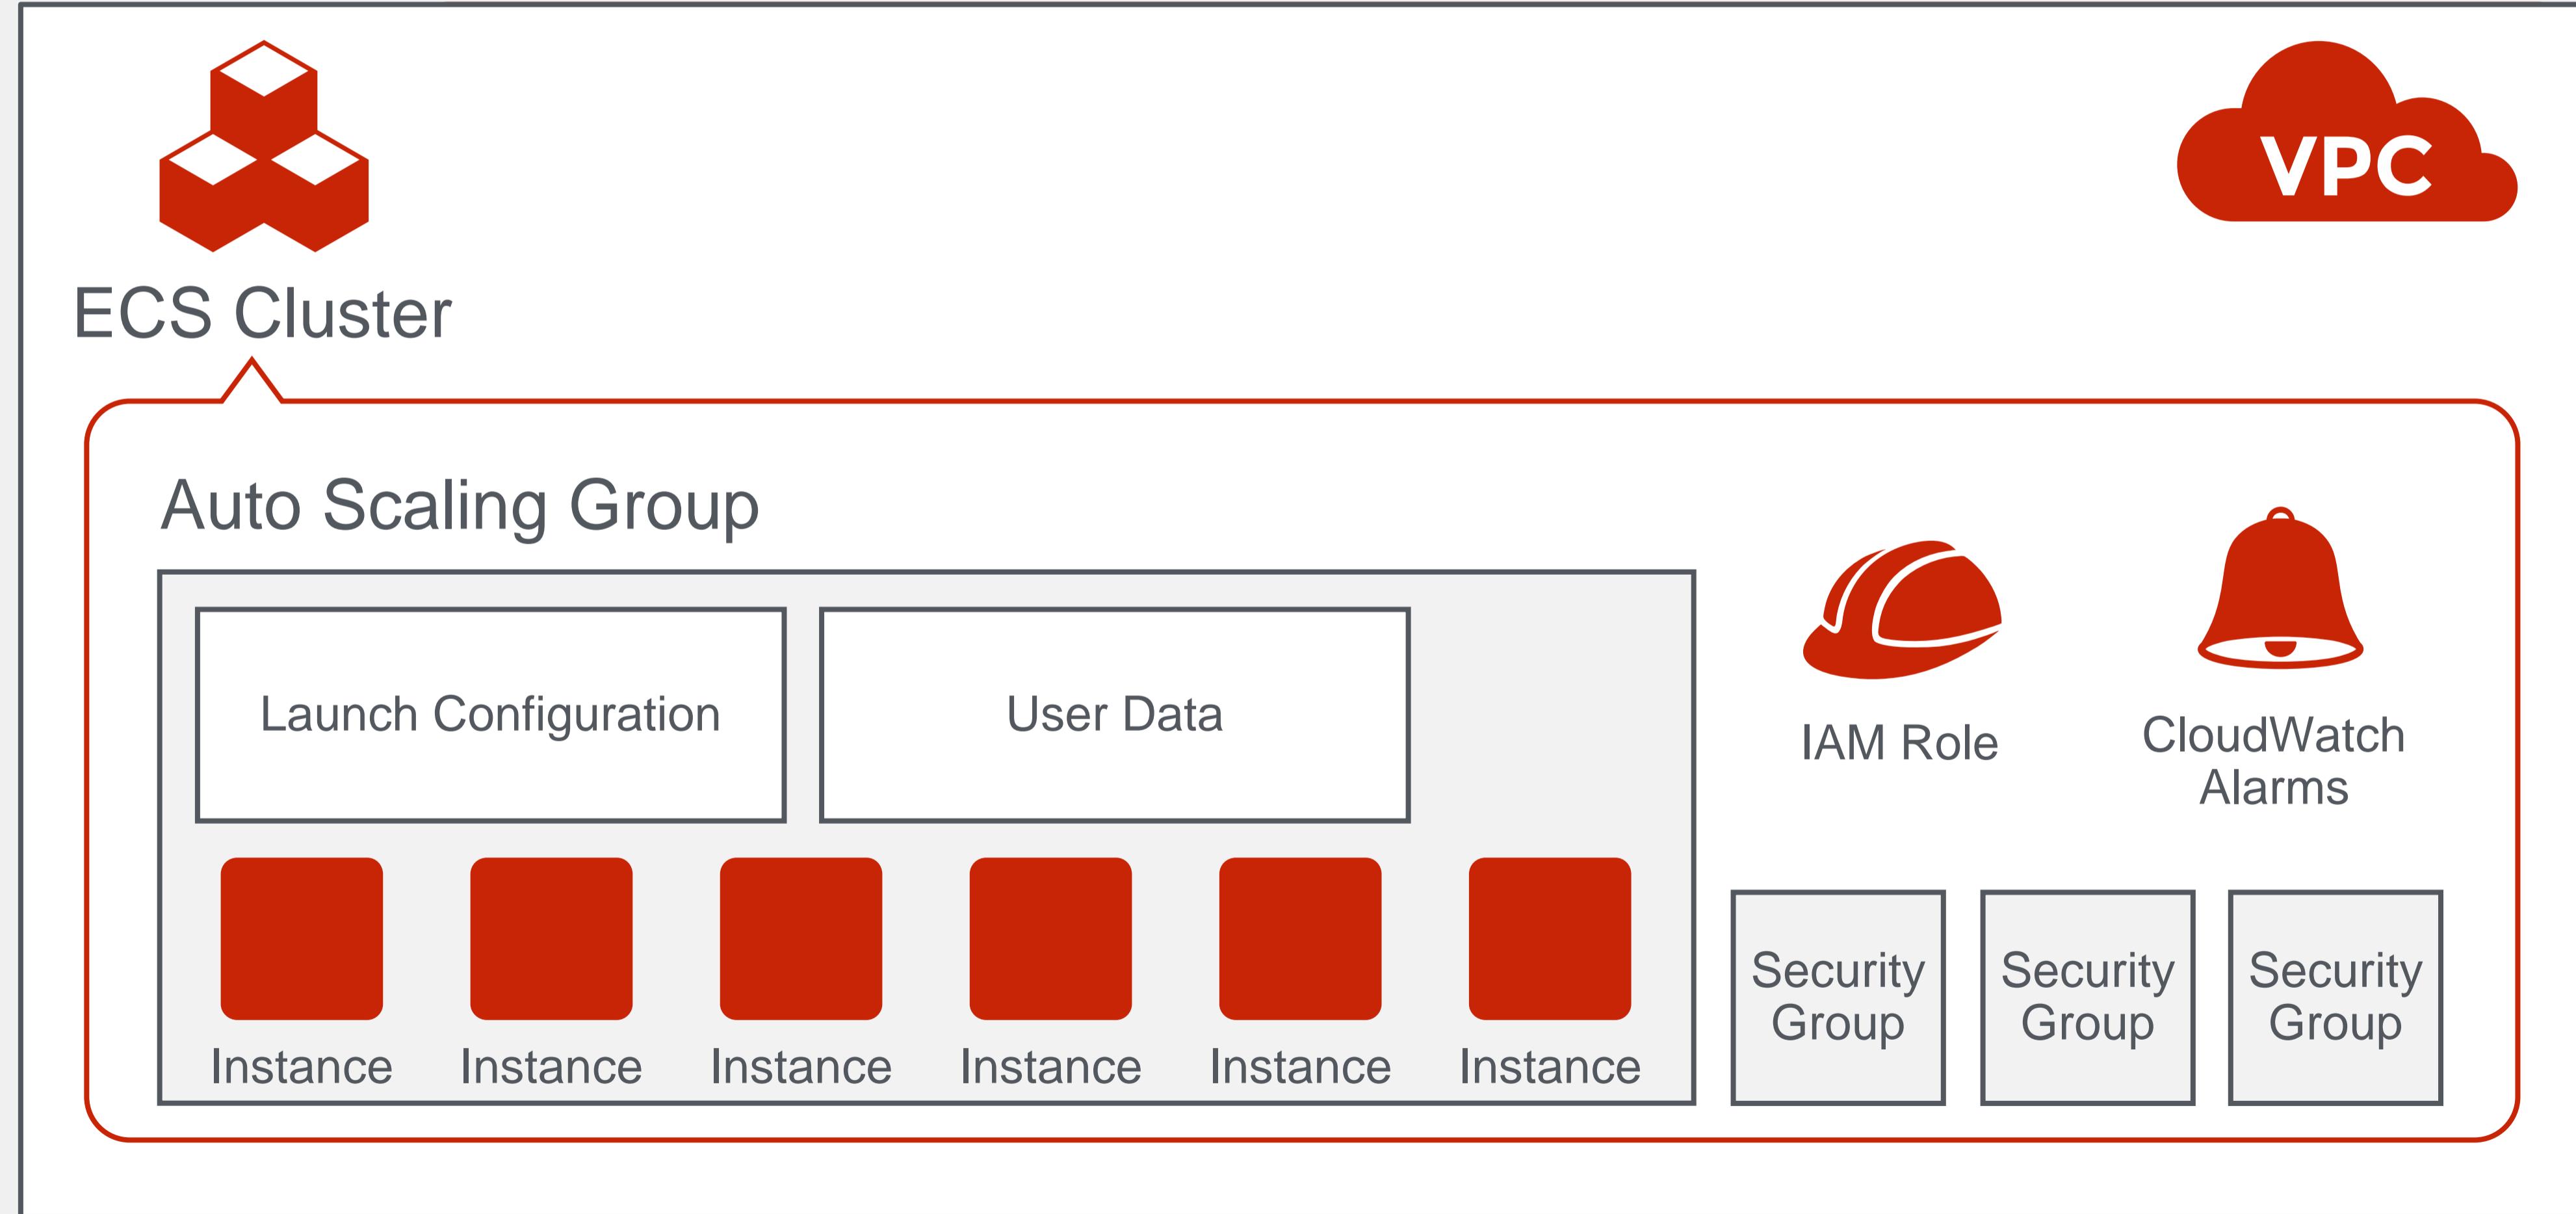


CASE STUDIES



ECS CLUSTER

TERRAFORM BUILDING BLOCKS



ECS CLUSTER

TERRAFORM BUILDING BLOCKS

Module 1

```
output "private_subnets" {  
  value = "${join("", aws_subnet.private.*.id)}"  
}
```

```
output "team_access_security_group" {  
  value = "${aws_security_group.team_access.id}"  
}
```

Module 2

Module 3

```
module "ecs-autoscaling" {  
  
  source = "./ecs-autoscaling"  
  cluster_name = "reinvent-demo-large-static-cluster"  
  
  # Tag values  
  owner = "arozumek"  
  team = "dps@riotgames.com"  
  
  # VPC Information  
  availability_zones = "${module.vpc.availability_zones}"  
  subnet_ids = "${module.vpc.private_subnet_ids}"  
  
  # Instance configuration  
  key_name = "reinvent-keypair"  
  instance_type = "r3.4xlarge"  
  security_group_ids = "${var.team_access_sg}"  
  iam_instance_profile = "Generic-ECS"  
  
  # ASG configuration  
  min_size = "3"  
  max_size = "3"  
  desired_capacity = "3"  
  
  # Docker registry  
  registry_url = "registry.rcluster.io"  
  registry_email = "${var.module_registry_email}"  
  registry_auth = "${var.module_registry_auth}"  
  
}
```

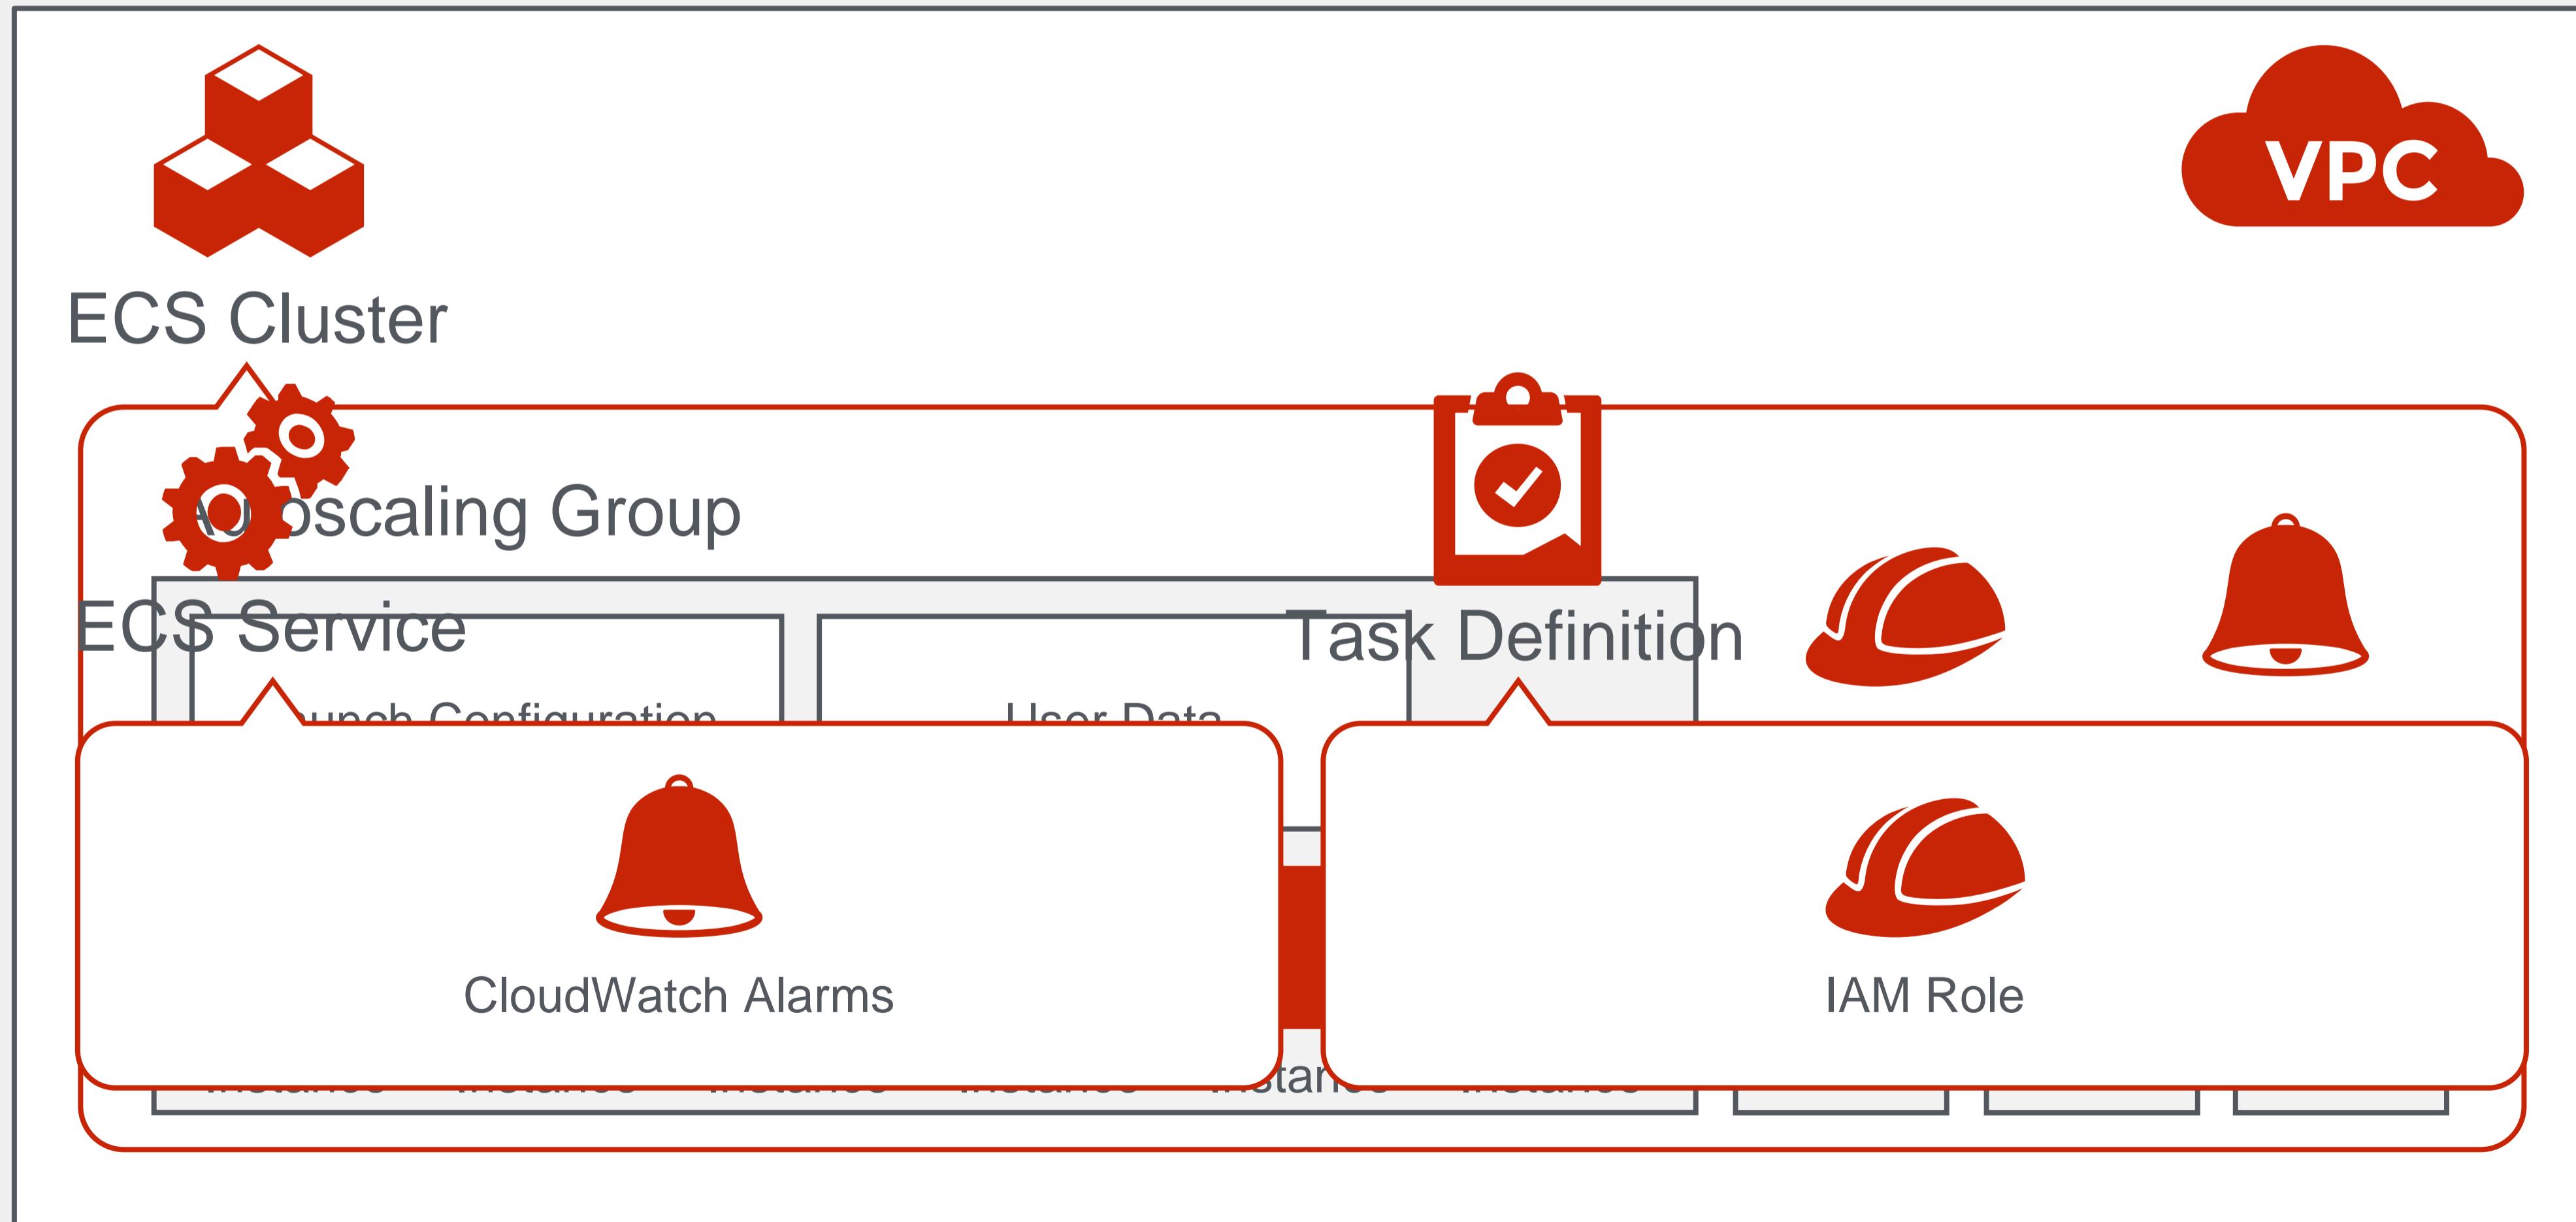
ECS CLUSTER

TERRAFORM BUILDING BLOCKS

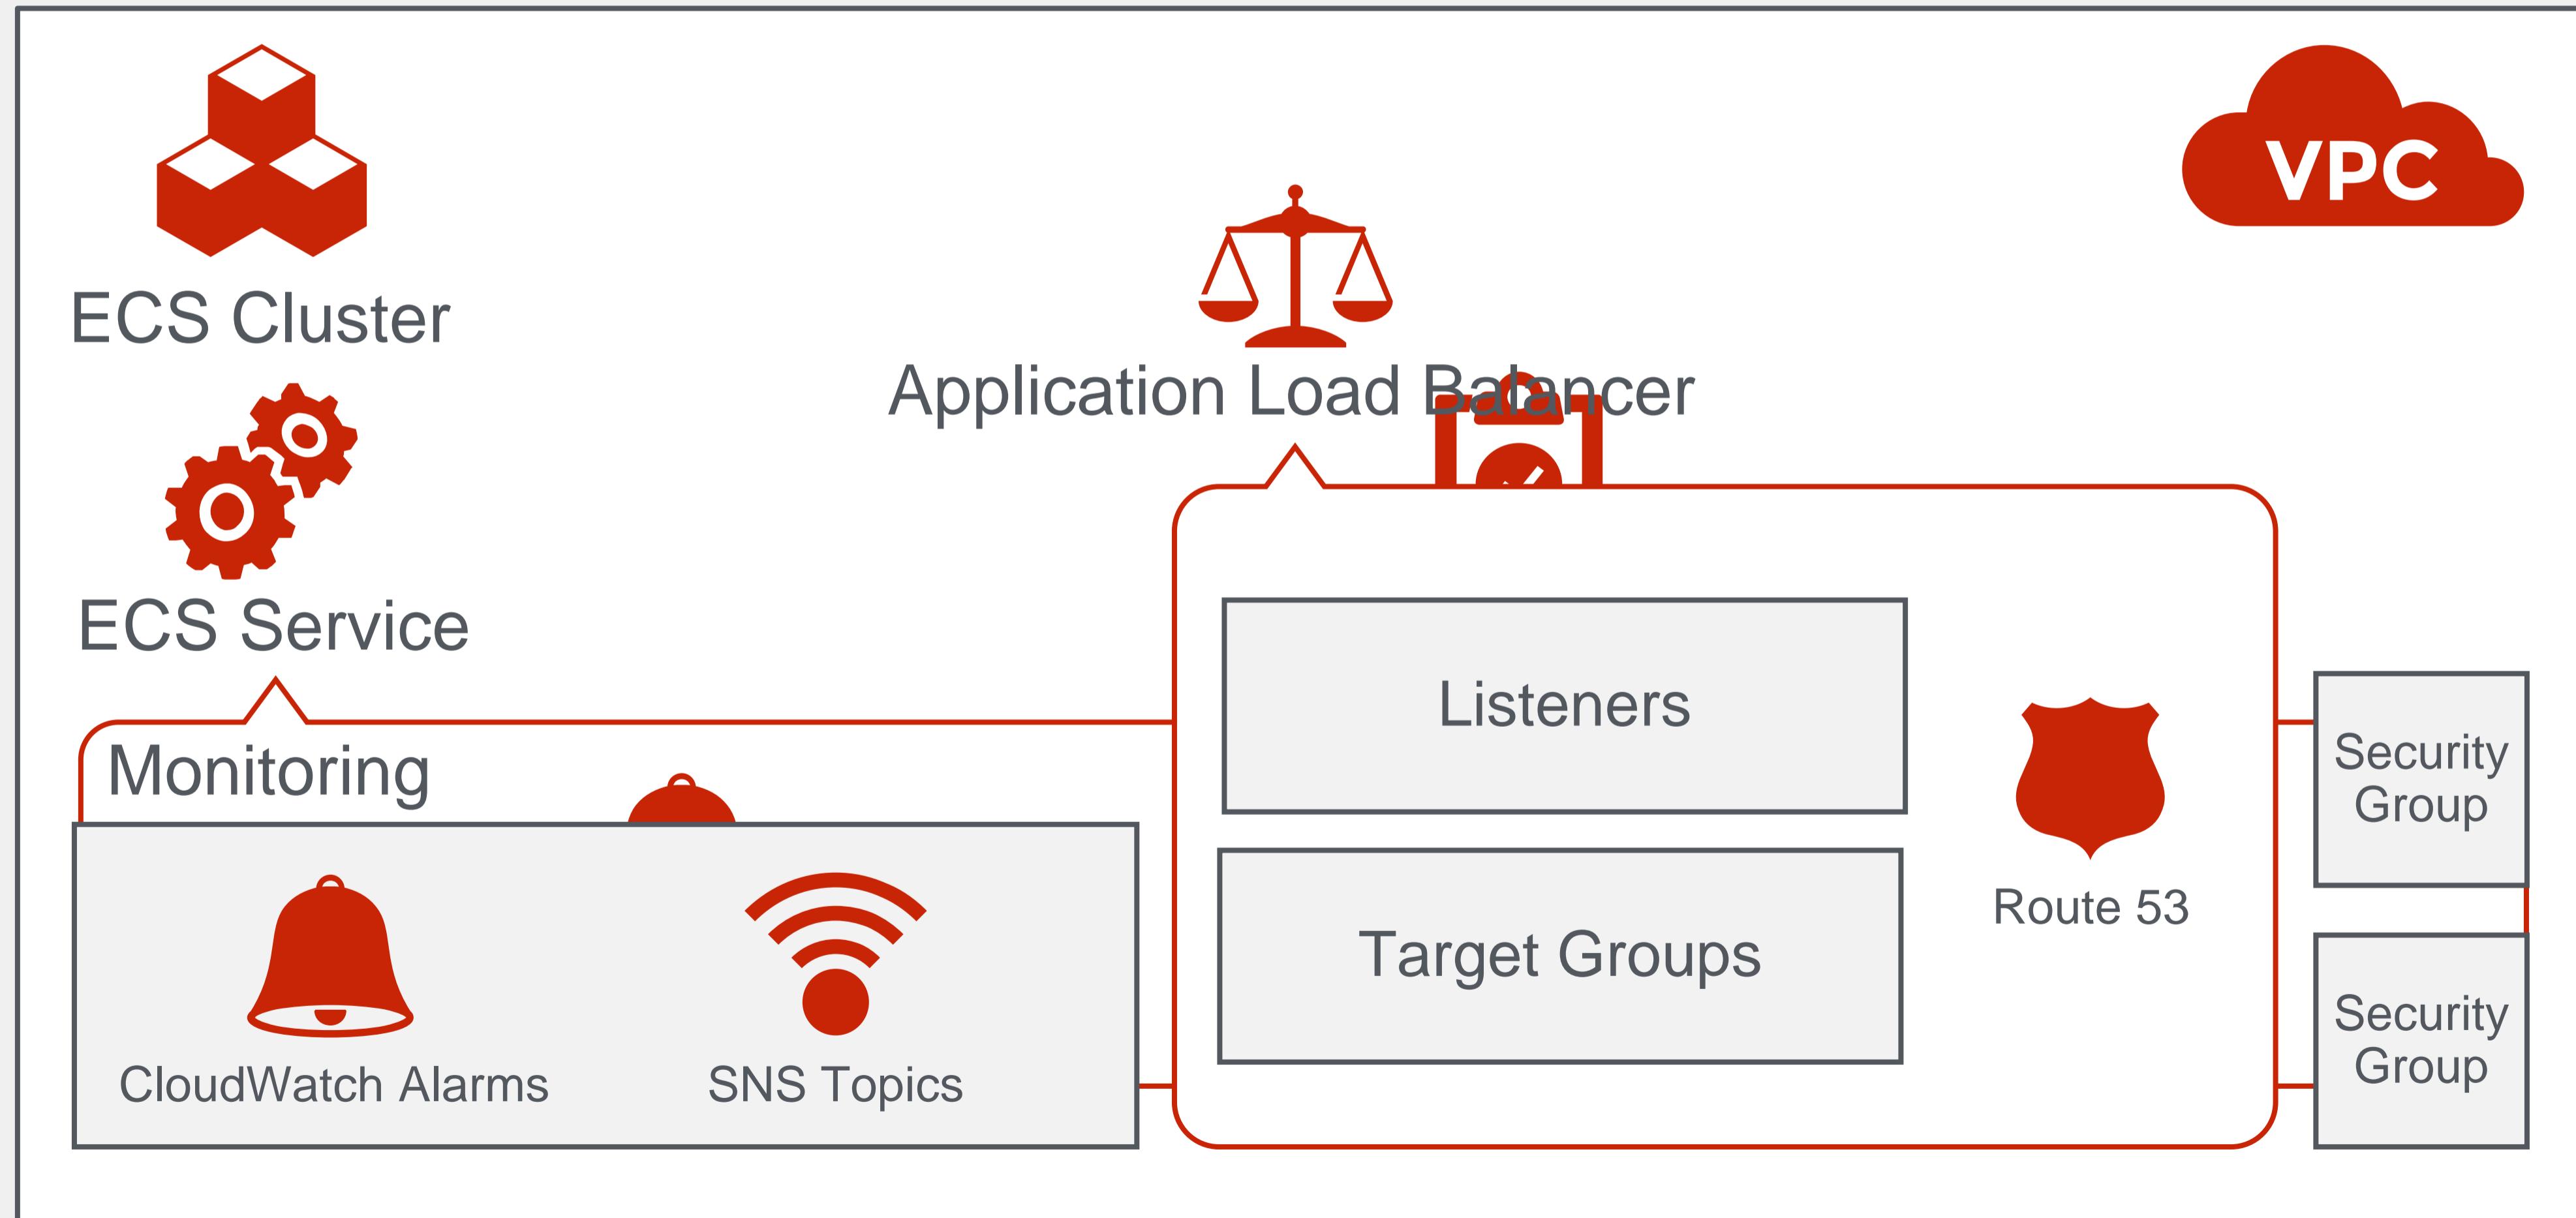
```
module "ecs-autoscaling" {  
  
    source = "./ecs-autoscaling"  
    cluster_name = "reinvent-demo"  
  
    # Tag values  
    owner = "arozumek"  
    team = "dps@riotgames.com"  
  
    # VPC Information  
    availability_zones = "${module.vpc.availability_zones}"  
    subnet_ids = "${module.vpc.private_subnet_ids}"  
  
    # Instance configuration  
    key_name = "reinvent-keypair"  
    instance_type = "m3.large"  
    security_group_ids="sg-d7b890b3"  
    iam_instance_profile = "Generic-ECS"  
  
    # ASG configuration  
    min_size = "3"  
    max_size = "10"  
    desired_capacity = "4"  
  
    # Docker registry  
    registry_url = "registry.rcluster.io"  
    registry_email = "${var.module_registry_email}"  
    registry_auth = "${var.module_registry_auth}"  
  
}
```

```
module "ecs-autoscaling" {  
  
    source = "./ecs-autoscaling"  
    cluster_name = "reinvent-demo"  
  
    # Tag values  
    owner = "arozumek"  
    team = "dps@riotgames.com"  
  
    # VPC Information  
    availability_zones = "${module.vpc.availability_zones}"  
    subnet_ids = "${module.vpc.private_subnet_ids}"  
  
    # Instance configuration  
    key_name = "reinvent-keypair"  
    instance_type = "r3.4xlarge"  
    security_group_ids="sg-d7b890b3"  
    iam_instance_profile = "Generic-ECS"  
  
    # ASG configuration  
    min_size = "3"  
    max_size = "3"  
    desired_capacity = "3"  
  
    # Docker registry  
    registry_url = "registry.rcluster.io"  
    registry_email = "${var.module_registry_email}"  
    registry_auth = "${var.module_registry_auth}"  
  
}
```

MICROSERVICES WITHOUT SERVICE ENDPOINTS

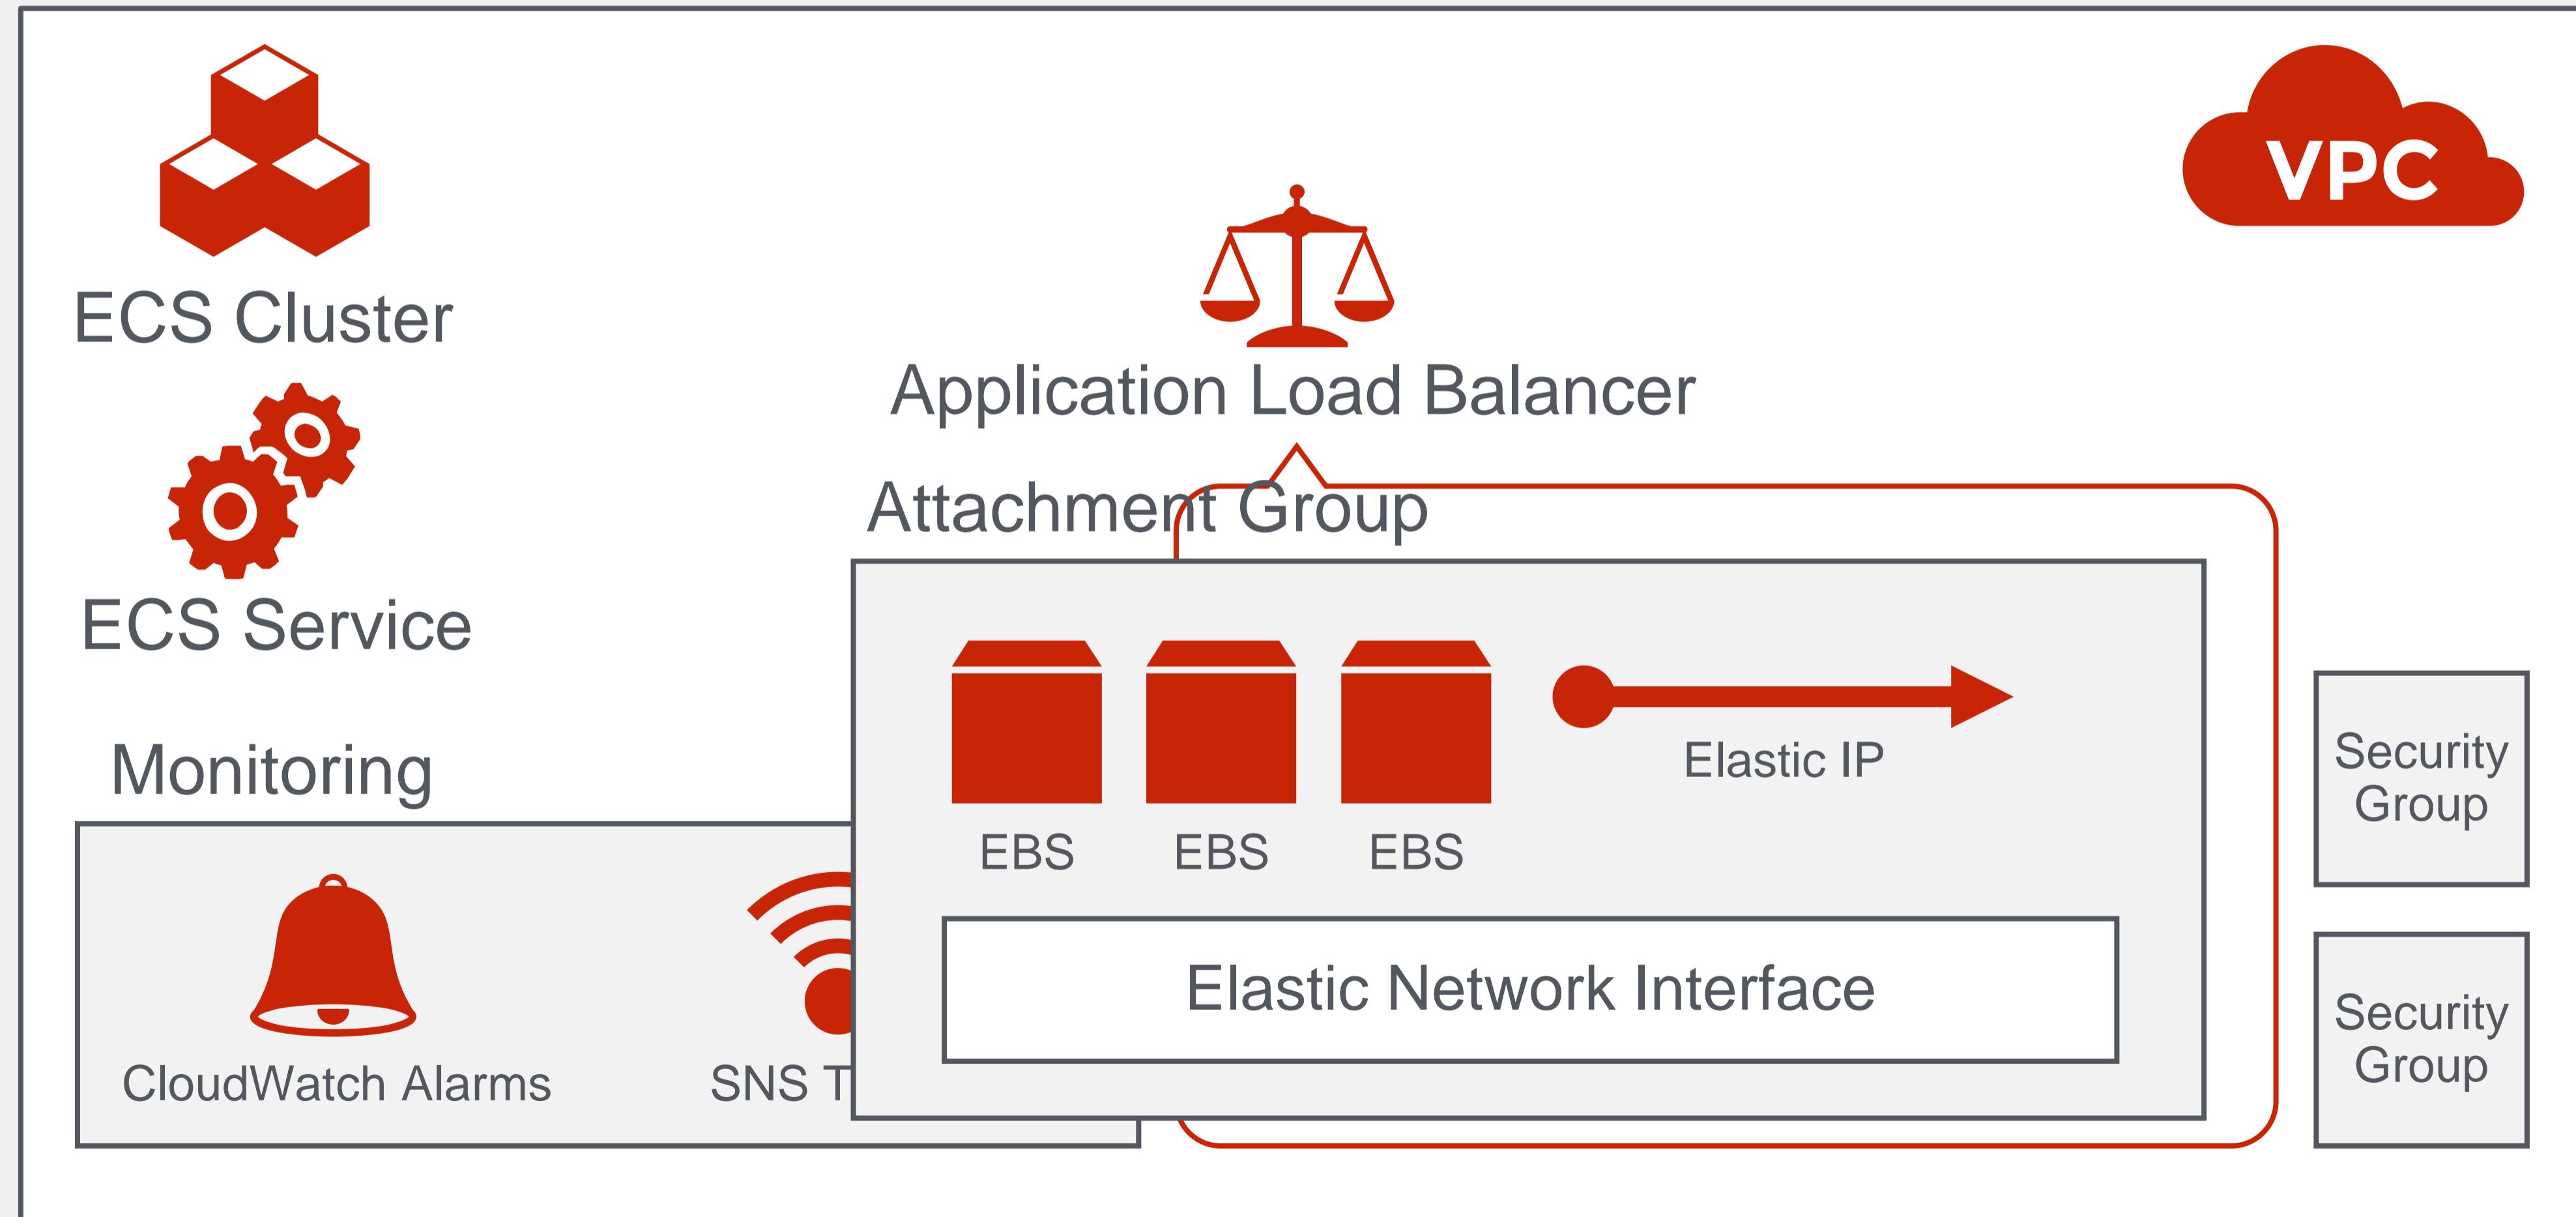


MICROSERVICES WITH SERVICE ENDPOINTS



PERSISTENT DATA

LOSE ECS HOSTS WITHOUT LOSING DATA





LESSONS LEARNED



LESSONS

WHAT WORKED FOR US

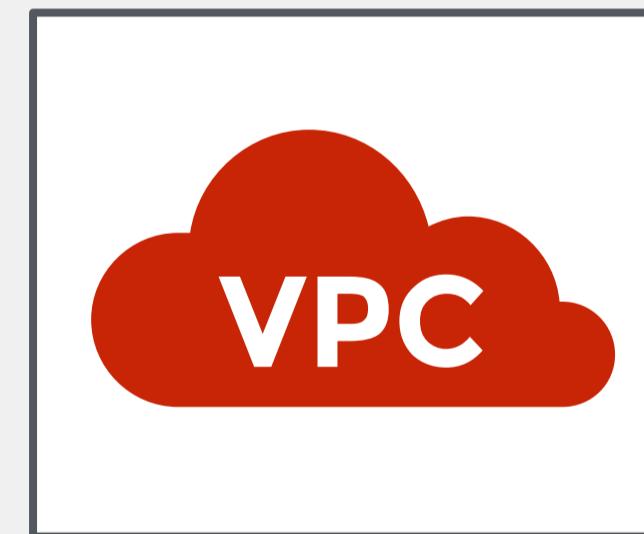
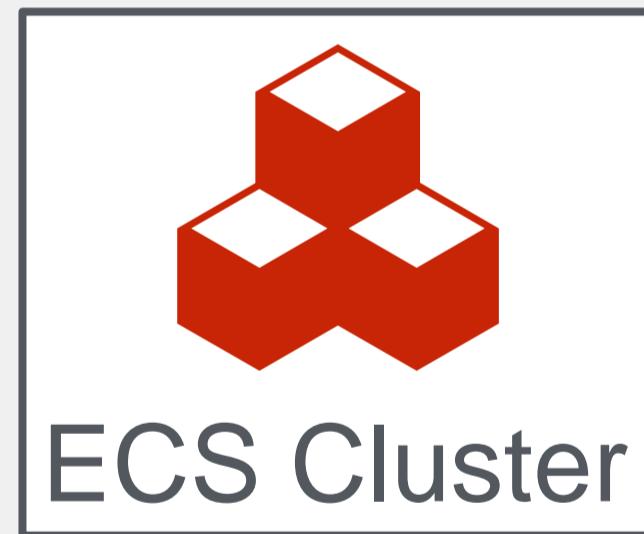
- Socialize your templates & take advantage of the community!
- Terraform community modules
- Terraform AWS blog



LESSONS

WHAT WORKED FOR US

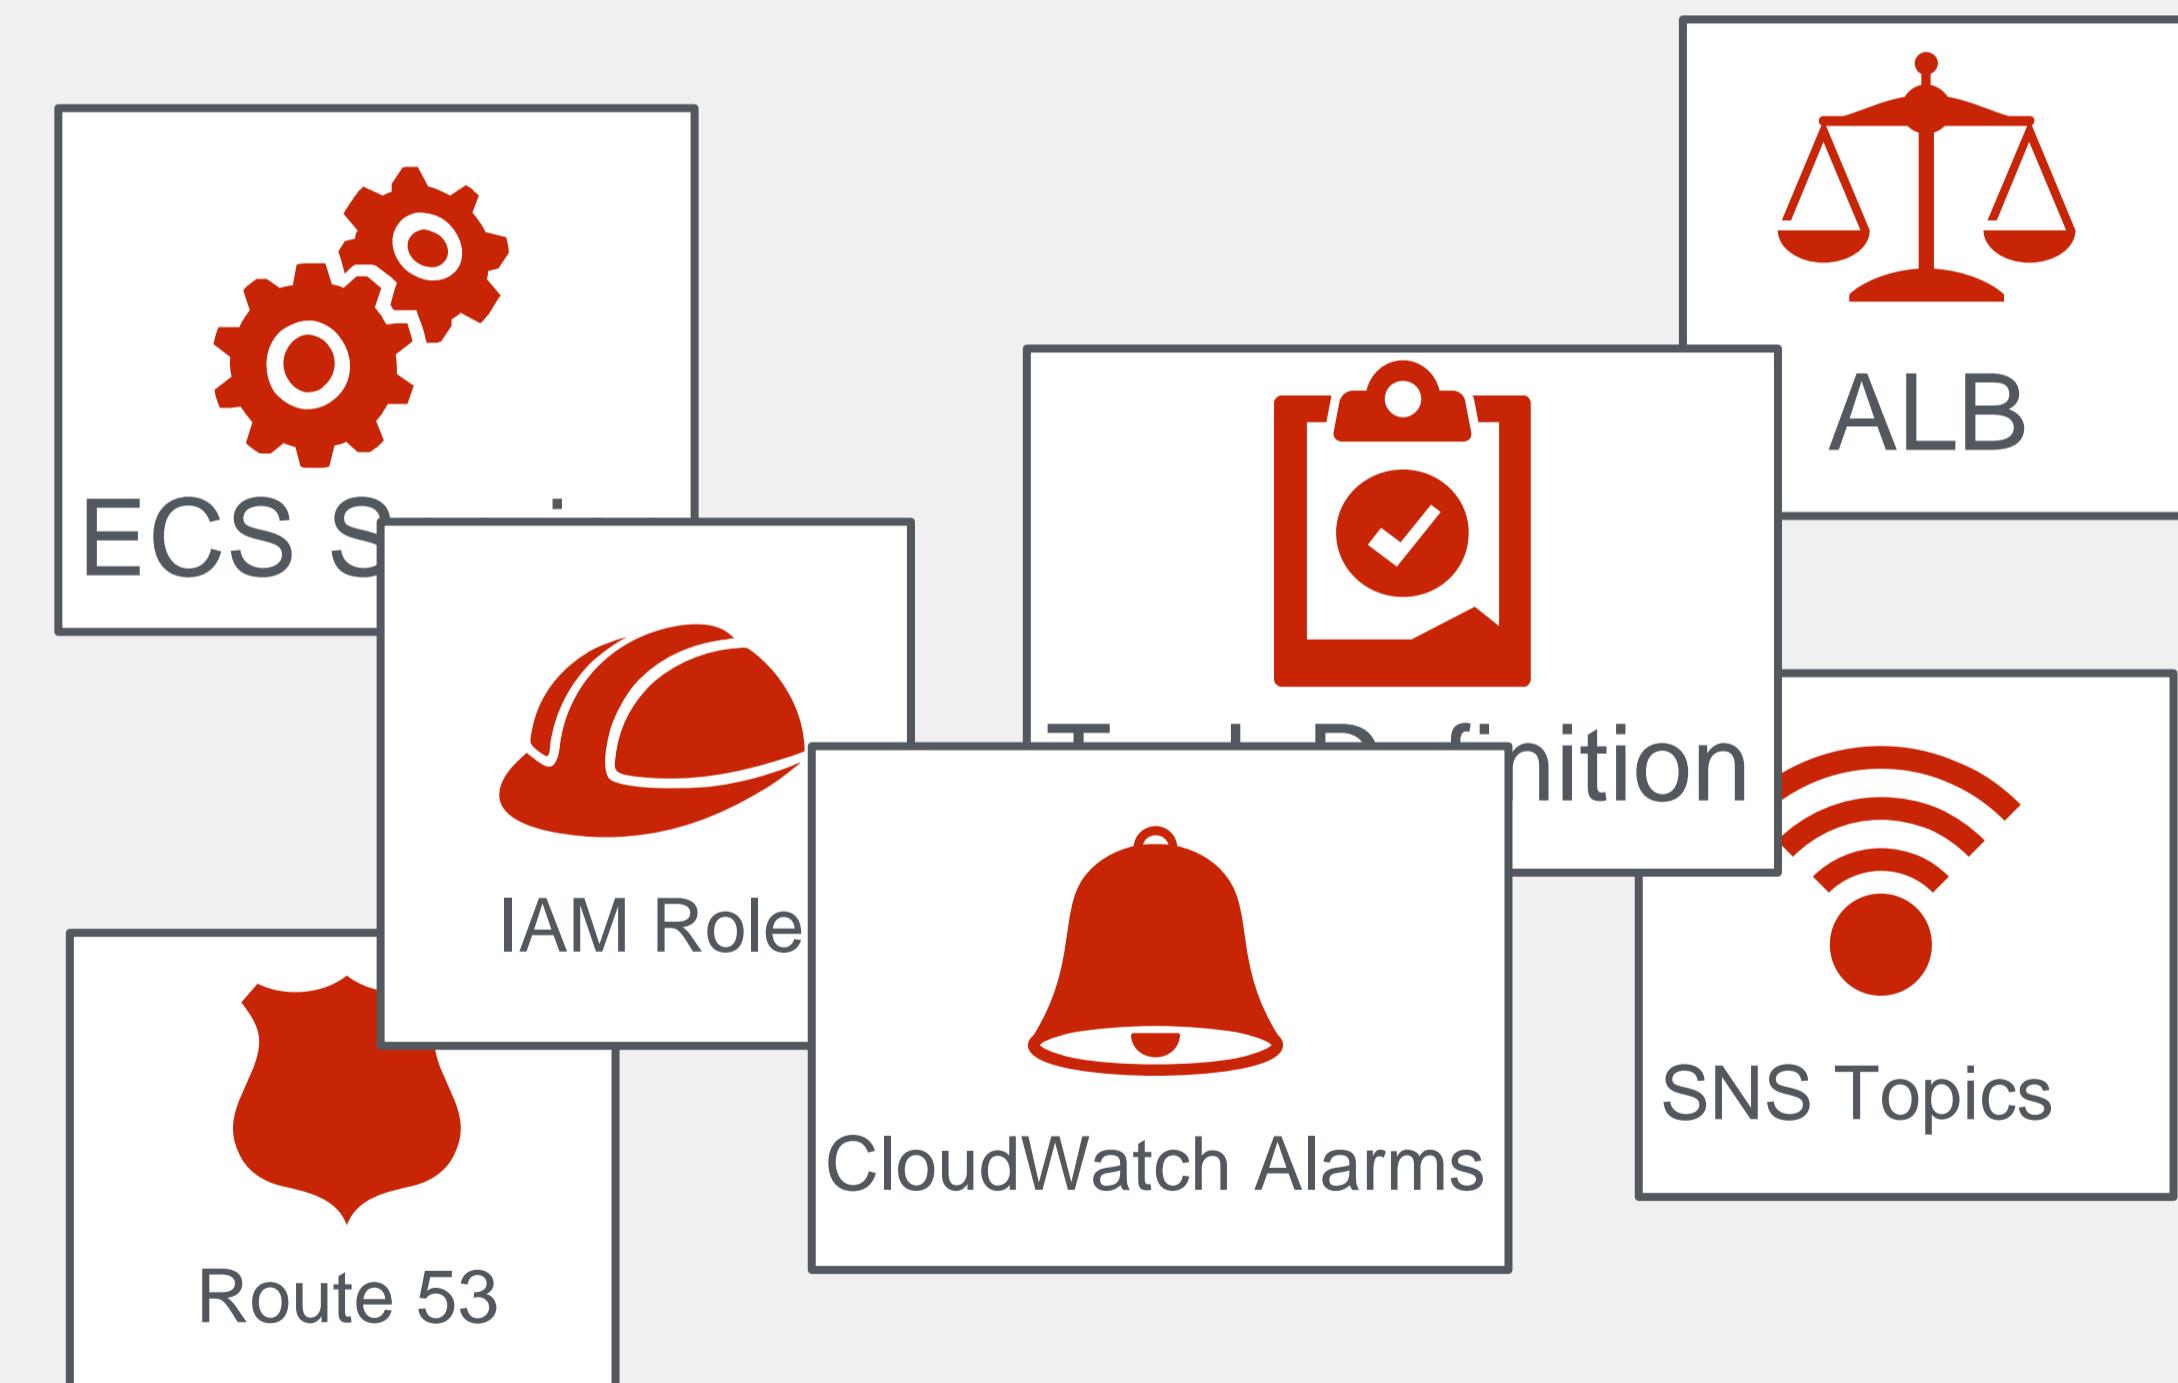
- Break apart your stacks



LESSONS

WHAT WORKED FOR US

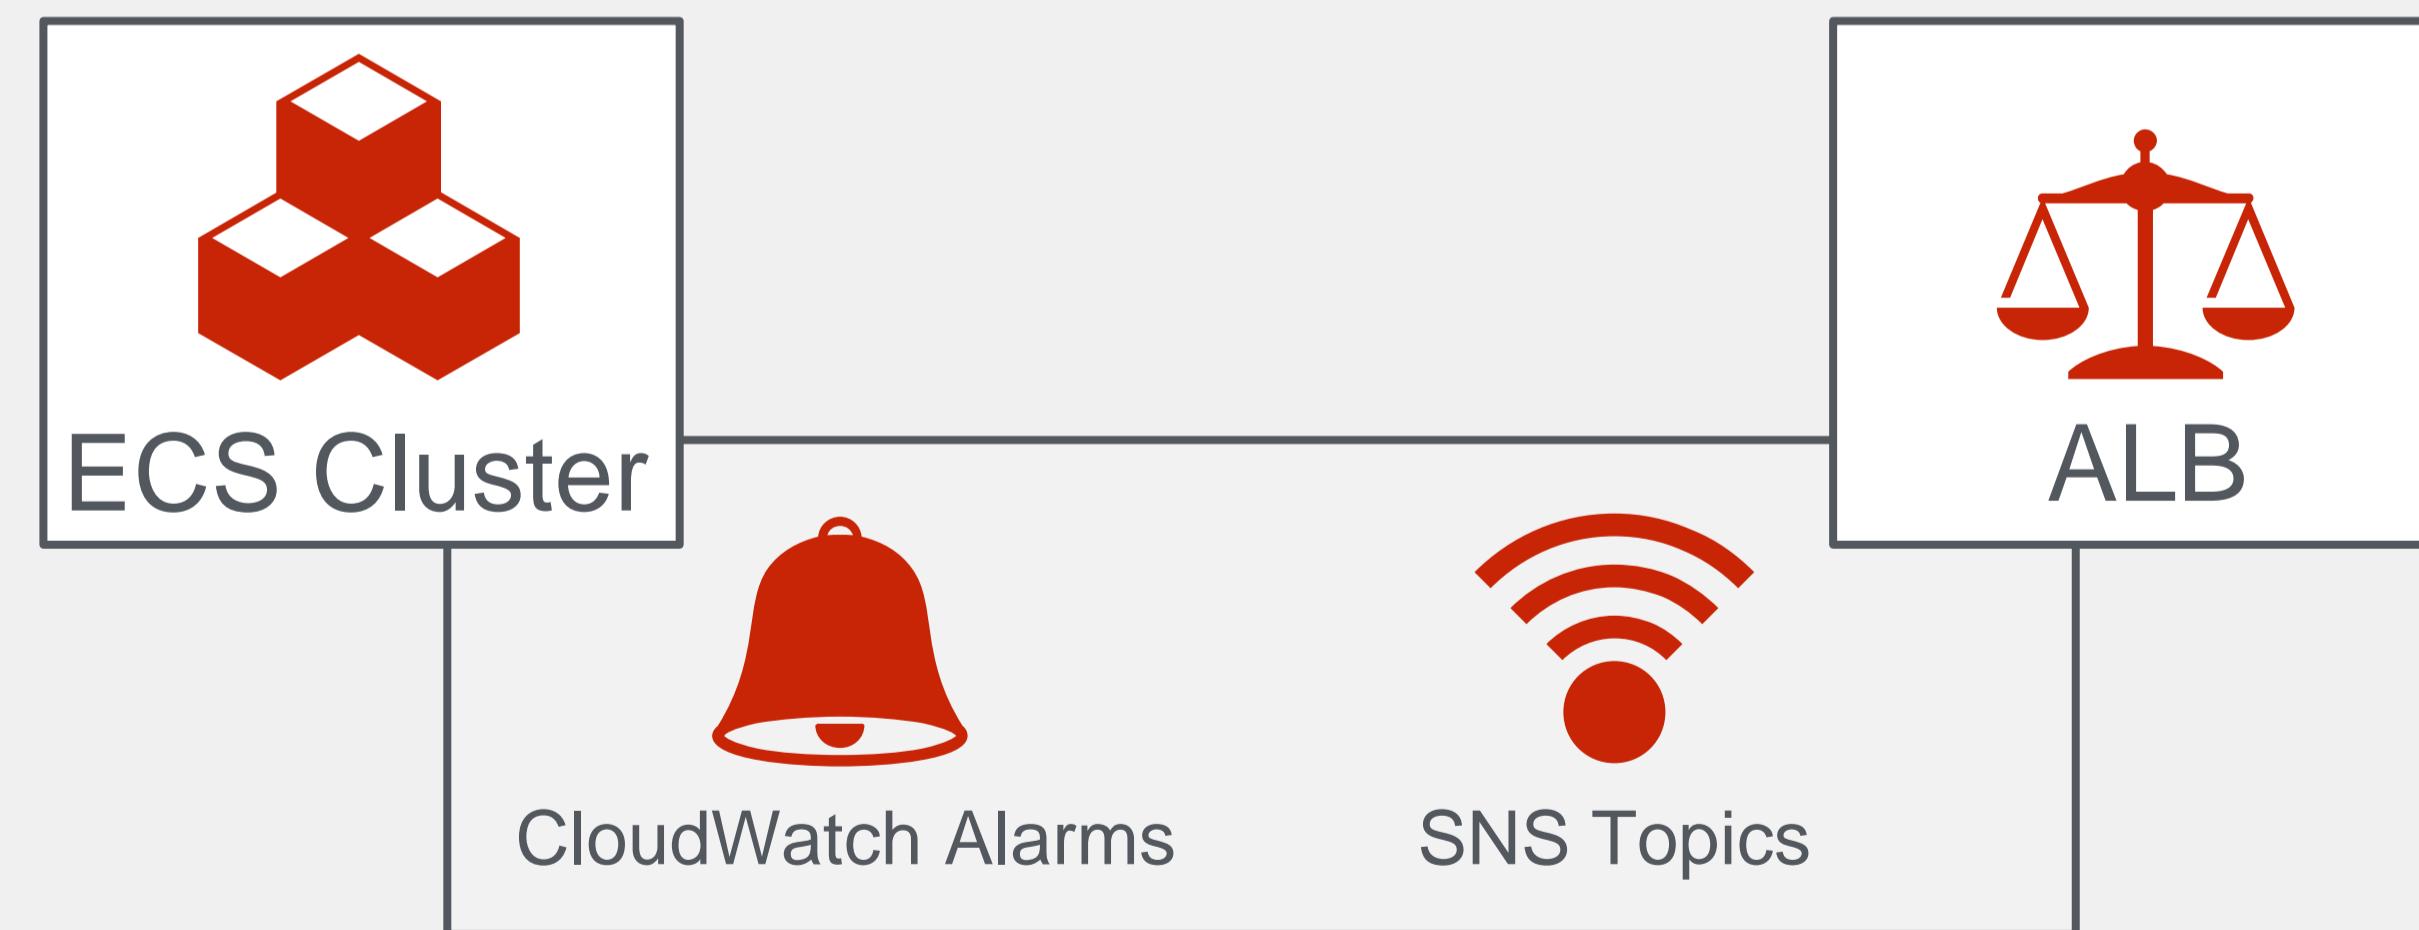
- Break apart your stacks, but don't overdo it



LESSONS

WHAT WORKED FOR US

- Break apart your stacks, but don't overdo it
- Use modules to reduce code duplication



LESSONS

WHAT WORKED FOR US

- Use remote state files

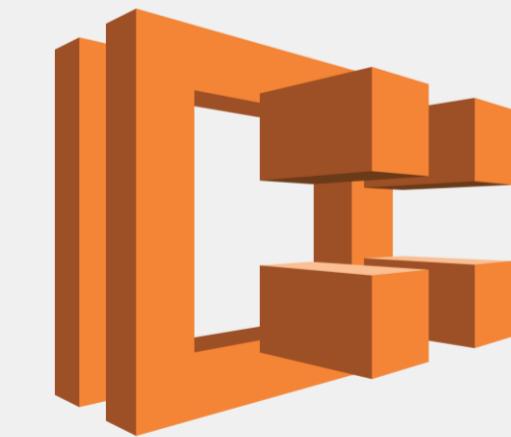
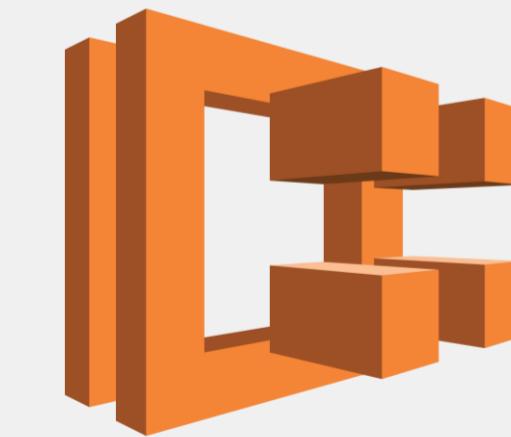
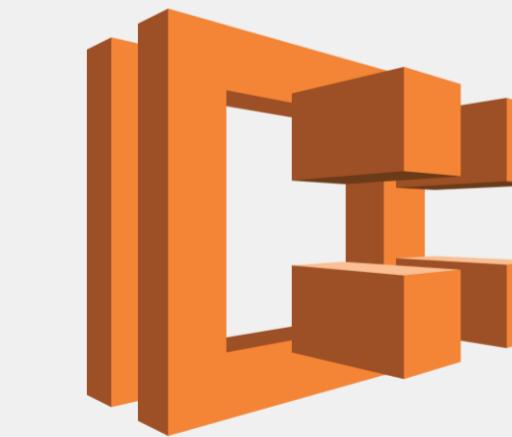
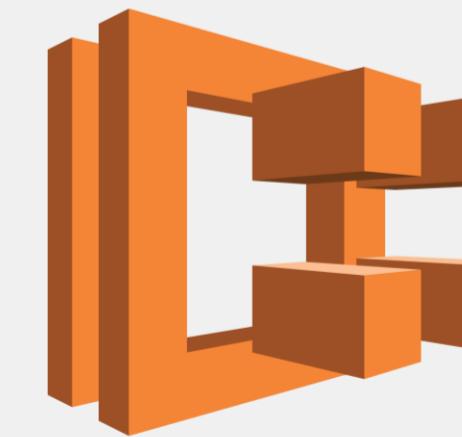
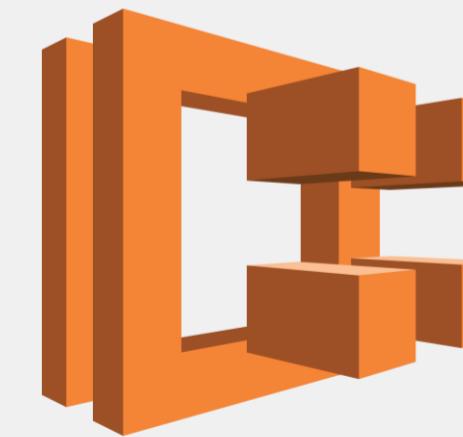
```
data "terraform_remote_state" "foo" {  
    backend = "s3"  
    config {  
        bucket = "terraform-state-prod"  
        key = "network/terraform.tfstate"  
        region = "us-east-1"  
    }  
}
```



LESSONS

WHAT WORKED FOR US

- Be liberal with your cluster provisioning
- Don't risk resource contention in production
- With good orchestration, additional clusters != additional operational overhead



LESSONS

WHAT WORKED FOR US

- Tag everything all of the time
- Keep your tags organized in your Terraform templates
- Have top level variables that get applied to every resource
- Create a tag for every dimension that is useful to your business



LESSONS

WHAT WORKED FOR US

- Centralize your logs



or



LESSONS

WHAT WORKED FOR US

- Stay up to date with release blogs and application updates
- ECS updates on the AWS blog
- Terraform's open source GitHub repository changelog

0.7.9 (November 4, 2016)

FEATURES:

- New Data Source: `aws_acm_certificate` (#8359)
- New Resource: `aws_autoscaling_attachment` (#9146)

LESSONS

WHAT WORKED FOR US

- Capture your AMI generation process

```
- hosts: "{{ env }}"
  remote_user: ec2-user
  become: yes
  become_user: root

vars:
  logstash_version: 1.5.5
  shield_logstash_username: logstash

tasks:
  - name: Include Vault secured variables
    include_vars: ../conf/es.vault.yml

  # Install / Configure NTP
  - name: Install / Configure NTP
    include: ../includes/support_modules/ntp_setup.yaml

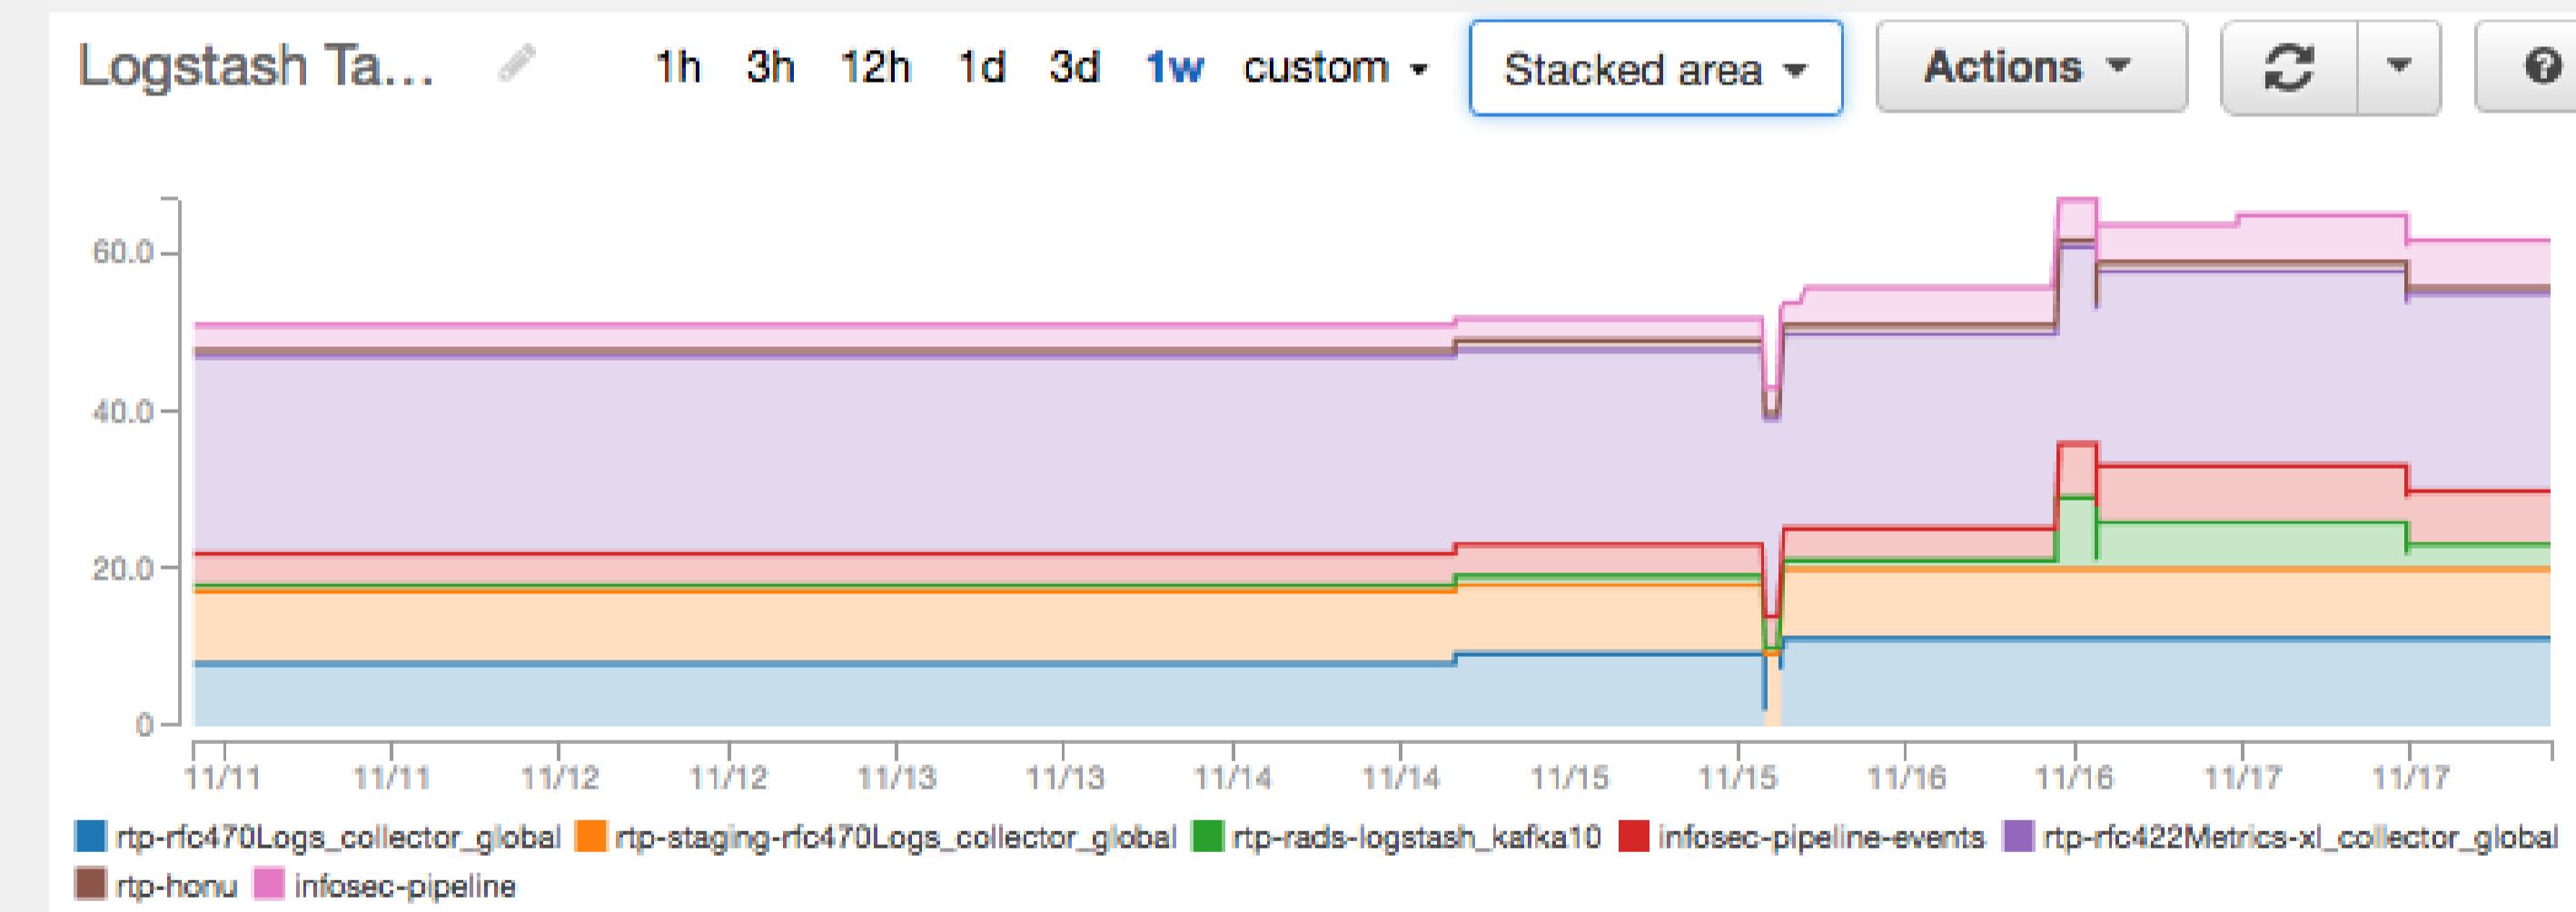
  # Perform Logstash Setup Operations
  - name: Perform Logstash Setup
    include: ../includes/logstash_setup.yaml
```



LESSONS

WHAT WORKED FOR US

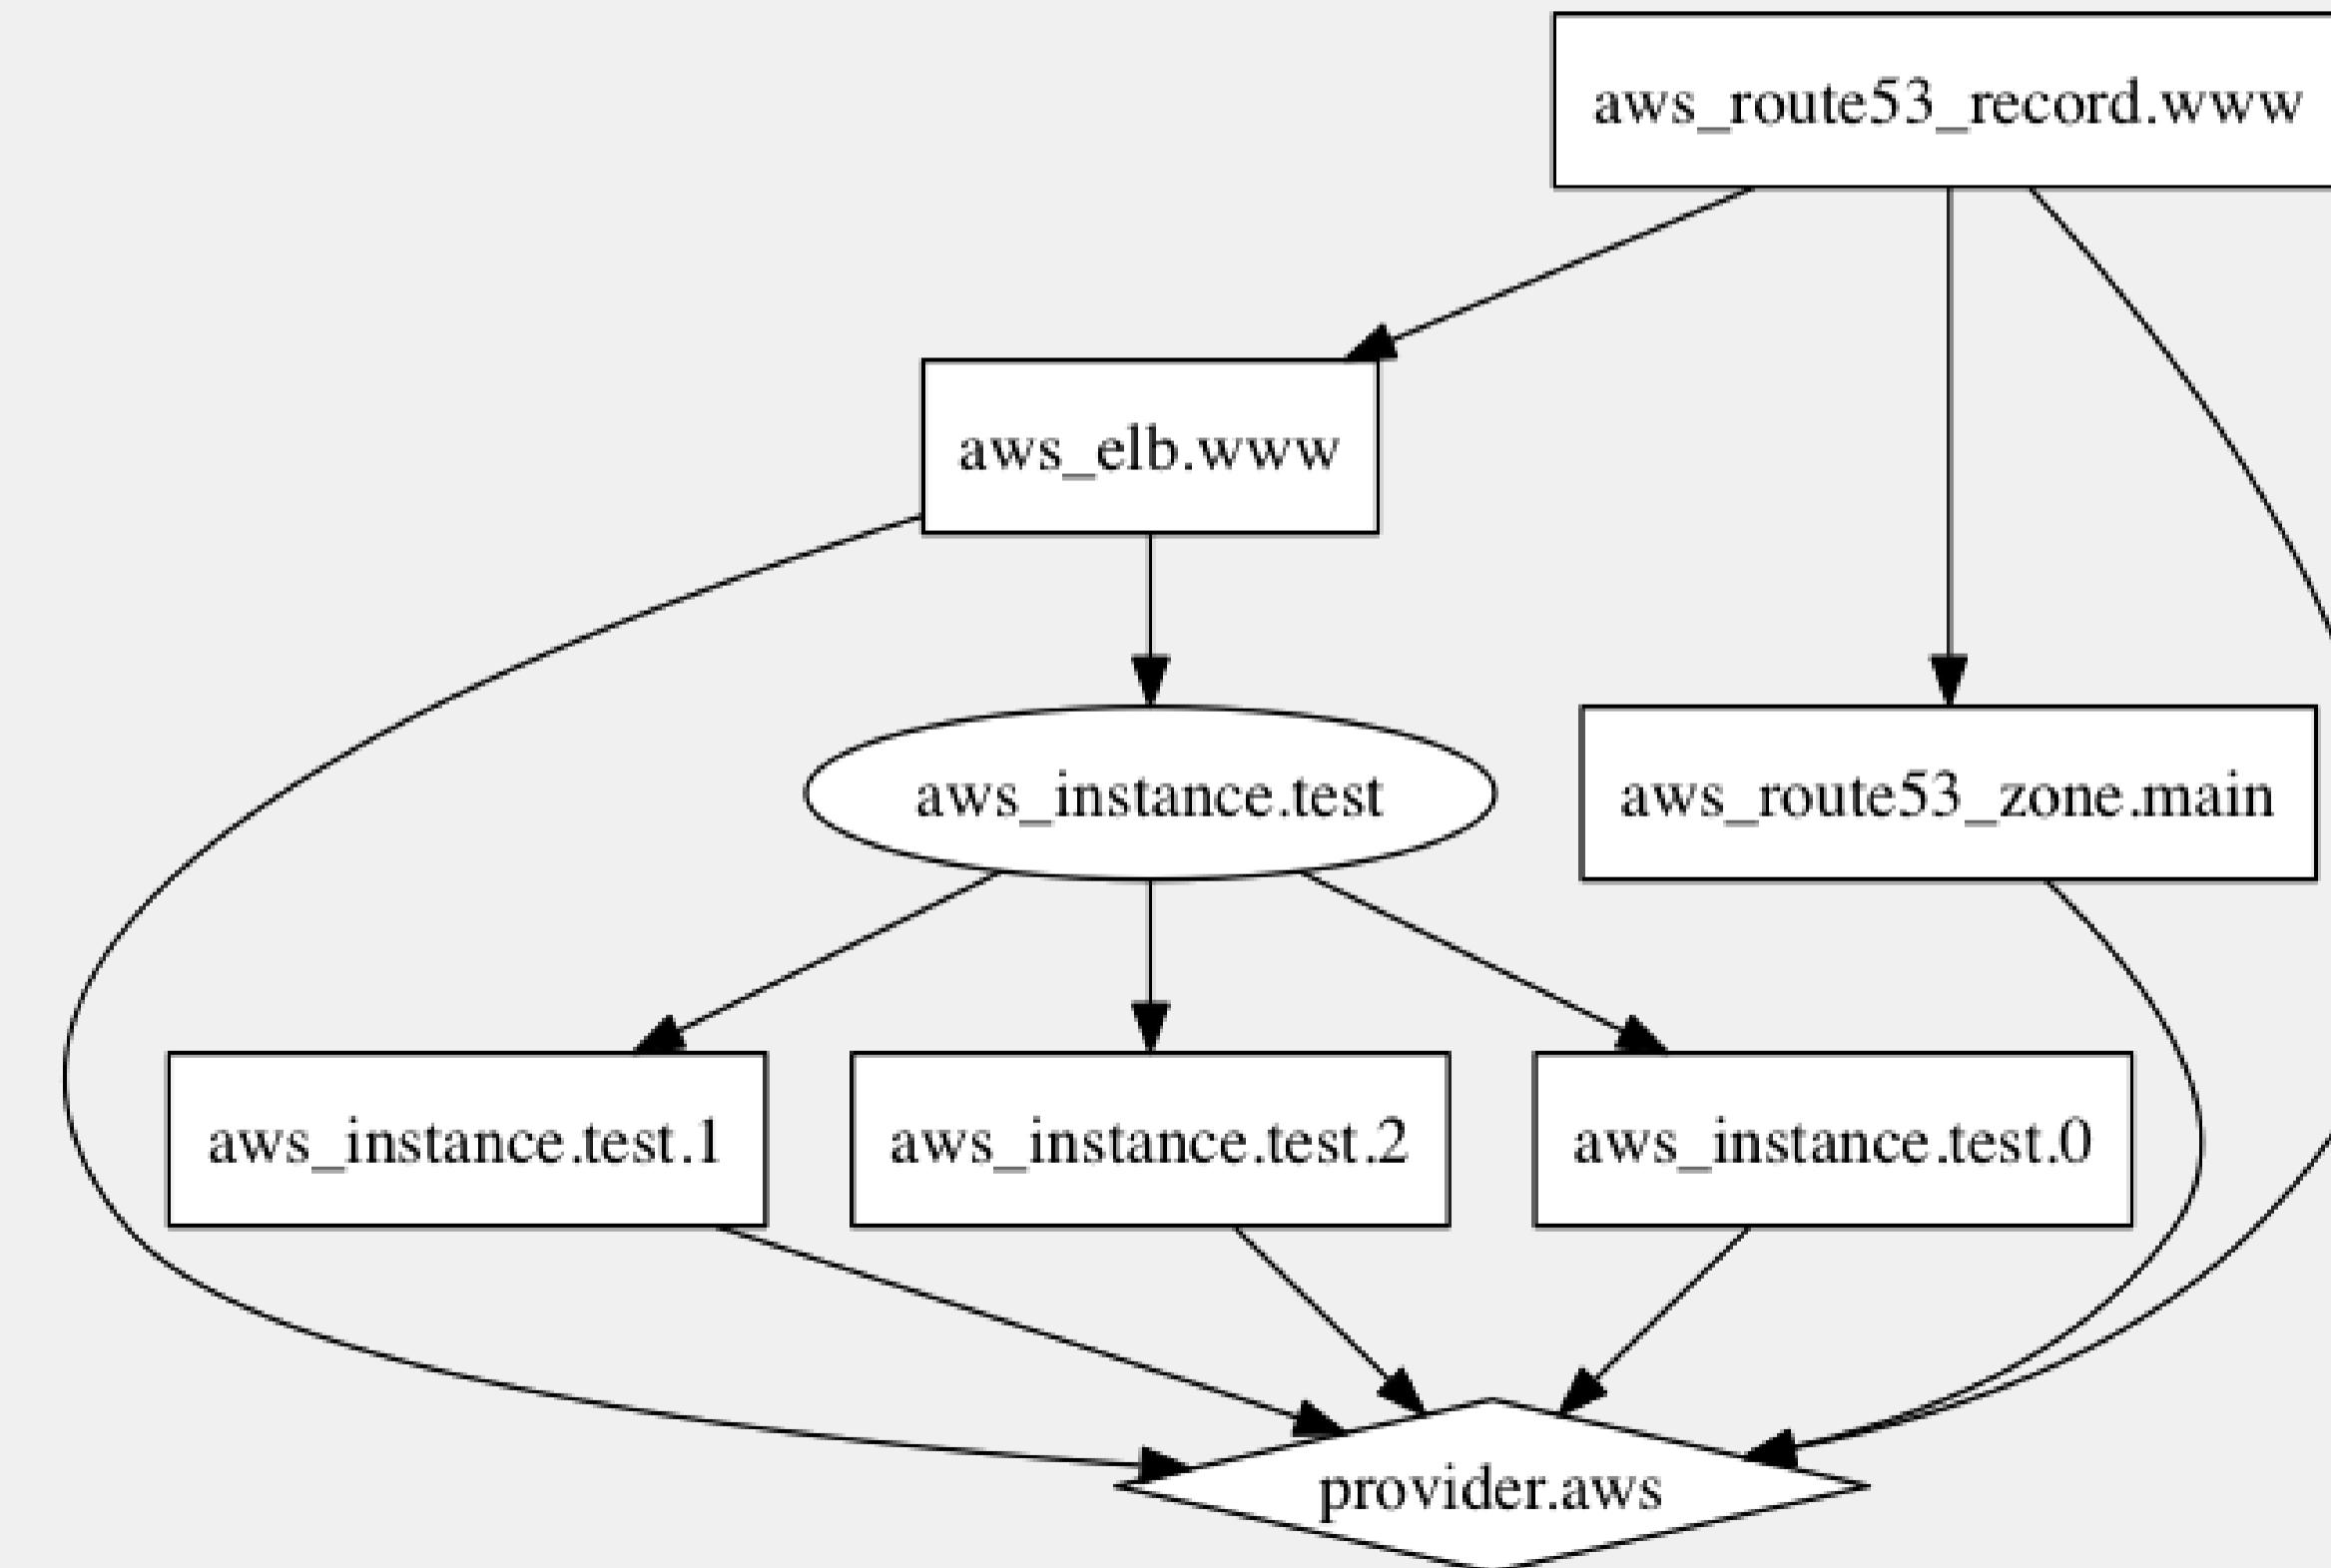
- Profile your memory requirements, monitor for scheduling issues



LESSONS

WHAT WORKED FOR US

- Take advantage of flow control in Terraform
- Use the Terraform graph tool to visualize your dependencies



RIOT ENGINEERING BLOG

THESE PROBLEMS AND MORE

Riot engineering



<https://engineering.riotgames.com/>

How we use data



CLAIRVOYANCE BLOG

<http://na.leagueoflegends.com/en/tag/insights>

ENJOY NOMS.

DRINK DRINKS.

PLAY GAMES.

**Tuesday
November 29th, 2016**

6:30 – 10:30PM

**Paiza Lounge
The Venetian
36th Floor**



JOIN US!



AWS
re:Invent

Thank you!



Remember to complete
your evaluations!