



flint

# Distributed Time Series Analysis Framework For Spark

Larisa Sawyer

Two Sigma



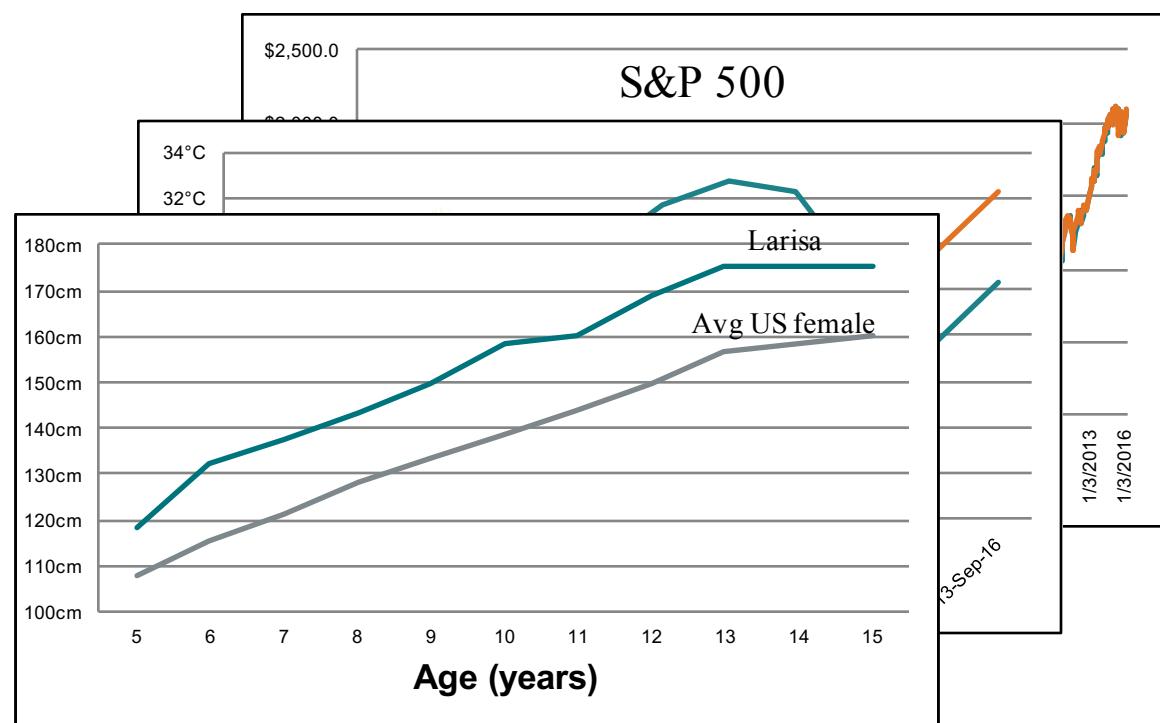
Larisa Sawyer



**TWO SIGMA**

## Time series examples

- ◆ Stock market prices
- ◆ Temperatures
- ◆ Height
- ◆ ...



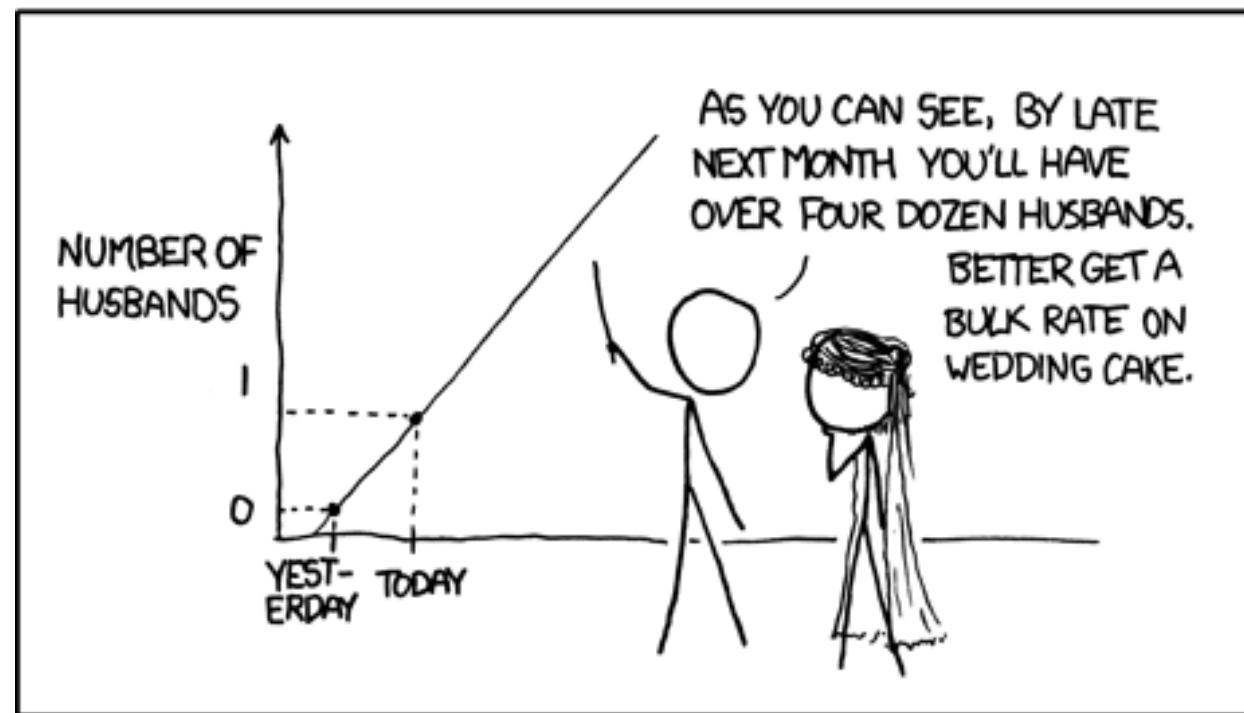
# What do we do with time series data?

- Forecast future values given past observations



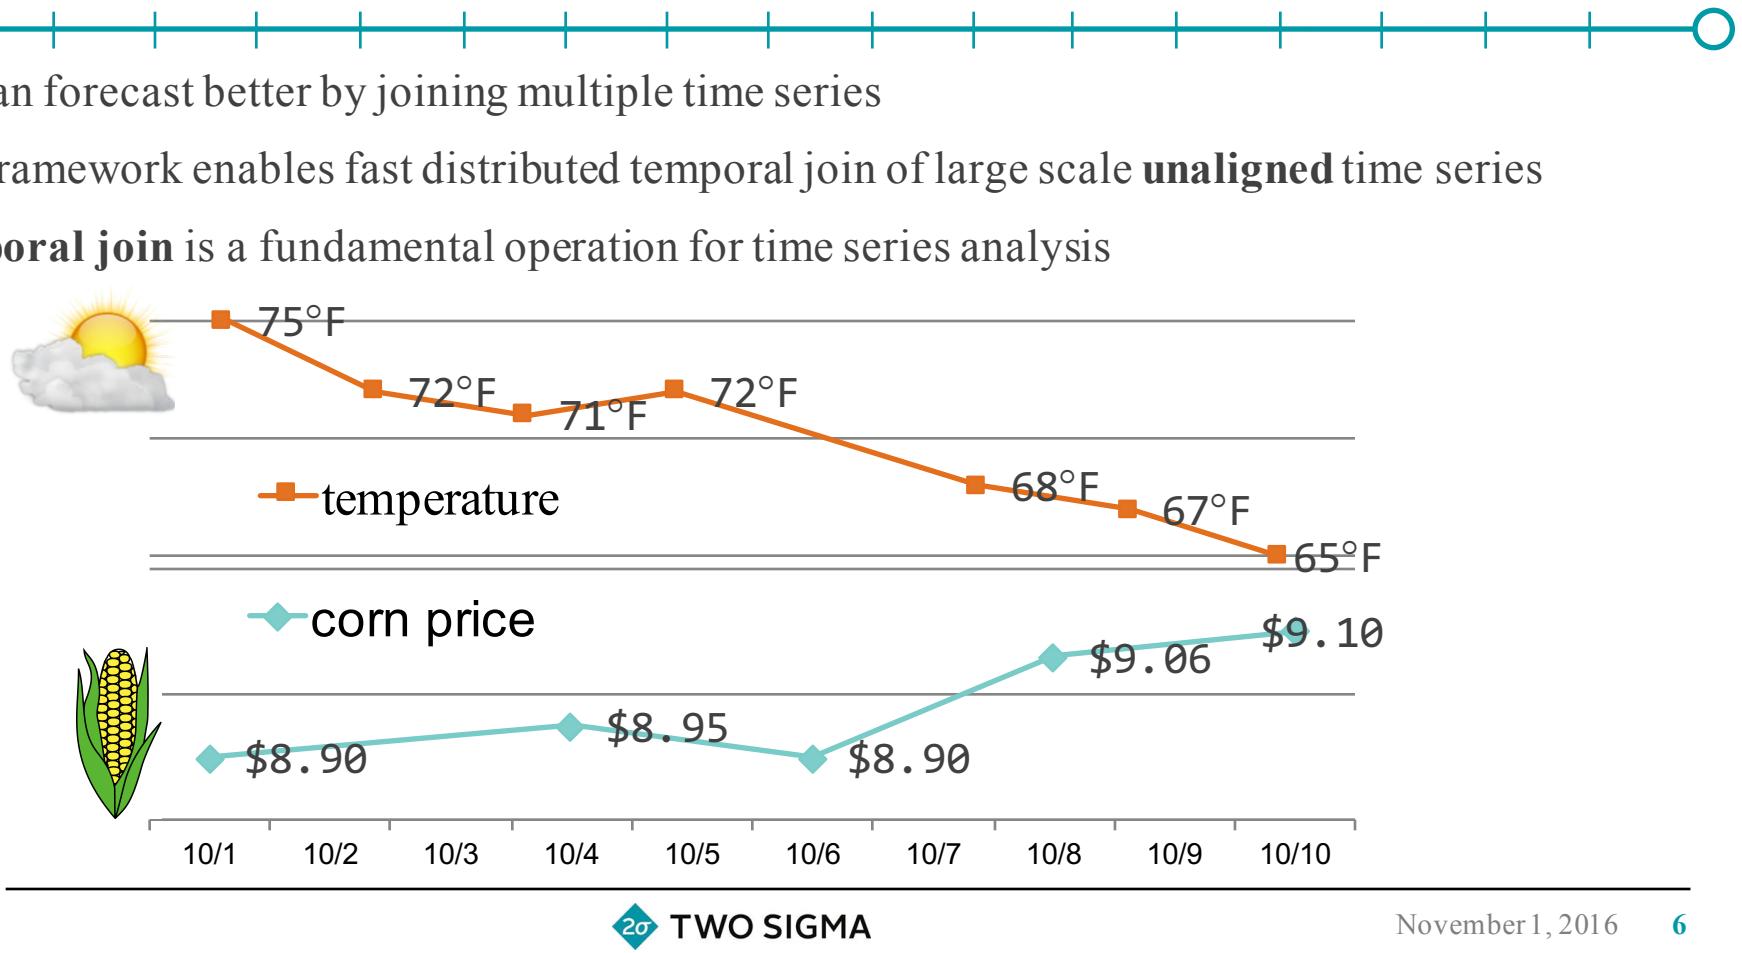
# Univariate time series

## MY HOBBY: EXTRAPOLATING



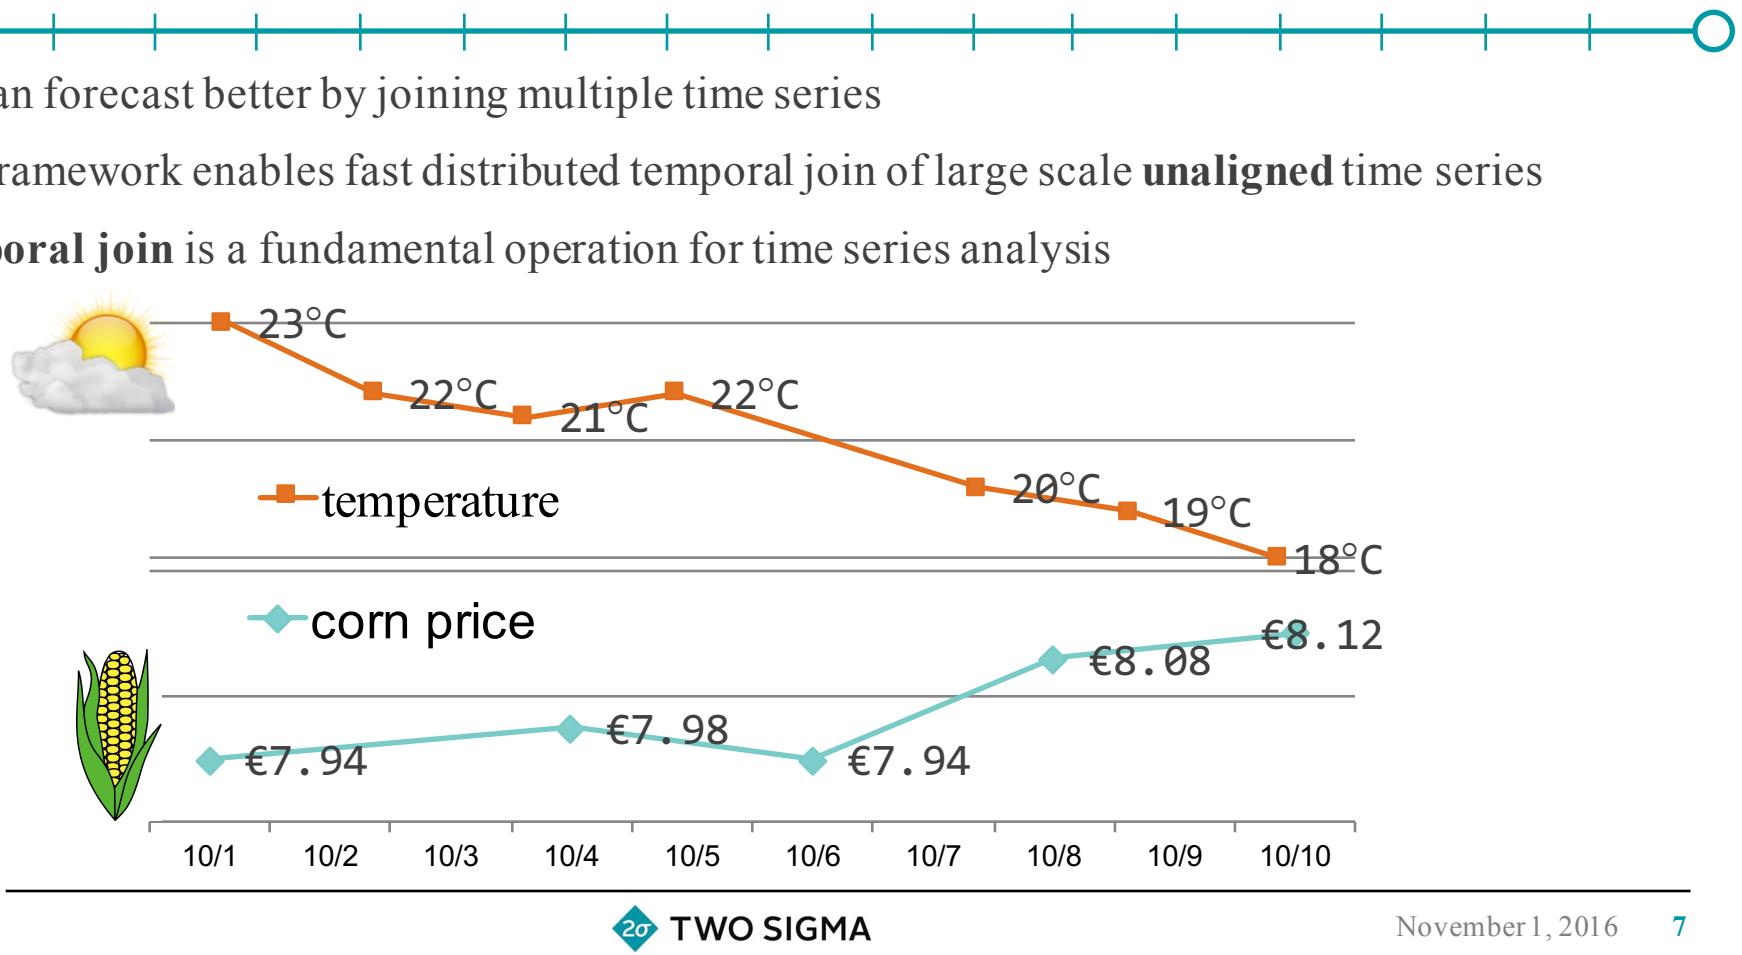
## Multivariate time series

- We can forecast better by joining multiple time series
- Our framework enables fast distributed temporal join of large scale **unaligned** time series
- **Temporal join** is a fundamental operation for time series analysis

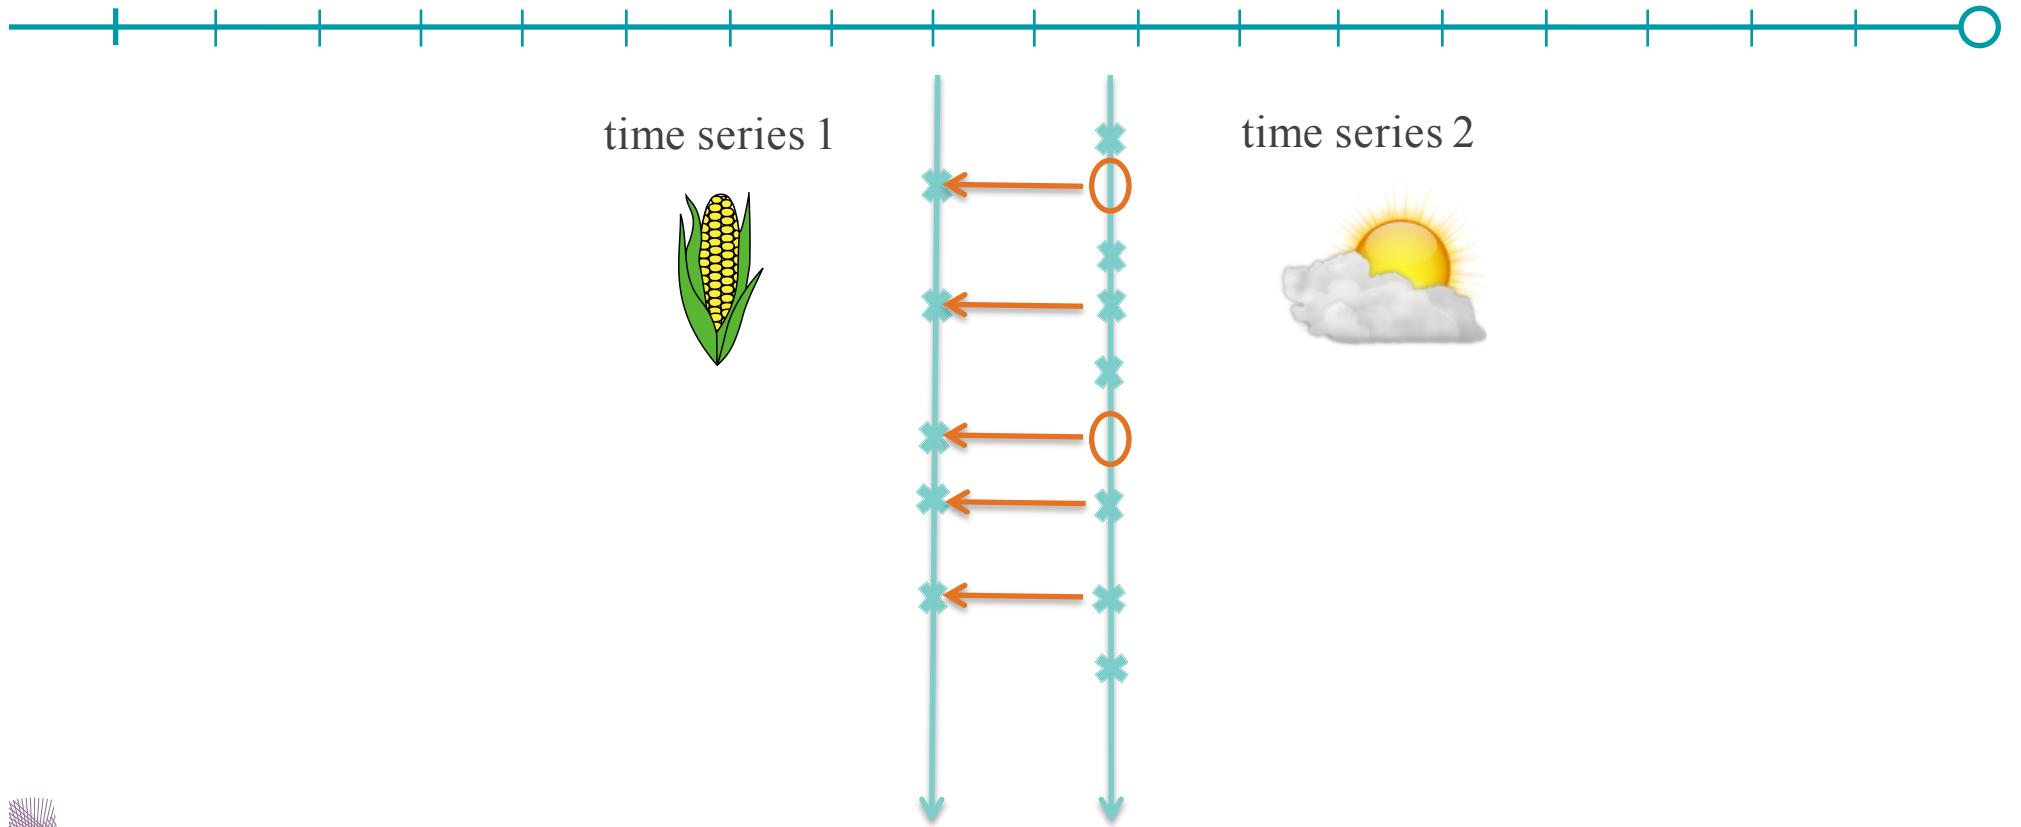


## Multivariate time series

- We can forecast better by joining multiple time series
- Our framework enables fast distributed temporal join of large scale **unaligned** time series
- **Temporal join** is a fundamental operation for time series analysis



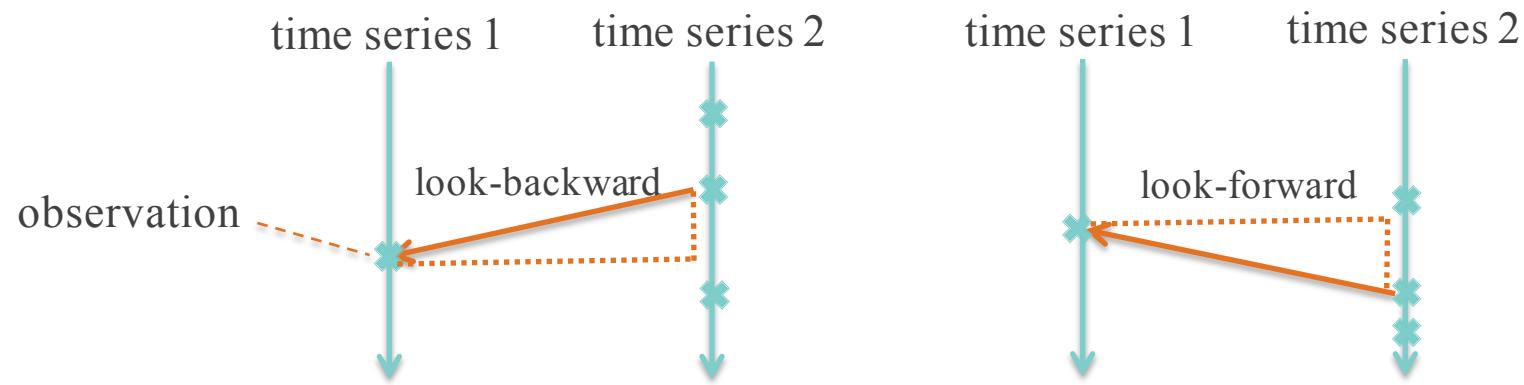
# What is a left join?



# What is temporal join?



- A particular **join** function defined by a matching criteria over **time**
- Examples of criteria
  - **look-backward**
  - **look-forward**



## Temporal join with look-backward criteria



time	tweets
08:00 AM	
10:00 AM	
12:00 PM	

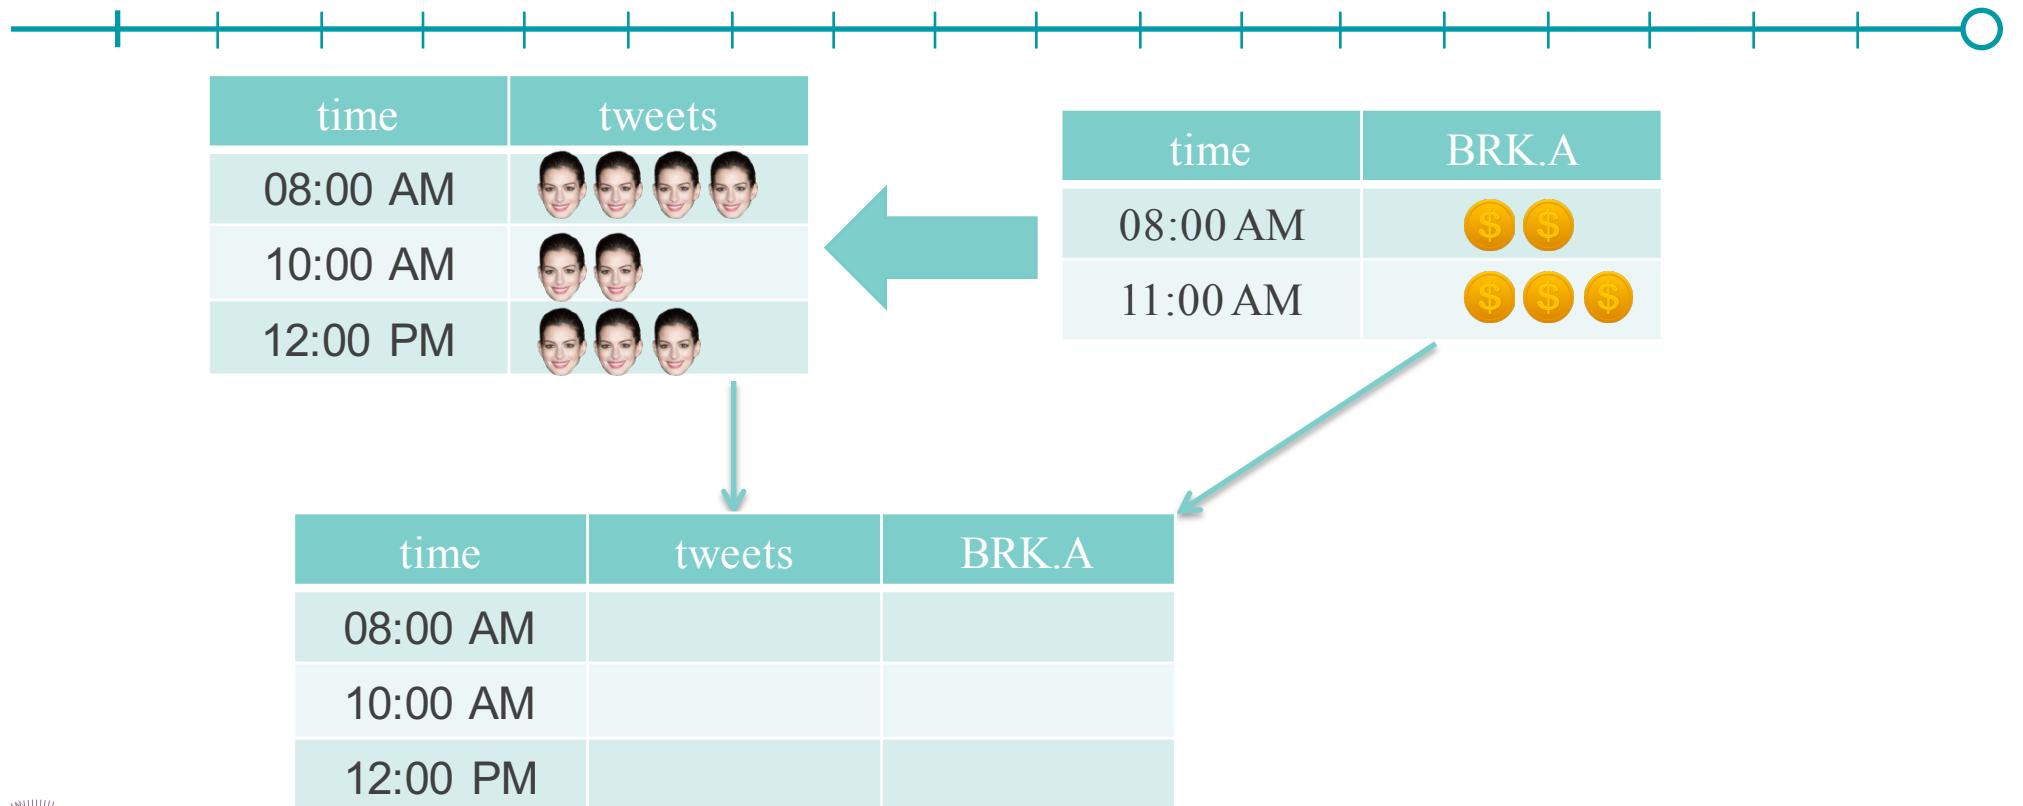
time	BRK.A
08:00 AM	
11:00 AM	

## Important Legal Information



The information presented here is offered for informational purposes only and should not be used for any other purpose (including, without limitation, the making of investment decisions). Examples provided herein are for illustrative purposes only and are not necessarily based on actual data. Nothing herein constitutes an offer to sell or the solicitation of any offer to buy any security or other interest. We consider this information to be confidential and not for redistribution or dissemination.

## Temporal join with look-backward criteria



## Temporal join with look-backward criteria

time	tweets
08:00 AM	
10:00 AM	
12:00 PM	

time	BRK.A
08:00 AM	
11:00 AM	

time	tweets	BRK.A
08:00 AM		
10:00 AM		
12:00 PM		

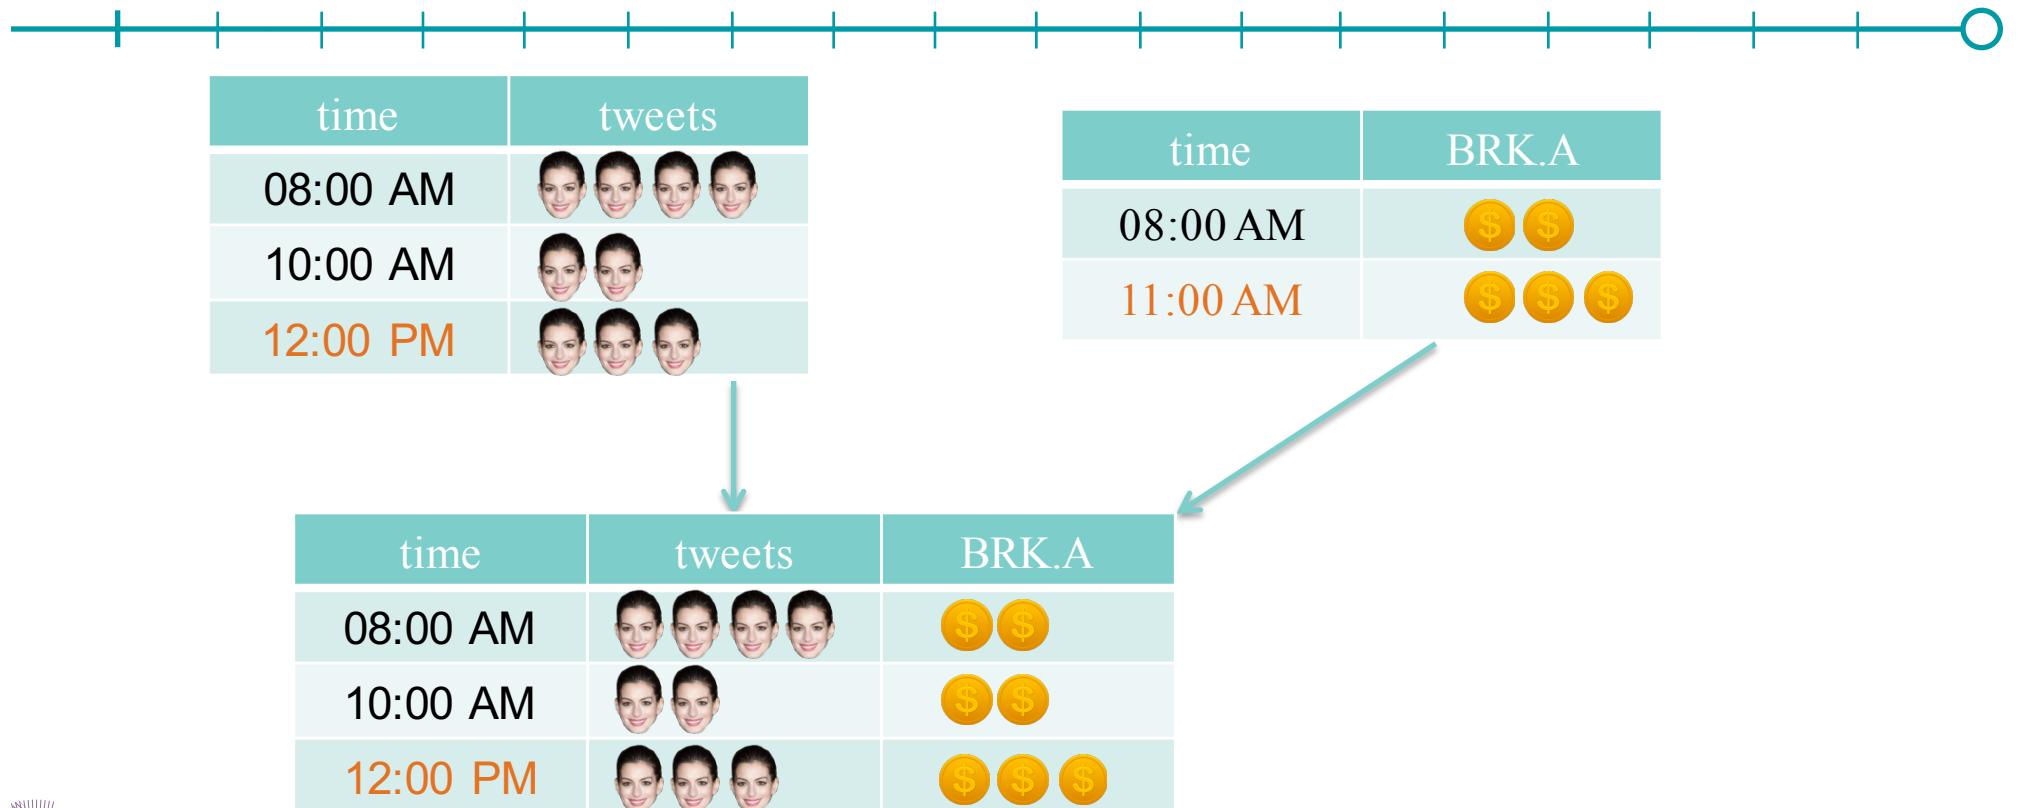
## Temporal join with look-backward criteria

time	tweets
08:00 AM	
10:00 AM	
12:00 PM	

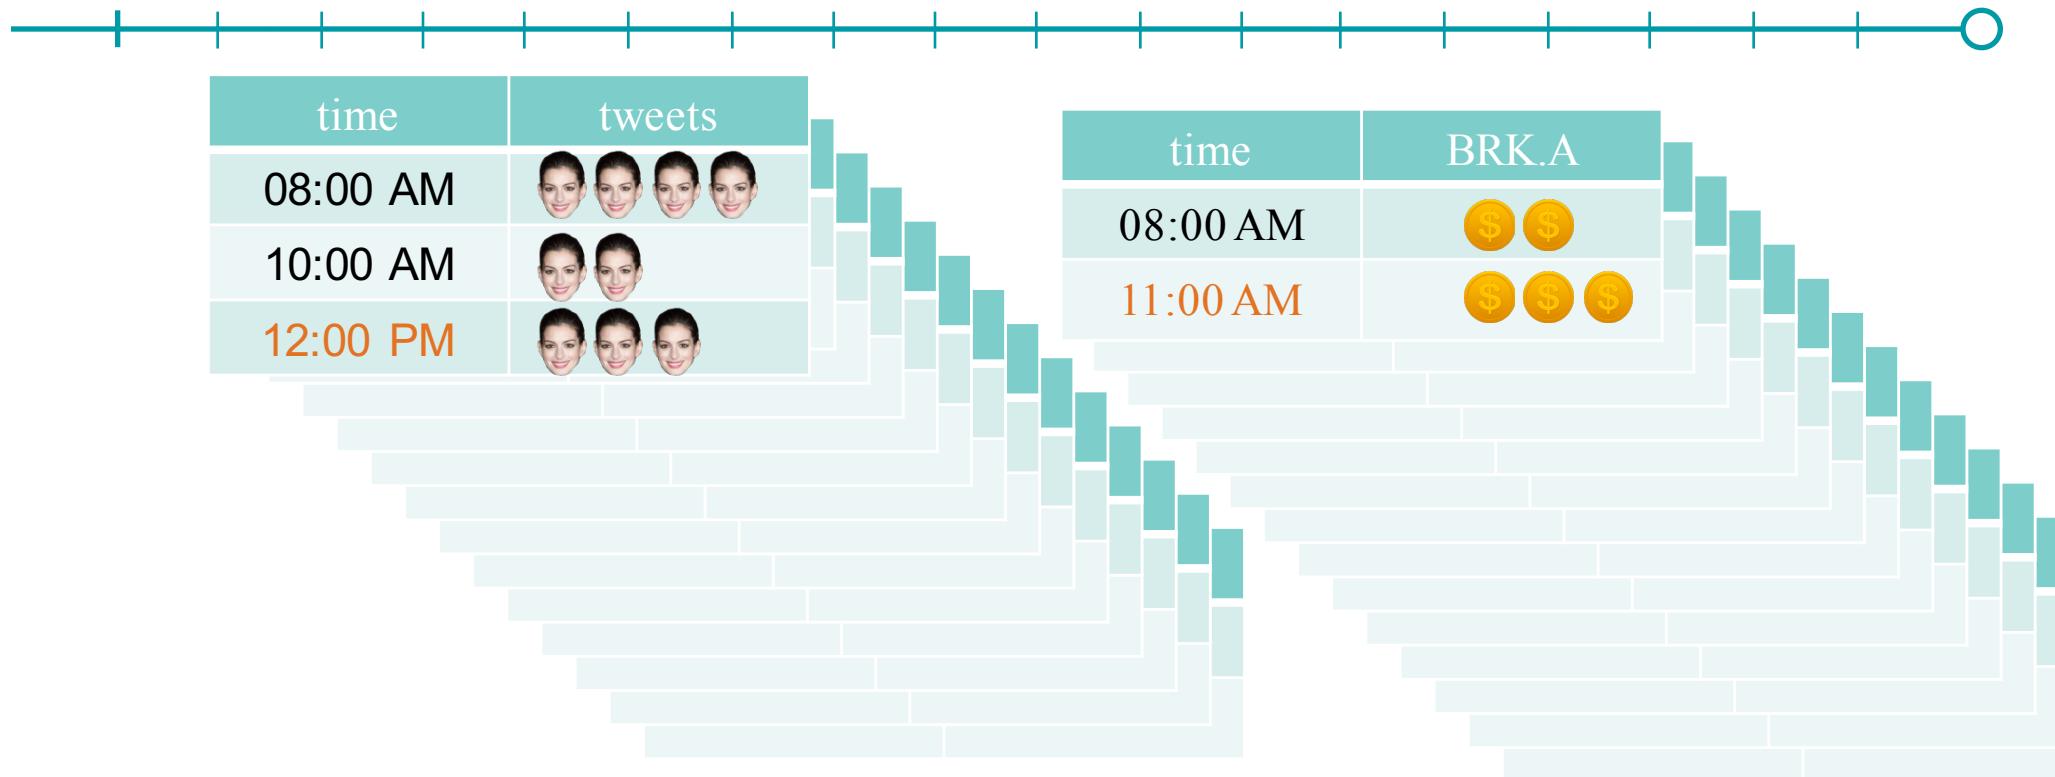
time	BRK.A
08:00 AM	
11:00 AM	

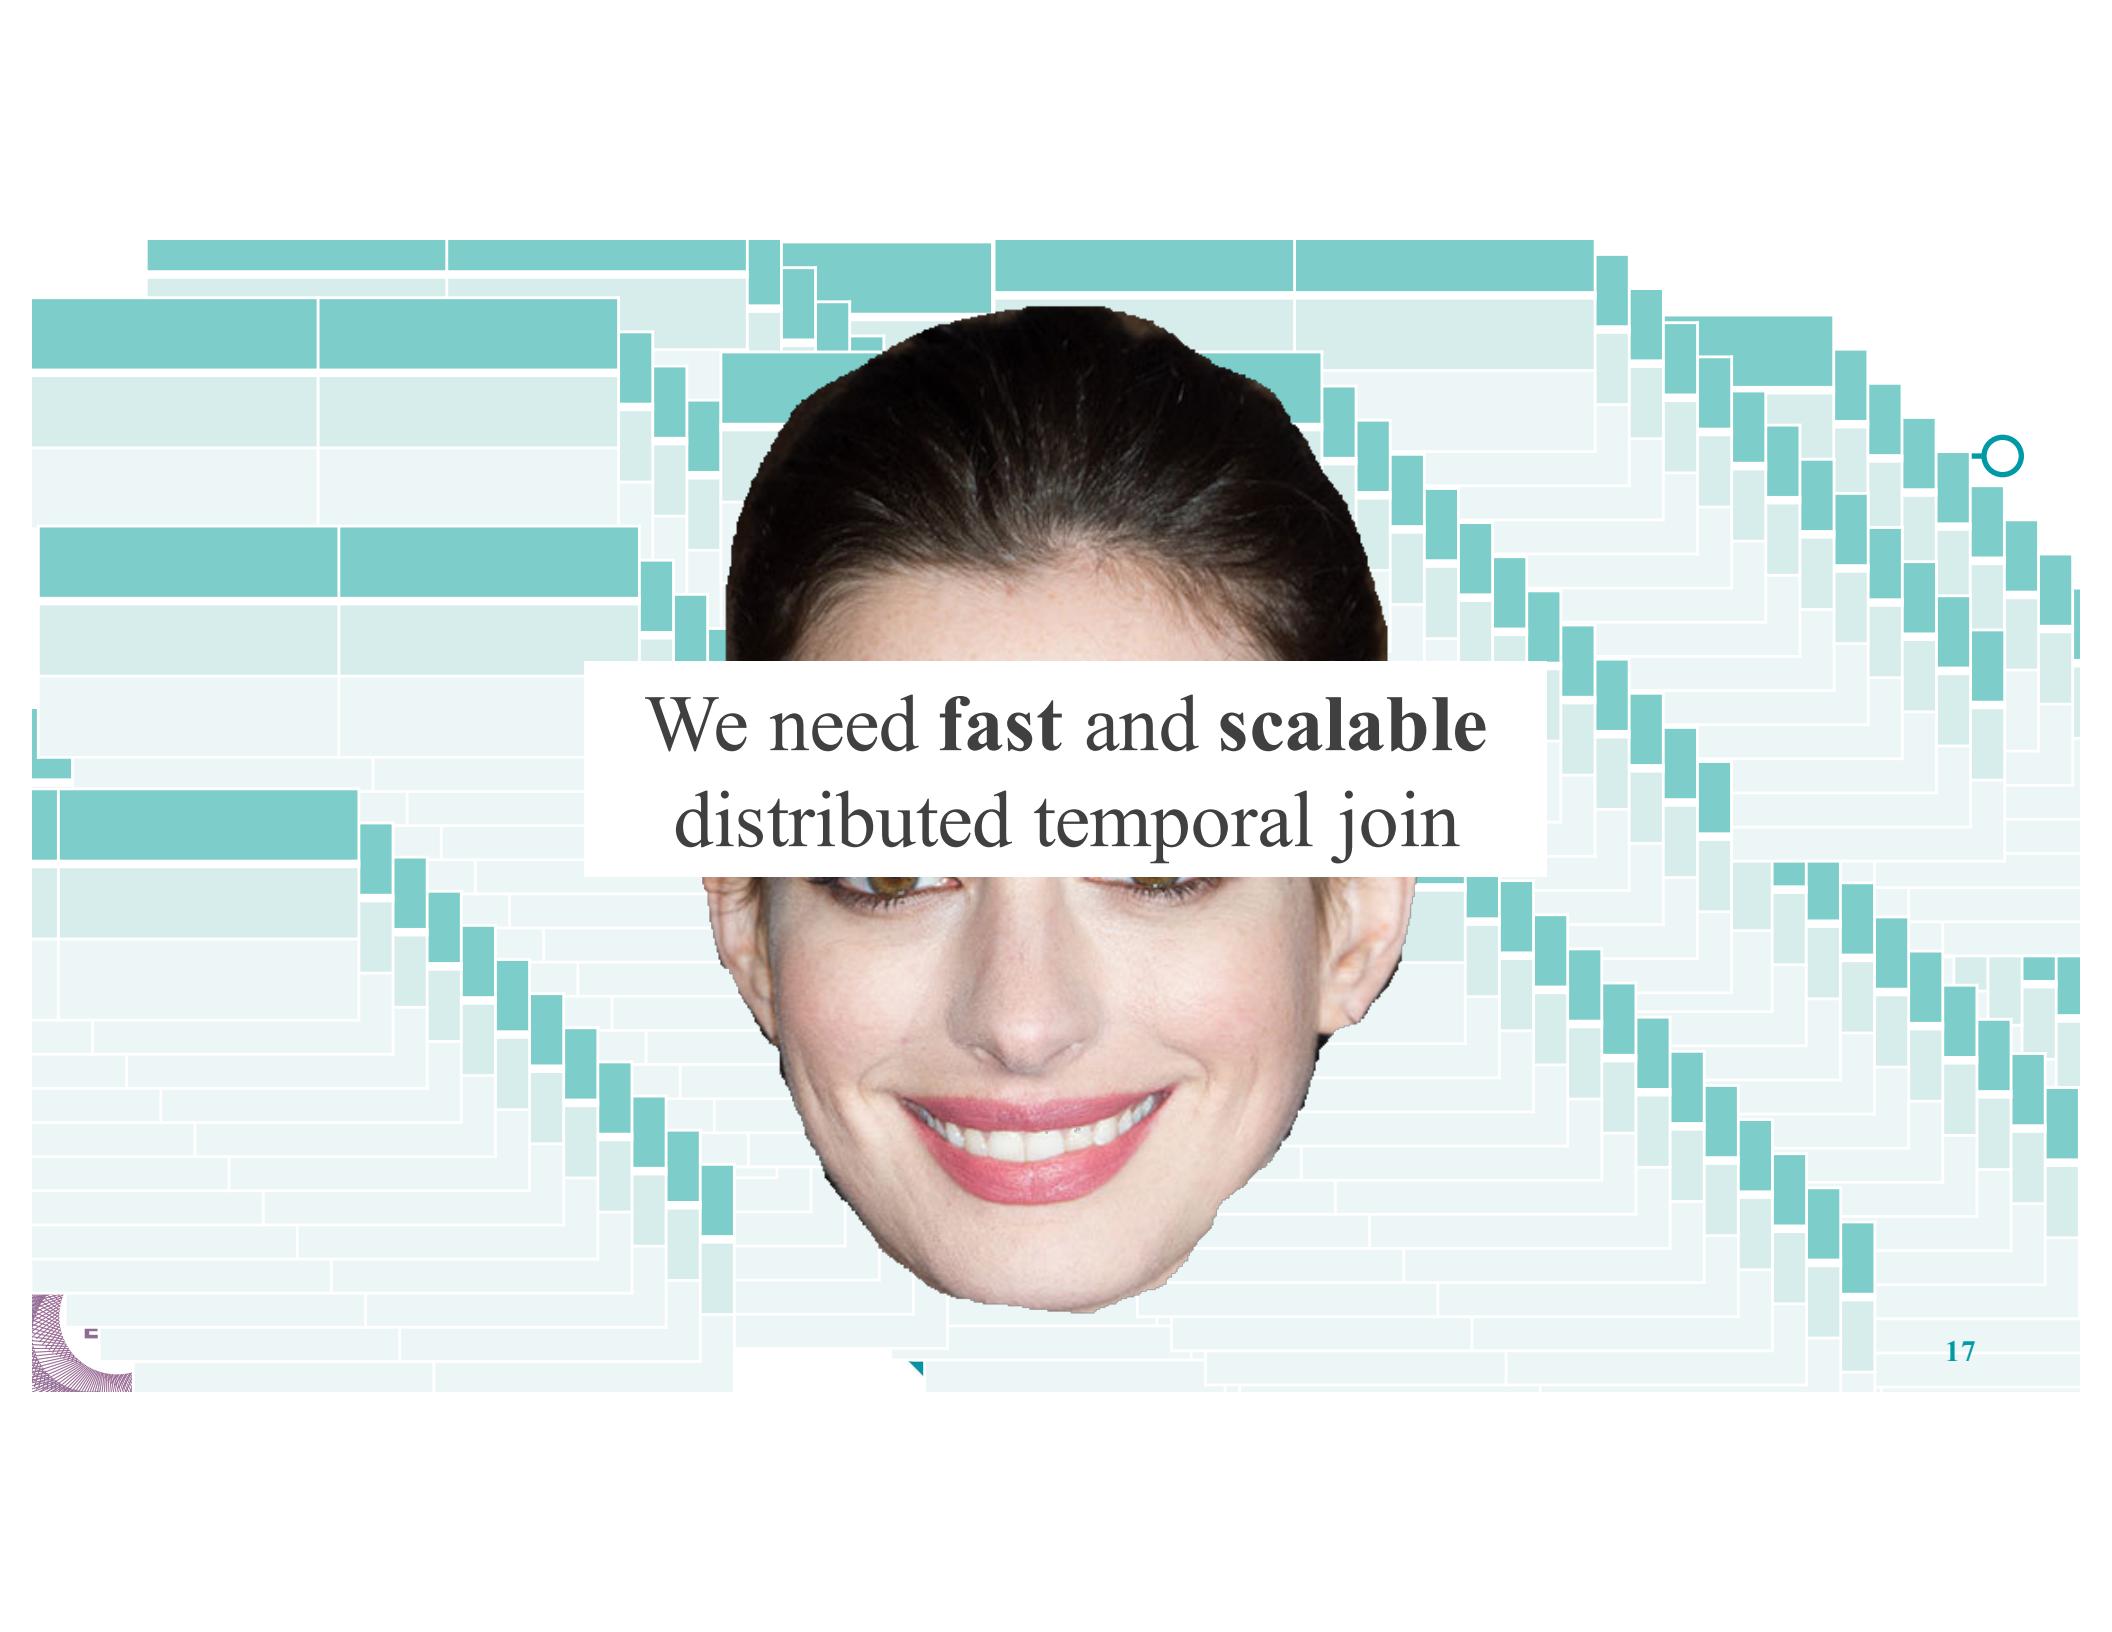
time	tweets	BRK.A
08:00 AM		
10:00 AM		
12:00 PM		

## Temporal join with look-backward criteria



## Temporal joins in practice





We need **fast** and **scalable**  
distributed temporal join

# Existing solutions



- ◆ Existing packages don't support temporal join or can't handle large time series
  - ◆ Pandas / R / Matlab
    - ◆ Limited to single machine
  - ◆ Spark
    - ◆ Does scale, but all data is unordered
  - ◆ spark-ts
    - ◆ Expects univariate time series to fit on single machine
    - ◆ Splits by col
    - ◆ Supports only snapshot data

# Flint: A new time series library for Spark



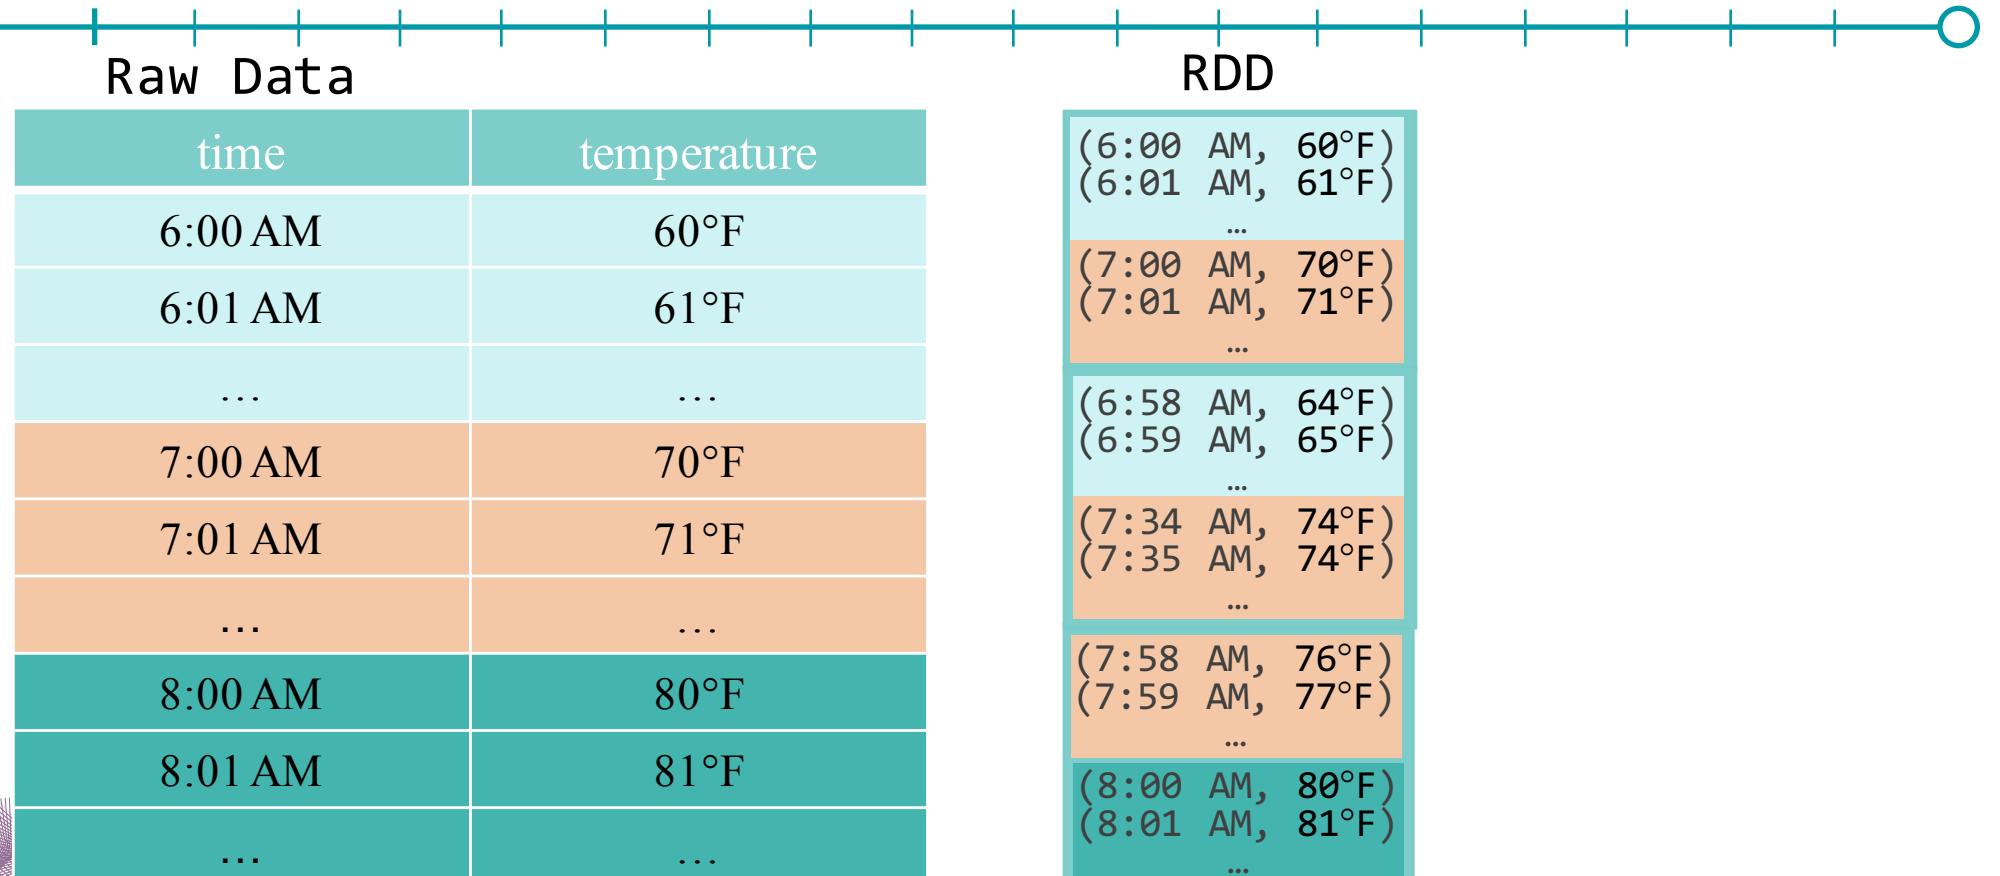
- ◆ Goal
  - ◆ Provide a collection of functions to **manipulate** and **analyze** time series at scale
    - ◆ Group, temporal join, summarize, aggregate ...
- ◆ How
  - ◆ Build a time series aware data structure
    - ◆ `TimeSeriesRDD` extends `RDD`
  - ◆ Optimize using temporal locality
    - ◆ Reduce shuffling
    - ◆ Reduce memory pressure by streaming



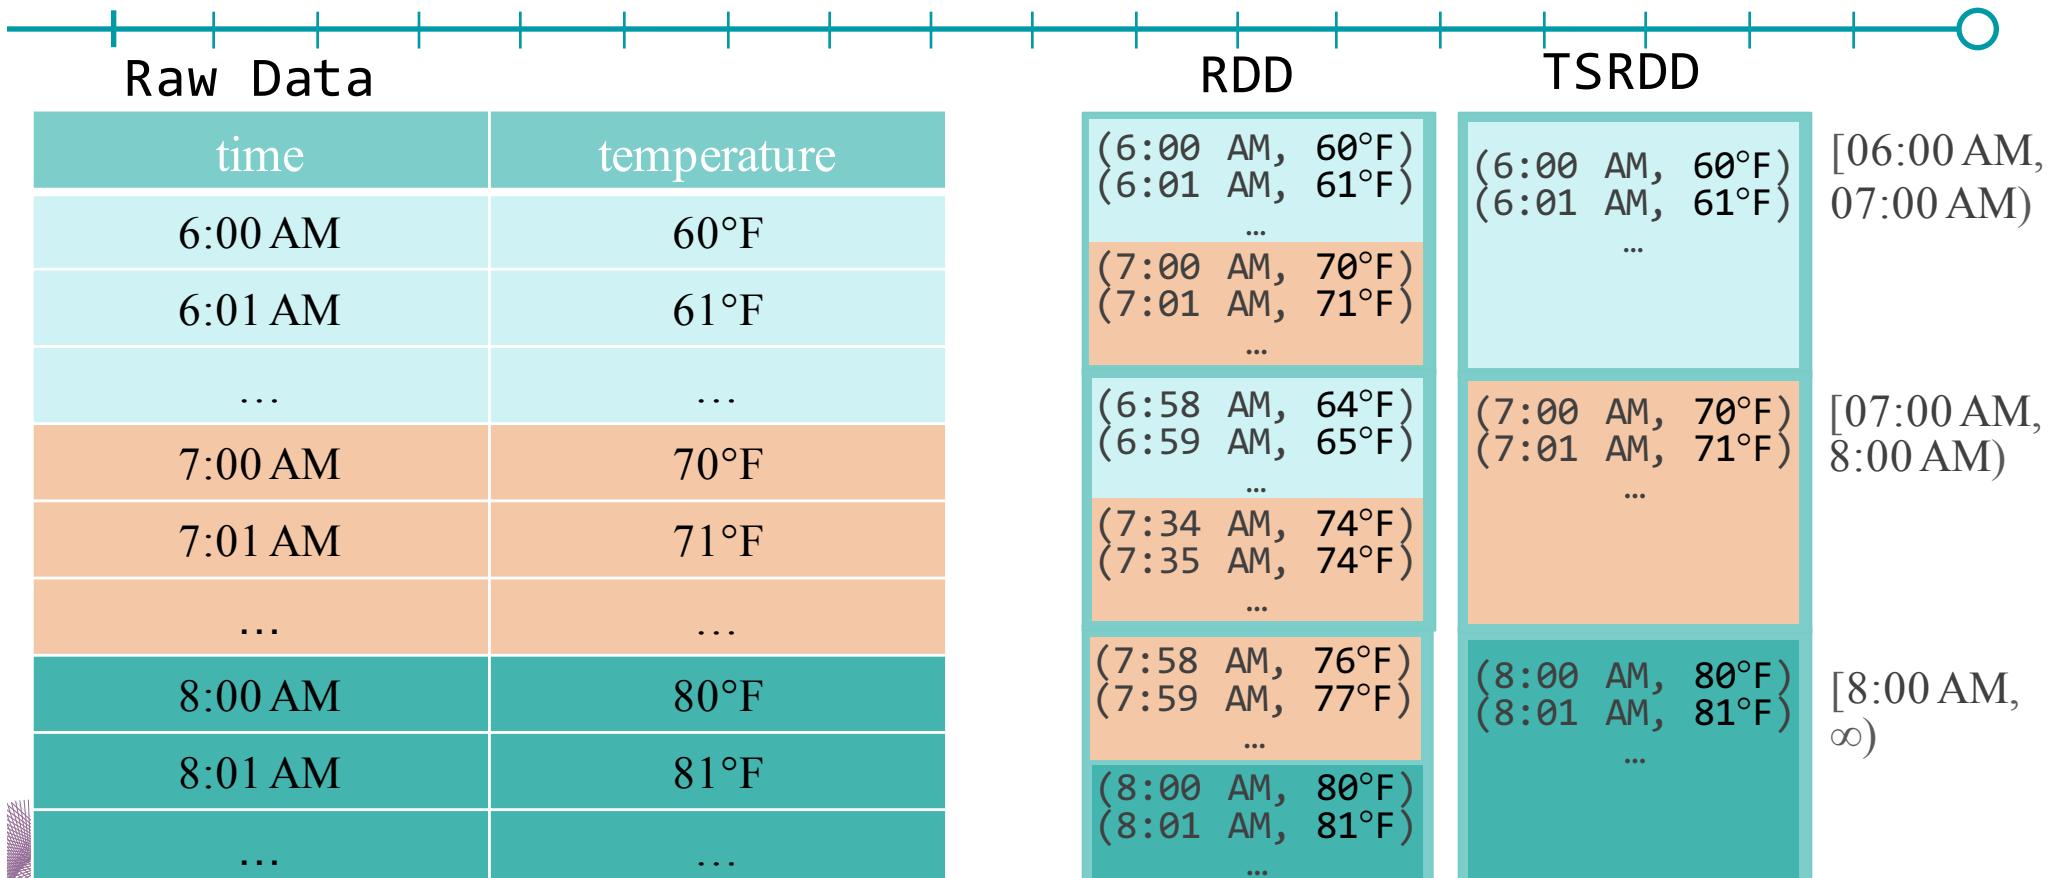
# What is a TimeSeriesRDD?

- 
- ◆ TimeSeriesRDD vs RDD
    - ◆ Associate time range on each partition
    - ◆ Track partition time-ranges
    - ◆ Preserve temporal order

# RDD



## TimeSeriesRDD



# Group function

- A **group** function groups rows with exactly the same timestamps

time	city	temperature
1:00 PM	New York	70°F
1:00 PM	Brussels	60°F
2:00 PM	New York	71°F
2:00 PM	Brussels	61°F
3:00 PM	New York	72°F
3:00 PM	Brussels	62°F
4:00 PM	New York	73°F
4:00 PM	Brussels	63°F

The diagram illustrates the grouping of data by timestamp. A horizontal timeline is shown with 12 vertical tick marks. Four orange curly braces group the data into four categories labeled group 1, group 2, group 3, and group 4. Group 1 contains the first two rows (1:00 PM). Group 2 contains the next two rows (2:00 PM). Group 3 contains the next two rows (3:00 PM). Group 4 contains the last two rows (4:00 PM).

# Group function

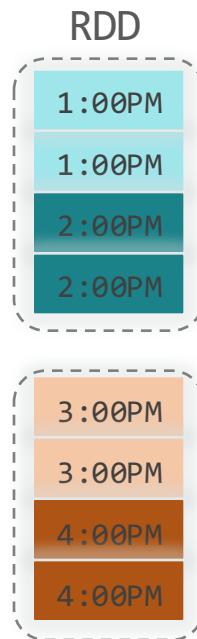
- A **group** function groups rows with nearby timestamps

time	city	temperature
1:00 PM	New York	70°F
1:00 PM	Brussels	60°F
2:00 PM	New York	71°F
2:00 PM	Brussels	61°F
3:00 PM	New York	72°F
3:00 PM	Brussels	62°F
4:00 PM	New York	73°F
4:00 PM	Brussels	63°F

The diagram illustrates a timeline with 13 vertical tick marks. Two groups of three data rows are highlighted: group 1 (rows 3-5) and group 2 (rows 6-8), both spanning from 2:00 PM to 4:00 PM.

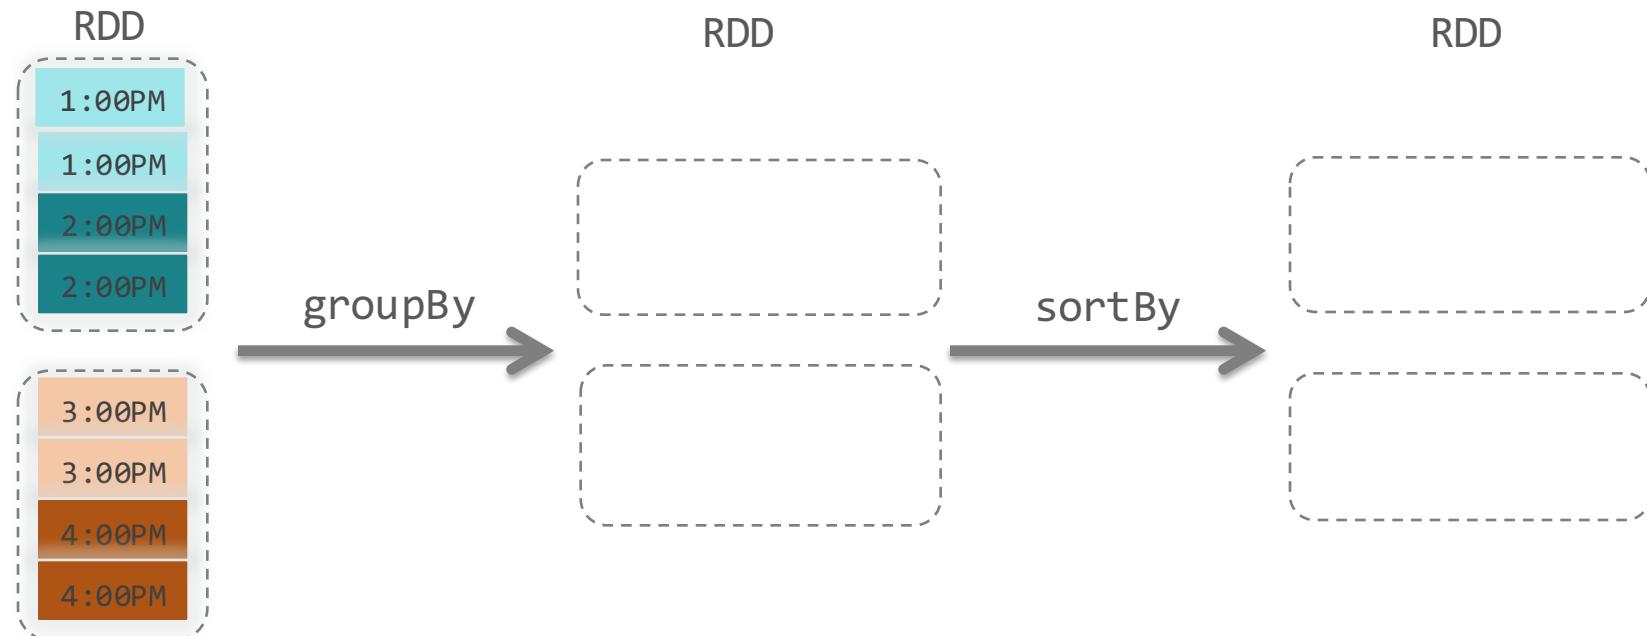
# Group in Spark

- Groups rows with exactly the same timestamps



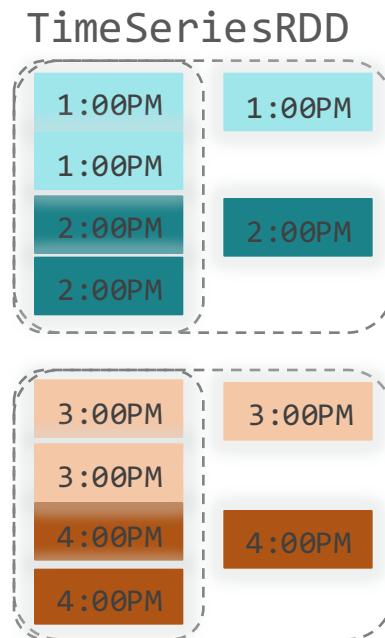
# Group in Spark

- Back to Temporal Order
- Data is shuffled and materialized on the workers



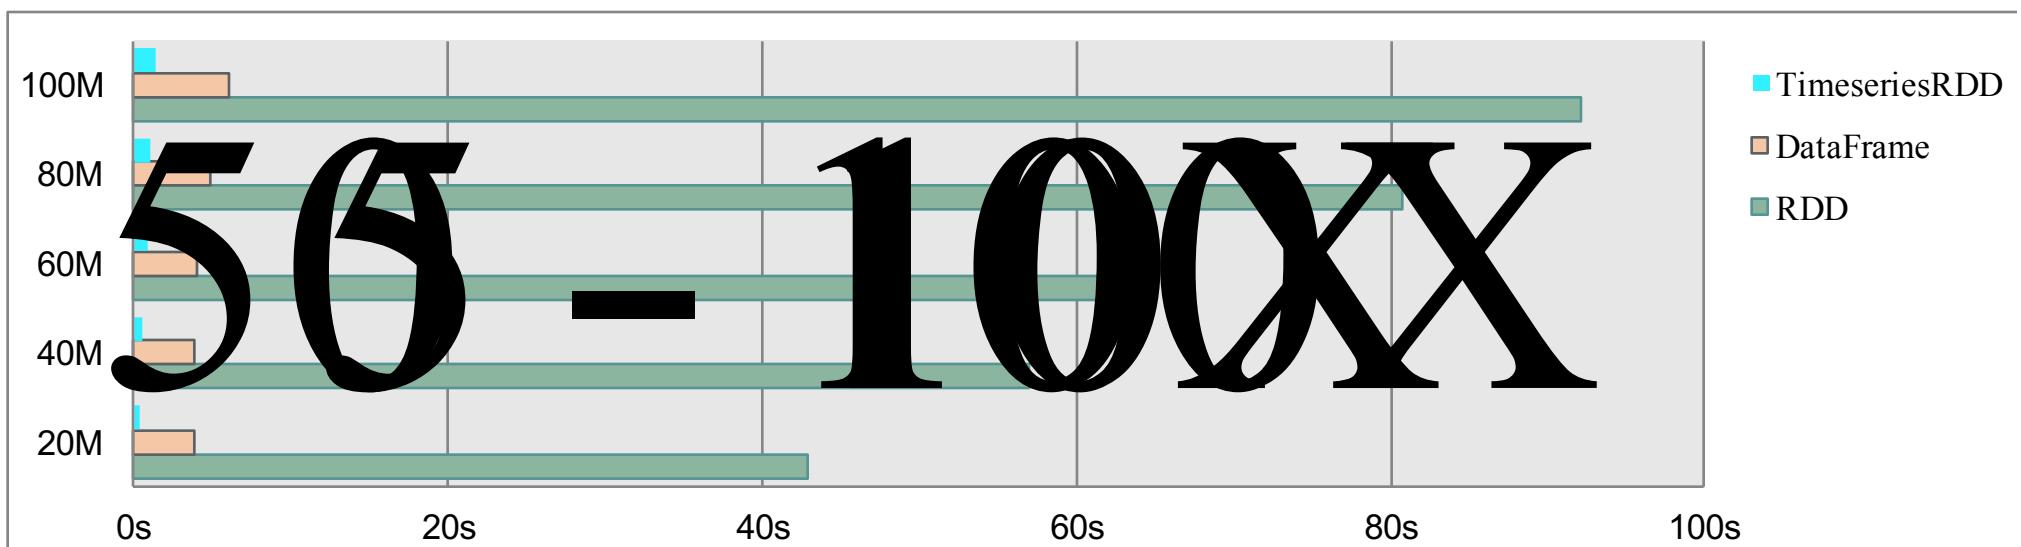
## Group in TimeSeriesRDD

- Data is grouped per partition locally as streams



## Benchmark for group + count

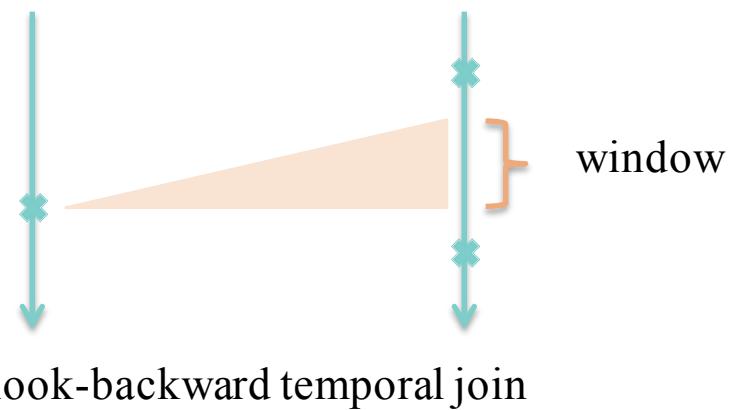
- Running time of **count** after **group**
  - 16 executors (10G memory and 4 cores per executor)
  - Data read from HDFS



## Temporal join



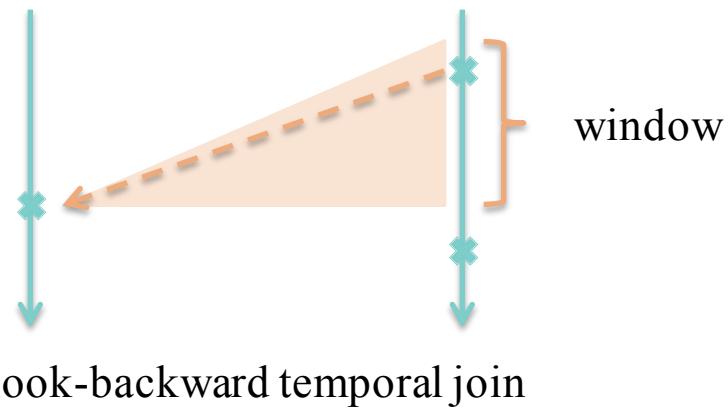
- A **temporal join** function is defined by a matching criteria over **time**
- A typical matching criteria has two parameters
  - **direction** – look-backward or look-forward
  - **window** – how much to look-backward or look-forward



## Temporal join



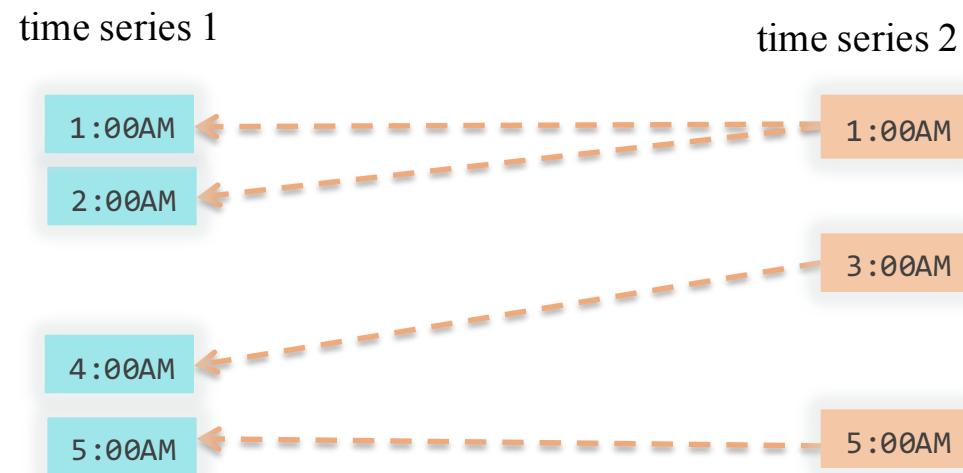
- A **temporal join** function is defined by a matching criteria over **time**
- A typical matching criteria has two parameters
  - **direction** – look-backward or look-forward
  - **window** – how much to look-backward or look-forward



## Temporal join

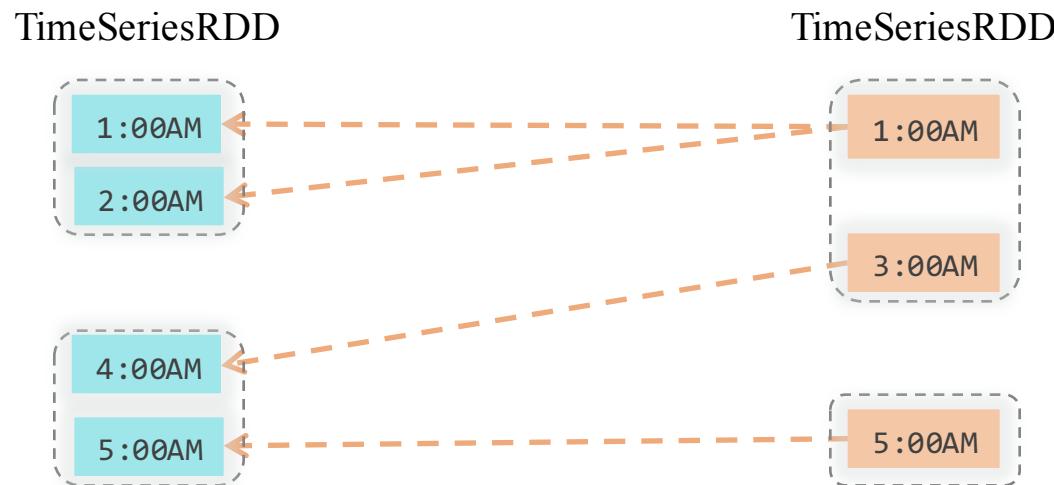


- Temporal join with criteria look-back and window of 1 hour



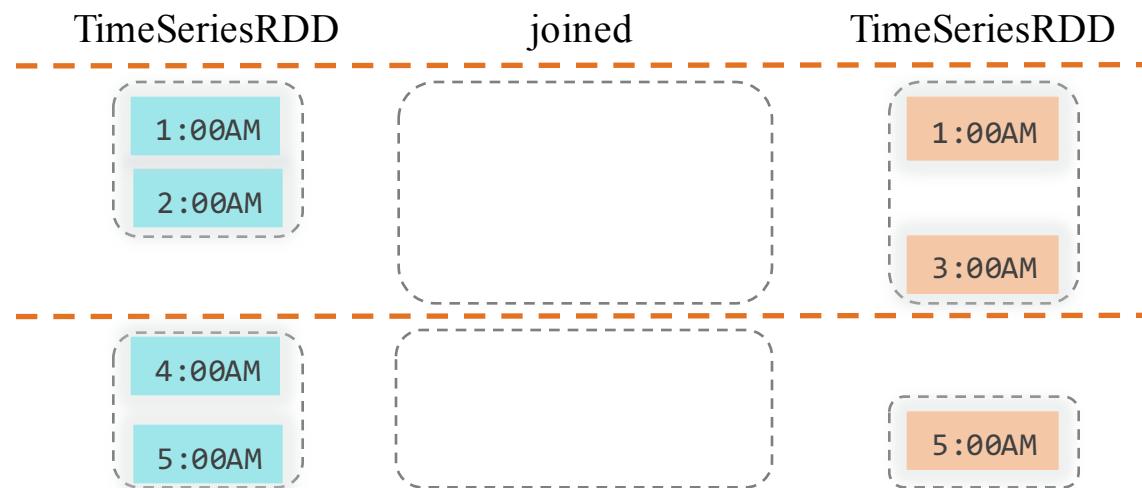
# Temporal join

- Temporal join with criteria look-back and window of 1 hour
  - How do we do temporal join in TimeSeriesRDD?



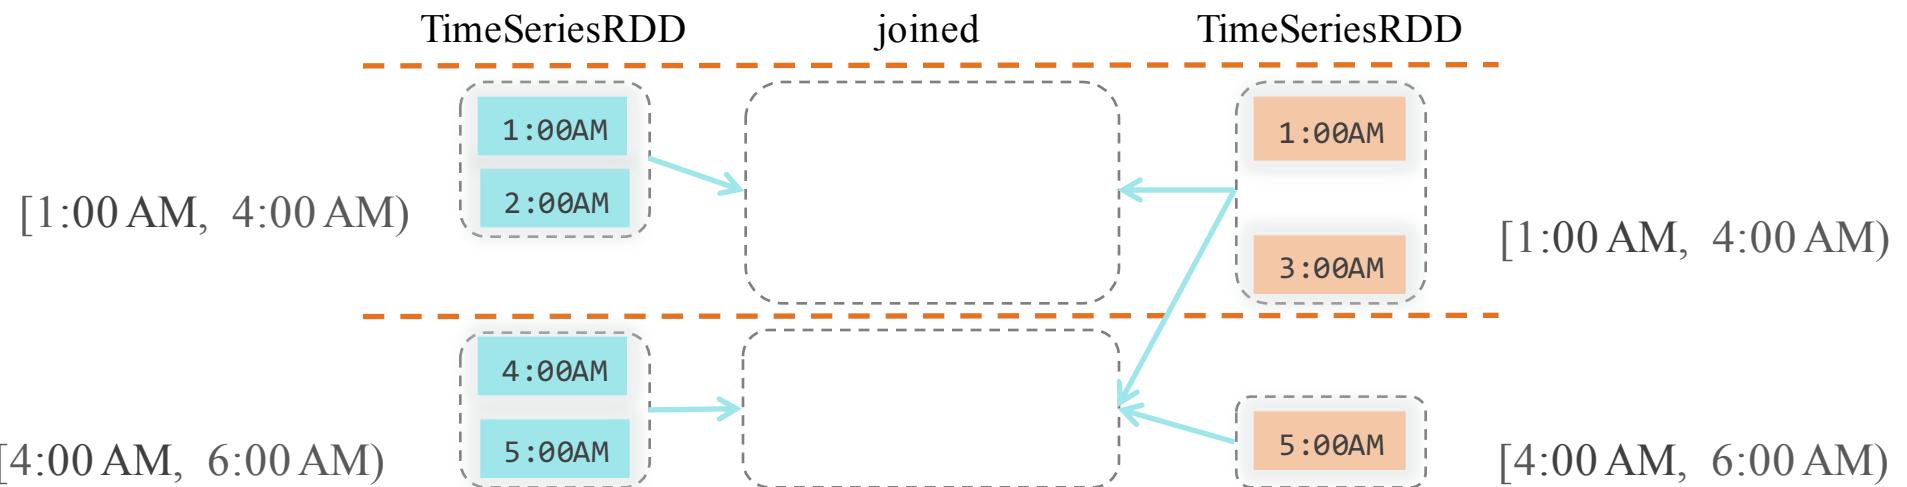
## Temporal join in TimeSeriesRDD

- Temporal join with criteria look-back and window of 1 hour
  - partition **time** space into disjoint intervals



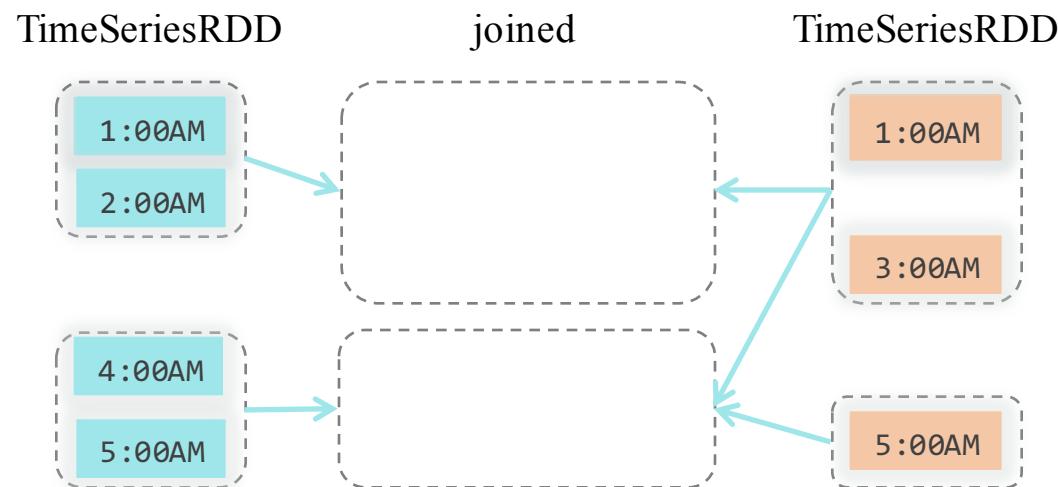
# Temporal join in TimeSeriesRDD

- Temporal join with criteria look-back and window of 1 hour
  - Build dependency graph for the joined TimeSeriesRDD



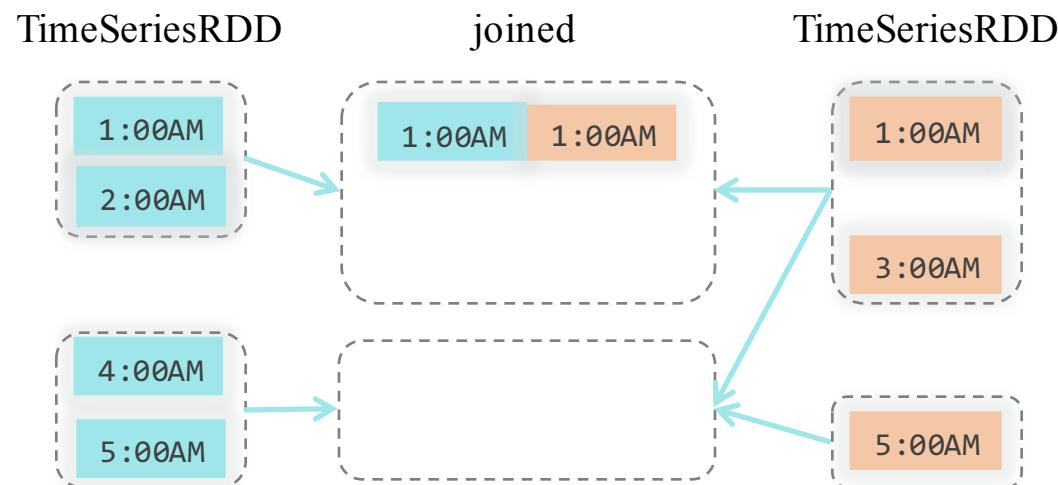
# Temporal join in TimeSeriesRDD

- Temporal join with criteria look-back and window of 1 hour
  - Join data as streams per partition



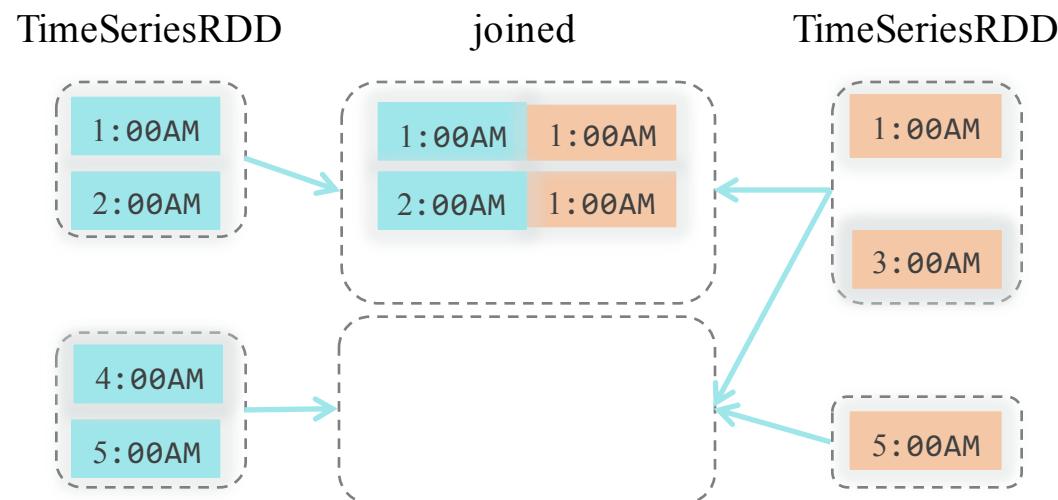
# Temporal join in TimeSeriesRDD

- Temporal join with criteria look-back and window of 1 hour
  - Join data as streams



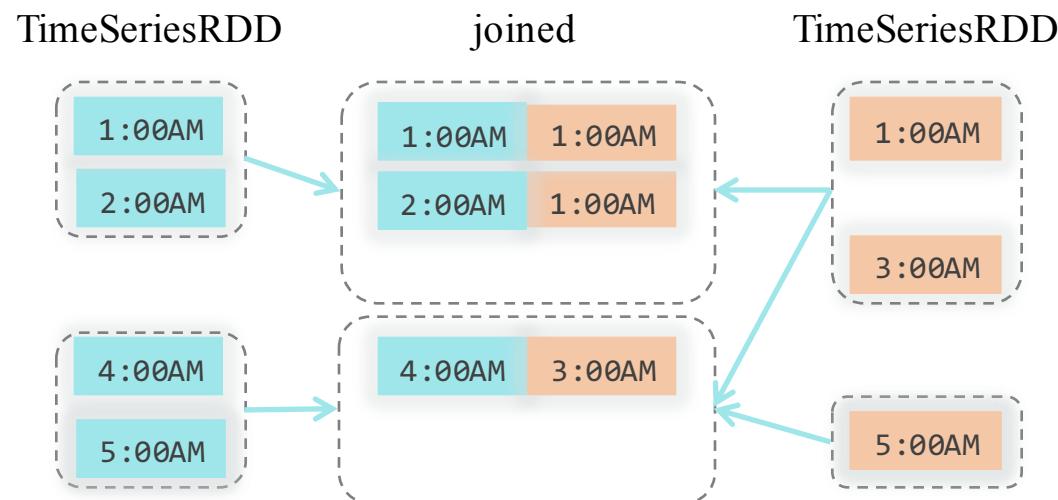
# Temporal join in TimeSeriesRDD

- Temporal join with criteria look-back and window of 1 hour
  - Join data as streams



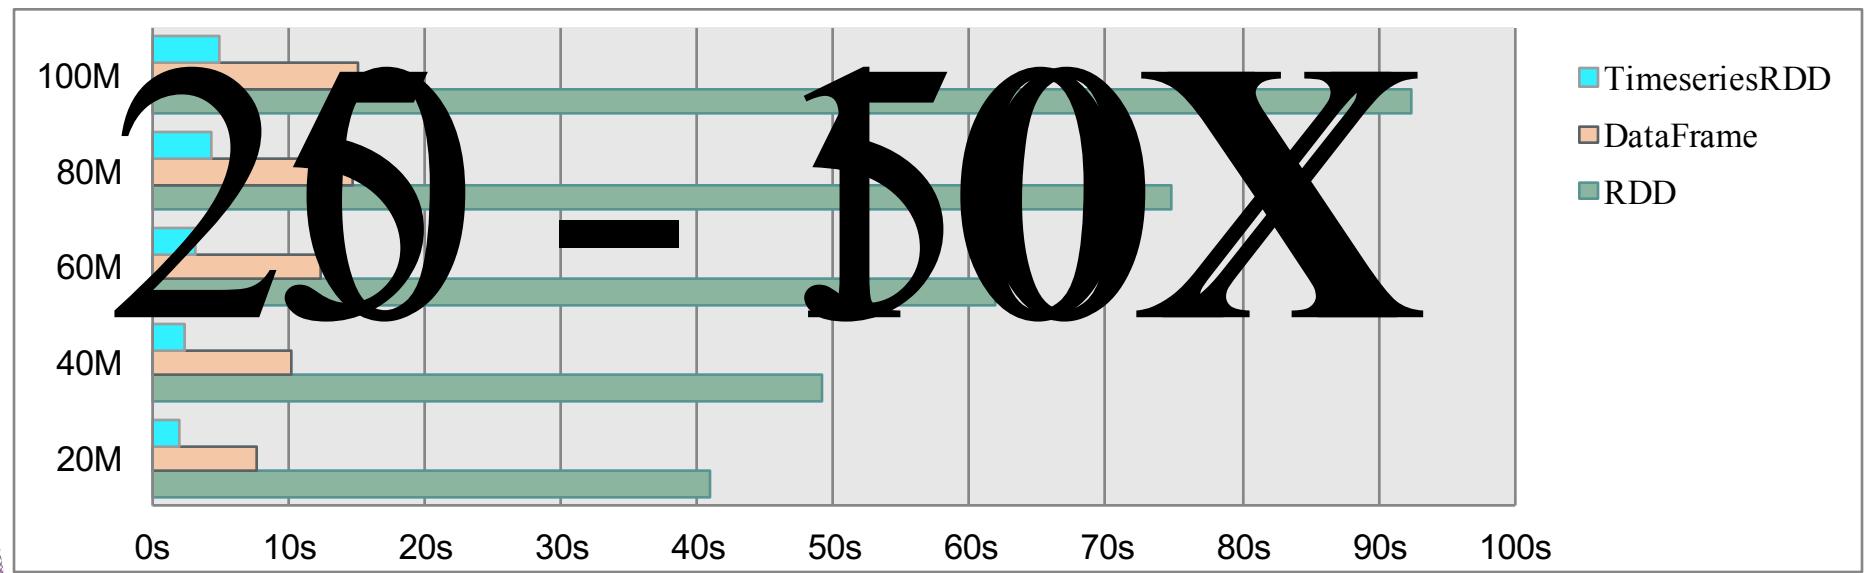
# Temporal join in TimeSeriesRDD

- Temporal join with criteria look-back and window of 1 hour
  - Join data as streams



## Benchmark for temporal join + count

- Running time of **count** after **temporal join**
  - 16 executors (10G memory and 4 cores per executor)
  - Data read from HDFS



## Functions over TimeSeriesRDD

- 
- ◆ Grouping functions
  - ◆ Temporal joins such as look-forward, look-backward etc.
  - ◆ Summarizers such as average, variance, z-score etc. over grouping functions

# Open Source

- ♦ True!
- ♦ <https://github.com/twosigma/flint>



**flint**

## What's next?

- 
- ◆ TimeSeriesDataframe / TimeSeriesDataset
  - ◆ Speed up
    - ◆ Richer APIs
  - ◆ Python bindings
  - ◆ Additional summarizers

## Key contributors

- 
- ◆ Christopher Aycock
  - ◆ Yuri Bogomolov
  - ◆ Jonathan Coveney
  - ◆ Li Jin
  - ◆ David Medina
  - ◆ Julia Meinwald
  - ◆ David Palaitis
  - ◆ Larisa Sawyer
  - ◆ Leif Walsh
  - ◆ Wenbo Zhao



**flint**

# Flint: Time Series For Spark

---

A library to solve for general time series analysis operations at massive scale

Anne Hathaway has nothing to do with Berkshire Hathaway

Check it out in open source, and contribute

<https://github.com/twosigma/flint>



**flint**

# THANK YOU.

Larisa.Sawyer@twosigma.com

