

# Spark @ Bloomberg: Dynamic Composable Analytics

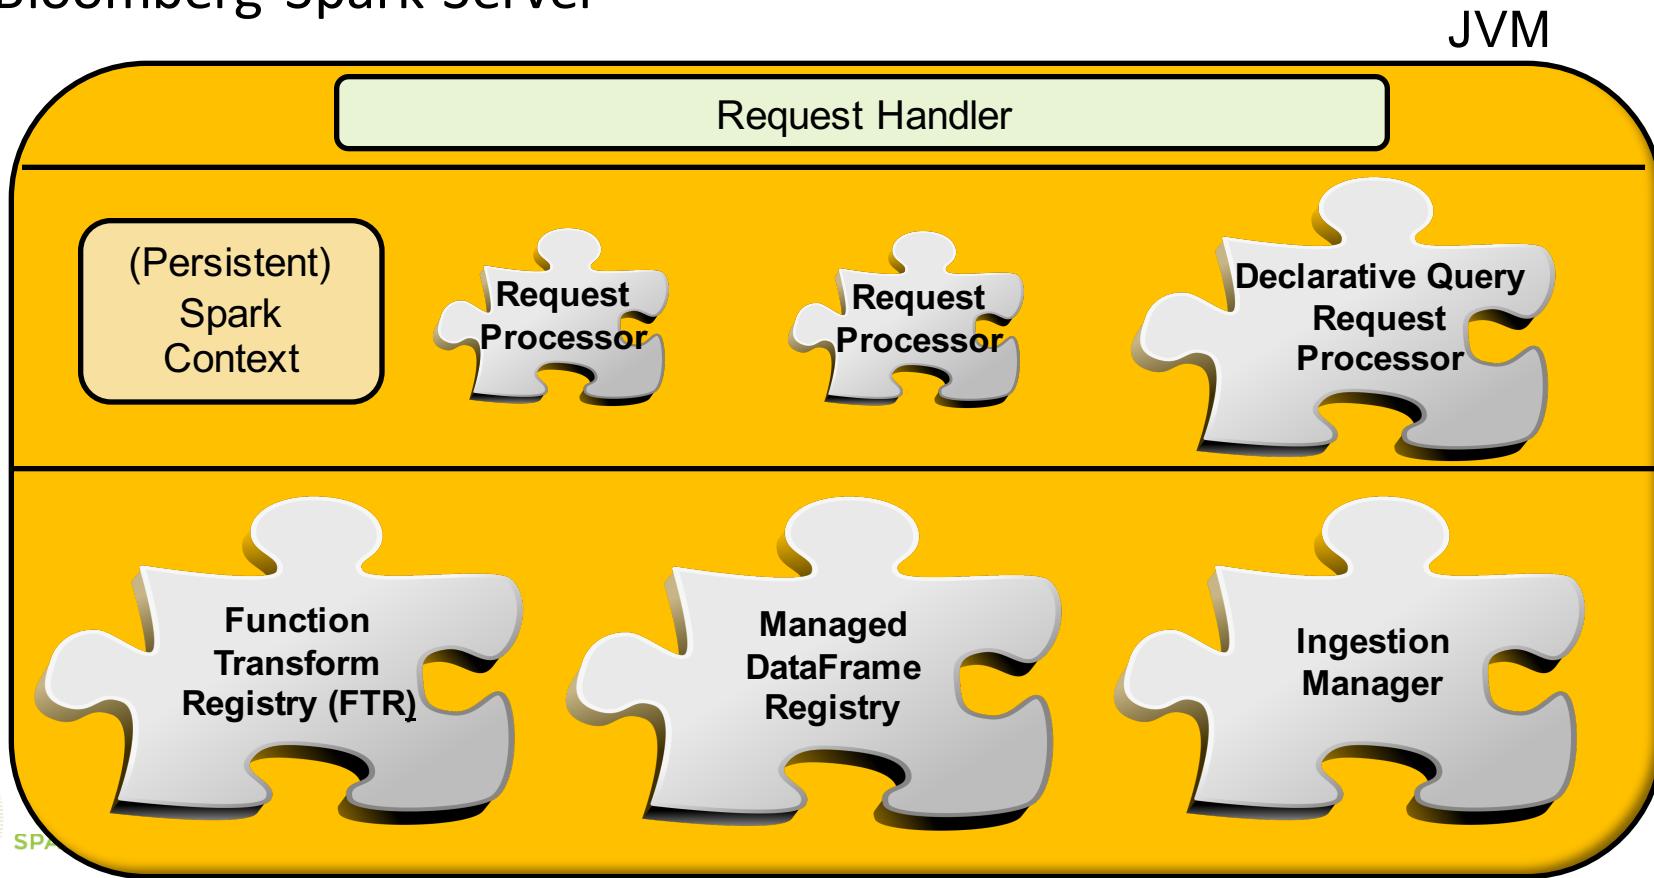
Partha Nageswaran  
Sudarshan Kadambi  
BLOOMBERG L.P.



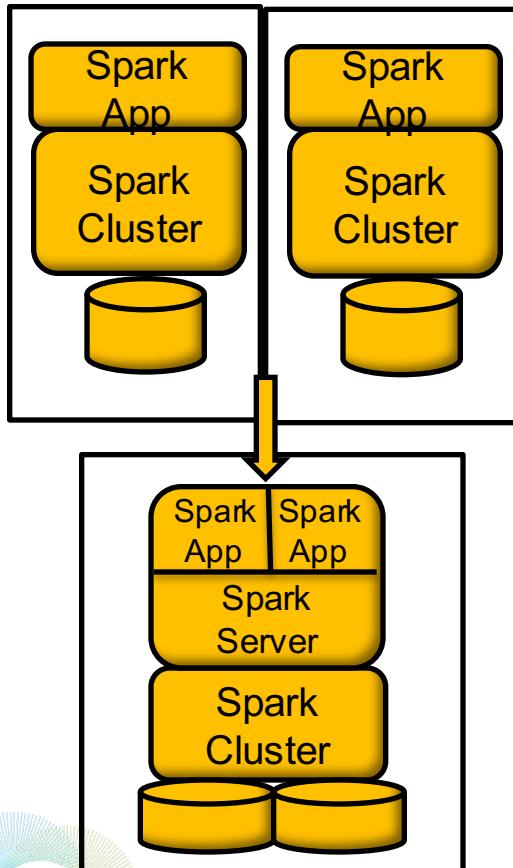
SPARK SUMMIT 2016  
DATA SCIENCE AND ENGINEERING AT SCALE  
JUNE 6-8, 2016 SAN FRANCISCO

# Bloomberg Spark Server

Spark Serverization at Bloomberg has culminated in the creation of the Bloomberg Spark Server



# Spark Serverization – Motivation



- Stand-alone Spark Apps on isolated clusters pose challenges:
  - Redundancy in:
    - » Crafting and Managing of RDDs/DFs
    - » Coding of the same or similar types of transforms/actions
  - Management of clusters, replication of data, etc.
  - Analytics are confined to specific content sets making Cross-Asset Analytics much harder
  - Need to handle Real-time ingestion in each App

# Dynamic Composable Analytics

- Compositional Analytics are common place in the Financial Sector

**Decile Rank the 14-day Relative Strength Index of Active Equity Stocks:**

```
DECILE(  
    RSI(  
        Price,  
        14,  
        ['IBM US Equity', 'VOD LN Equity', ... ]  
    )  
)
```

- Price is data abstracted as a Spark Data Frame
- RSI, DECILE are building block analytics, expressible as Spark transforms and actions



# Dynamic Composable Analytics

- Another usecase may want to compose Percentile with RSI

**Percentile Rank the 14-day Relative Strength Index of Active Equity Stocks:**

```
PERCENTILE(  
    RSI(  
        Price,  
        14,  
        ['IBM US Equity', 'VOD LN Equity', ... ]  
    )  
)
```

- Or Percentile with ROC, etc. And the compositions maybe arbitrarily complex

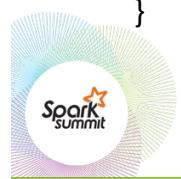


# Dynamic Composable Analytics

```
def RSI(df: DataFrame, period:Int=14) : DataFrame = {
    val smmaCoeff = udf( i:Double ) => scala.math.pow(period-1, i-1)/scala.math.pow(period,i) )
    val rsi_from_rs = udf( n:Double, d:Double ) => 100 - 100*d/(d+n) )
    val rsi_window = Window.partitionBy('id).orderBy('date.desc)

    df.withColumn("weight", smmaCoeff(row_number.over(rsi_window)))
        .withColumn("diff", 'value - lead('value,1).over(rsi_window))
        .withColumn("U", when('diff > 0, 'diff).otherwise(0))
        .withColumn("D", when('diff < 0, abs('diff)).otherwise(0))
        .groupBy('id).agg(rsi_from_rs( sum('U * 'weight), sum('D * 'weight) )as 'value)
}

def Decile(df: Dataframe) : DataFrame ={
    df.withColumn("value", ntile(10).over( Window.orderBy('value.desc) ))
}
```

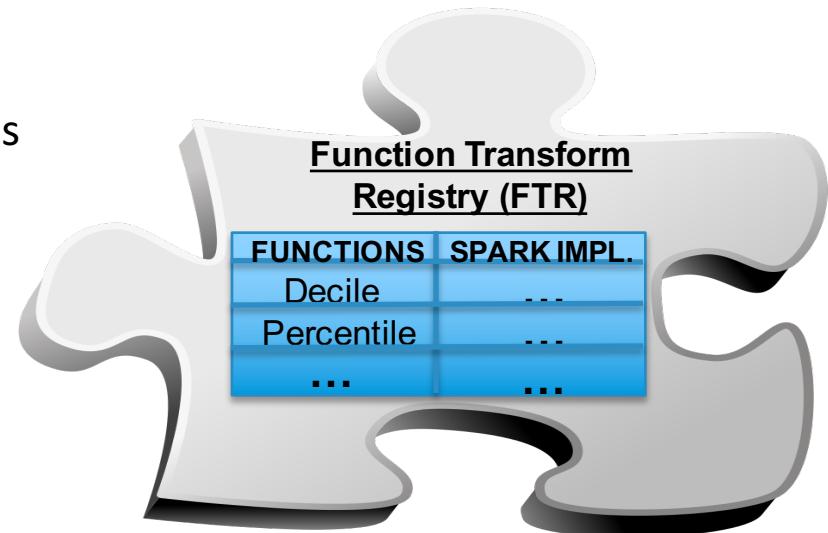


SPARK SUMMIT 2016

Ack: Andrew Foster

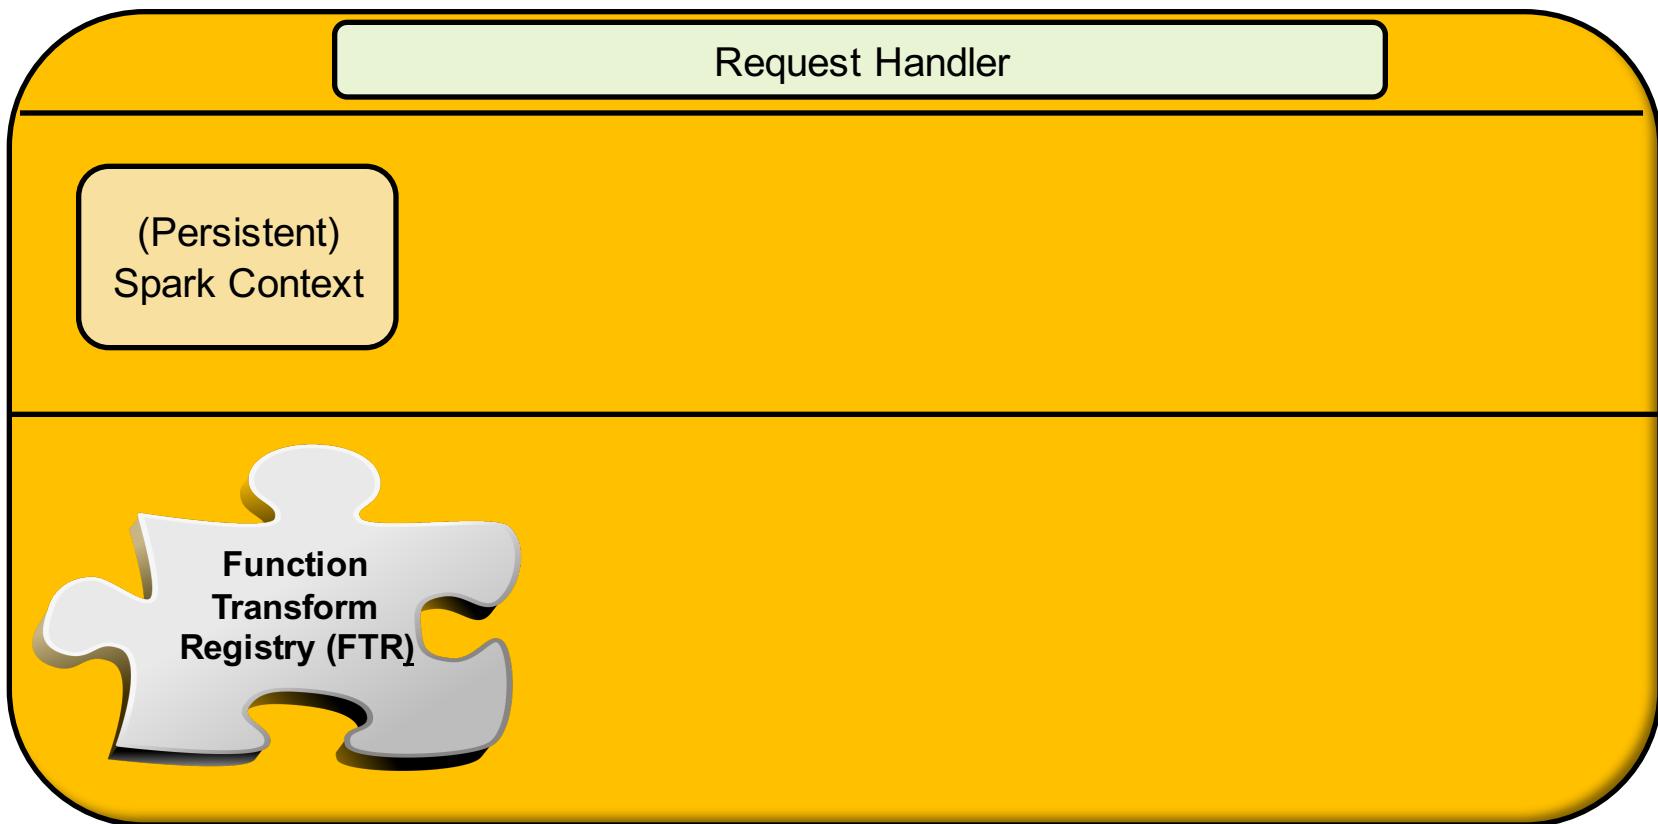
# Function Transform Registry

- Maintain a Registry of Analytic functions (FTR) with functions expressed as Parametrized Spark Transforms and Actions
- Functions can compose other functions, along with additional transforms, within the Registry
- Registry supports 'bind' and 'lookup' operations



# Bloomberg Spark Server

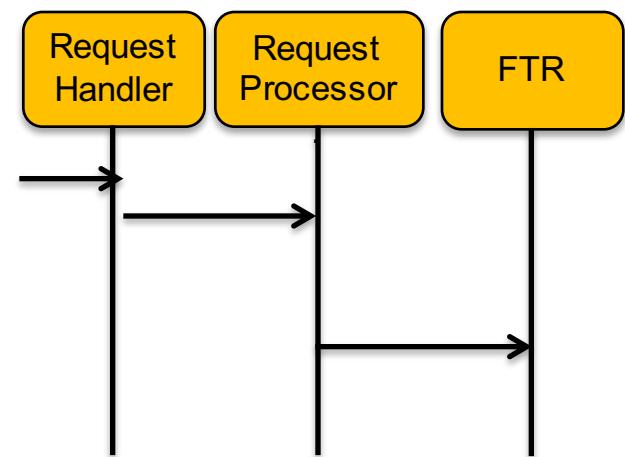
JVM



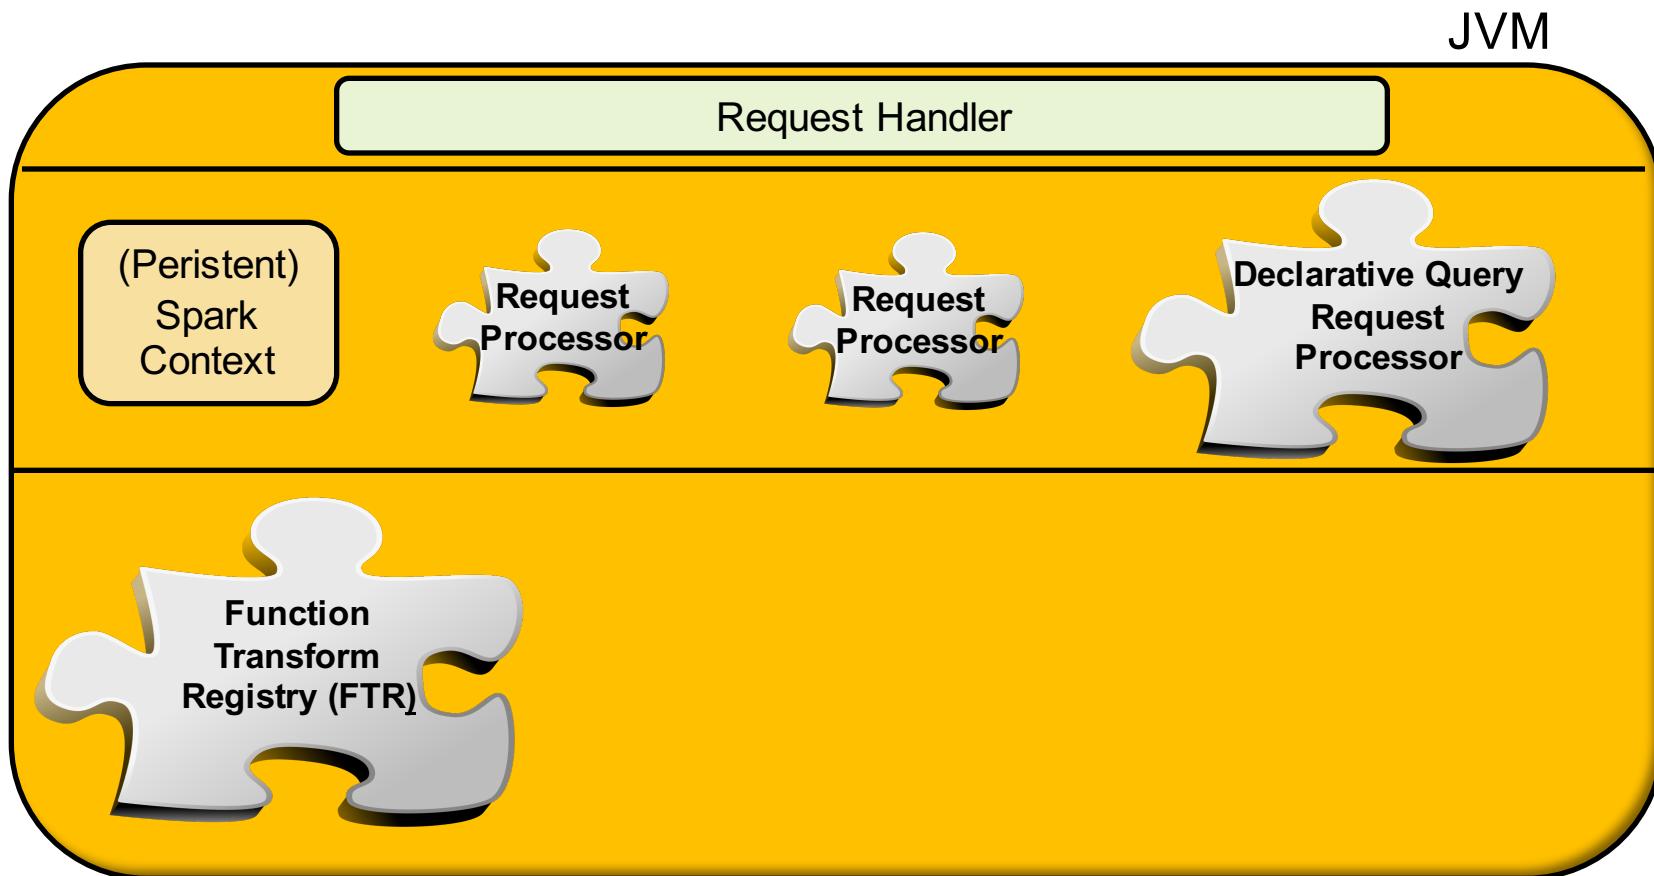
SPARK SUMMIT 2016

# Request Processor

- Request Processors (RPs) are spark applications that orchestrate composition of analytics on Data Frames
  - RPs comply with a specification that allows them to be hosted by the Bloomberg Spark Server
  - Each request (such as: compute the Decile Rank of the RSI) is handled by a Request Processor that looks up functions from the FTR, Composes them and applies them to Data Frames



# Bloomberg Spark Server



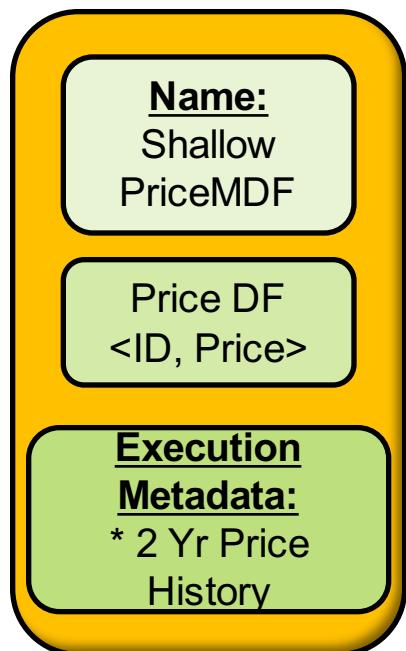
# Managed Data Frames

- Besides locating functions from the FTR, Request Processors have to pass in Data Frames to these functions as inputs
- Rather than instantiate Data Frames, lookup Data Frames from a Data Frames Registry
  - Such Data Frames are called Managed Data Frames (**MDF**)
  - The Registry that Manages these Data Frames is the Managed Data Frame Registry (**MDF Registry**)



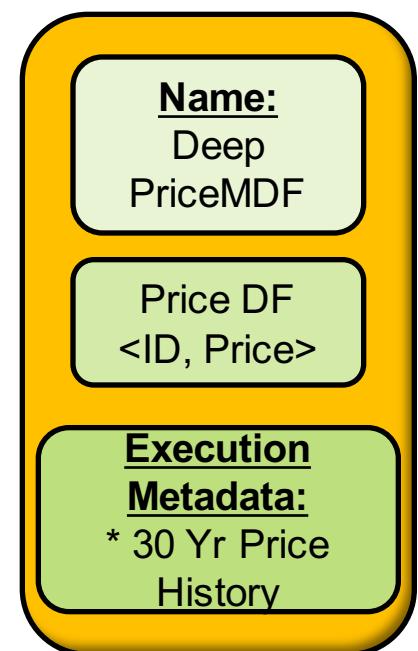
# Introducing Managed DataFrames (MDFs)

MDF



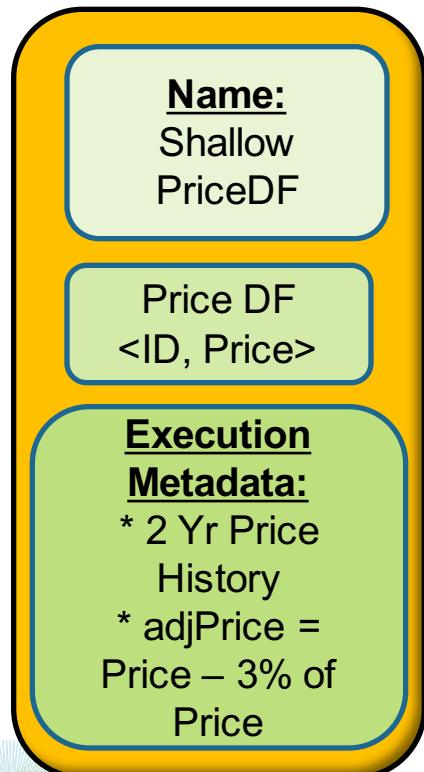
- A Managed DataFrame (MDF) is a named DataFrame, optionally combined with **Execution Metadata**
  - MDFs can be located by name OR by any Column Name defined in the Schema of the corresponding DF
- Execution Metadata includes:
  - **Data Distribution** metadata captures information about the data depth, histogram information, etc.
  - E.g.: A managed DataFrame for pricing of stocks, representing 2 years of historical data and another for representing 30 years of historical data

MDF

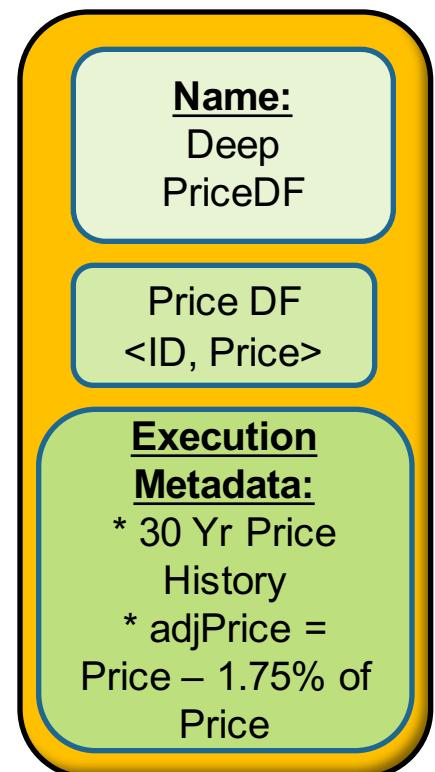


# Managed DataFrames

MDF



MDF



- **Data Derivation** metadata which are mathematical expressions that define how additional columns can be synthesized from existing columns in the schema
- E.g.: adjPrice is a derived Column, defined in terms of the base Price column
- In essence, an MDF with data derivation metadata have a Schema that is a union of the contained DF and the derived columns



SPARK SUMMIT 2016

# The MDF Registry

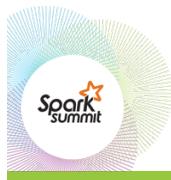
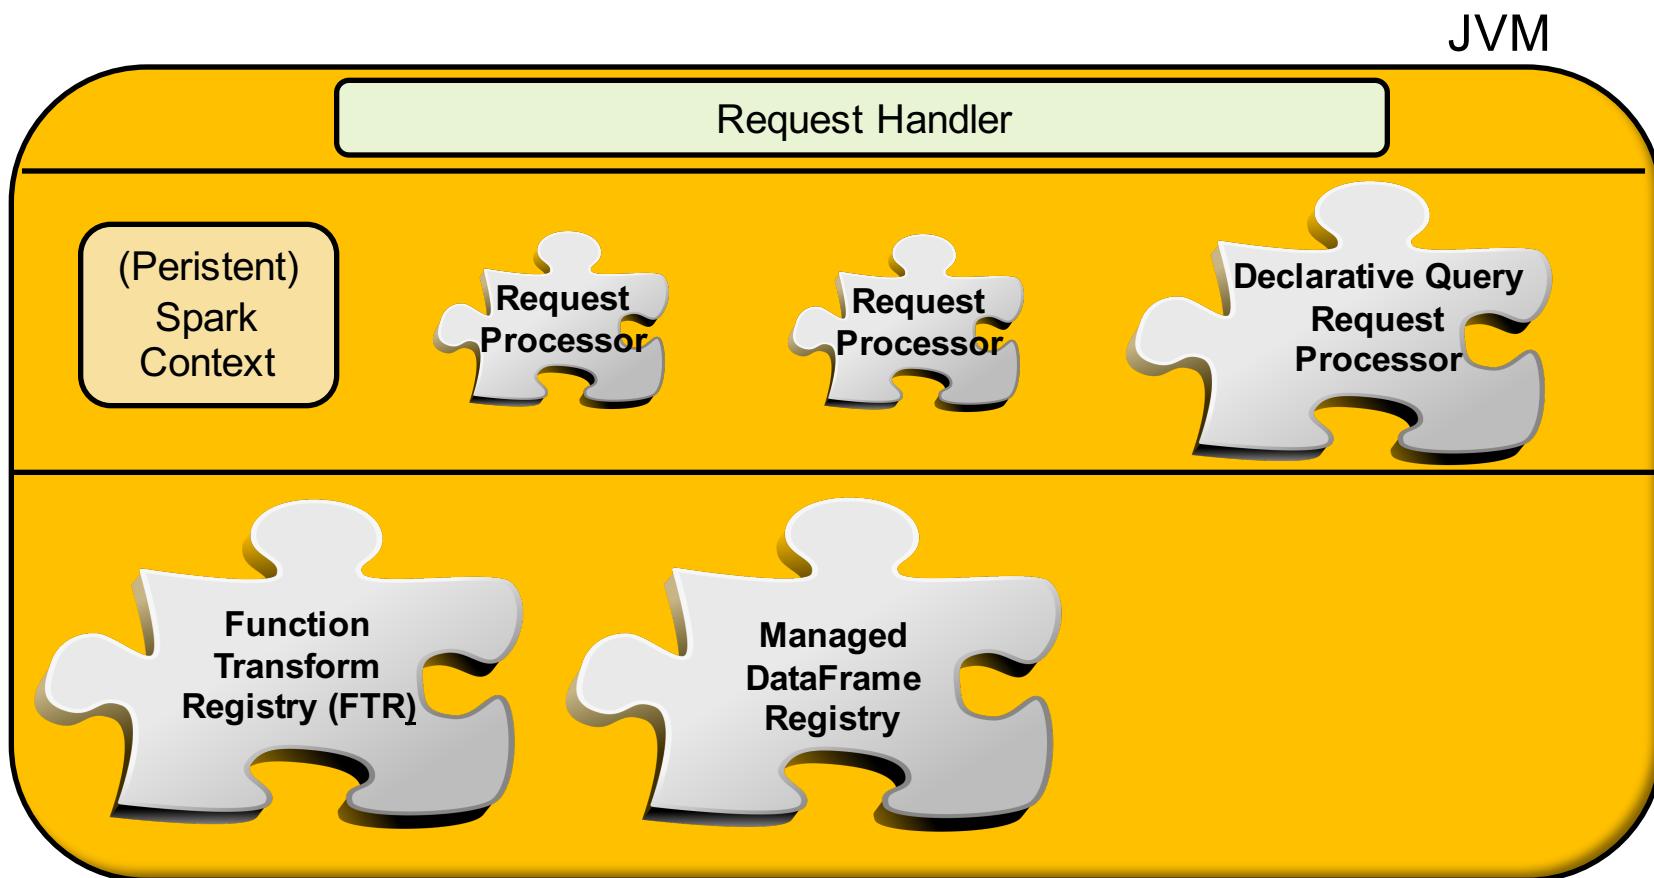
## MDF Registry

Name	Columns	DF Ref.	Meta Data	...
Shallow Price DF	Price, adjPrice	○	...	...
Deep Price DF	Price, adjPrice	○	...	...

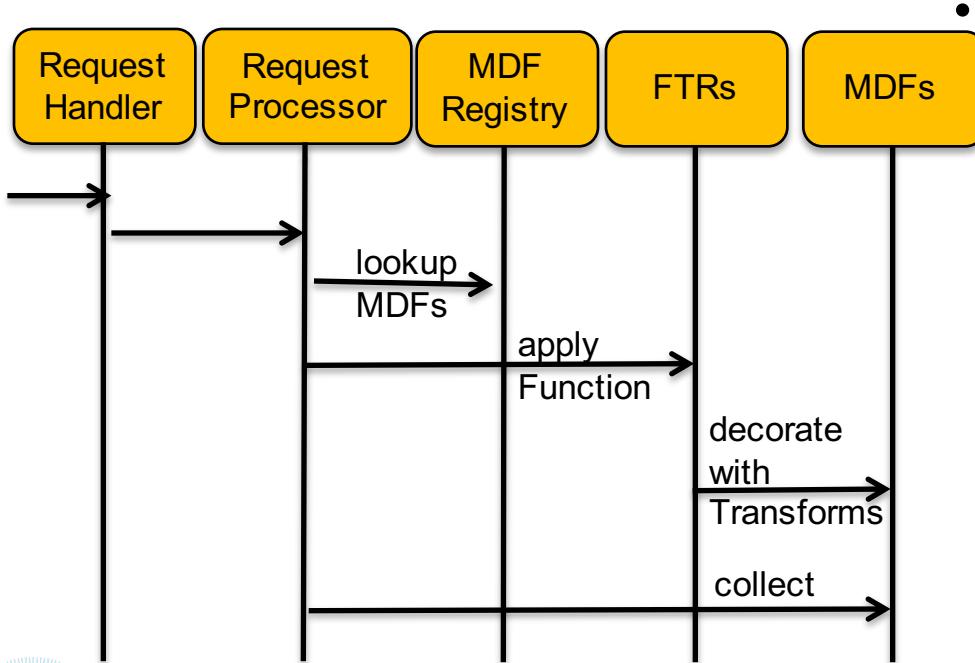
- The MDF Registry within the Bloomberg Spark Server provides support for:
  - Binding MDFs by Name
  - Looking up MDFs by Name
  - Looking up MDF by a Column Name (an element of the MDF Schema), etc.
- The MDF Registry maintains a 'table' that associates the Name of the MDF with the DF reference and Columns in the DF



# Bloomberg Spark Server



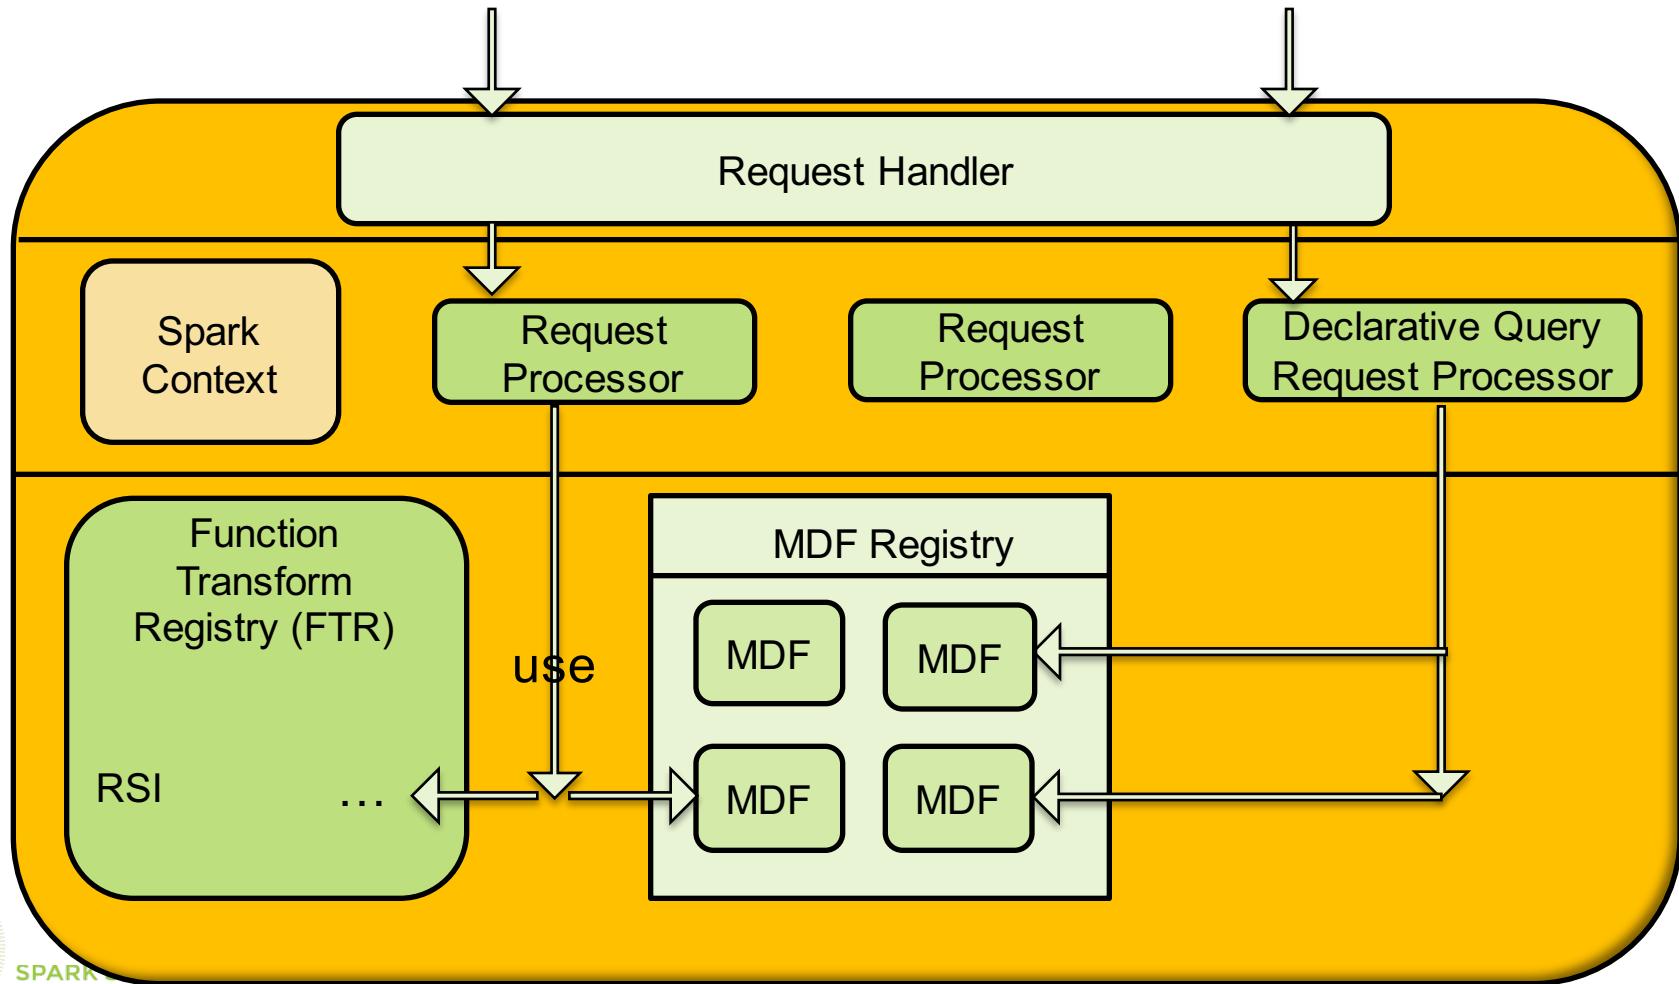
# Flow of Requests



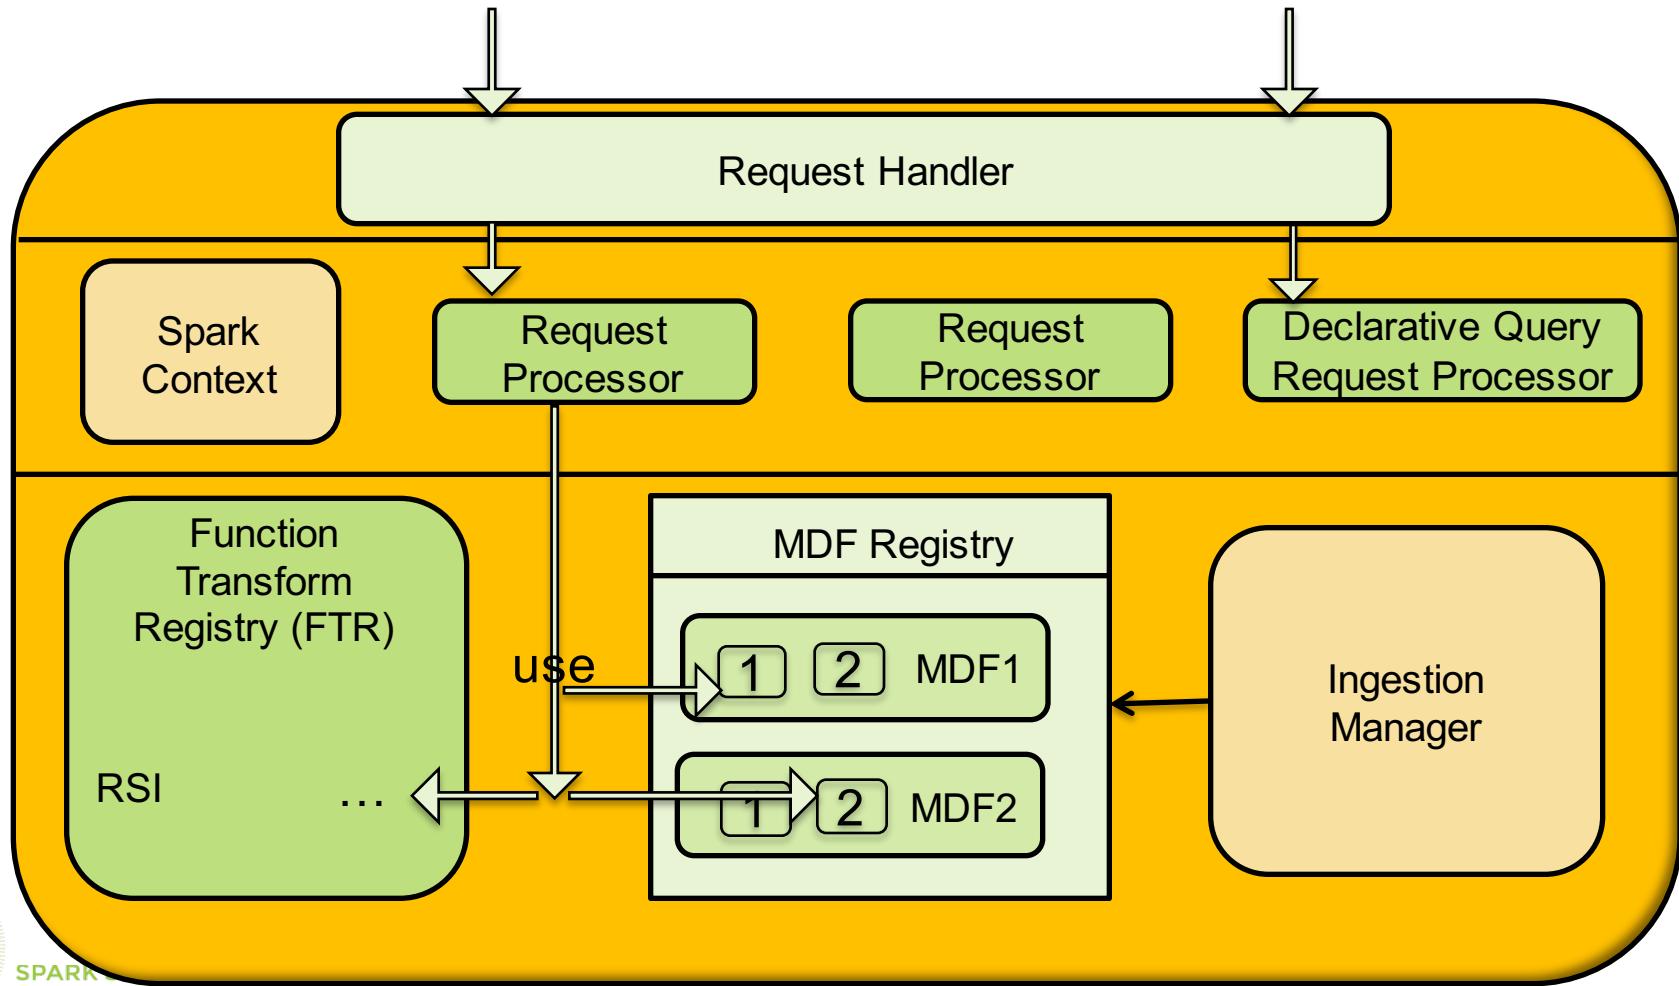
- Request Processors within the Spark Server orchestrate analytics
  - These RPs have access to the Registry and FTRs
  - Are responsible for composing transforms and actions on one or more MDFs
  - May dynamically bind additional MDFs (materialized or otherwise) for use by other Apps



# Bloomberg Spark Server



# Bloomberg Spark Server



# Schema Repository

- Enterprise-wide data pipeline
- External (to Spark) schema repository and service
- Enables MDF lookup by a dataset schema element
  - Analytic expressions can now be composed over data elements

# Execution Metadata

- Dataset Source Connection Identifiers
  - Backing Stores
  - Real-time Topics
- Storage Level & Refresh Rate
- Subset Predicate, etc.

# Ad-hoc Cross-Domain Analytics

- Registration of pre-materialized DataFrames
  - Collaborative analytics between application workflows
- Dynamic creation of Managed DataFrames
  - Spark Servers have data pertaining to a single domain materialized
  - Ad-hoc cross-domain analytics requires capability to synthesize MDFs on demand



# Content Subsetting

- High value data sub-setted within Spark
  - Reduce cost of querying external datastore
- Specified as a filter predicate at time of registration
  - E.g. Member companies of popular indices [Dow 30, S&P 500,...] have records placed within Spark

# Content Subsetting

- Seamless unification of data in Spark ( $DF_{subset}$ ) and backing store ( $DF_{subset'}$ )  
$$(DF_{subset} \cup DF_{subset'}).filter(query) = DF_{subset}.filter(query) \cup DF_{subset'}.filter(query)$$
- Dataset owners provided knobs for cost vs performance.
- LRU cache like mechanism planned in the future
- **Make sense as a capability native to Spark dataframes**



## Ingestion: Periodic Refresh

- Periodic data pull into Spark from the backing store
- Subset criteria applied during data retrieval
- Used when a dataset has a backing store, but no real time update stream that we can tap into
- Dataset owners have control over storage level of the dataframes created within a given MDF

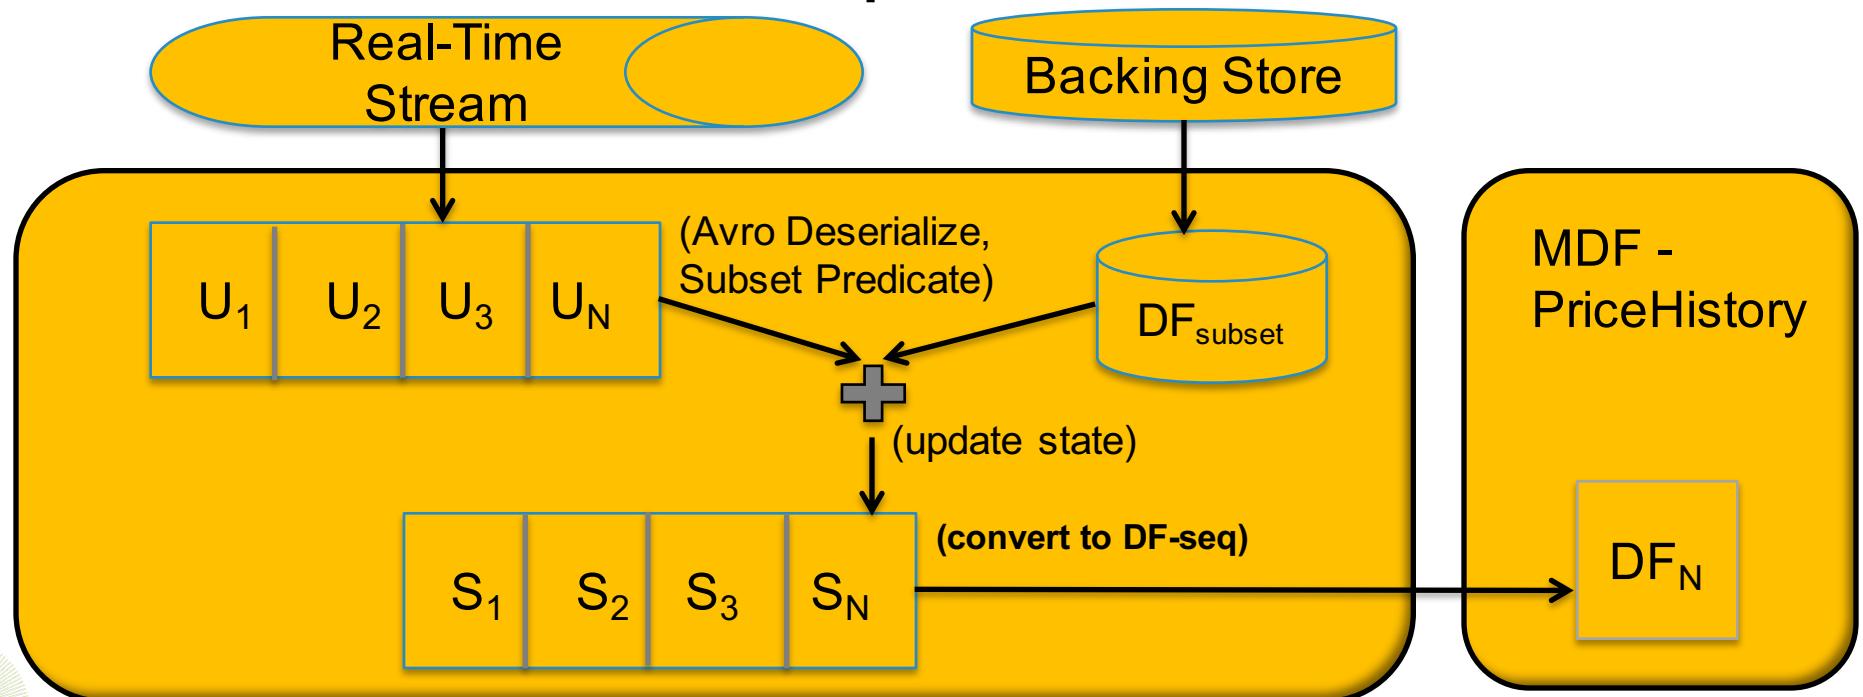


# Ingestion: Stream Reconciliation

- Analytics needs to be low-latency with respect to queries, but also data freshness
- Since data is being sub-setted within Spark, need to keep the subset up to date
- Datasets published to different Kafka topics.
  - 1:1 mapping between datasets, topics and DStreams.

# Ingestion: Stream Reconciliation

Similar intent as Structured Streaming, to be introduced in Spark 2.0



# Ingestion: Data Transformation

- Data in backing stores may need representation transforms before being used in queries
  - Data in multiple tables denormalized into a single DF within Spark
  - Or, quickly see effect of different storage representations on performance, without changing the representation in the backing store
- Implemented via. user transforms associated with a given MDF



SPARK SUMMIT 2016

# Spark Server: Memory Management

- An MDF contains multiple generation of DFs, being generated and destroyed
- Multiple generations operated upon by RPs at given point in time
- Reference counting to keep track of what DFs are being used and by whom
- Long running queries aborted for forced reclamation



SPARK SUMMIT 2016

# Query Consistency

- Multiple queries need to operate on same snapshot of data
- How to achieve, if data constantly changing underneath?
- Each DF within MDF associated with time epoch
- Registry lookup with a reference time
- Time-align sub-setted dataframes with data in backing store

# Spark for Online Analytics

- High Availability of Spark Driver
  - High bootstrap cost to reconstructing cluster and cached state
  - Naïve HA models (such as multiple active clusters) surface query inconsistency
- High Availability of RDD Partitions
  - With subset or universe cached, lost RDD partitions kill query performance
- Performance Consistency
  - Performance gated by slowest executor
  - *High Availability and Low Tail Latency closely related*
- Interactions effects between low-latency queries and low-latency updates
  - No to Minimal sandboxing between jobs sharing executor JVMs

**First Bloomberg contribution: SPARK-15352**



SPARK SUMMIT 2016

# Spark Server Acknowledgements



Andrew Foster



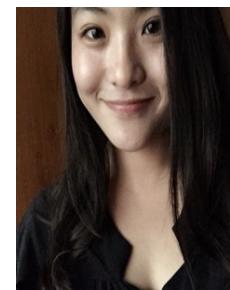
Joe Davey



Shubham Chopra



Nimbus Goehausen



Tracy Liang

# THANK YOU.

[pnageswaran@bloomberg.net](mailto:pnageswaran@bloomberg.net)

[skadambi@bloomberg.net](mailto:skadambi@bloomberg.net)



SPARK SUMMIT 2016  
DATA SCIENCE AND ENGINEERING AT SCALE  
JUNE 6-8, 2016 SAN FRANCISCO