

Re-Architecting Apache Spark for Performance Understandability

Kay Ousterhout

Joint work with Christopher Canel, Max Wolfson,
Sylvia Ratnasamy, Scott Shenker



About Me

PhD candidate at UC Berkeley

Thesis work on performance of large-scale distributed systems

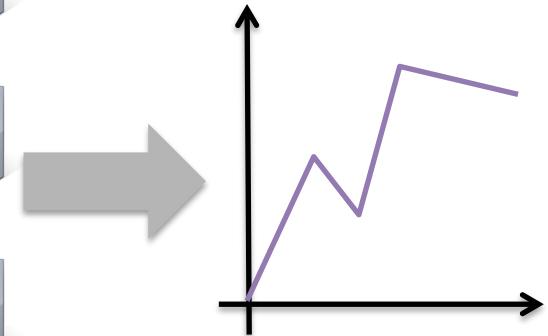
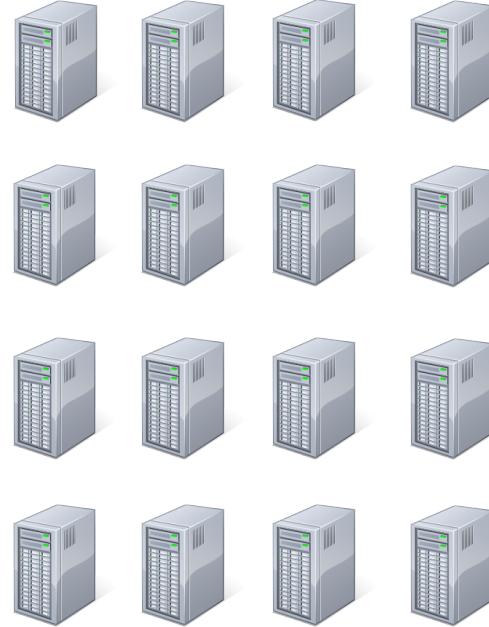
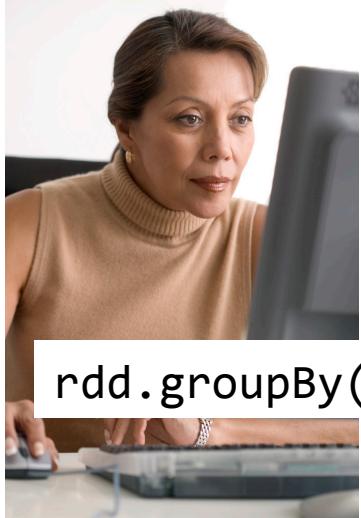
Apache Spark PMC member

About this talk

Future architecture for systems like Spark

Implementation is API-compatible with Spark

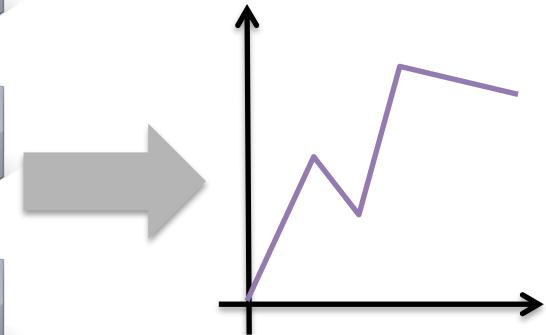
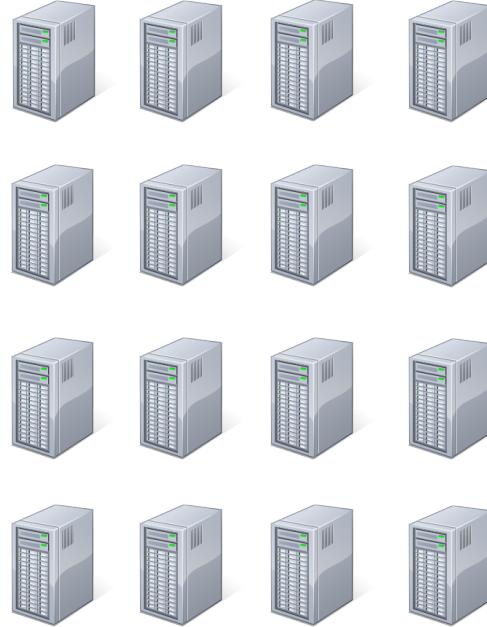
Major change to Spark's internals (~20K lines of code)



Spark cluster is a black box, runs the job fast



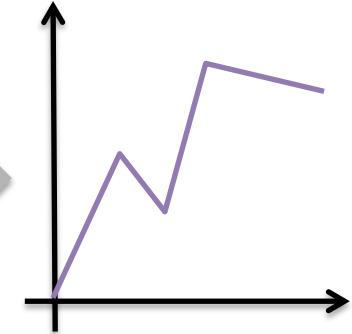
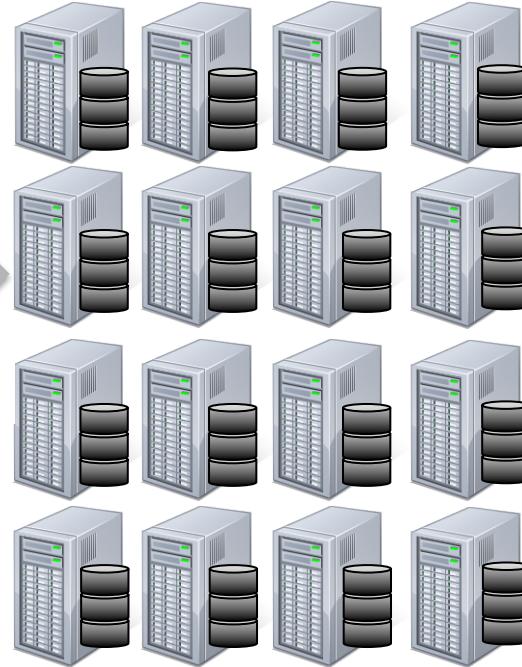
`rdd.groupBy(...)`



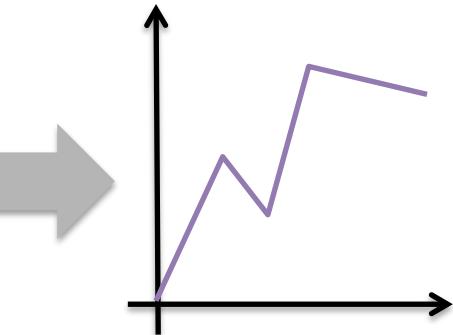
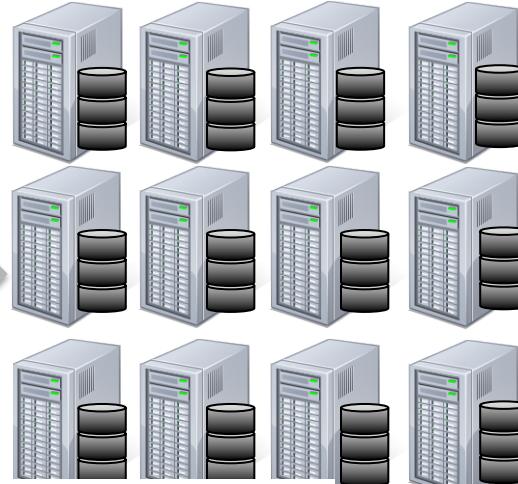
Idealistic view: Spark cluster is a black box, runs the job fast



rdd.groupBy(...)



Idealistic view: Spark cluster is a black box, runs the job fast



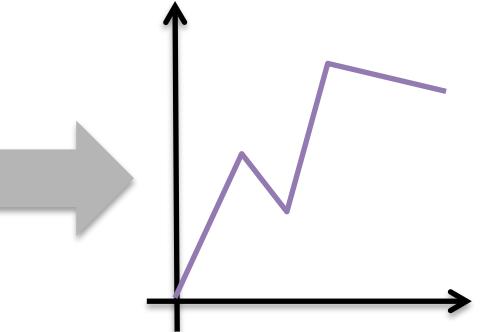
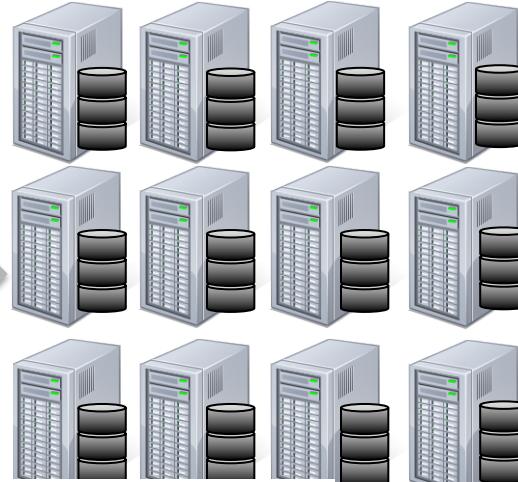
~~rdd.groupBy(...)~~
rdd.reduceByKey(...)

Configuration:

spark.serializer KryoSerializer
spark.executor.cores 8



Idealistic view: Spark cluster is a black box, runs the job fast



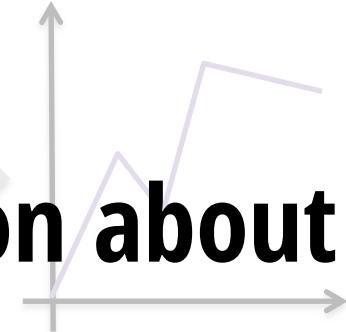
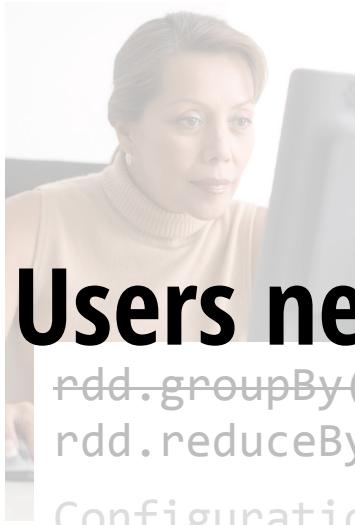
~~rdd.groupBy(...)~~
rdd.reduceByKey(...)

Configuration:

spark.serializer KryoSerializer
spark.executor.cores 8



Realistic view: user uses performance characteristics to tune job, configuration, hardware, etc.



Users need to be able to reason about performance

~~rdd.groupBy(...)~~
~~rdd.reduceByKey(...)~~

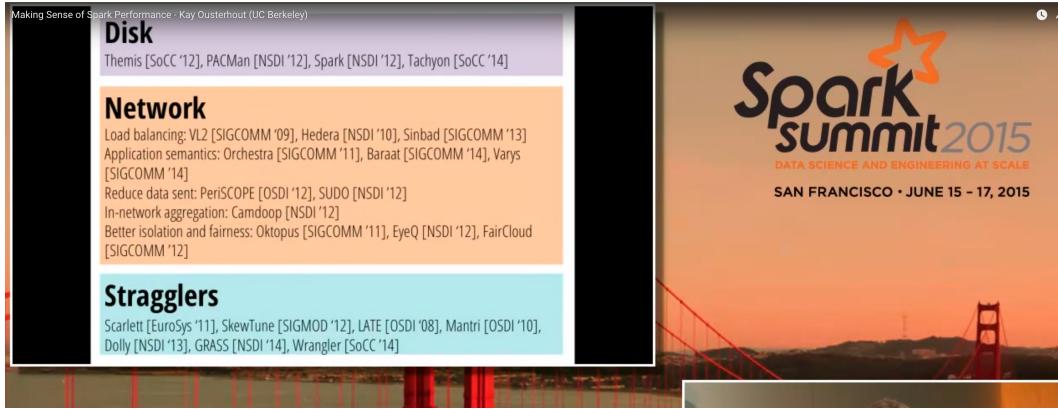
Configuration:

spark.serializer KryoSerializer
spark.executor.cores 8



Realistic view: user uses performance characteristics to tune job, configuration, hardware, etc.

Reasoning about Spark Performance

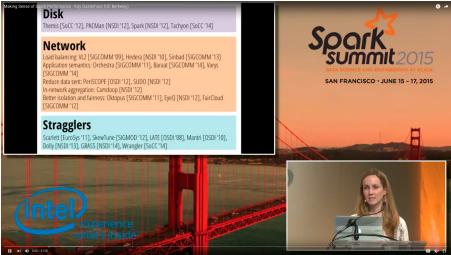


Widely accepted that network and disk I/O are bottlenecks

CPU (not I/O) typically the bottleneck

network optimizations can improve job completion time by at most 2%

Reasoning about Spark Performance



Spark Summit 2015:
CPU (not I/O) often
the bottleneck

Project Tungsten:
**initiative to optimize Spark's CPU use,
driven in part by our measurements**

 [databricks](#) [PRODUCT](#) [SPARK](#) [RESOURCES](#) [COMPANY](#) [BLOG](#)

COMPANY

All Posts  

Partners

Events

Press Releases

DEVELOPER

All Posts

Spark

Spark SQL

Spark Streaming

MLlib

Project Tungsten: Bringing Spark Closer to Bare Metal

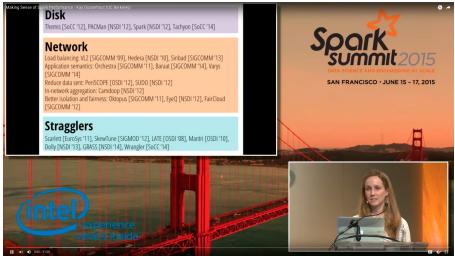
April 28, 2015 | by Reynold Xin and Josh Rosen

In a previous [blog post](#), we looked back and surveyed performance improvements made to Spark in the past year. In this post, we look forward and share with you the next chapter, which we are calling *Project Tungsten*. 2014 witnessed Spark setting the world record in large-scale sorting and saw major improvements across the entire engine from Python to SQL to machine learning. Performance optimization, however, is a never ending process.

Project Tungsten will be the largest change to Spark's execution engine since the project's inception. It focuses on substantially [improving the efficiency of memory and CPU for Spark applications](#), to push performance closer to the limits of modern hardware. This effort includes three initiatives:

1. *Memory Management and Binary Processing*: leveraging application semantics to manage memory explicitly and eliminate the overhead of JVM object model and garbage collection
2. *Caching more computations, algorithms and data structures to exploit memory hierarchy*

Reasoning about Spark Performance



Spark Summit 2015:
CPU (not I/O) often
the bottleneck

A screenshot of a Databricks blog post titled "Project Tungsten: Bringing Spark Closer to Bare Metal". The post discusses performance improvements made to Spark's execution engine since its inception. It highlights three main areas of focus: Memory Management and Binary Pruning, Cache-aware computation, and Code-generation. The post is dated April 04, 2015, and is written by Bryan Fann and Ash Narayan.

Project Tungsten:
initiative to optimize
Spark's CPU use

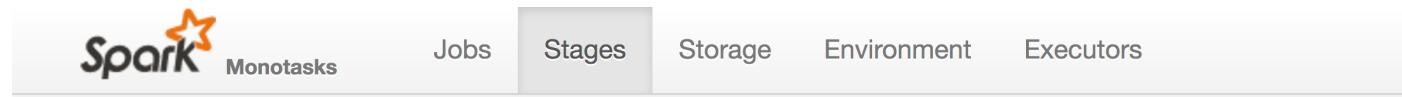
Spark 2.0:
Some evidence that I/O
is again the bottleneck
[HotCloud '16]

Users need to understand performance to extract the best runtimes

Reasoning about performance is currently difficult

**Software and hardware are constantly evolving,
so performance is always in flux**

Vision: jobs report high-level performance metrics



Details for Stage 17

Total task time across all tasks: 13 min

Shuffle read: 2.5 GB / 31589120

Performance Information

Bottleneck: Disk (if disk bandwidth were increased by 23% or more, network would become the bottleneck)

Vision: jobs report high-level performance metrics



Details for Stage 17

Total task time across all tasks: 13 min

Shuffle read: 2.5 GB / 31589120

Performance Information

Bottleneck: Disk (if disk bandwidth were increased by 23% or more, network would become the bottleneck)

Non-bottlenecks: Network (could reduce network bandwidth by up to 30% slower without impacting runtime), CPU (could increase CPU time by up to 2x without impacting runtime)

Vision: jobs report high-level performance metrics



Details for Stage 17

Total task time across all tasks: 13 min

Shuffle read: 2.5 GB / 31589120

Performance Information

Bottleneck: Disk (if disk bandwidth were increased by 23% or more, network would become the bottleneck)

Non-bottlenecks: Network (could reduce network bandwidth by up to 30% slower without impacting runtime), CPU (could increase CPU time by up to 2x without impacting runtime)

Benefit of caching: Storing input in-memory would reduce job completion time by 42%

Vision: jobs report high-level performance metrics



Details for Stage 17

Total task time across all tasks: 13 min

Shuffle read: 2.5 GB / 31589120

Performance Information

Bottleneck: Disk (if disk bandwidth were increased by 23% or more, network would become the bottleneck)

Non-bottlenecks: Network (could reduce network bandwidth by up to 30% slower without impacting runtime), CPU (could increase CPU time by up to 2x without impacting runtime)

Benefit of caching: Storing input in-memory would reduce job completion time by 42%

Job Runtime Predictor: (enter in properties of different cluster to estimate job's runtime)

Number of machines:

CPU cores per machine:

Network bandwidth per machine: Gbps

I/O bandwidth per machine: MB/s

Calculate new runtime

Predicted new job runtime:

How can we achieve this vision?

Spark overview

Reasoning about Spark's performance: why it's hard

New architecture: monotasks

Reasoning about monotasks performance: why it's easy

Monotasks in action (results)

Example Spark Job:
Read remote data
Filter records
Write result to disk

Task 1:
Read and filter **block 1**
write result to disk

Task 2:
Read and filter **block 2**
write result to disk

Task 3:
Read and filter **block 3**
write result to disk

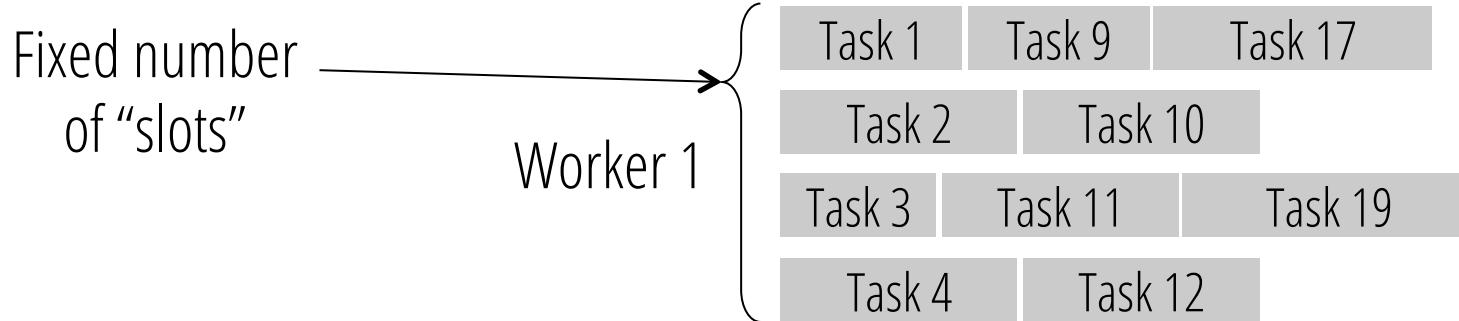
Task 4:
Read and filter **block 4**
write result to disk

Task 5:
Read and filter **block 5**
write result to disk

Task 6:
Read and filter **block 6**
write result to disk

⋮

Task n:
Read and filter **block n**
write result to disk



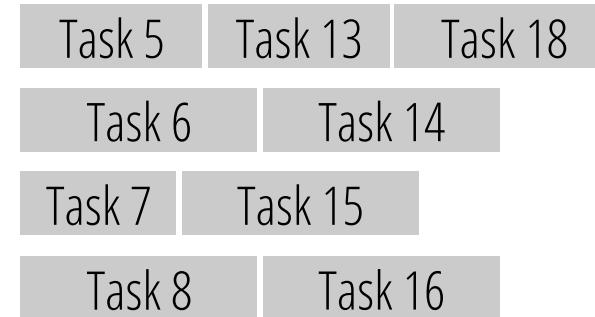
Example Spark Job:

Read remote data

Filter records

Write result to disk

Worker 2



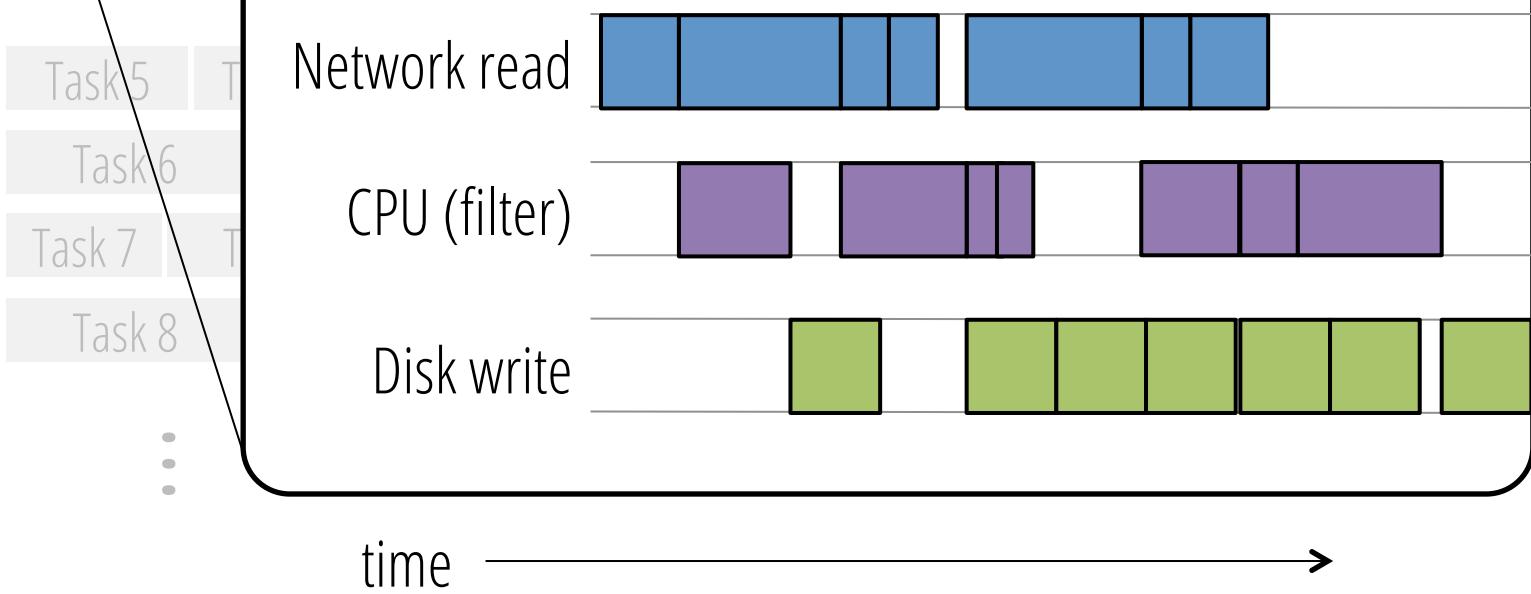
⋮

time

Worker 1



Worker 2



How can we achieve this vision?

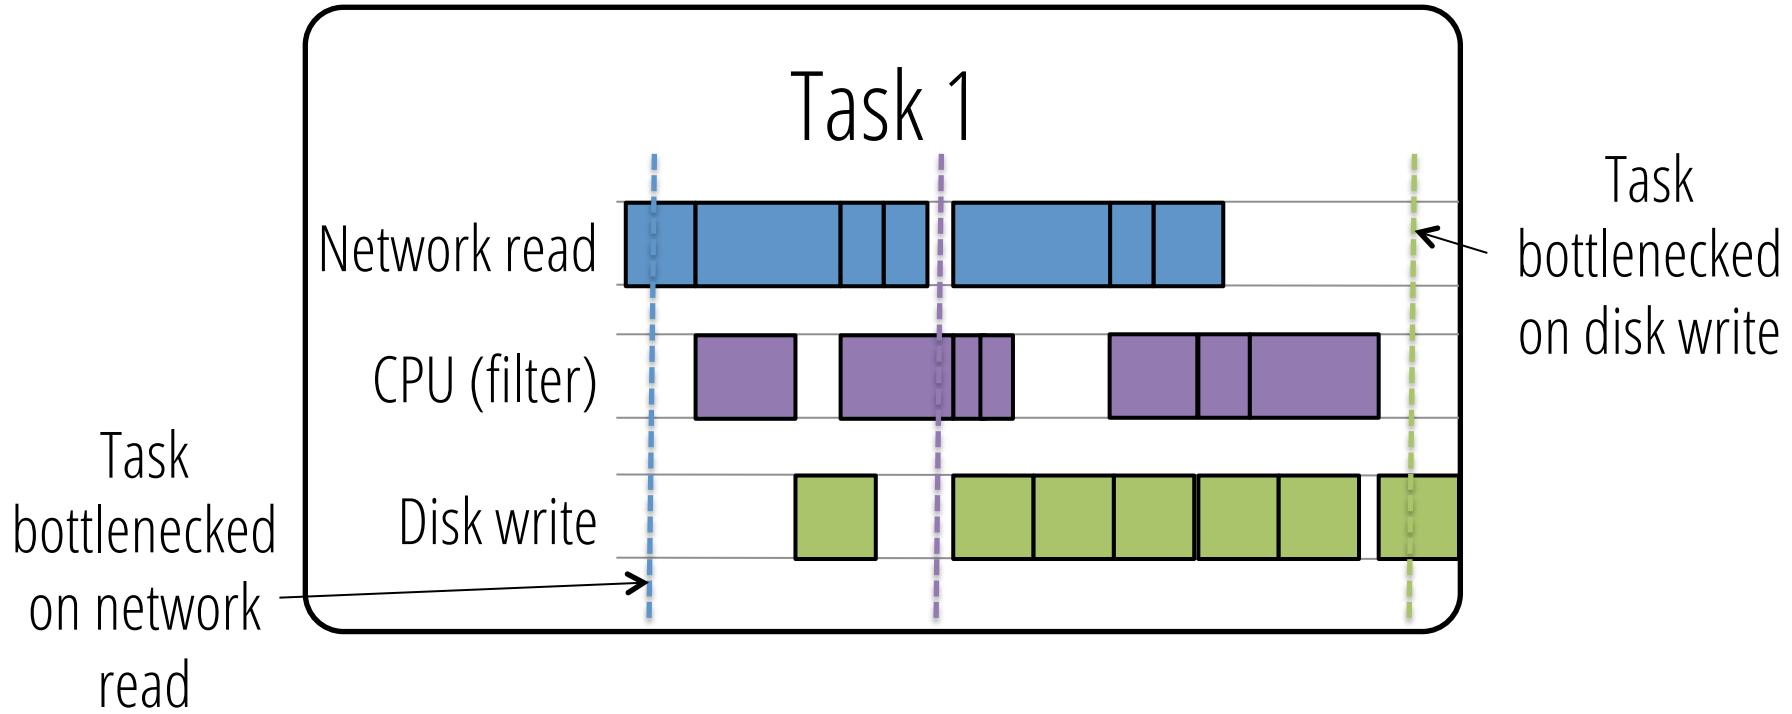
Spark overview

Reasoning about Spark's performance: why it's hard

New architecture: monotasks

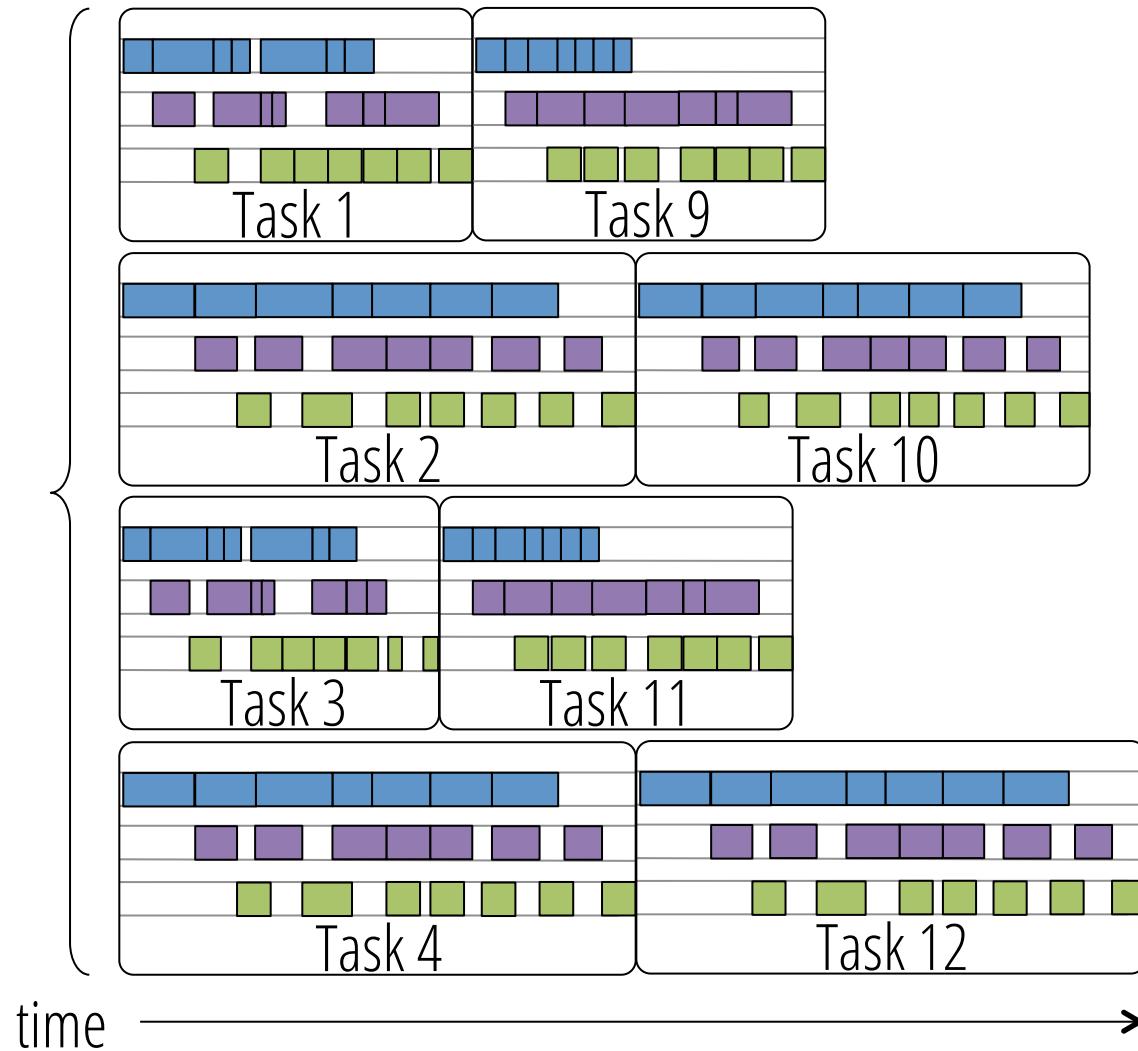
Reasoning about monotasks performance: why it's easy

Monotasks in action (results)

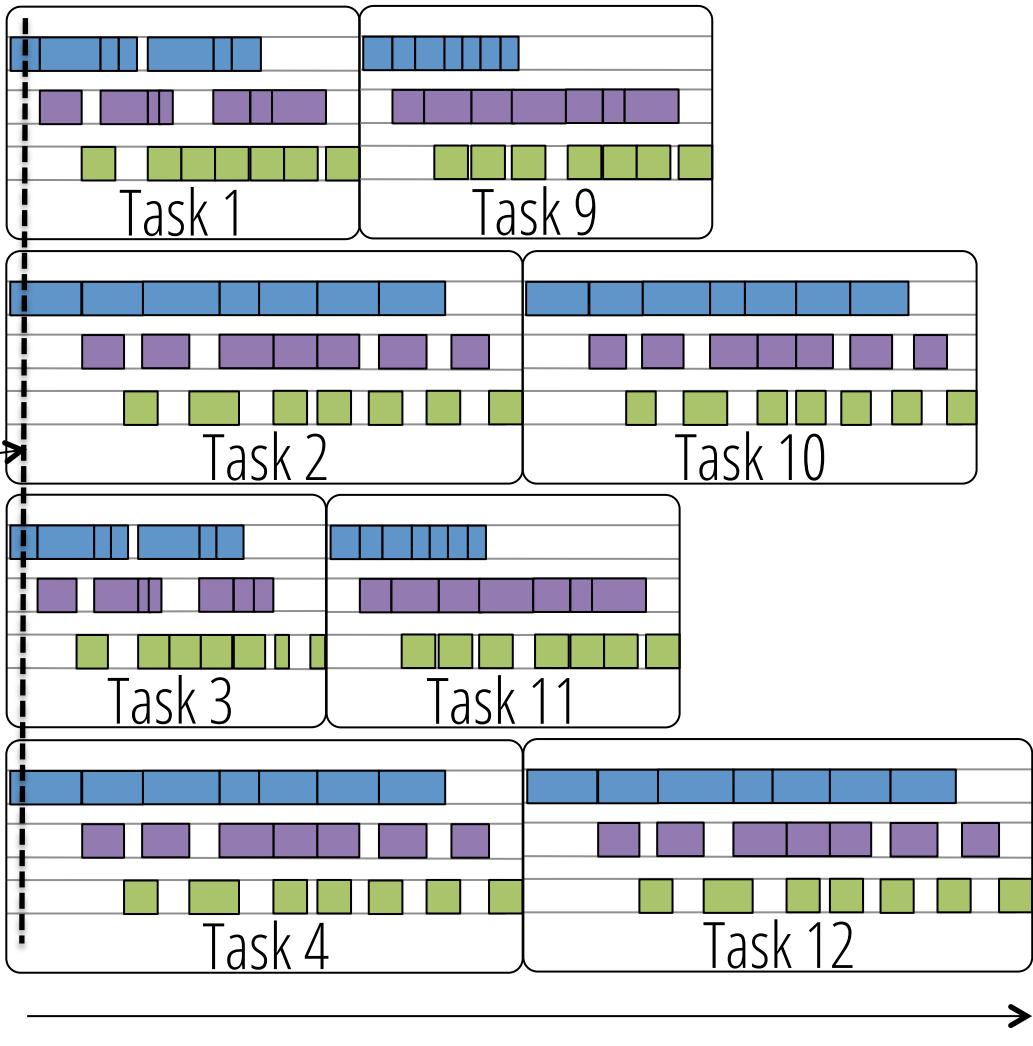


Challenge 1: Task pipelines multiple resources, bottlenecked on different resources at different times

4 concurrent tasks
on Worker 1



Challenge 2:
Concurrent tasks may
contend for
the same resource
(e.g., network)



Spark Summit 2015:

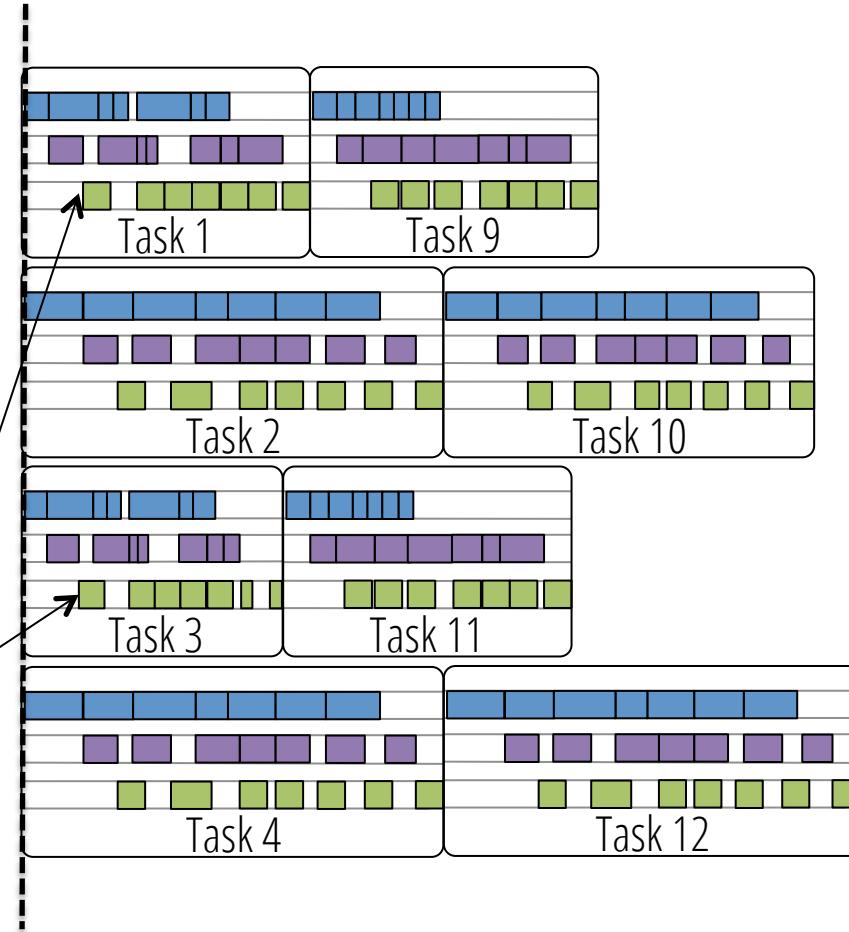
Blocked time analysis: how quickly
could a job have completed if a resource
were infinitely fast? (upper bound)

Result of ~1 year of adding metrics to Spark!

How much faster
would the job be with
2x disk throughput?

How would runtimes for these disk writes change?

How would that change timing of
(and contention for) other resources?



Challenges to reasoning about performance

Tasks bottleneck on different resources at different times

Concurrent tasks on a machine may contend for resources

No model for performance

How can we achieve this vision?

Spark overview

Reasoning about Spark's performance: why it's hard

New architecture: monotasks

Reasoning about monotasks performance: why it's easy

Monotasks in action (results)

Spark:

Tasks
bottleneck on
different
resources

Concurrent
tasks may
contend

No model for
performance

Spark:

Tasks bottleneck on different resources

Concurrent tasks may contend

No model for performance

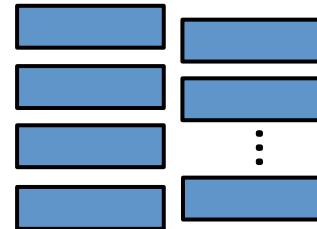
Monotasks: Each task uses one resource

Example Spark Job:

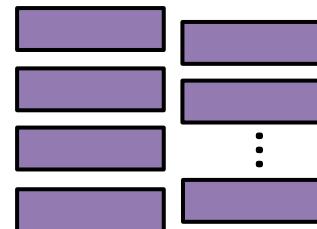
Read remote data

Filter records

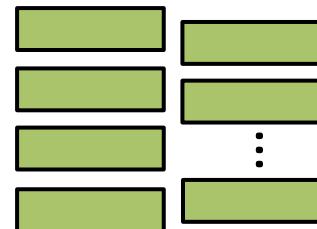
Write result to disk



Network monotasks:
Each read one remote block



CPU monotasks:
Each filter one block,
generate serialized output



Disk monotasks:
Each writes one block to disk

Spark:

Tasks bottleneck on different resources

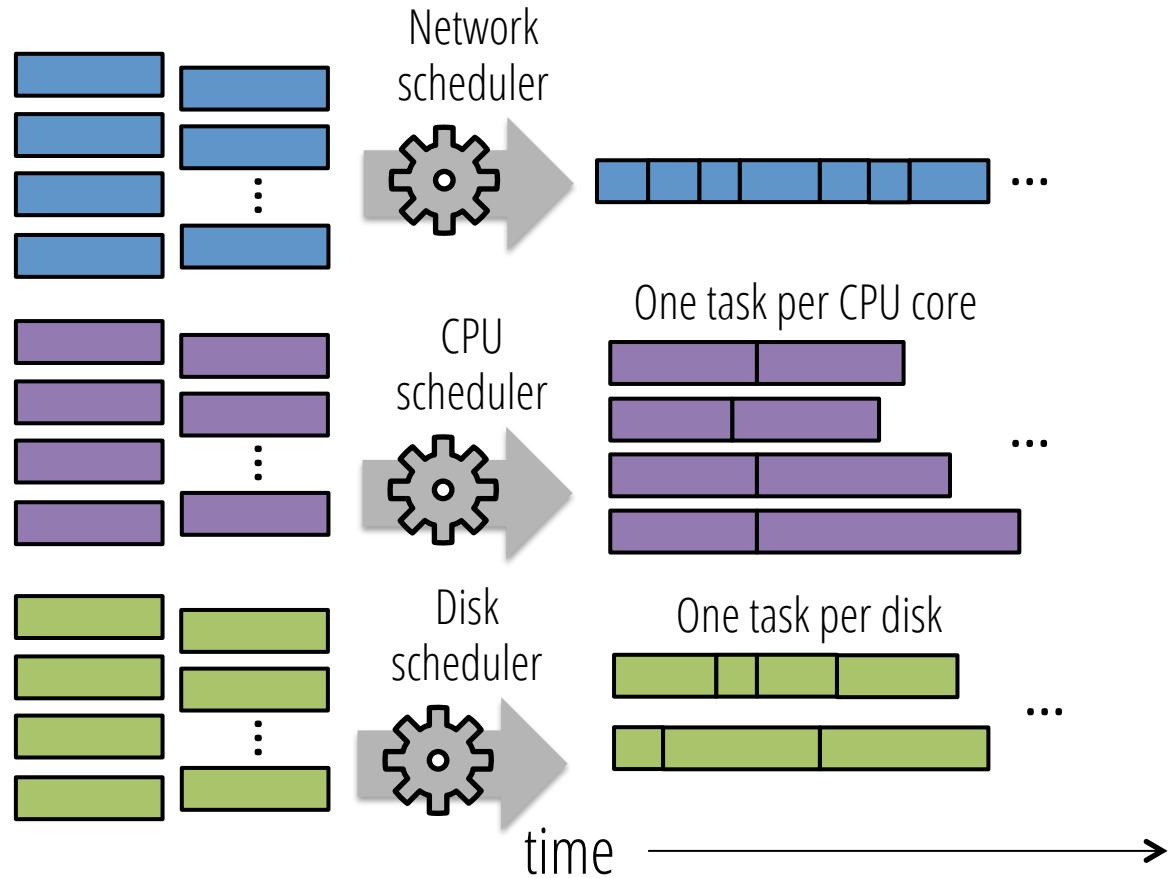
Concurrent tasks may contend

No model for performance

Monotasks:

Each task uses one resource

Dedicated schedulers control contention



Spark:

Tasks bottleneck on different resources

Concurrent tasks may contend

No model for performance

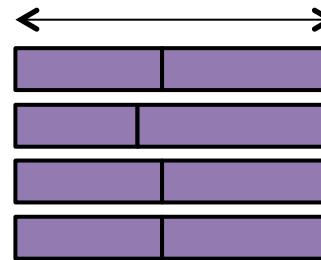
Monotasks:

Each task uses one resource

Dedicated schedulers control contention

Monotask times can be used to model performance

Ideal CPU time: total CPU monotask time / # CPU cores



Spark:

Tasks bottleneck on different resources

Concurrent tasks may contend

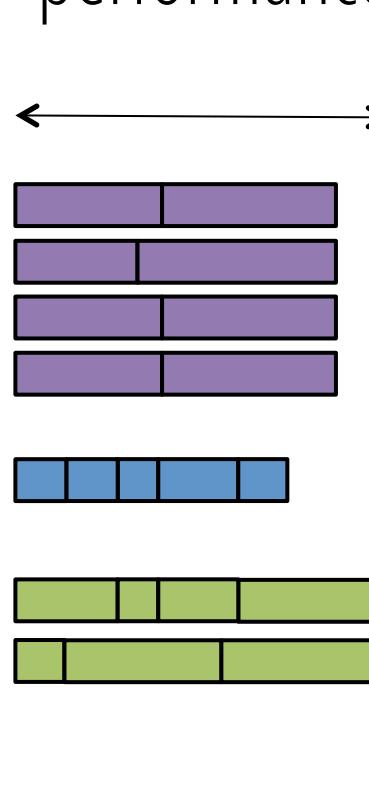
No model for performance

Monotasks:

Each task uses one resource

Dedicated schedulers control contention

Monotask times can be used to model performance



Ideal CPU time:
total CPU monotask
time / # CPU cores

Ideal network runtime

Ideal disk runtime

Job runtime:
max of ideal times

Spark:

Tasks bottleneck on different resources

Concurrent tasks may contend

No model for performance

Monotasks:

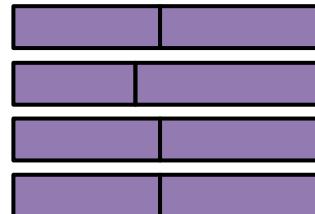
Each task uses one resource

Dedicated schedulers control contention

Monotask times can be used to model performance

How much faster would the job be with 2x disk throughput?

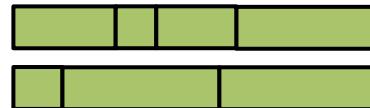
Ideal CPU time:
total CPU monotask
time / # CPU cores



Ideal network runtime



Ideal disk runtime



Spark:

Tasks bottleneck on different resources

Concurrent tasks may contend

No model for performance

Monotasks:

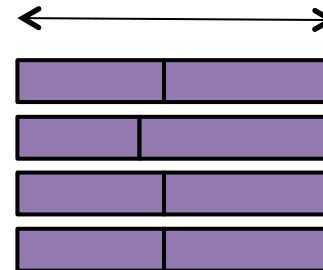
Each task uses one resource

Dedicated schedulers control contention

Monotask times can be used to model performance

How much faster would the job be with 2x disk throughput?

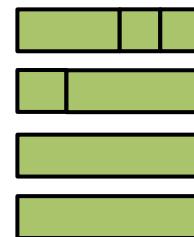
Ideal CPU time:
total CPU monotask
time / # CPU cores



Ideal network runtime



Ideal disk runtime
(2x disk concurrency)



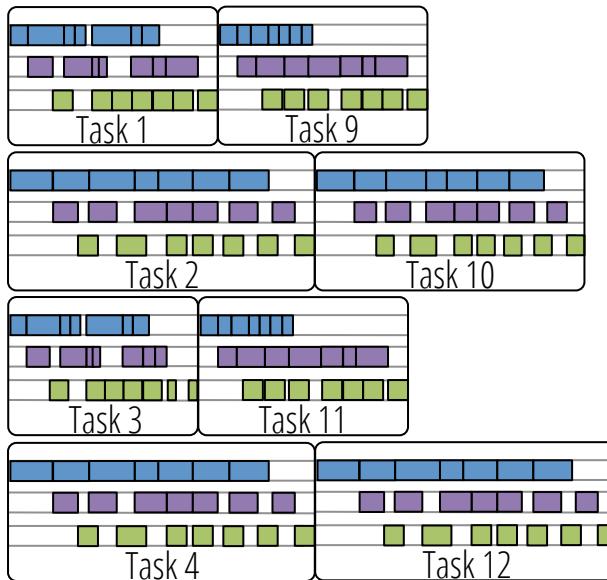
New job runtime

Spark:

Tasks bottleneck on different resources

Concurrent tasks may contend

No model for performance

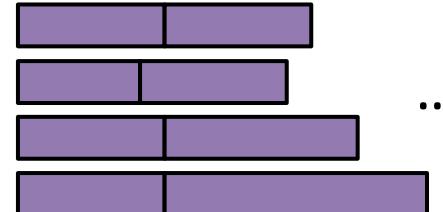


How does this decomposition work?

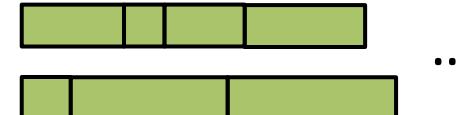
Network monotasks



Compute monotasks:
one task per CPU core



Disk monotasks: one
per disk



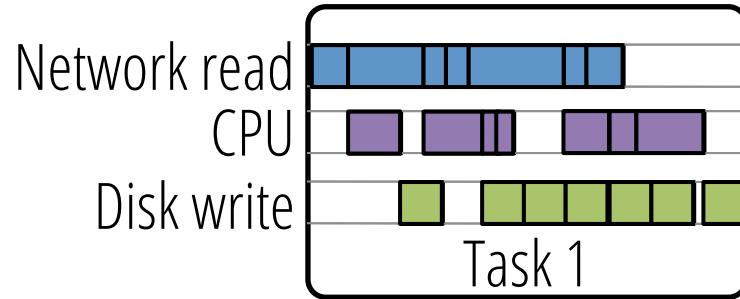
Monotasks:

Each task uses one resource

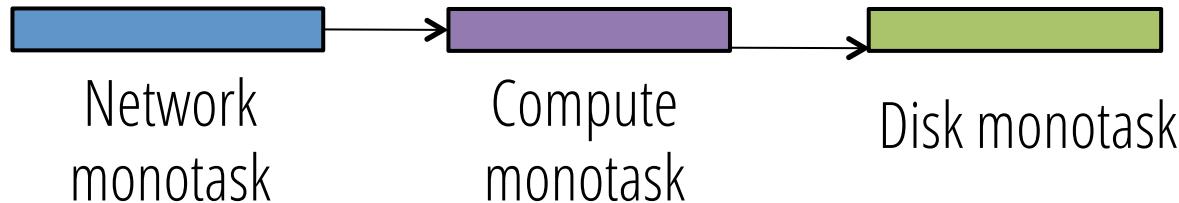
Dedicated schedulers control contention

Monotask times can be used to model performance

How does this decomposition work?



Monotasks



Implementation

API-compatible with Apache Spark

Workloads can be run on monotasks without re-compiling

Monotasks decomposition handled by Spark internals

Monotasks works at the application level

No operating system changes

How can we achieve this vision?

Spark overview

Reasoning about Spark's performance: why it's hard

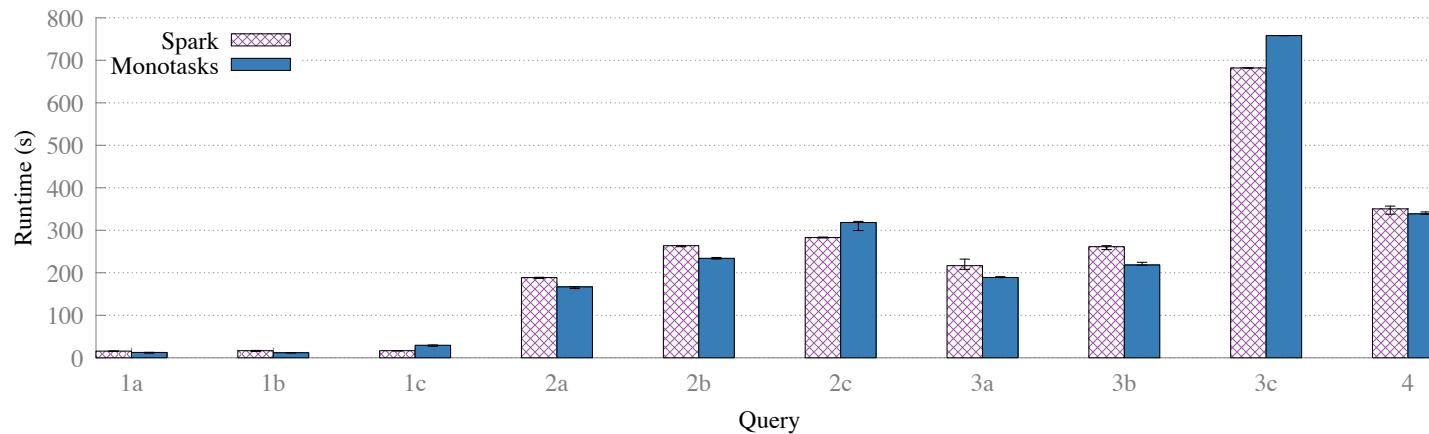
New architecture: monotasks

Reasoning about monotasks performance: why it's easy

Monotasks in action (results)

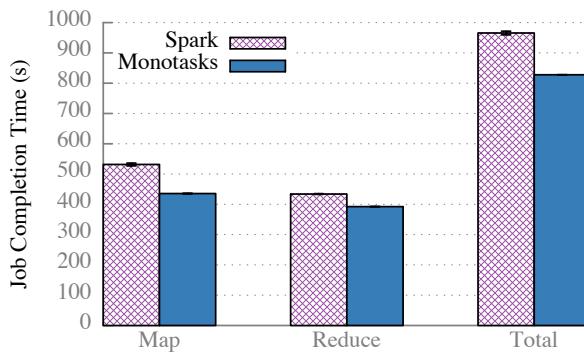
Performance on-par with Apache Spark

Big data benchmark
(SQL workload)

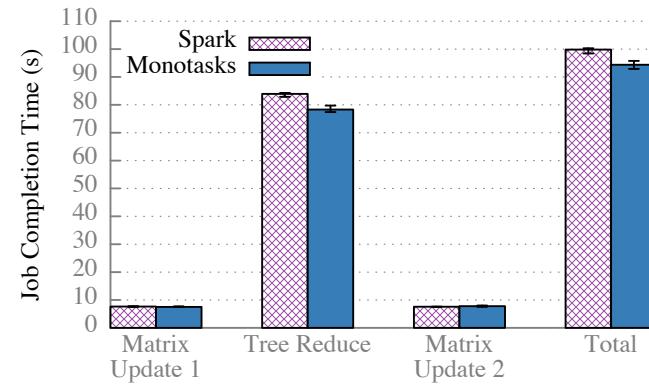


Performance on-par with Apache Spark

Sort
(600 GB, 20 machines)



Block coordinate descent
(Matrix workload used in ML applications)
16 machines

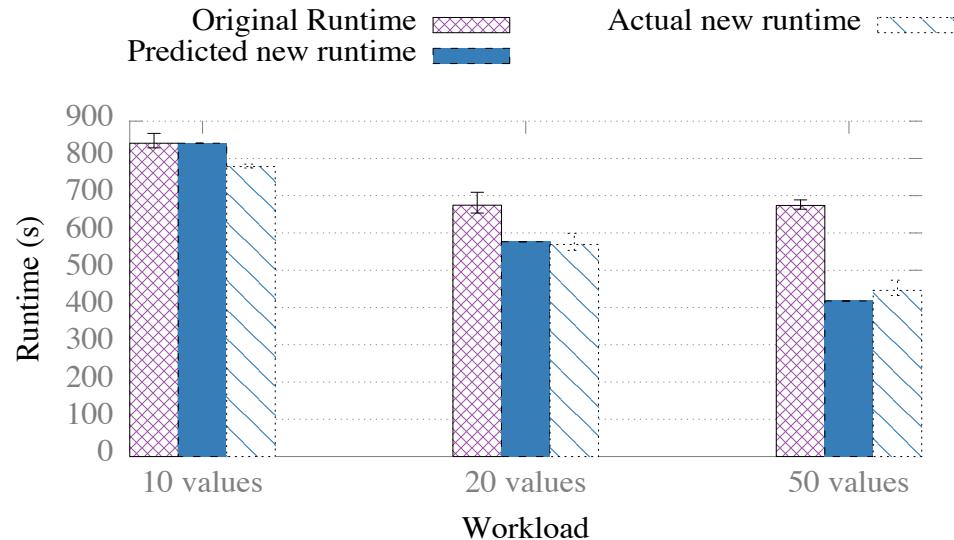


Monotasks in action

Modeling performance

Leveraging performance clarity to optimize performance

How much faster would jobs run if each machine had 2 disks instead of 1?



Predictions for different hardware
within 10% of the actual runtime

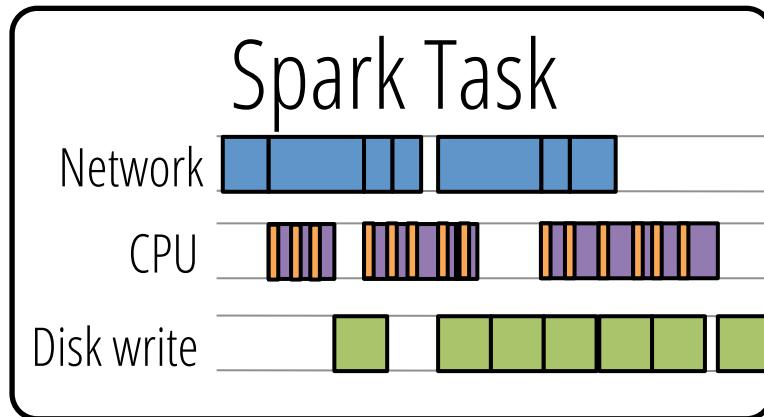
How much faster would job run if data were de-serialized and in memory?

Eliminates disk time to read input data

Eliminates CPU time to de-serialize data

How much faster would job run if data were de-serialized and in memory?

Measuring (de) serialization time with Spark



(De) serialization pipelined with other processing *for each record*

Application-level measurement incurs high overhead

■ : (de)serialization time

How much faster would job run if data were de-serialized and in memory?

Measuring (de) serialization time with Monotasks



Original compute monotask



Un-rolled monotask

Eliminating fine-grained pipelining enables measurement!

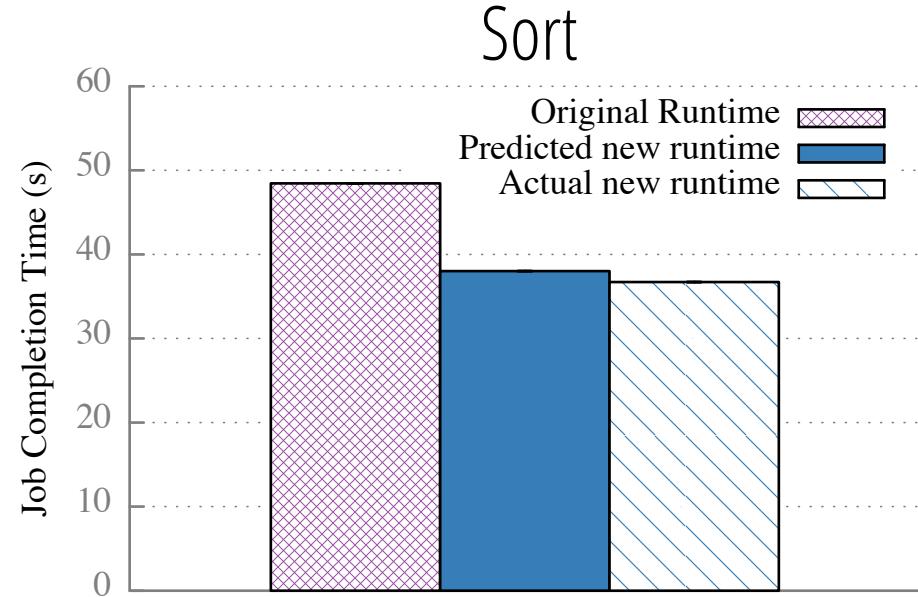
■ : (de)serialization time

How much faster would job run if data were de-serialized and in memory?

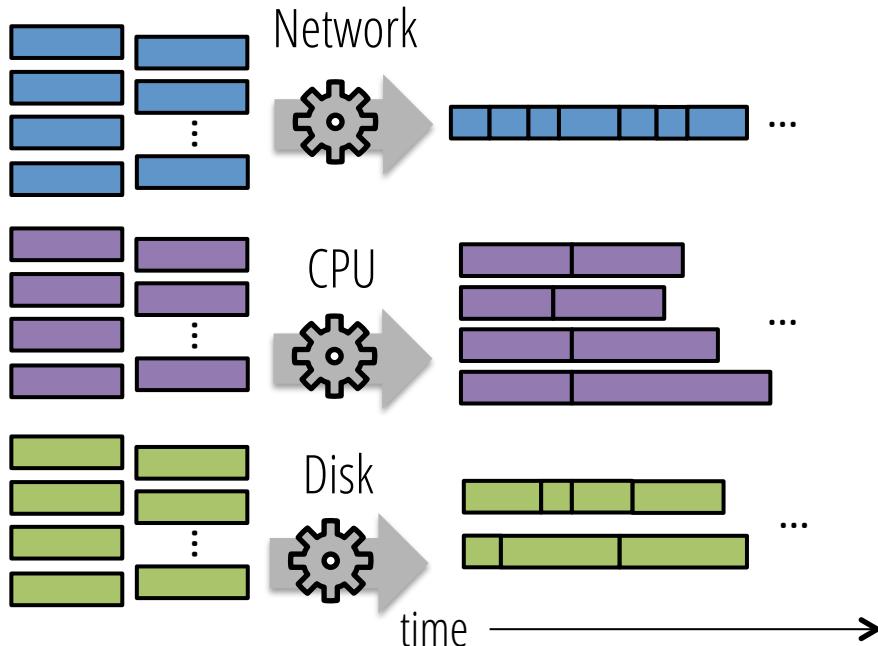
Eliminate disk monotask time

Eliminate CPU monotask time
spent (de)serialiazing

Re-run model



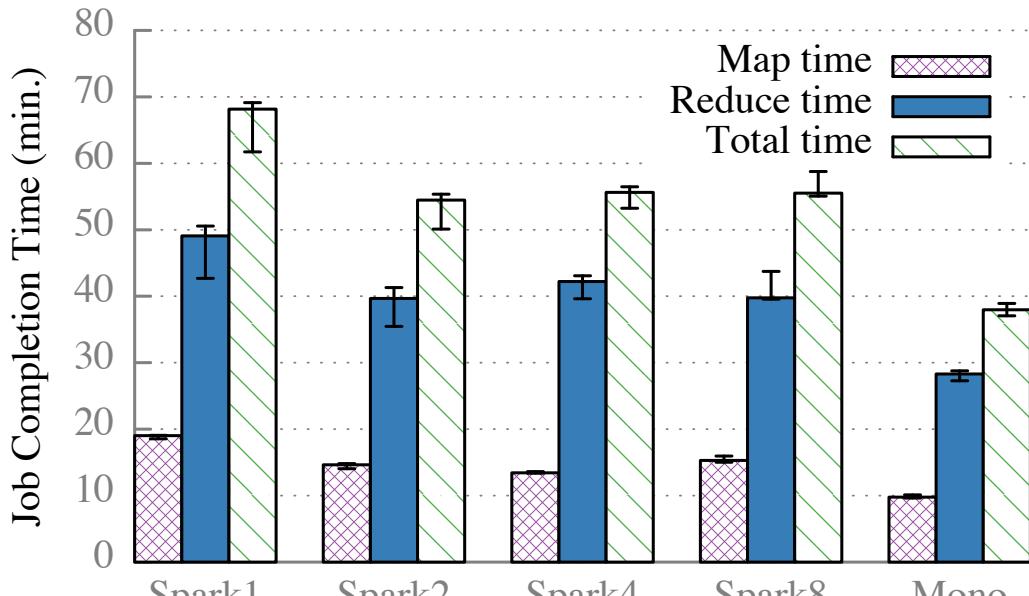
Leveraging Performance Clarity to Improve Performance



Schedulers have complete visibility over resource use

Framework can configure for best performance

Configuring the number of concurrent tasks



Spark with different numbers of concurrent tasks

Monotasks better than any configuration:
per-resource schedulers automatically schedule with the ideal concurrency

Future Work

Leveraging performance clarity to improve performance

- Use resource queue lengths to dynamically adjust job

Automatically configuring for multi-tenancy

- Don't need jobs to specify resource requirements
- Can achieve higher utilization: no multi-resource bin packing

Details for Stage 17

Total task time across all tasks: 13 min

Shuffle read: 2.5 GB / 31589120

Performance Information

Bottleneck: Disk (if disk bandwidth were increased by 23% or more, network would become the bottleneck)

Non-bottlenecks: Network (could reduce network bandwidth by up to 30% slower without impacting runtime, increase CPU time by up to 2x without impacting runtime)

Benefit of caching: Storing input in-memory would reduce job completion time by 42%

Job Runtime Predictor: (enter in properties of different cluster to estimate job's runtime)

Number of machines:

CPU cores per machine:

Network bandwidth per machine:

Gbps

Vision:

Spark always reports bottleneck information

Challenging with existing architecture

Monotasks:

Network monotasks

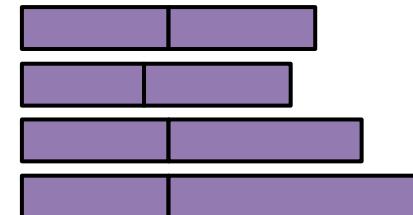


Each task uses one resource

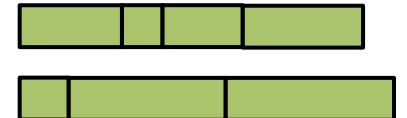
Dedicated schedulers control contention

Monotask times can be used to model performance

Compute monotasks:
one task per CPU core



Disk monotasks: one per disk



Interested? Have a job whose performance you can't figure out? Email me:
keo@cs.berkeley.edu