



Efficient state management with Spark 2.0 and scale-out databases

Jags Ramnarayan

Co-founder, CTO, SnappyData

Barzan Mozafari

Strategic Advisor @ Snappydata


Assistant Professor,

Univ. of Michigan, Ann Arbor



SPARK SUMMIT 2016
DATA SCIENCE AND ENGINEERING AT SCALE
JUNE 6-8, 2016 SAN FRANCISCO

SnappyData's focus – DB on Spark



- New Spark-based open source project started by Pivotal GemFire founders+engineers
- Decades of in-memory data management experience
- Focus on real-time, operational analytics: Spark inside an OLTP+OLAP database



Agenda

- Mixed workloads are important but complex today
- State management challenges with such workloads in spark Apps
- The SnappyData solution
- Approximate query processing
- How do we extend spark for real time, mixed workloads?
- Q&A

Mixed Workloads are increasingly common



IoT Sensor Streams


Anomaly detection – score against models

Correlate current temperature trend with history....

- Map sensors to tags
- Monitor temperature
- Send alerts

Interact using dynamic queries

Mixed workload Architecture is Complex (Lambda)



Lambda Architecture is Complex

- Complexity: learn and master multiple products → data models, disparate APIs, configs
- Slower
- Wasted resources



Can we simplify & optimize?

Perhaps a single clustered DB that can manage stream state, transactional data and run OLAP queries?



SPARK SUMMIT 2016


Deeper Look into Managing State in Spark Applications



SPARK SUMMIT 2016

Deeper Look – Ad Impression Analytics

Ad Network Architecture – Analyze log impressions in real time



Ad Impression Analytics

To simplify, let's consider that impression logs are in this format:

timestamp	publisher	advertiser	website	geo	bid	cookie
2013-01-28 13:21:12	pub1	adv10	abc.com	NY	0.0001	1214
2013-01-28 13:21:13	pub1	adv10	abc.com	NY	0.0005	1214
2013-01-28 13:21:14	pub2	adv20	xyz.com	CA	0.0003	4321
2013-01-28 13:21:15	pub2	adv20	xyz.com	CA	0.0001	5675

Our goal is to aggregate these logs by publisher and geo, and compute the average bid, the number of impressions and the number of uniques by minute. So the aggregation will look something like:

timestamp	publisher	geo	avg_bid	imps	uniques
2013-01-28 13:21:00	pub1	NY	0.0003	256	104
2013-01-28 13:21:00	pub2	CA	0.0002	121	15
2013-01-28 13:22:00	pub1	NY	0.0001	190	98
2013-01-28 13:22:00	pub2	CA	0.0007	137	19



Bottlenecks in the write path

```
val input :DataFrame= sqlContext.read  
  .options(kafkaOptions).format(..)  
  .stream("Kafka url")
```

Stream micro batches in parallel from Kafka to each Spark executor

```
val result :DataFrame = input  
  .where("geo != 'UnknownGeo'")  
  .groupBy(  
    window("event-time", "1min"),  
    "Publisher", "geo")  
  .agg(avg("bid"), ....)
```

Filter and collect event for 1 minute.
Reduce to 1 event per Publisher, Geo every minute

```
val query = result.write.format("My Favorite NoSQLDB")  
  .outputMode("append")  
  .startStream("dest-path")
```


Execute GROUP BY ... Expensive Spark Shuffle ...

Shuffle again in DB cluster ...
data format changes ...
serialization costs




Bottlenecks in the Write Path

Shuffle Costs (Copying, Serialization)




- Aggregations – GroupBy, MapReduce
- Joins with other streams, Reference data

Excessive copying in Java based Scale out stores




Avoid bottleneck - Localize process with State



SPARK SUMMIT 2016

New State Management API in Spark 2.0 – KV Store

- Preserve state across streaming aggregations across multiple batches
- Fetch, store KV pairs
- Transactional
- Supports versioning (batch)
- Built in store that persists to HDFS



Source: Databricks presentation



SPARK SUMMIT 2016

Impedance mismatch with KV stores?

We want to run interactive “scan”-intensive queries:

- Find total *uniques* for Ads grouped on geography
- Impression trends for Advertisers(group by query)

Two alternatives: row-stores(**all KV Stores**) vs. column stores

Fast key-based lookup

But, too slow to run
aggregations, scan based
interactive queries

Emp_no	Dept_id	Hire_date	Emp_fn	Emp_ln
1	1	2001-01-01	Smith	Bob
2	1	2002-02-01	Jones	Jim
3	1	2002-05-01	Young	Sue
4	2	2003-02-01	Sternle	Bill
5	2	1999-06-15	Aurora	Jack
6	3	2000-08-15	Jung	Laura

Fast scans, aggregations

Updates and random
writes are very difficult

Row-Oriented Database

1	1	2001-01-01	Smith	Bob
2	1	2002-02-01	Jones	Jim
3	1	2002-05-01	Young	Sue

Consume too
much memory


Column-Oriented Database

1	2	3	4	5
1	1	1	2	2
2001-01-01	2002-02-01	2002-02-01	2002-02-01	2002-02-01



Why columnar storage in-memory?

Hardware Changes: The Memory Wall



Source: MonetDB




SPARK SUMMIT 2016

But, are in-memory column-stores fast enough?

AWS c4.2xlarge ; 4 x (8 cores, 15GB)

Column Table: AdImpressions; 450million rows

```
select count(*) from adImpressions  
group by geo  
order by count desc  
limit 20;
```




Single User, 10+seconds. Is this Interactive speed?




SPARK SUMMIT 2016

Use statistical techniques to shrink data!

- Most apps happy to tradeoff 1% accuracy for 200x speedup!
 - Can usually get a 99.9% accurate answer by only looking at a tiny fraction of data!
- Often can make perfectly accurate decisions with imperfect answers!
 - A/B Testing, visualization, ...
- The data itself is usually noisy
 - Processing entire data doesn't necessarily mean exact answers!
- Inference is probabilistic anyway



Original (2 billion points), 71 mins



VAS (1M), 3 secs



SPARK SUMMIT 2016


Our Solution: SnappyData

Open Sourced @ <https://github.com/SnappyDataInc/snappydata>

SnappyData: A New Approach



A Single Unified Cluster: OLTP + OLAP + Streaming
for real-time analytics




Vision: Drastically reduce the cost and complexity in modern big data



SPARK SUMMIT 2016

Unified In-memory DB for Streams, Txn, OLAP queries



Real-time operational Analytics – TBs in memory

Features

- Deeply integrated database for Spark
 - 100% compatible with Spark
 - Extensions for Transactions (updates), SQL stream processing
 - Extensions for High Availability
 - Approximate query processing for interactive OLAP
- OLTP+OLAP Store
 - Replicated and partitioned tables
 - Tables can be Row or Column oriented (in-memory & on-disk)
 - SQL extensions for compatibility with SQL Standard
 - create table, view, indexes, constraints, etc



Approximate Query Processing Features


- Uniform random sampling
- Stratified sampling
 - Solutions exist for stored data (BlinkDB)
 - SnappyData works for infinite streams of data too
- Support for synopses
 - Top-K queries, heavy hitters, outliers, ...
- Exponentially decaying windows over time




SPARK SUMMIT 2016

Interactive-Speed Analytic Queries – Exact or Approximate

Select avg(Bid), Advertiser from T1 group by Advertiser → 100 secs
Select avg(Bid), Advertiser from T1 group by Advertiser with error 0.1 → 2 secs




Airline On-time performance Analytics – Demo




SPARK SUMMIT 2016

Revisiting AdAnalytics – Spark with colocated store




Revisiting AdAnalytics – Spark with colocated store



SPARK SUMMIT 2016

Concurrent Ingest + Query Performance




- AWS 4 c4.2xlarge instances - 8 cores, 15GB mem
- Each node parallelly ingests stream from Kafka
- Parallel batch writes to store (32 partitions)
- Only few cores used for Stream writes as most of CPU reserved for OLAP queries

<https://github.com/SnappyDataInc/snappy-poc>



SPARK SUMMIT 2016

Concurrent Ingest + Query Performance



Sample “scan” oriented OLAP query(Spark SQL) performance executed while ingesting data

<https://github.com/SnappyDataInc/snappy-poc>




SPARK SUMMIT 2016

How SnappyData Extends Spark




SPARK SUMMIT 2016

Unified Cluster Architecture




How do we extend Spark for Real Time?



- Spark Executors are long running. **Driver failure** doesn't shutdown Executors
- **Driver HA** – Drivers run “Managed” with standby secondary
- **Data HA** – Consensus based clustering integrated for eager replication



How do we extend Spark for Real Time?



- By pass scheduler for low latency SQL
- Deep integration with **Spark Catalyst(SQL)** – collocation optimizations, indexing use, etc
- Full SQL support – Persistent Catalog, Transaction, DML



Unified OLAP/OLTP/streaming with Spark

- Far fewer resources: TB problem becomes GB.
 - CPU contention drops
- Far less complex
 - single cluster for stream ingestion, continuous queries, interactive queries and machine learning
- Much faster
 - compressed data managed in distributed memory in columnar form reduces volume and is much more responsive



SPARK SUMMIT 2016

SnappyData is Open Source

- Ad Analytics example/benchmark -
<https://github.com/SnappyDataInc/snappy-poc>
- <https://github.com/SnappyDataInc/snappydata>



Join SnappyData for pizza and drinks, tonight at *Tradition!*

JUNE 7, 2016 @ 7:30 PM

441 JONES ST, SAN FRANCISCO, CA
(JUST AROUND THE CORNER)



THANK YOU.

Drop by our booth to learn more.



SPARK SUMMIT 2016
DATA SCIENCE AND ENGINEERING AT SCALE
JUNE 6-8, 2016 SAN FRANCISCO

Snappy Spark Cluster Deployment topologies

Unified Cluster



- Snappy store and Spark Executor share the JVM memory
- Reference based access – zero copy

Split Cluster



- SnappyStore is isolated but use the same COLUMN FORMAT AS SPARK for high throughput

Simple API – Spark Compatible

- Access Table as DataFrame
Catalog is automatically recovered

```
val impressionLogs: DataFrame = context.table(colTable)
val campaignRef: DataFrame = context.table(rowTable)

val parquetData: DataFrame = context.table(parquetTable)
<... Now use any of DataFrame APIs ... >
```

- Store RDD[T]/DataFrame can be stored in SnappyData tables
- Access from Remote SQL clients
- Additional API for updates, inserts, deletes

```
//Save a dataFrame using the Snappy or spark context ...
context.createExternalTable("T1", "ROW", myDataFrame.schema,
props );

//save using DataFrame API
dataDF.write.format("ROW").mode(SaveMode.Append).options(prop
s).saveAsTable("T1");
```



Extends Spark

```
CREATE [Temporary] TABLE [IF NOT EXISTS] table_name
(
    <column definition>
) USING 'JDBC | ROW | COLUMN'
OPTIONS (
    COLOCATE_WITH 'table_name',          // Default none
    PARTITION_BY 'PRIMARY KEY | column name', // will be a replicated table, by default
    REDUNDANCY    '1',                  // Manage HA
    PERSISTENT "DISKSTORE_NAME ASYNCHRONOUS | SYNCHRONOUS",
                // Empty string will map to default disk store.
    OFFHEAP "true | false"
    EVICTION_BY "MEMSIZE 200 | COUNT 200 | HEAPPERCENT",
    ....
    [AS select_statement];
```



Simple to Ingest Streams using SQL

