

Spark Uber Development Kit

Kelvin Chu, Hadoop Platform, Uber
Gang Wu, Hadoop Platform, Uber

Spark Summit 2016
June 07, 2016



Uber Mission

**“Transportation as reliable as
running water, everywhere,
for everyone”**

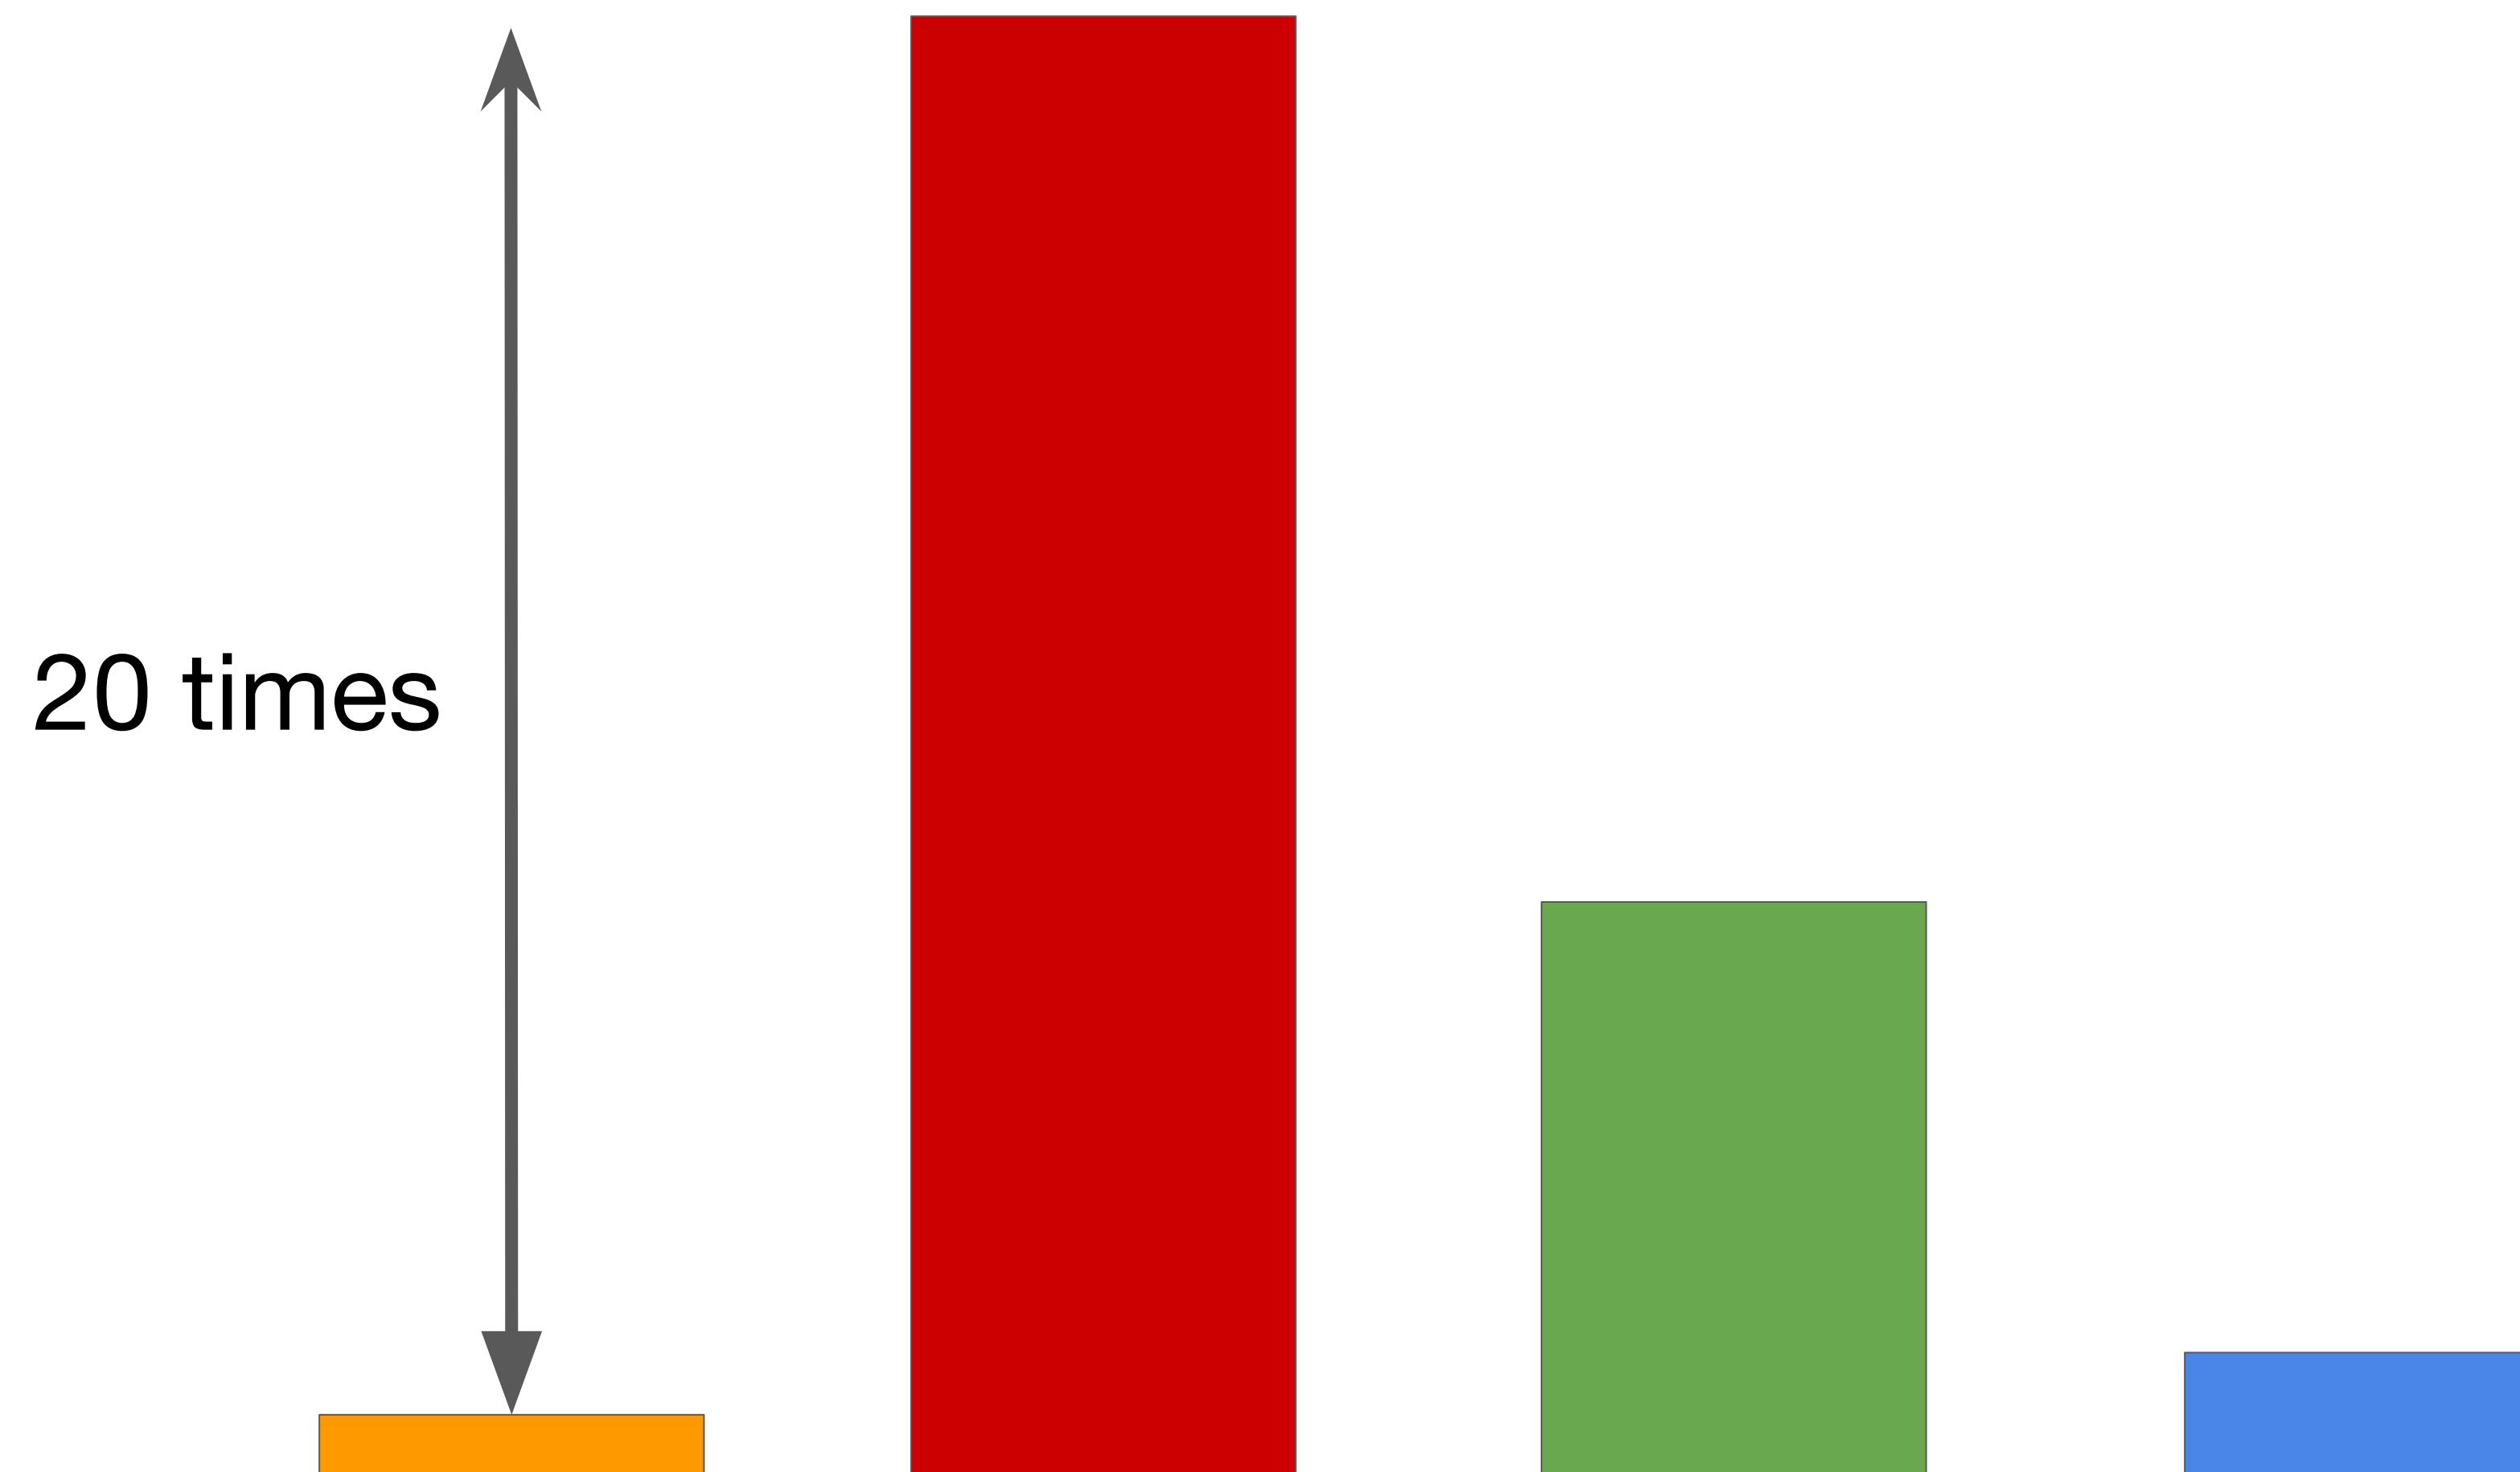
About Us

- Hadoop team of Data Infrastructure at Uber
- Schema systems
- HDFS data lake
- Analytics engines on Hadoop
- Spark computing framework and toolings

Execution Environment

Complexity

Cluster Sizes



YARN

Mesos

Docker

JVM

Parquet

ORC

Sequence

Text

Home Built Services

Hive

Kafka

ELK

Consequence:

Pretty hard for beginners, sometimes hard for experienced users too.

Goals:

Multi-Platform: Abstract out environment

Self-Service: Create and run Spark jobs super easily

Reliability: Prevent harm to infrastructure systems

Engineers

SRE

API

Tools

API

Tools

Engineers

Easy
Self-Service
Multi-Platform

SRE

No Harm
Reliability

API

Engineers

SRE

Tools

- SCBuilder
- Kafka dispersal

- SparkPlug

SCBuilder

Encapsulate cluster environment details

- Builder Pattern for SparkContext
- Incentive for users:
 - performance optimized (default can't pass 100GB)
 - debug optimized (history server, event logs)
 - don't need to ask around YARN, history servers, HDFS configs
- Best practices enforcement:
 - SRE approved CPU and memory settings
 - resource efficient serialization

Kafka Dispersal

Kafka as data sink of RDD result

`publish(data: RDD, topic: String, schemald: Int, appId: String)`

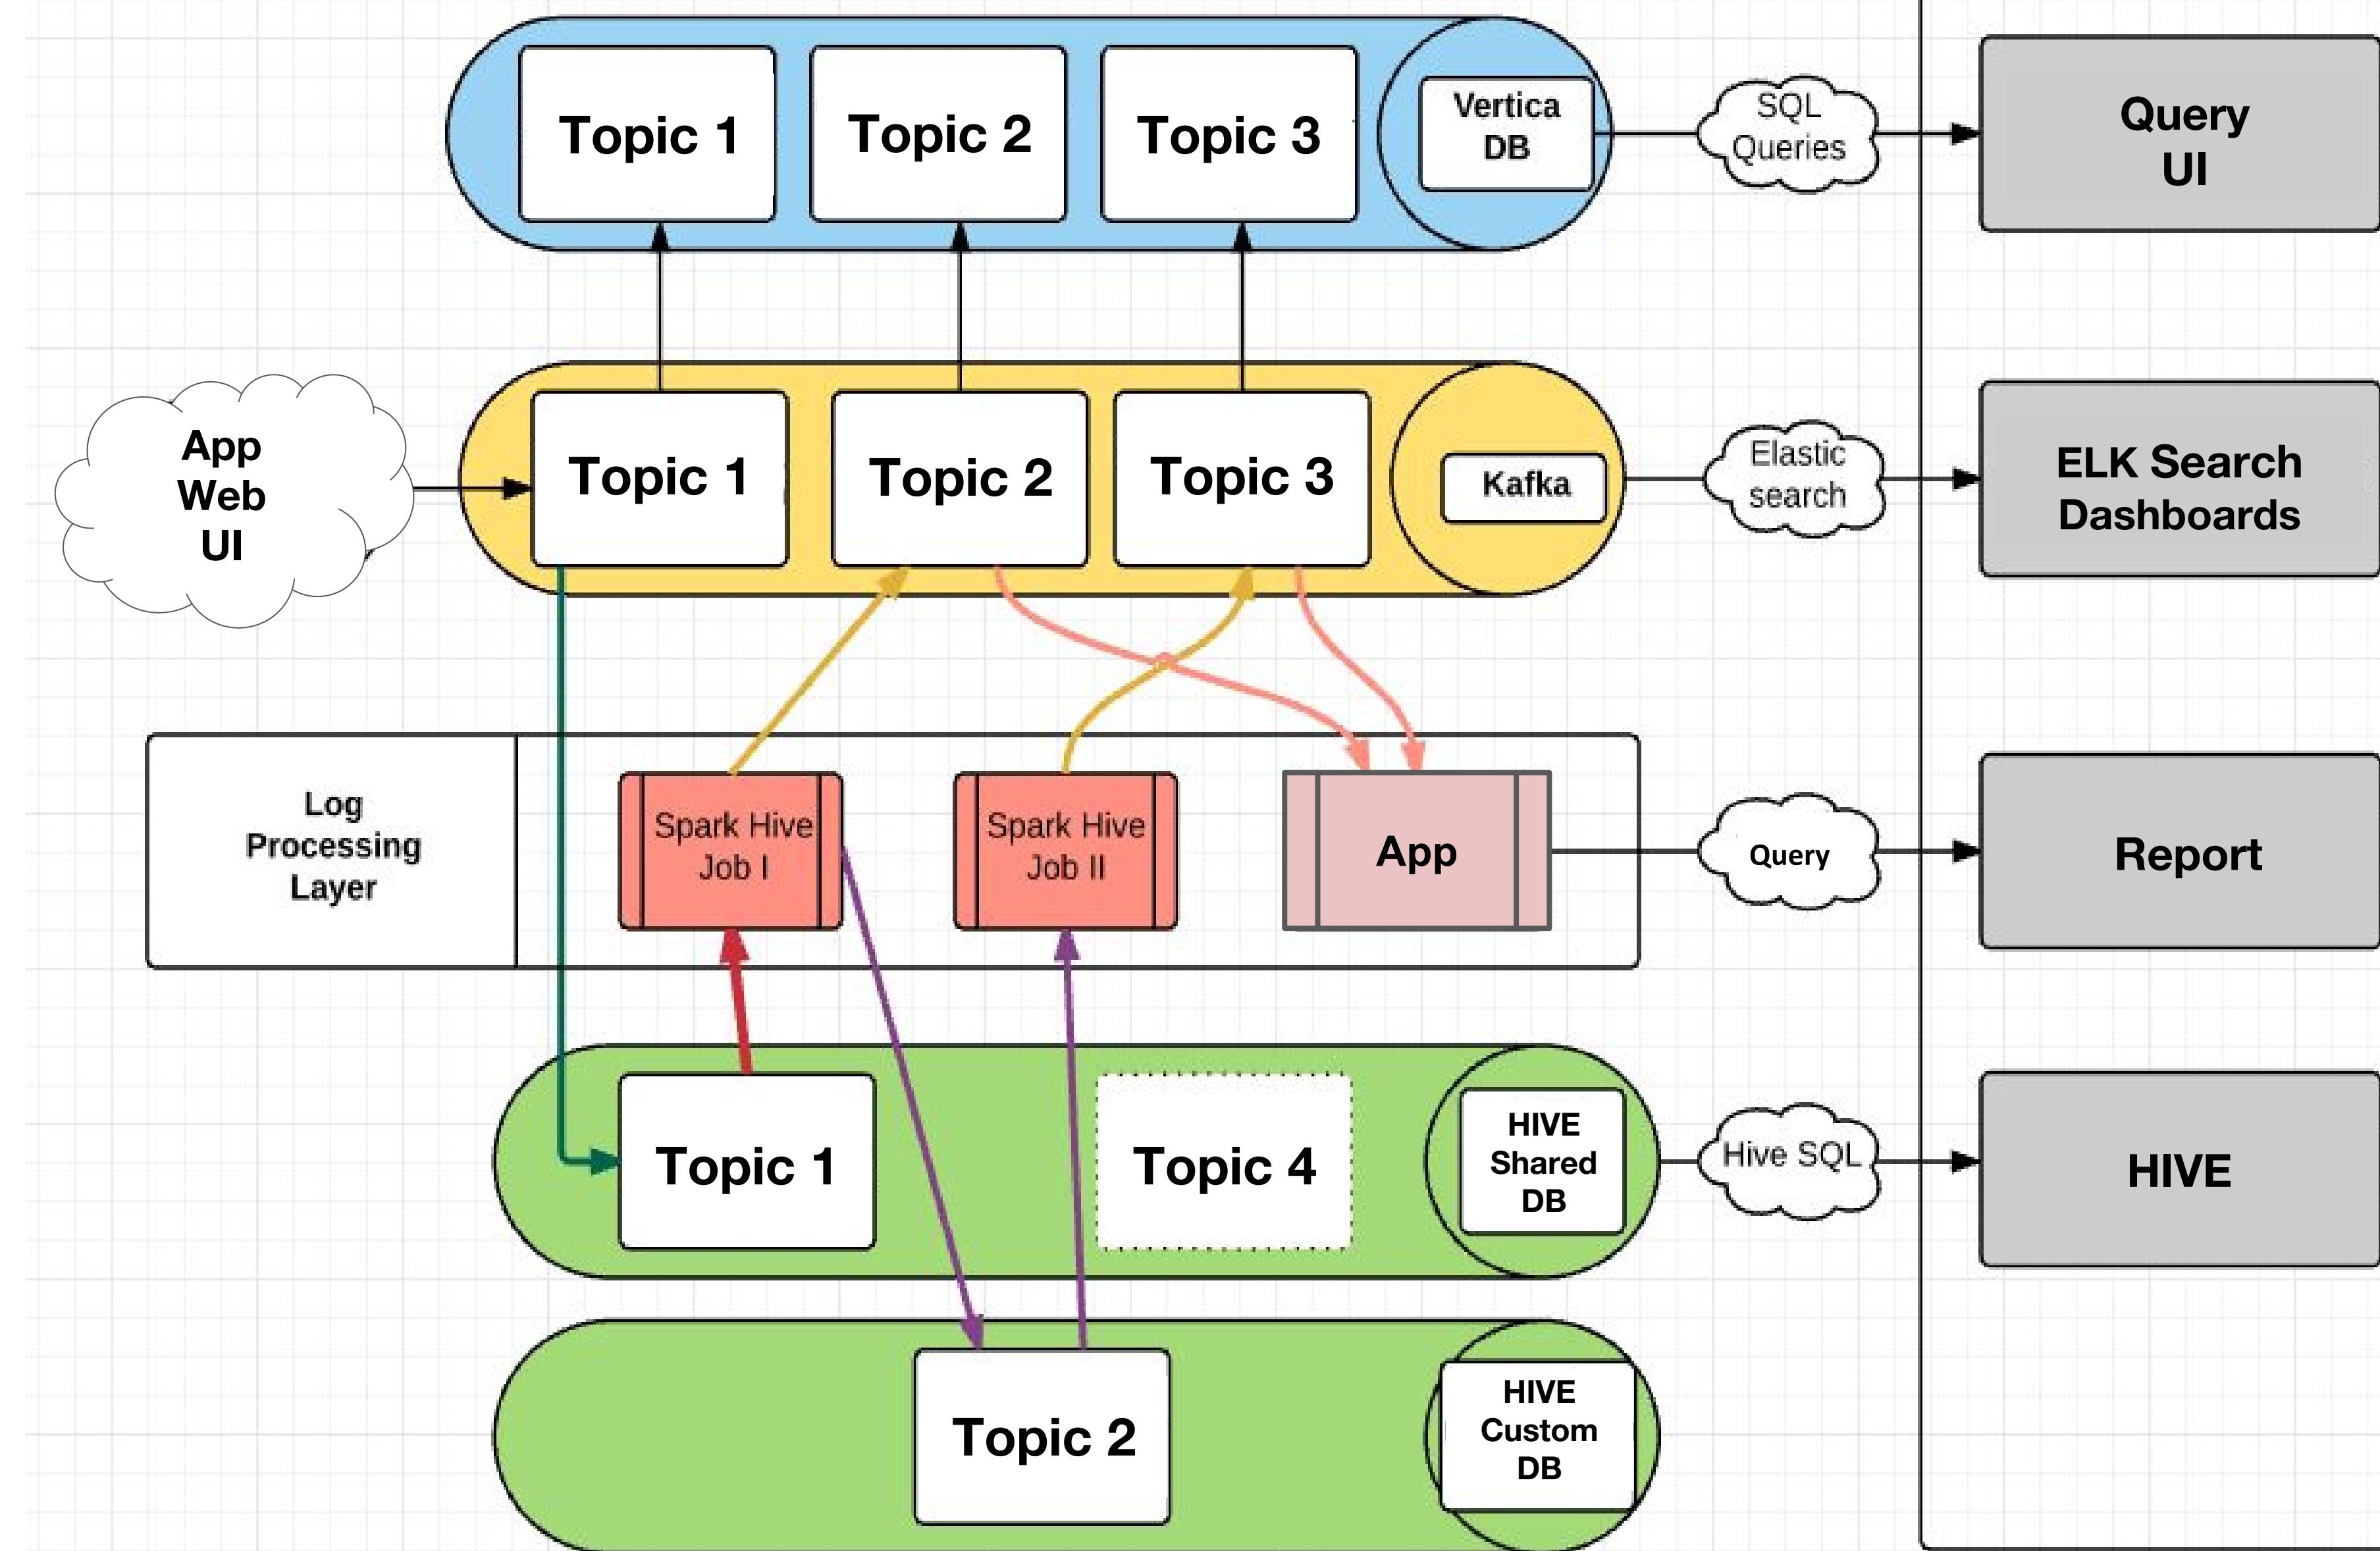
- Incentive for users:
 - RDD as first class citizen => parallelization
 - built-in HA
- Best practices enforcement:
 - rate limiting
 - message integrity by schema
 - bad messages tracking

SparkPlug

Kickstart job development

- A collection of popular job templates
 - Two commands to run the first job in Dev
- One use case per template
 - e.g. Ozzie + SparkSQL + Incremental processing
 - e.g. Incremental processing + Kafka dispersal
- Best Practices
 - built-in unit tests, test coverage, Jenkins
 - built-in Kafka, HDFS mocks

Example Spark Application Data Flow Chart





API

Engineers

SRE

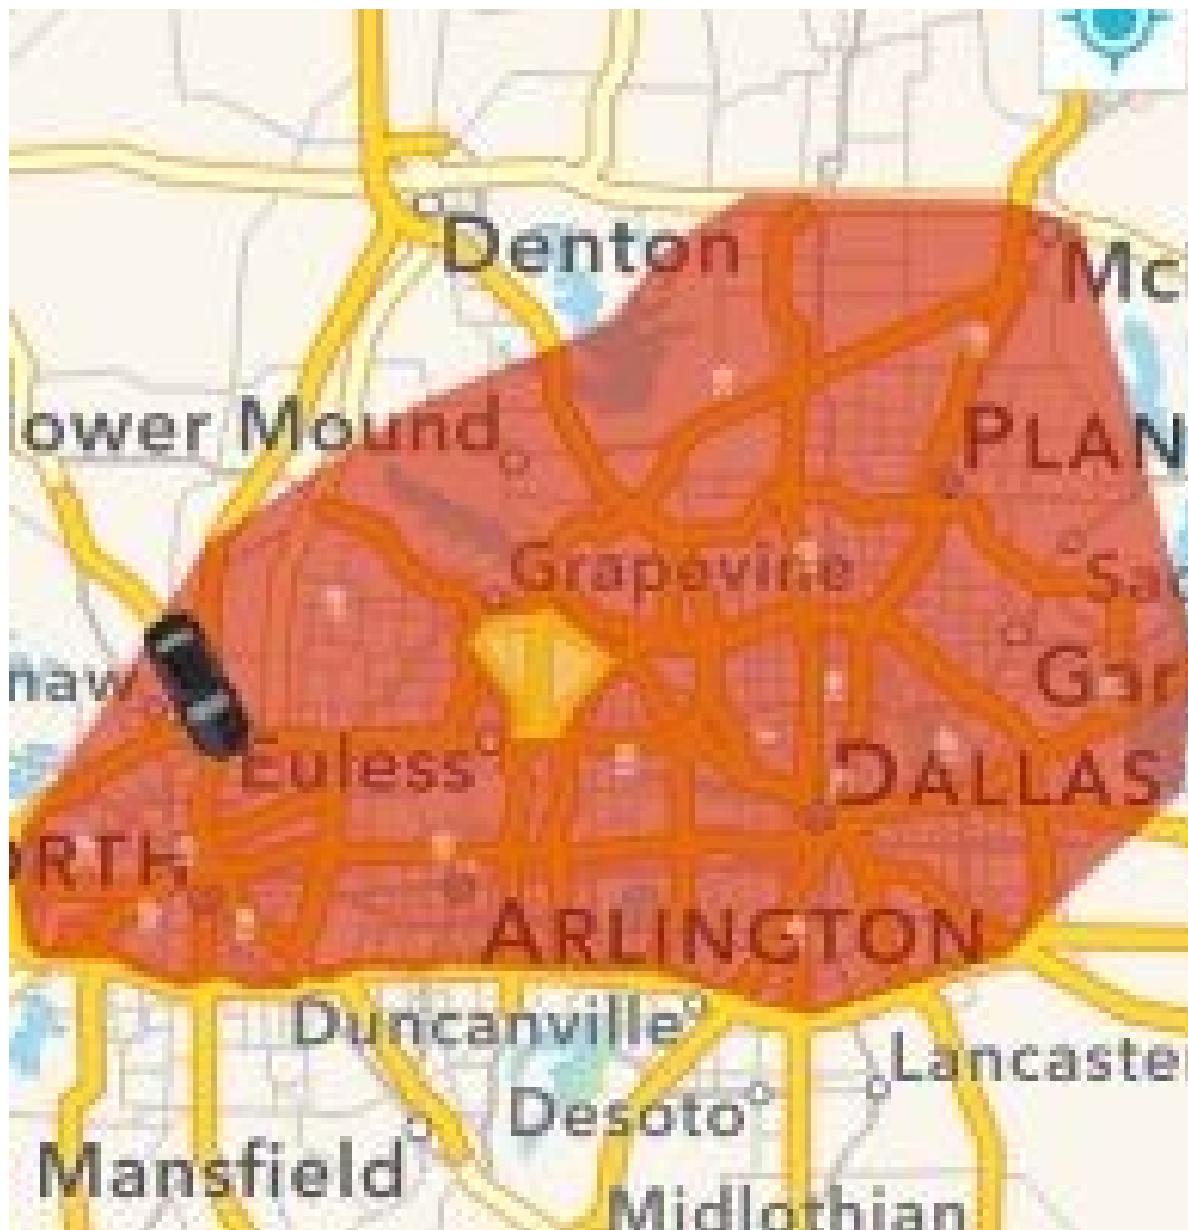
Tools

- Geo-spatial processing
- SCBuilder
- Kafka dispersal

- SparkPlug

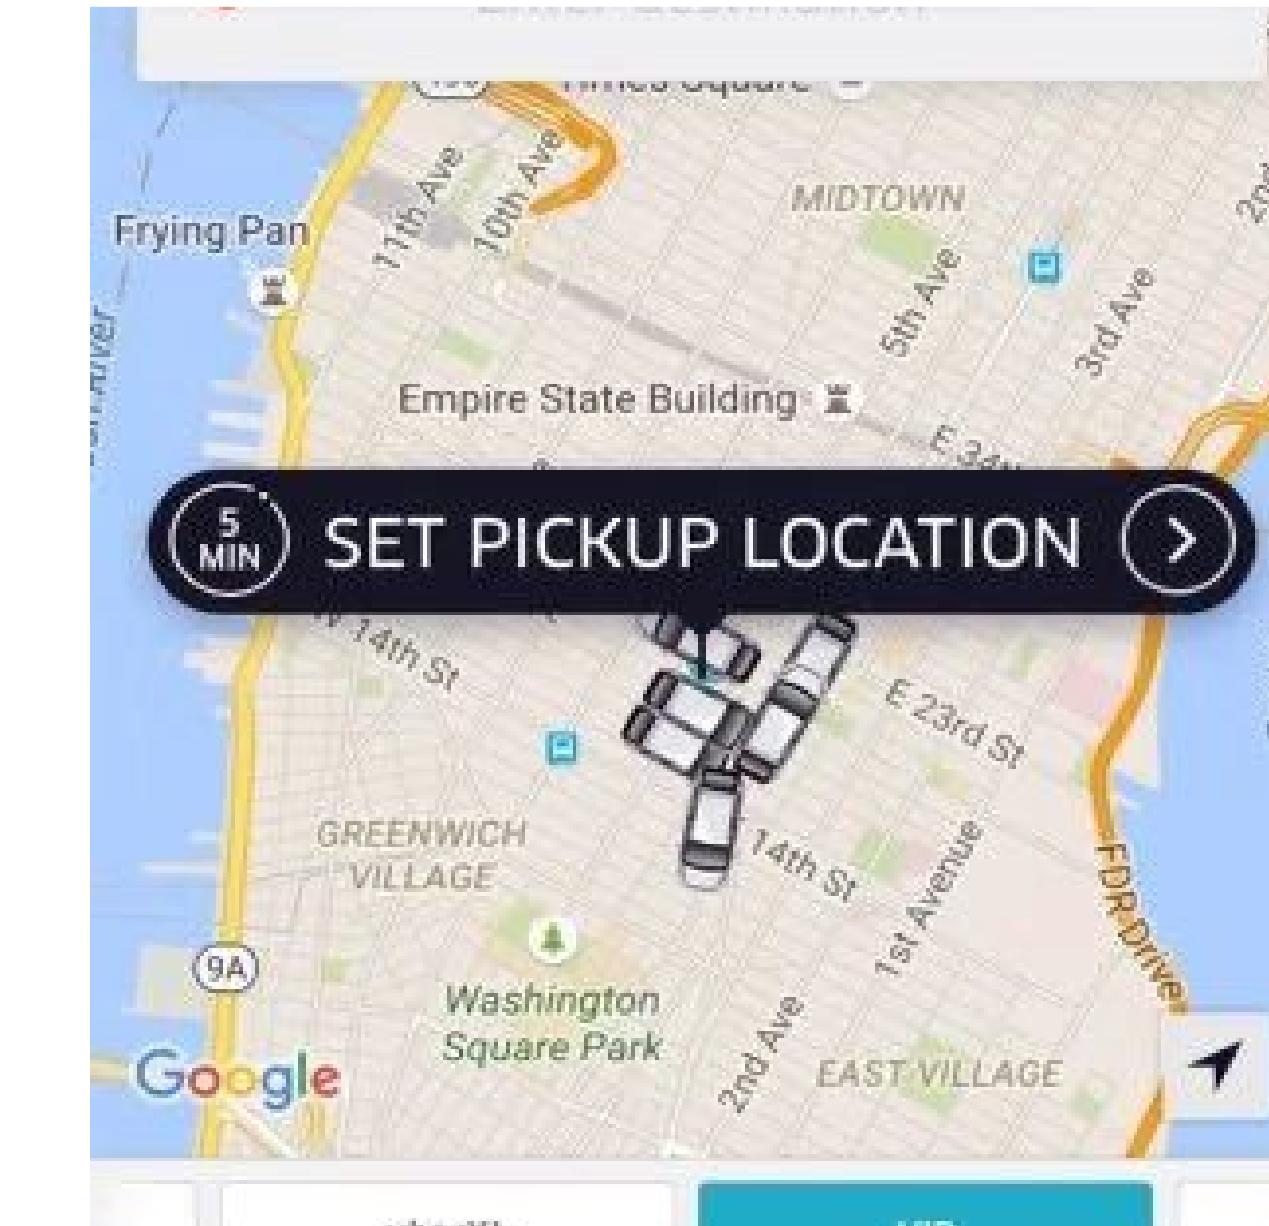
GeoSpatial UDF

Commonly used UDFs



within(trip_location, city_shape)

Find if a car is inside a city

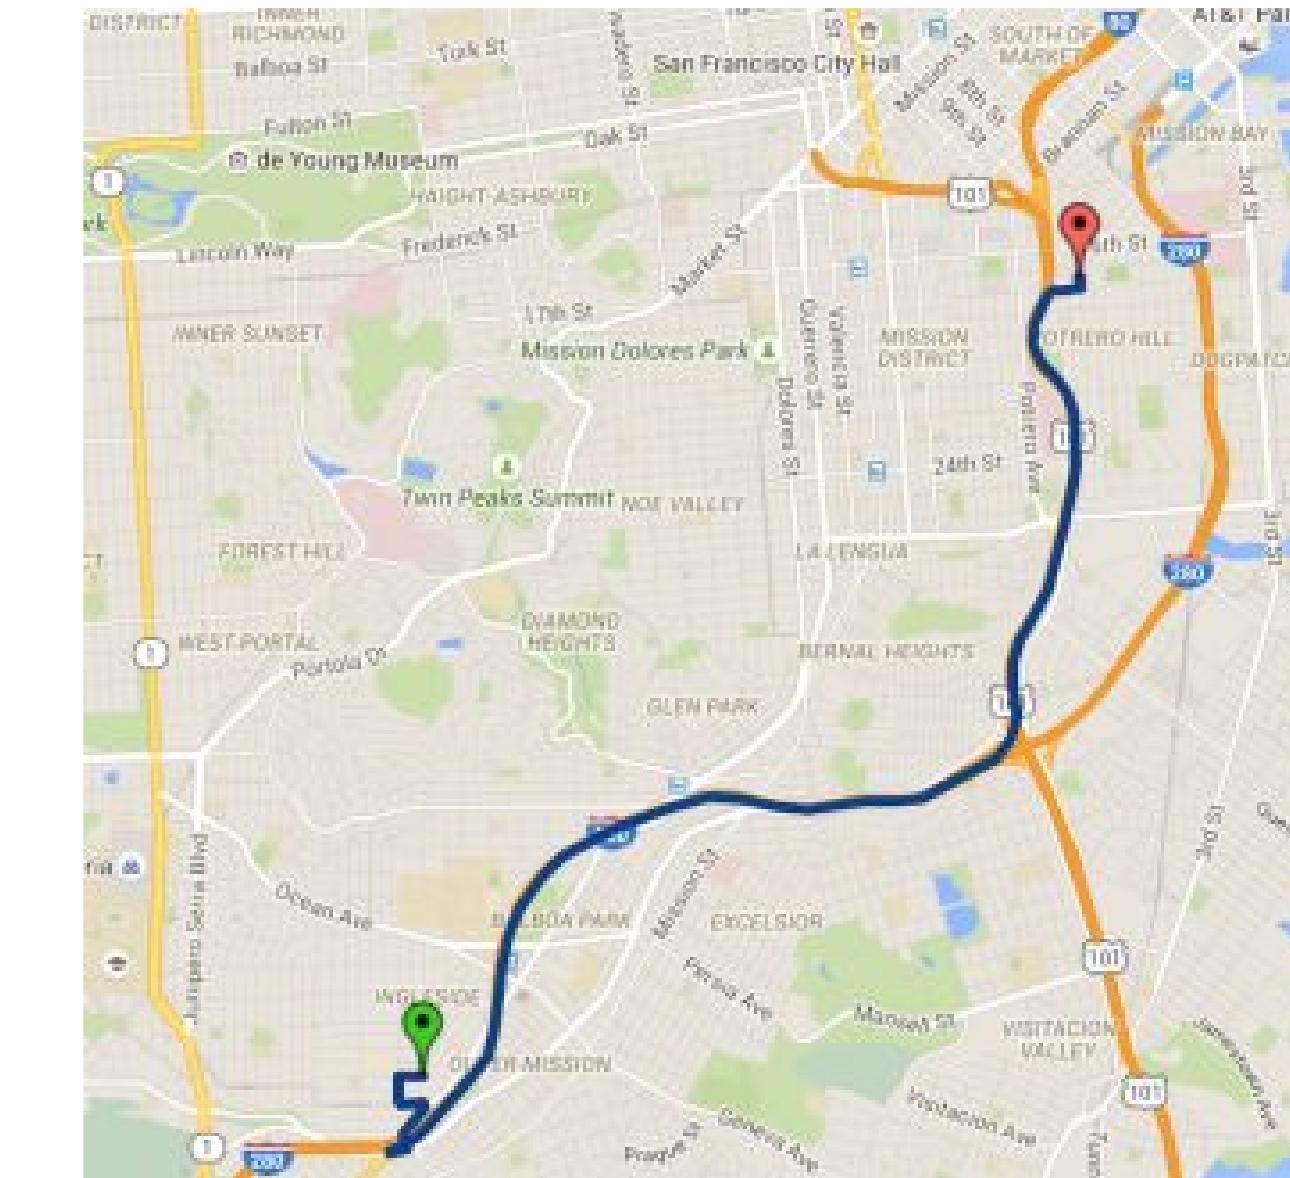
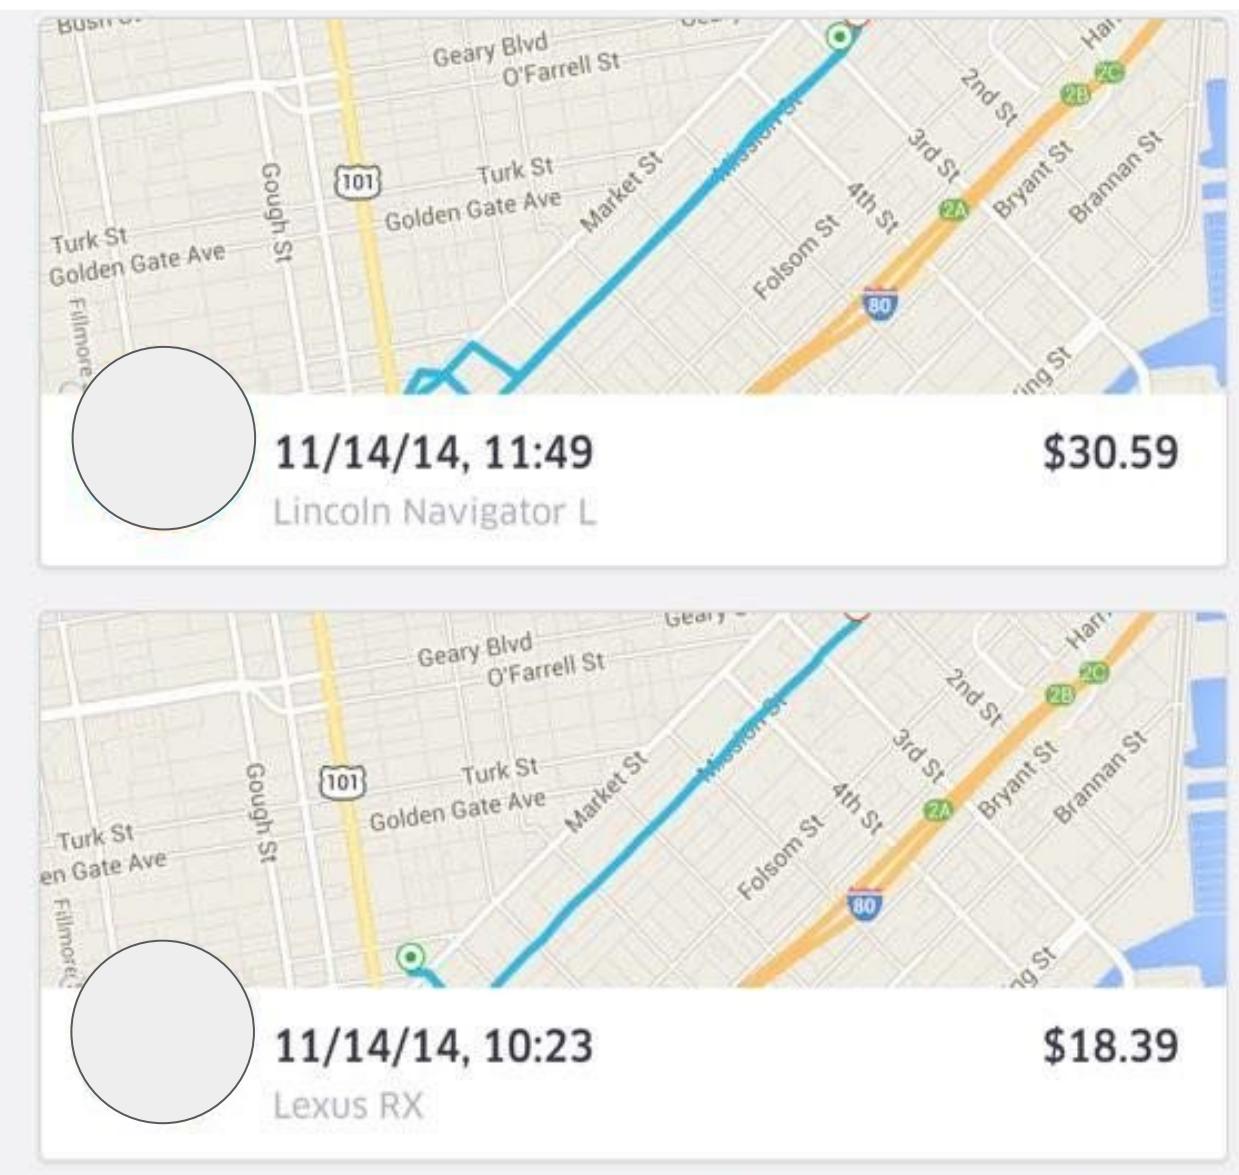


contains(geofence, auto_location)

Find all autos in one area

GeoSpatial UDF

Commonly used UDFs



overlaps(trip1, trip2)

Find trips that have similar routes

intersects(trip_location, gas_location)

Find all gas stations a trip route has passed by

Spatial Join

Common query at Uber

Objective: associate all trips with city_id for a single day.

SELECT trip.trip_id, city.city_id

FROM trip **JOIN** city

WHERE contains(city.city_shape, trip.start_location)

AND trip.datestr = '2016-06-07'

Spatial Join

Problem

It takes nearly ONE WEEK to run at Uber's data scale.

1. Spark does not have broadcast join optimization for non-equation join.
2. Not scalable, only one executor is used for cartesian join.

Spatial Join

Build a UDF to broadcast geo-spatial index

Spatial Join

Runtime Index Generation

1. Build Index

Index data is small but change often (city table)

Get fields from geo tables (city_id and city_shape)

Build QuadTree or RTree index at Spark Driver

Spatial Join

Executor Execution

2. Broadcast Index

UDF code is part of the Spark UDK jar.

⇒ `get_city_id(location)`, returns city_id of a location

Use the broadcasted spatial index for fast spatial retrieval

Spatial Join

Runtime UDF Generation

3. Rewrite Query (**2 mins only!** compared to 1 week before)

SELECT

trip_id, **get_city_id**(start_location)

FROM

trip

WHERE

datestr = '2016-06-07'

API

Engineers

SRE

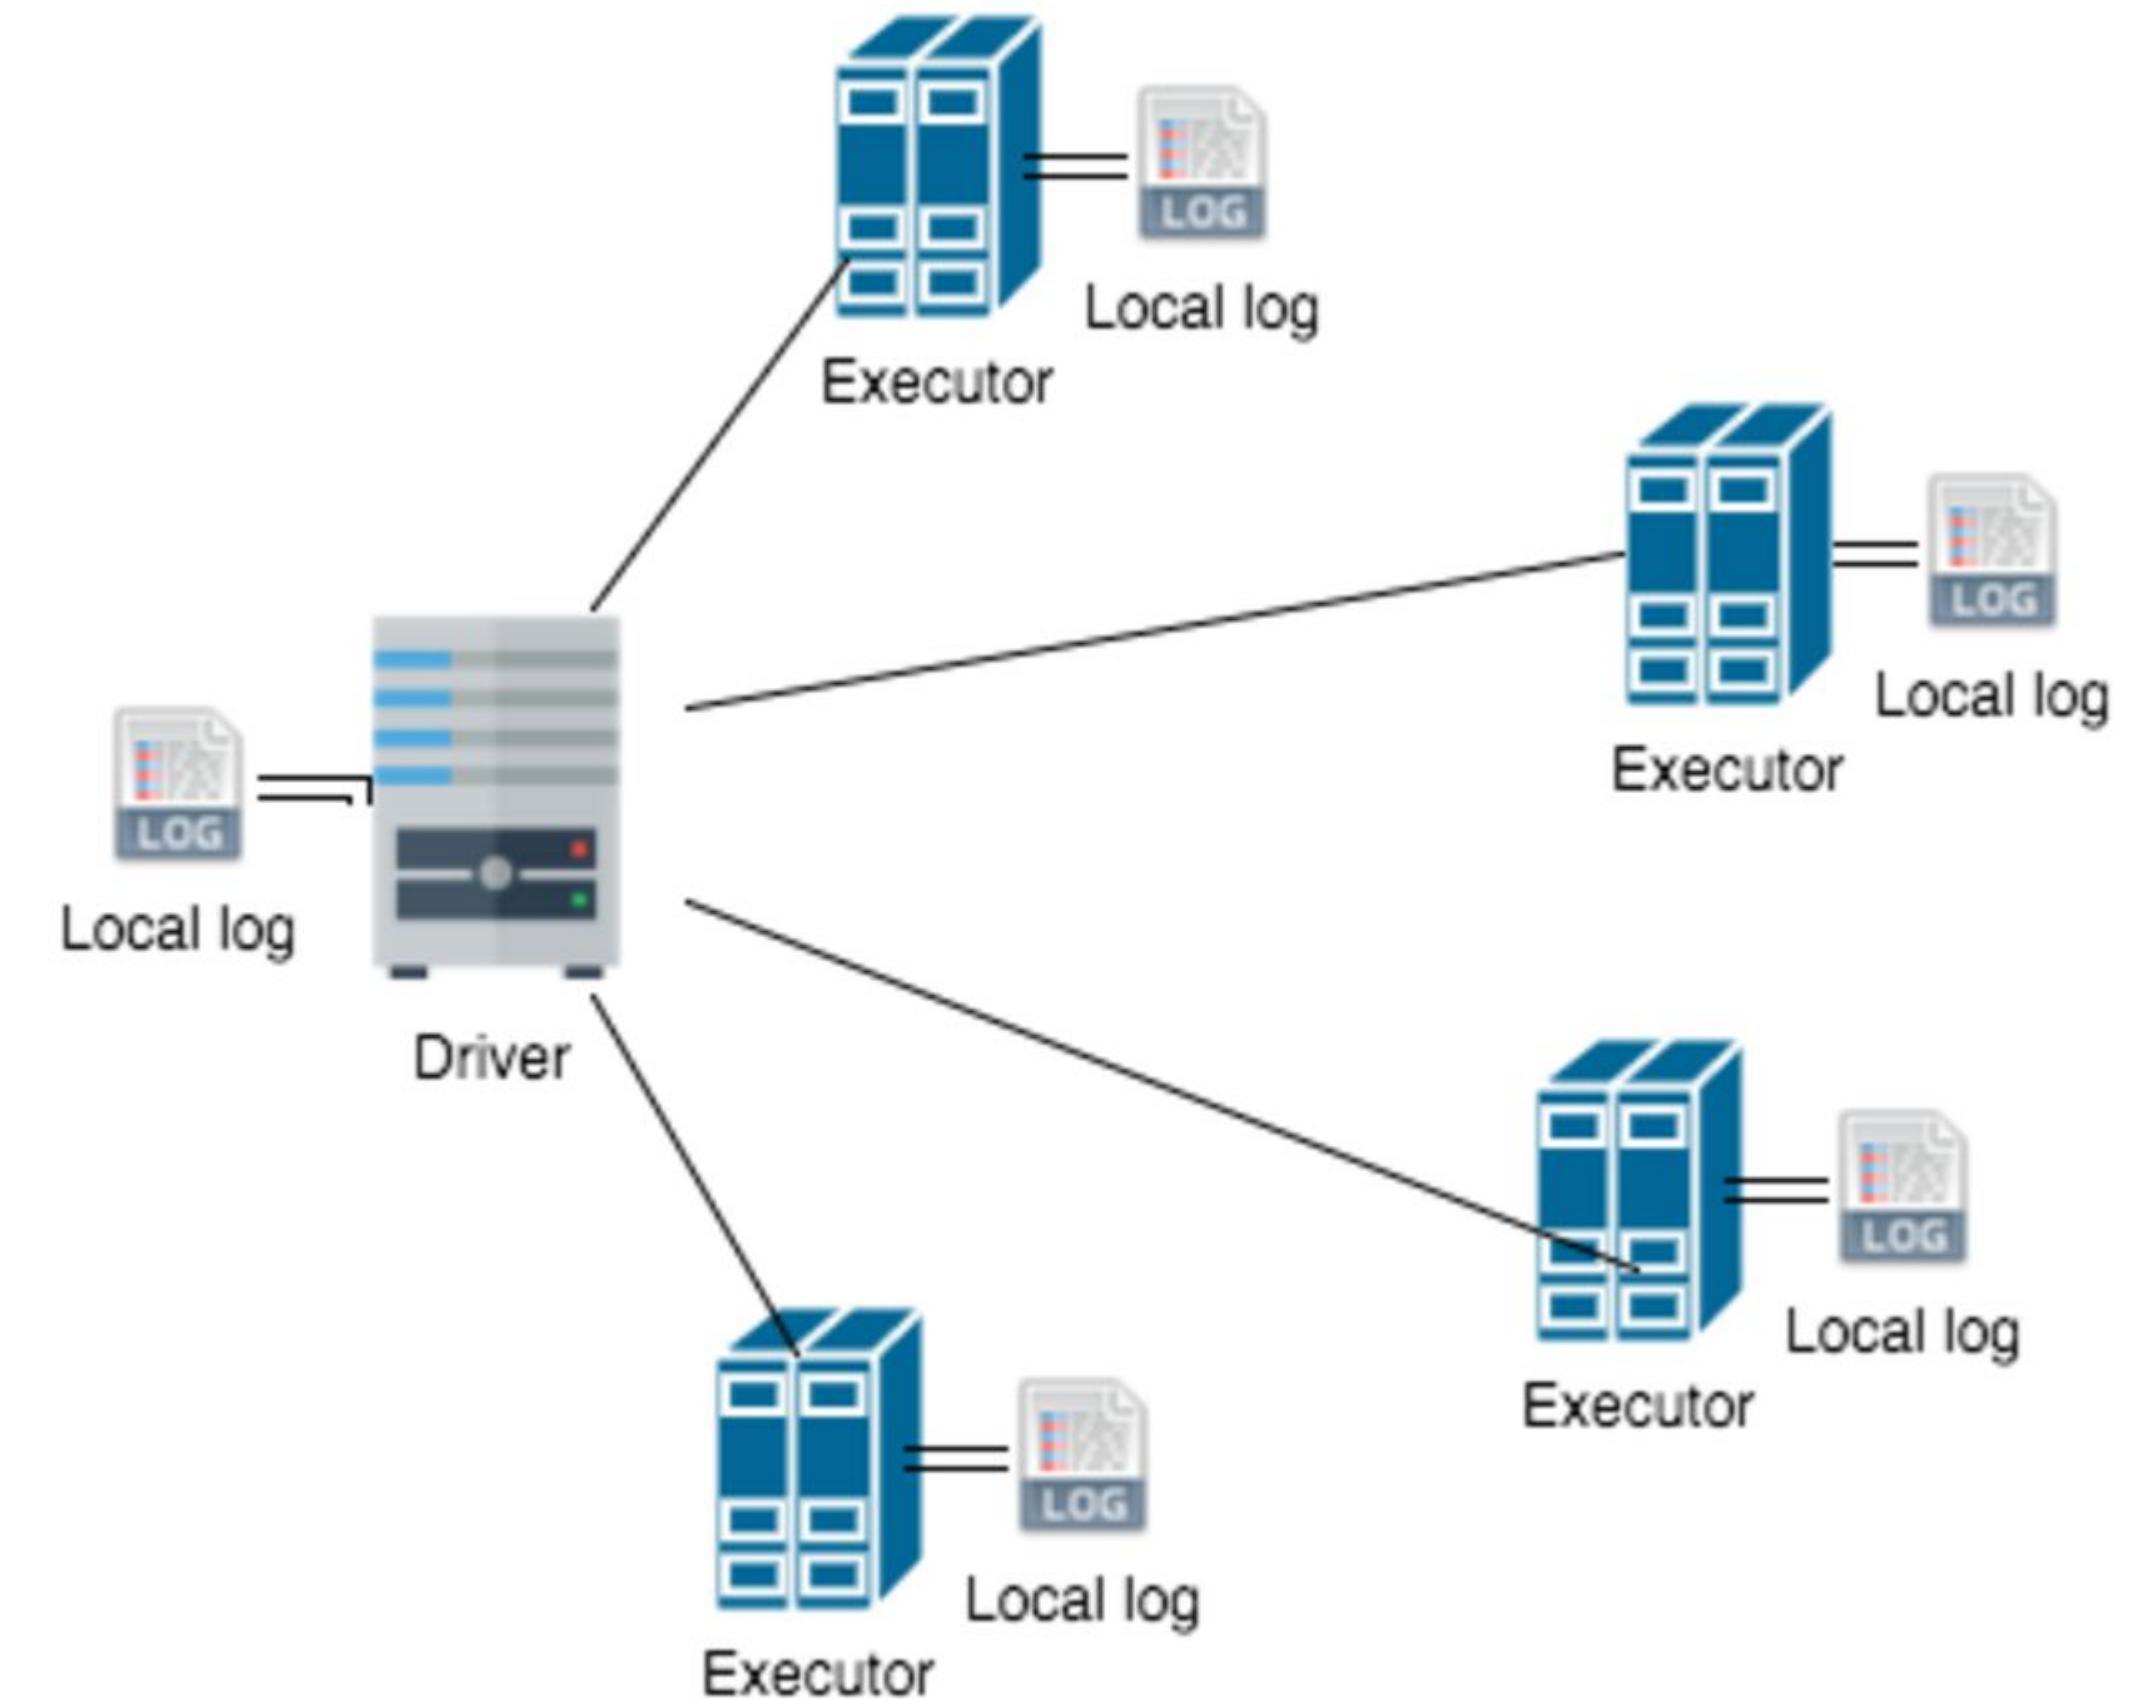
Tools

- Geo-spatial processing
- SCBuilder
- Kafka dispersal

- SparkChamber
- SparkPlug

Spark Debugging

- 1. Tons of local log files across many machines.**
- 2. Overall file size is huge and difficult to be handled by a single machine.**
- 3. Painful for debugging, which log is useful?**



Spark Chamber

Distributed Log Debugger for Spark

Interactive

Extend Spark Shell by Hooks.

Easy to adopt for Spark developers.

Spark Chamber Session

Welcome to



version 0.1

Using Scala version 2.10.5 (OpenJDK 64-Bit Server VM, Java 1.7.0_101)

Spark Chamber beta version 0.1

Maintained by Hadoop Compute Team, HipChat room: @Spark

```
scala> username
```

Username:

```
scala> applicationID
```

Application ID:

scala>

```
scala> allApplicationIds
```

Recent Spark Applications:

[]:	application_1463943621508_147799	InspectorGadget	2016-05-24T17:42:56.237GMT
[1]:	application_1463943621508_148157	InspectorGadget	2016-05-24T17:42:56.773GMT
[2]:	application_1463943621508_147498	InspectorGadget	2016-05-24T17:28:01.696GMT
[3]:	application_1463943621508_148089	InspectorGadget	2016-05-24T17:42:57.023GMT
[4]:	application_1463943621508_147842	InspectorGadget	2016-05-24T17:42:59.347GMT
[5]:	application_1463943621508_147798	InspectorGadget	2016-05-24T17:43:01.423GMT
[6]:	application_1463943621508_147589	InspectorGadget	2016-05-24T17:28:36.430GMT
[7]:	application_1463943621508_147805	InspectorGadget	2016-05-24T17:42:57.561GMT
[8]:	application_1463943621508_147845	InspectorGadget	2016-05-24T17:42:56.266GMT
[9]:	application_1463943621508_147937	InspectorGadget	2016-05-24T17:42:59.013GMT
[10]:	application_1463943621508_148051	InspectorGadget	2016-05-24T17:42:55.401GMT
[11]:	application_1463943621508_148117	InspectorGadget	2016-05-24T17:42:55.583GMT
[12]:	application_1463943621508_148160	InspectorGadget	2016-05-24T17:42:59.121GMT
[13]:	application_1463943621508_147984	InspectorGadget	2016-05-24T17:43:00.954GMT
[14]:	application_1463943621508_147962	InspectorGadget	2016-05-24T17:43:00.844GMT
[15]:	application_1463943621508_148114	InspectorGadget	2016-05-24T17:42:56.078GMT
[16]:	application_1463943621508_148142	InspectorGadget	2016-05-24T17:43:11.641GMT
[17]:	application_1463943621508_148097	InspectorGadget	2016-05-24T17:43:00.286GMT
[18]:	application_1463943621508_147777	InspectorGadget	2016-05-24T17:28:38.652GMT
[19]:	application_1463943621508_148106	InspectorGadget	2016-05-24T17:42:55.861GMT
[20]:	application_1463943621508_148045	InspectorGadget	2016-05-24T17:42:59.113GMT
[21]:	application_1463943621508_147989	InspectorGadget	2016-05-24T17:42:56.853GMT
[22]:	application_1463943621508_148002	InspectorGadget	2016-05-24T17:43:02.013GMT

```
scala> setApplicationId("application_1463943621508_147842")
```

```
scala> hosts
```

Hosts (ExecutorId, HostName, LogWebUI):

[0]: driver	hadoopworker348-sjc1	http://hadoopworker348-sjc1.pr...
[1]: executor(1)	hadoopworker347-sjc1	http://hadoopworker347-sjc1.pr...
[2]: executor(10)	hadoopworker313-sjc1	http://hadoopworker313-sjc1.pr...
[3]: executor(11)	hadoopcompute093-sjc1	http://hadoopcompute093-sjc1.pr...
[4]: executor(12)	hadoopworker490-sjc1	http://hadoopworker490-sjc1.pr...
[5]: executor(13)	hadoopworker373-sjc1	http://hadoopworker373-sjc1.pr...
[6]: executor(14)	hadoopworker255-sjc1	http://hadoopworker255-sjc1.pr...
[7]: executor(15)	hadoopworker542-sjc1	http://hadoopworker542-sjc1.pr...
[8]: executor(16)	hadoopworker549-sjc1	http://hadoopworker549-sjc1.pr...
[9]: executor(17)	hadoopworker311-sjc1	http://hadoopworker311-sjc1.pr...
[10]: executor(18)	hadoopworker437-sjc1	http://hadoopworker437-sjc1.pr...
[11]: executor(19)	hadoopworker288-sjc1	http://hadoopworker288-sjc1.pr...
[12]: executor(2)	hadoopcompute034-sjc1	http://hadoopcompute034-sjc1.pr...
[13]: executor(20)	hadoopcompute032-sjc1	http://hadoopcompute032-sjc1.pr...
[14]: executor(21)	hadoopcompute018-sjc1	http://hadoopcompute018-sjc1.pr...

scala>

```
scala> search("data count:")
```

Search result for data count:

```
[0]: (309)@hadoopcompute050-sjc1/executor(4) 16/05/28 01:49:05 INFO SparkChamberExample: Partition: 2, data count: 5980
[1]: (311)@hadoopcompute050-sjc1/executor(4) 16/05/28 01:49:05 INFO SparkChamberExample: Partition: 3, data count: 654245
[2]: (314)@hadoopworker275-sjc1/executor(1) 16/05/28 01:49:05 INFO SparkChamberExample: Partition: 4, data count: 4539
[3]: (316)@hadoopworker275-sjc1/executor(1) 16/05/28 01:49:05 INFO SparkChamberExample: Partition: 5, data count: 6777
[4]: (321)@hadoopworker305-sjc1/executor(2) 16/05/28 01:49:05 INFO SparkChamberExample: Partition: 0, data count: 6434
[5]: (323)@hadoopworker305-sjc1/executor(2) 16/05/28 01:49:05 INFO SparkChamberExample: Partition: 1, data count: 5432
[6]: (328)@hadoopworker419-sjc1/executor(3) 16/05/28 01:49:05 INFO SparkChamberExample: Partition: 6, data count: 4887
[7]: (330)@hadoopworker419-sjc1/executor(3) 16/05/28 01:49:05 INFO SparkChamberExample: Partition: 7, data count: 6012
[8]: (342)@hadoopworker275-sjc1/executor(1) 16/05/28 01:49:55 INFO SparkChamberExample: Partition: 2, data count: 5980
[9]: (344)@hadoopworker275-sjc1/executor(1) 16/05/28 01:49:55 INFO SparkChamberExample: Partition: 3, data count: 654245
[10]: (408)@hadoopworker419-sjc1/executor(3) 16/05/28 01:50:44 INFO SparkChamberExample: Partition: 2, data count: 5980
[11]: (410)@hadoopworker419-sjc1/executor(3) 16/05/28 01:50:44 INFO SparkChamberExample: Partition: 3, data count: 654245
[12]: (480)@hadoopworker305-sjc1/executor(2) 16/05/28 01:51:45 INFO SparkChamberExample: Partition: 2, data count: 5980
[13]: (482)@hadoopworker305-sjc1/executor(2) 16/05/28 01:51:45 INFO SparkChamberExample: Partition: 3, data count: 654245
[14]: (923)@hadoopcompute094-sjc1/executor(2) 16/05/28 01:52:47 INFO SparkChamberExample: Partition: 6, data count: 4887
[15]: (925)@hadoopcompute094-sjc1/executor(2) 16/05/28 01:52:47 INFO SparkChamberExample: Partition: 7, data count: 6012
[16]: (941)@hadoopworker303-sjc1/executor(4) 16/05/28 01:52:47 INFO SparkChamberExample: Partition: 0, data count: 6434
[17]: (943)@hadoopworker303-sjc1/executor(4) 16/05/28 01:52:47 INFO SparkChamberExample: Partition: 1, data count: 5432
[18]: (952)@hadoopworker456-sjc1/executor(1) 16/05/28 01:52:47 INFO SparkChamberExample: Partition: 2, data count: 5980
[19]: (954)@hadoopworker456-sjc1/executor(1) 16/05/28 01:52:47 INFO SparkChamberExample: Partition: 3, data count: 654245
[20]: (961)@hadoopworker490-sjc1/executor(3) 16/05/28 01:52:47 INFO SparkChamberExample: Partition: 4, data count: 4539
[21]: (963)@hadoopworker490-sjc1/executor(3) 16/05/28 01:52:47 INFO SparkChamberExample: Partition: 5, data count: 6777
[22]: (983)@hadoopworker303-sjc1/executor(4) 16/05/28 01:54:36 INFO SparkChamberExample: Partition: 2, data count: 5980
[23]: (985)@hadoopworker303-sjc1/executor(4) 16/05/28 01:54:36 INFO SparkChamberExample: Partition: 3, data count: 654245
[24]: (1059)@hadoopworker490-sjc1/executor(3) 16/05/28 01:56:25 INFO SparkChamberExample: Partition: 2, data count: 5980
[25]: (1061)@hadoopworker490-sjc1/executor(3) 16/05/28 01:56:25 INFO SparkChamberExample: Partition: 3, data count: 654245
[26]: (1139)@hadoopworker341-sjc1/executor(6) 16/05/28 01:57:11 INFO SparkChamberExample: Partition: 2, data count: 5980
[27]: (1141)@hadoopworker341-sjc1/executor(6) 16/05/28 01:57:11 INFO SparkChamberExample: Partition: 3, data count: 654245
```

Spark Chamber

Distributed Log Debugger for Spark

Features

1. Get all recent Spark Application IDs.
2. Get first exception, all exceptions grouped by types sorted by time, etc.
3. Display CPU, memory, I/O metrics.
4. Dive into a specific driver/executor/machine
5. Search

Spark Chamber

Distributed Log Debugger for Spark

Security

Developer mode: debug developer's own Spark job.

SRE mode: view and check all users' Spark job information.

Spark Chamber

Enable Yarn Log Aggregation

Home / tmp / logs / username / logs

	Name	Size	User	Group
	..		username	hadoop
	application_1464368688010_12439		username	hadoop
	application_1464368688010_12462		username	hadoop
	application_1464368688010_13128		username	hadoop
	application_1464368688010_14241		username	hadoop
	application_1464368688010_14323		username	hadoop

Home / tmp / logs / username / logs / application_1464368688010_12462

	Name	Size	User	Group
	..		username	hadoop
	hadoopworker176-sjc1	12.5 KB	username	hadoop
	hadoopworker230-sjc1	15.6 KB	username	hadoop
	hadoopworker255-sjc1	58.7 KB	username	hadoop
	hadoopworker302-sjc1	12.5 KB	username	hadoop
	hadoopworker437-sjc1	12.5 KB	username	hadoop

YARN aggregates log files on HDFS

All application IDs of the same user are under same place.

Files are named after host names

One machine has one log file, regardless of # executors on that machine.

Spark Chamber

[Use Spark to debug Spark](#)

Extend the Spark Shell by Hooks:

1. For ONE application Id, distribute log files to different executors.
2. Extract each lines and save into DataFrame.
3. Sort log dataframe by time and hostname.
4. Retrieve target log via SparkSQL DataFrame APIs.

API

Engineers

SRE

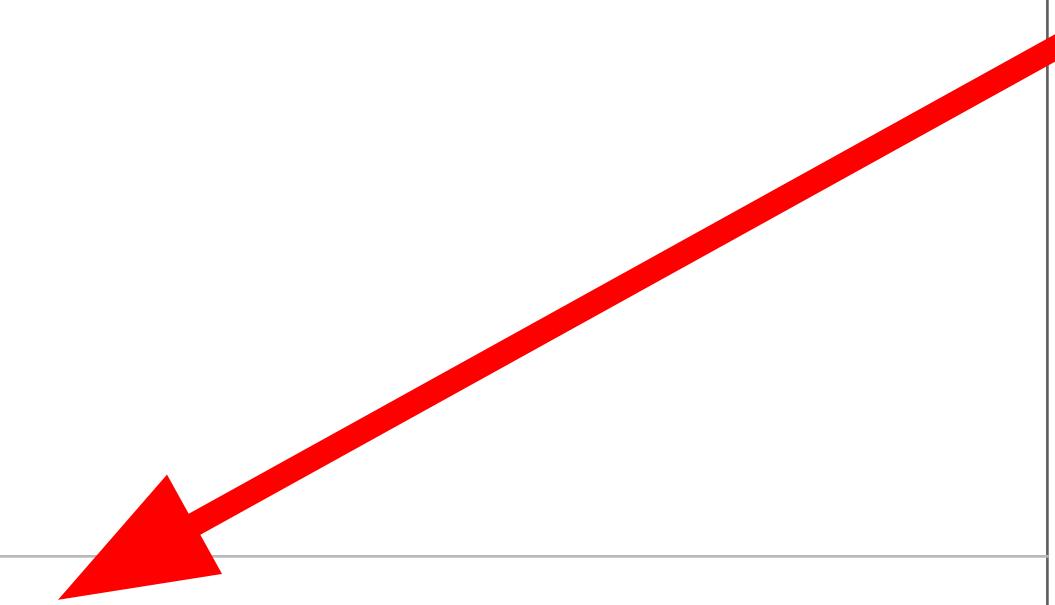
Tools

Future
Work

- Geo-spatial processing
- SCBuilder
- Kafka dispersal

- SparkChamber
- SparkPlug

- SparkChamber



Spark Chamber

SRE version - Cluster wide insights

- Dimensions - Jobs
 - All
 - Single team
 - Single engineer
- Dimensions - Time
 - Last month, week, day
- Dimensions - Hardware
 - Specific rack, pod

Spark Chamber

SRE version - Analytics and Machine Learning

- Analytics
 - Resource auditing
 - Data access auditing
- Machine Learning
 - Failures diagnostics
 - Malicious jobs detection
 - Performance optimization

Future Work

Engineers

SRE

API

Tools

- Geo-spatial processing
- SCBuilder
- Kafka dispersal
- Hive table registration (Didn't cover today)
- Incremental processing (Didn't cover today)
- Debug logging
- Metrics
- Configurations
- Data Freshness

- SparkChamber
- SparkPlug
- Unit testing (Didn't cover today)
- Oozie integration (Didn't cover today)

- Resource usage

- SparkChamber
- Resource usage auditing
- Data access auditing
- Machine learning on jobs

SPARK: INTERACTIVE TO PRODUCTION

Today, Tuesday, June 7

4:50 PM – 5:20 PM

Room: Ballroom B

Dara Adib, Uber

Locality Sensitive Hashing by Spark

Tomorrow, Wednesday, June 8

5:25 PM – 5:55 PM

Room: Imperial

Alain Rodriguez, Fraud Platform, Uber

Kelvin Chu, Hadoop Platform, Uber

Thank you

UBER