



Apache Ignite and Apache Spark

Where Fast Data Meets the IoT

Denis Magda
GridGain Product Manager
Apache Ignite PMC

<http://ignite.apache.org>



#apacheignite



#denismagda

GridGain

Agenda

- IoT Demands to Software
- IoT Software Stack
 - Device OS/RTOS
 - Data Collection and Enrichment
 - HTAP Platform
 - Application APIs
- Demo

IoT Demands to Software

Apache®, Apache Ignite, Ignite®, and the Apache Ignite logo are either registered trademarks or trademarks of the Apache Software Foundation in the United States and/or other countries.



IoT Demands to Software

- Real-time ingestion
- Real-time processing
- Time-series support
- Flexible Querying APIs
 - SQL
 - Full-text search
 - Geo-spatial
- Analytics
 - BI
 - Machine Learning
- High-availability
- Simple scalability



IoT Software Stack



Apache®, Apache Ignite, Ignite®, and the Apache Ignite logo are either registered trademarks or trademarks of the Apache Software Foundation in the United States and/or other countries.

IoT Software Stack

Application APIs

HTAP Platform

Data Collection and Enrichment

Device OS / Real-Time OS

Apache IoT Software Stack

Application APIs



HTAP Platform



Data Collection and Enrichment



Device OS / Real-Time OS



Device OS/RTOS

Apache®, Apache Ignite, Ignite®, and the Apache Ignite logo are either registered trademarks or trademarks of the Apache Software Foundation in the United States and/or other countries.



Apache MyNewt

- Open Source RTOS
 - Cortex M0-M4
 - MIPS & RISC-V
- Networking
 - Bluetooth Low Energy
 - Wi-Fi
 - TCP/IP & UDP
- Build & Package Management
- Secure bootloader and signed images
- Remote Firmware Upgrade



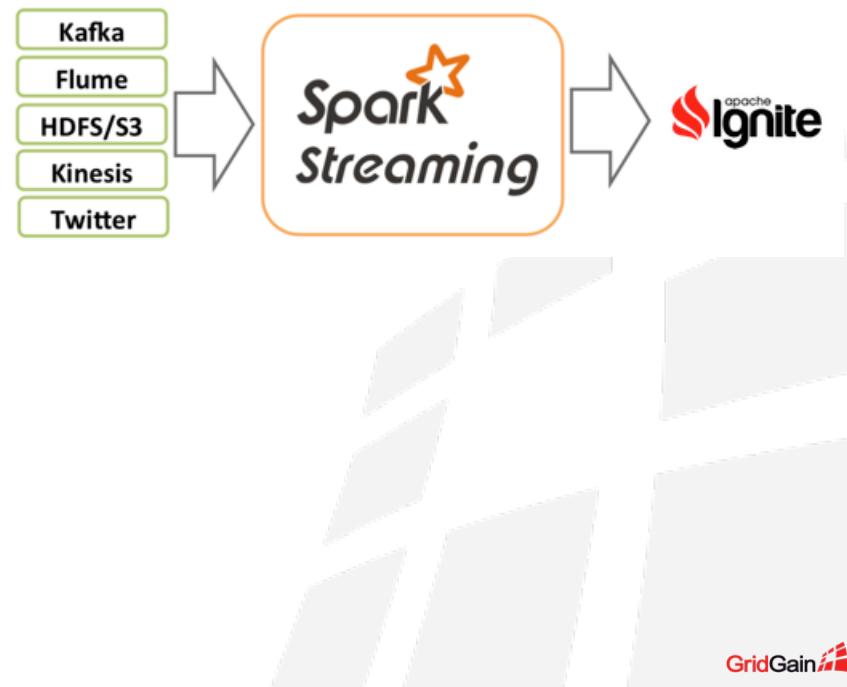
Data Collection and Enrichment

Apache®, Apache Ignite, Ignite®, and the Apache Ignite logo are either registered trademarks or trademarks of the Apache Software Foundation in the United States and/or other countries.



Data Collection and Enrichment

- Spark Streaming
 - Fault-Tolerant Streams Processing
- Data Collection
 - Sockets
 - Kafka
 - Flume
- Data Enrichment
 - Spark API
- Data Storage
 - Apache Ignite



HTAP Platform



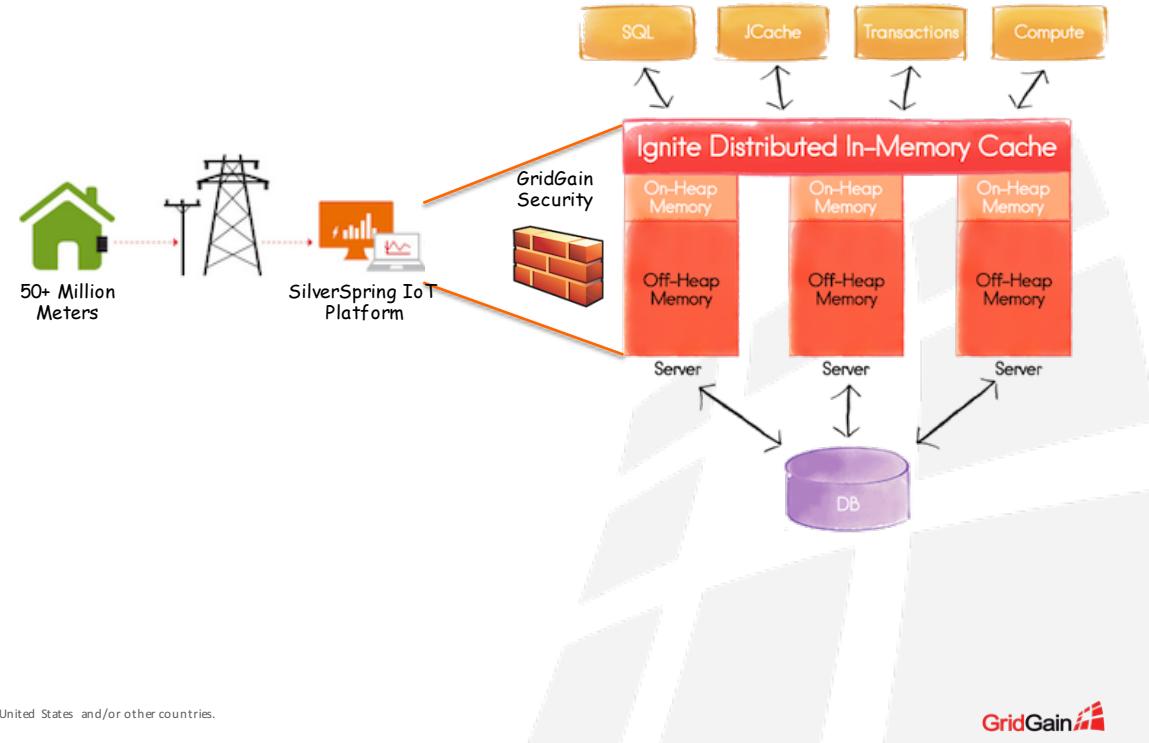
Apache®, Apache Ignite, Ignite®, and the Apache Ignite logo are either registered trademarks or trademarks of the Apache Software Foundation in the United States and/or other countries.



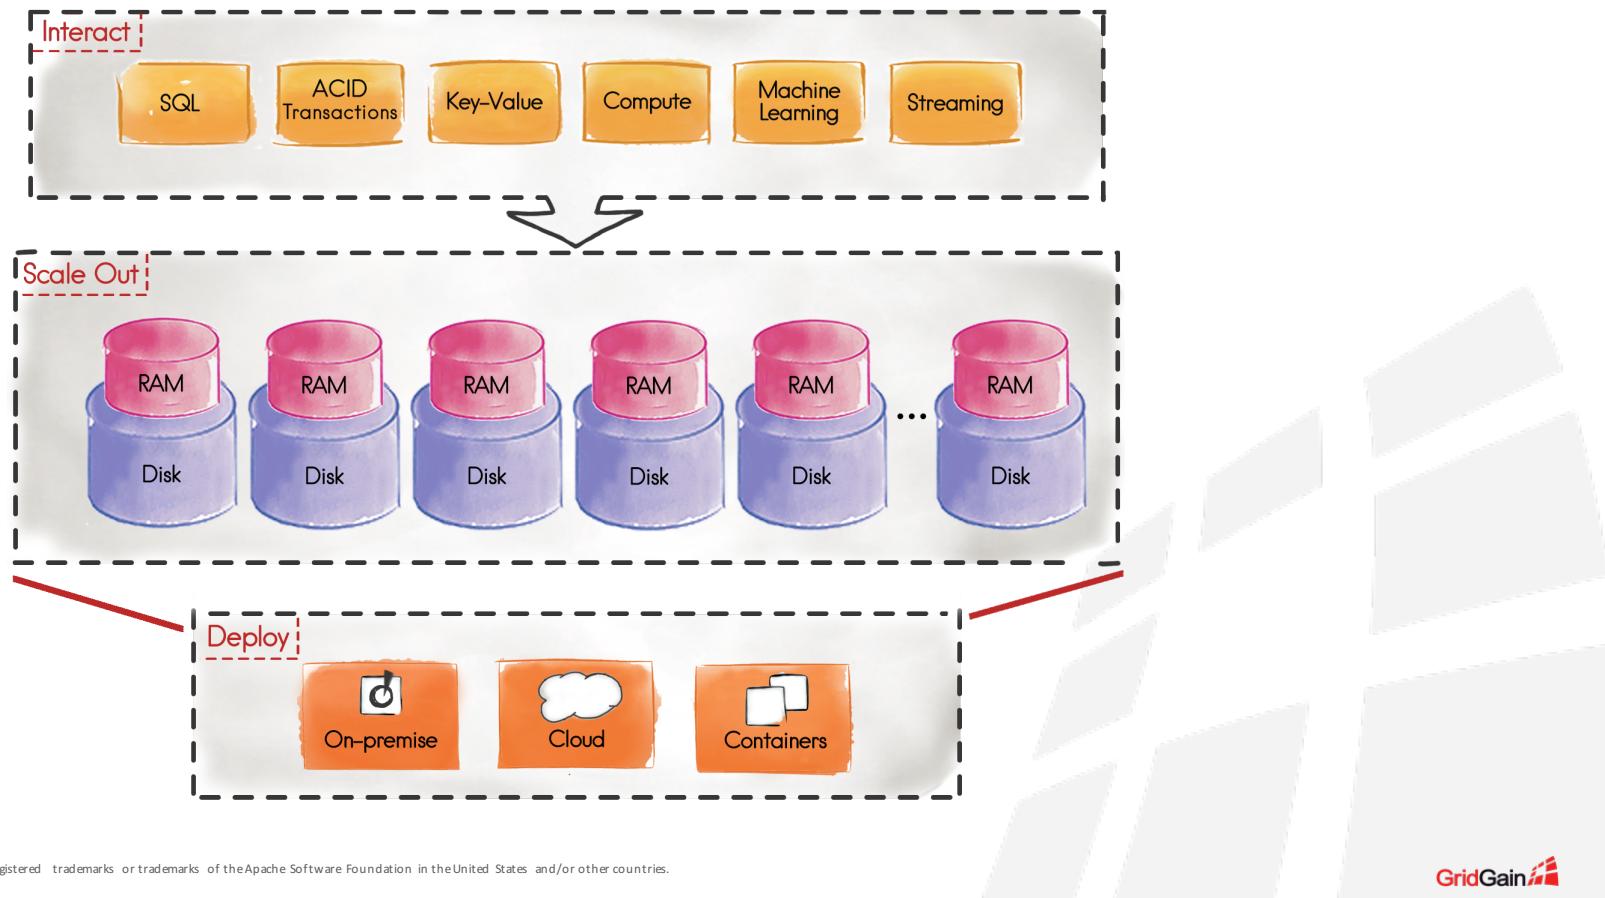
Smart Metering and Utilities – delivers a comprehensive IoT platform

Use Case:

- SilverSpring Requirements
 - Migrate to in-memory processing
 - Add scalability & elasticity
 - Use open source technologies
- SilverSpring Solution
 - SaaS Architecture
 - Multi-Tenancy
 - Advanced Security
 - Strong compute capabilities
 - Co-located in-memory processing
 - Demonstrated best
 - On-demand elasticity & scalability
 - ANSI-99 SQL Support
 - Transactional consistency



Apache Ignite HTAP Platform

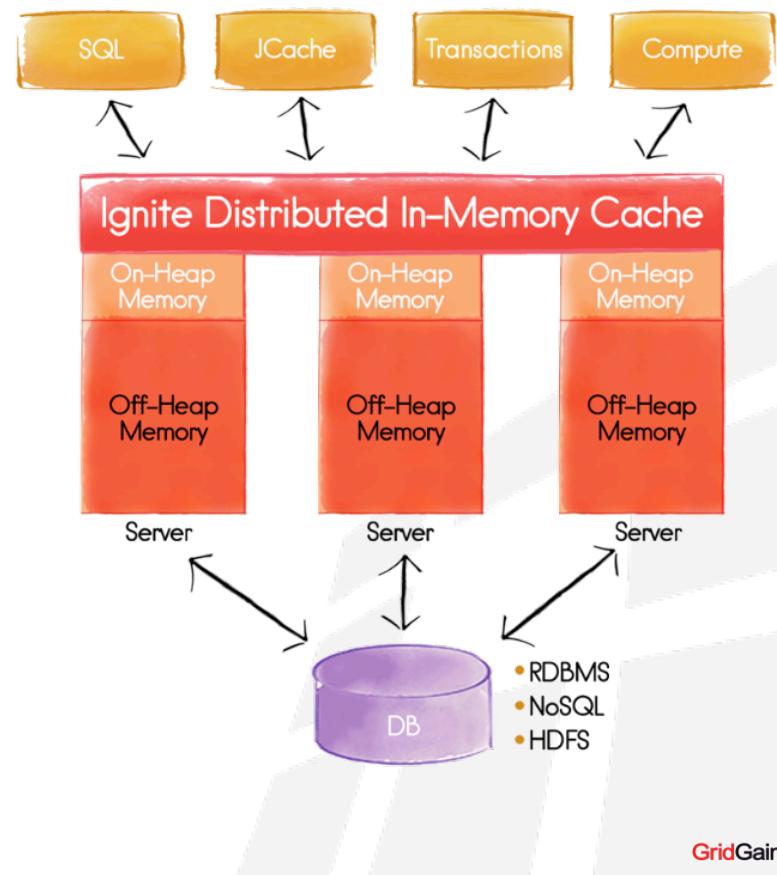


Apache®, Apache Ignite, Ignite®, and the Apache Ignite logo are either registered trademarks or trademarks of the Apache Software Foundation in the United States and/or other countries.

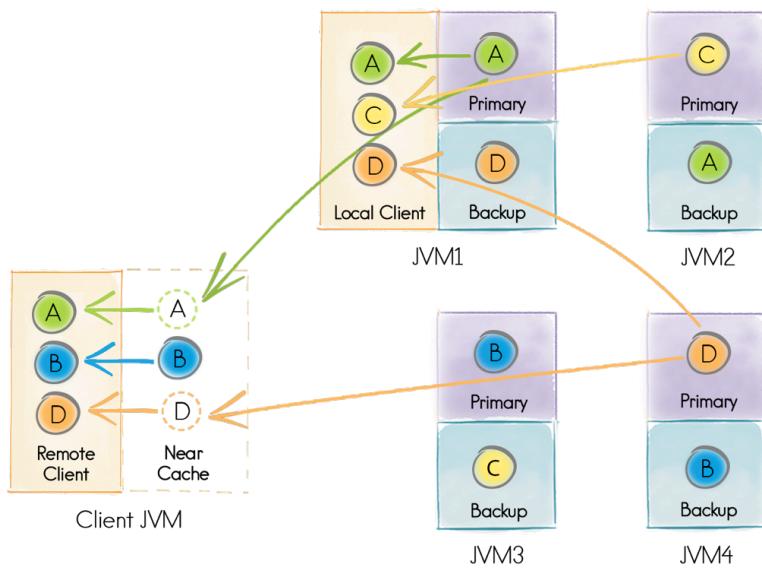
GridGain

Key-Value Data Grid

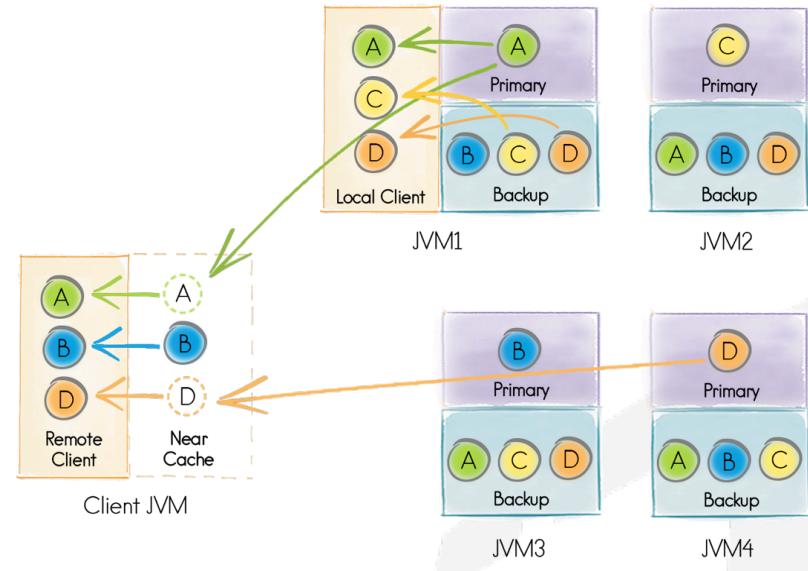
- Distributed Key-Value Data Store
- Data Reliability
- High-Availability
 - Active replicas, automatic failover
- Data Consistency
 - ACID distributed transactions



In-Memory Data Grid: Partitioning



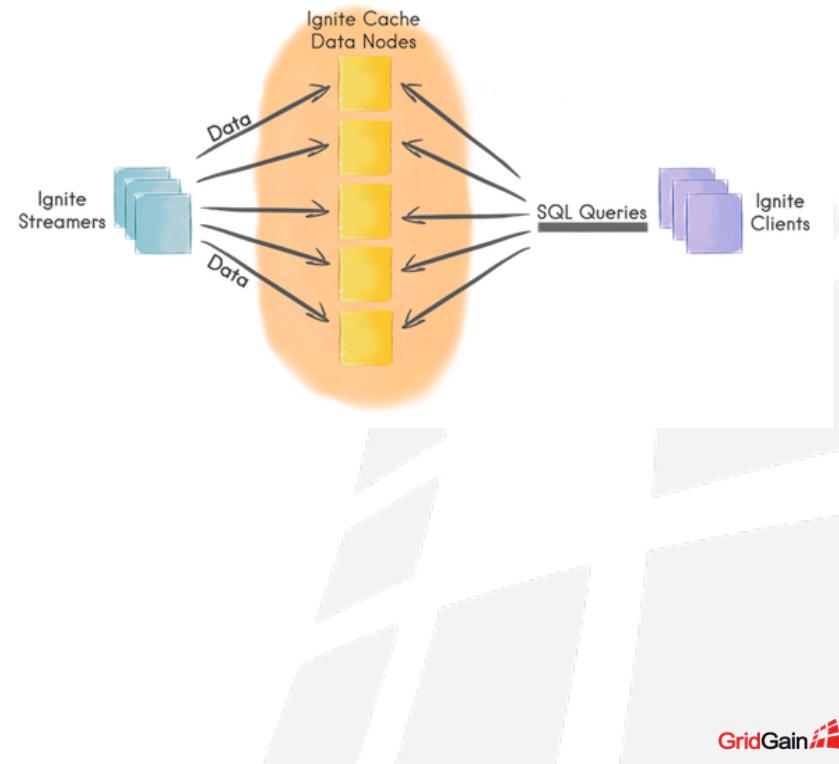
Partitioned Cache



Replicated Cache

Streaming to Ignite

- Ignite Data Streamer
 - Fastest Data Injection
 - Automatic data partitioning
- Stream Receivers
 - Custom logic execution
 - Additional data transformation
- Stream Adapter
 - Socket
 - Kafka
 - Flink
 - RocketMQ
 - etc.



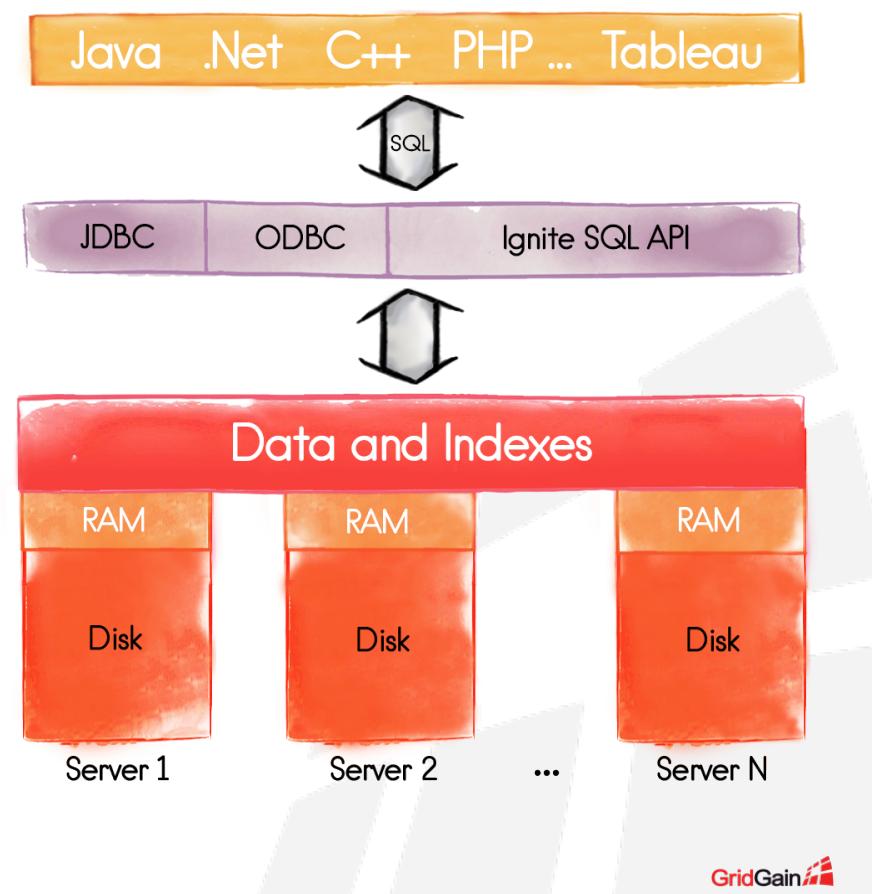
Application APIs

Apache®, Apache Ignite, Ignite®, and the Apache Ignite logo are either registered trademarks or trademarks of the Apache Software Foundation in the United States and/or other countries.



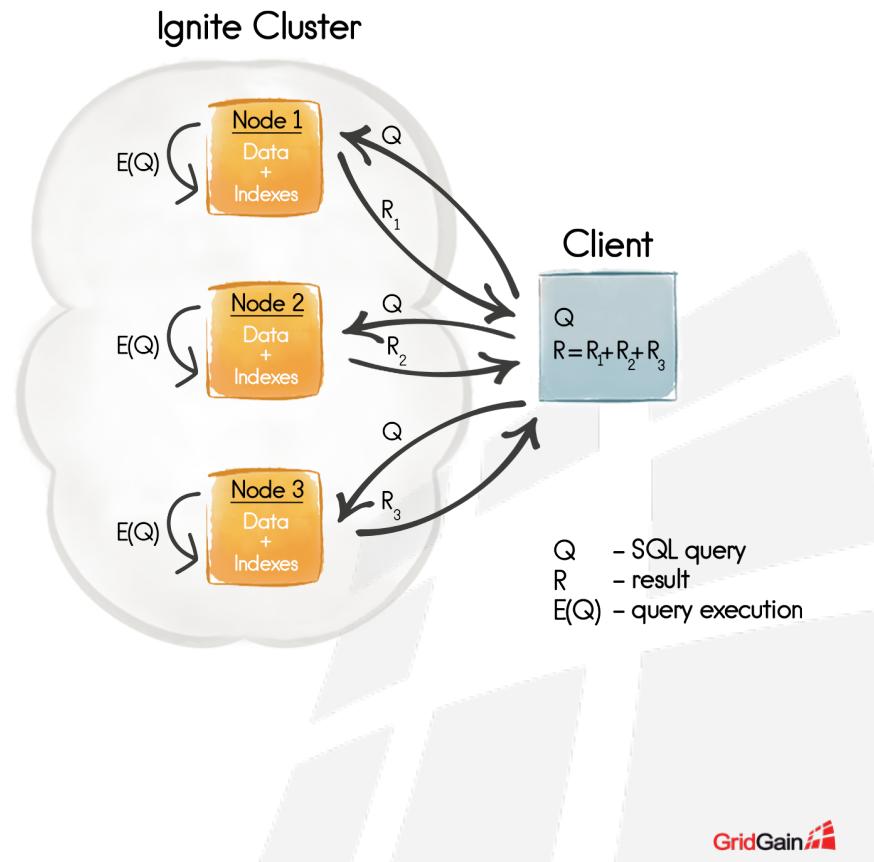
Apache Ignite SQL APIs

- ANSI-99 SQL
- Geo-spatial Queries
- Full-text Search
- Always Consistent
- Fault Tolerant
- Cross-Platform
 - JDBC and ODBC drivers
 - DML (INSERT, UPDATE, etc.)
 - DDL (CREATE, DROP, etc.)



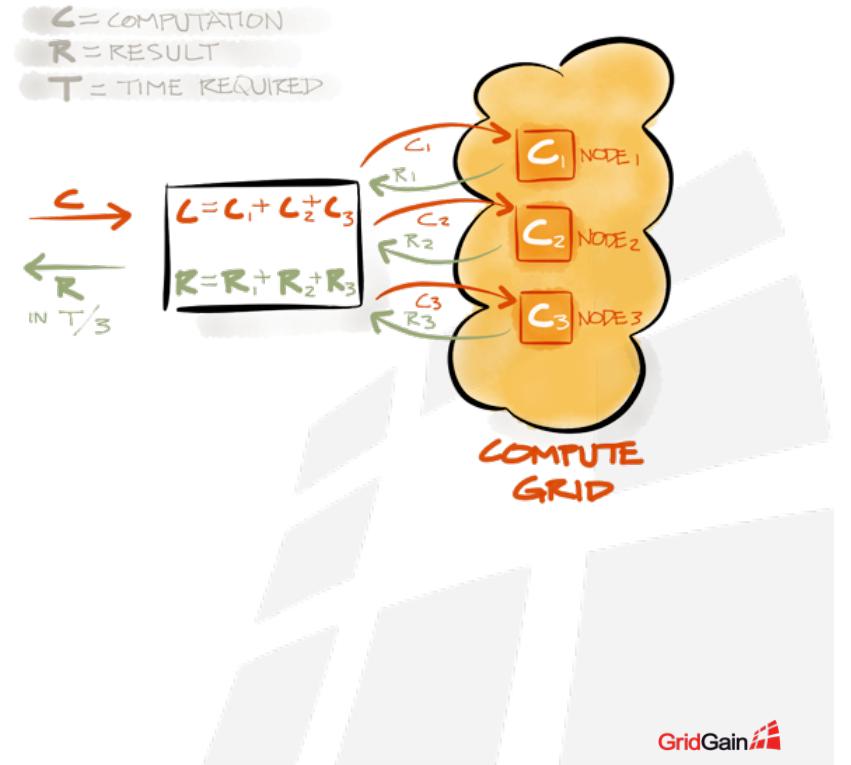
Apache Ignite SQL Grid: Queries

- Distributed Joins
- Automatic Group By, Aggregations, Sorting
- Cross-Cache Joins, Unions, etc.
- Ad-Hoc SQL Support



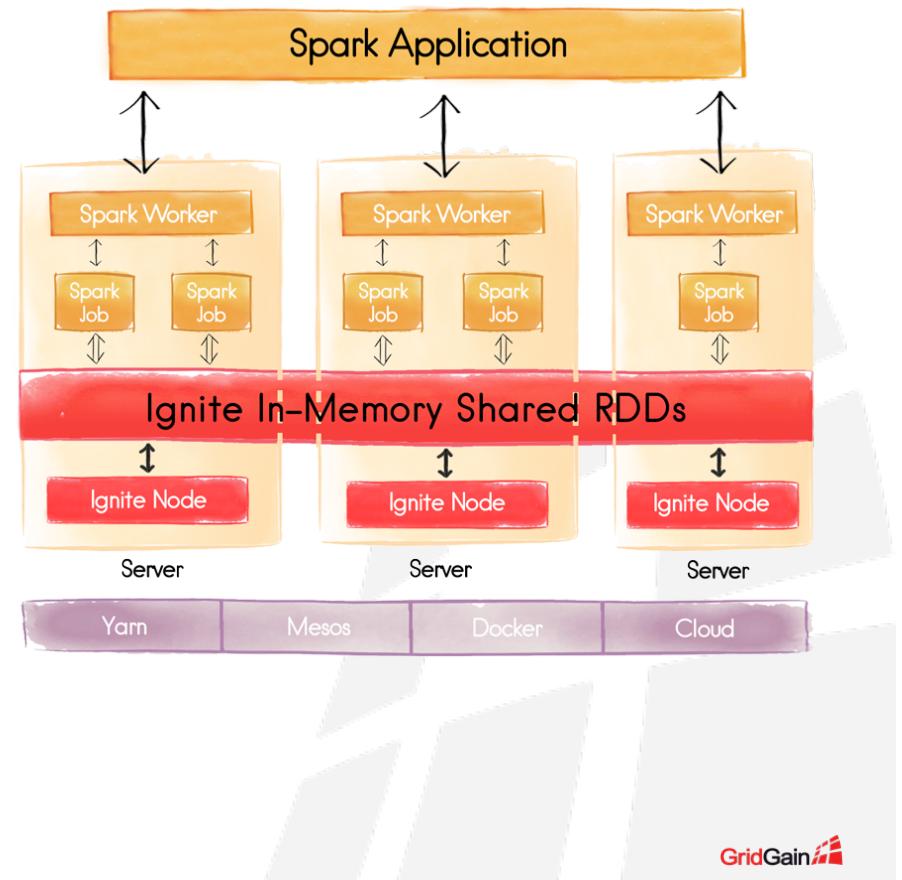
Apache Ignite Compute Grid

- Direct API for MapReduce
- Direct API for ForkJoin
- Zero Deployment
- State Checkpoints
- Load Balancing
- Automatic Failover

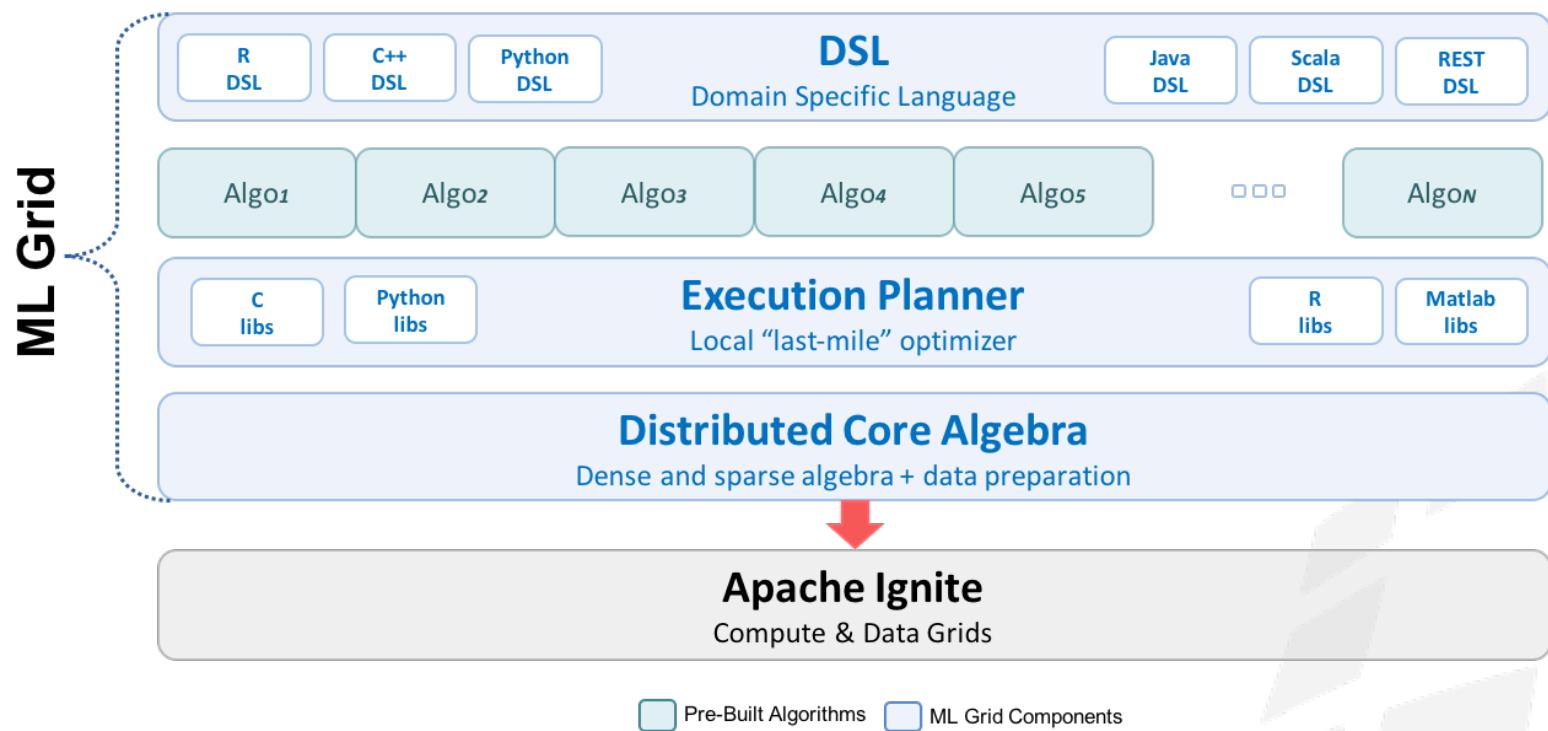


Spark Shared RDDs

- IgniteRDD
 - Share RDD across jobs on the host
 - Share RDD across jobs in the application
 - Share RDD globally
- Faster SQL
 - In-Memory Indexes
 - SQL on top of Shared RDD



Machine Learning Grid



Apache®, Apache Ignite, Ignite®, and the Apache Ignite logo are either registered trademarks or trademarks of the Apache Software Foundation in the United States and/or other countries.



Demo

Apache®, Apache Ignite, Ignite®, and the Apache Ignite logo are either registered trademarks or trademarks of the Apache Software Foundation in the United States and/or other countries.



Resources

- Documentation:
 - Apache Ignite: <https://apacheignite.readme.io/docs>
 - Apache Spark Streaming: <http://spark.apache.org/streaming/>
 - Ignite and Spark Integration: <https://apacheignite-fs.readme.io/docs/ignite-for-spark>
 - Apache MyNewt: <http://mynewt.apache.org>
- Demo Source Code:
 - <https://github.com/dmagda/IgniteSparkIoT>



ANY QUESTIONS?

Thank you for joining us. Follow the conversation.

<http://ignite.apache.org>



#apacheignite



#denismagda